Oracle Berkeley DB

Berkeley DB
API Reference
for C++

Release 4.8

ORACLE

BERKELEY DB

Legal Notice

This documentation is distributed under an open source license. You may review the terms of this license at:
http://www.oracle.com/technology/software/products/berkeley-db/htdocs/oslicense.html

Oracle, Berkeley DB, and Sleepycat are trademarks or registered trademarks of Oracle. All rights to these marks are reserved. No third-party use is permitted without the express prior written consent of Oracle.

To obtain a copy of this document's original source code, please submit a request to the Oracle Technology Network forum at:
http://forums.oracle.com/forums/forum.jspa?forumID=271

Published 4/12/2010

Table of Contents

Preface

Welcome to Berkeley DB (DB). This document describes the C++ API for DB, version 4.8. It is intended to describe the DB API, including all classes, methods, and functions. As such, this document is intended for C++ developers who are actively writing or maintaining applications that make use of DB databases.

Conventions Used in this Book

The following typographical conventions are used within in this manual:

Class names are represented in `monospaced font`, as are `method names`. For example: "`Db::open()` is a `Db` class method."

Variable or non-literal text is presented in *italics*. For example: "Go to your *DB_INSTALL* directory."

Program examples are displayed in a `monospaced font` on a shaded background. For example:

```
typedef struct vendor {
    char name[MAXFIELD];            // Vendor name
    char street[MAXFIELD];          // Street name and number
    char city[MAXFIELD];            // City
    char state[3];                  // Two-digit US state code
    char zipcode[6];                // US zipcode
    char phone_number[13];          // Vendor phone number
} VENDOR;
```

☞   Finally, notes of interest are represented using a note block such as this.

For More Information

Beyond this manual, you may also find the following sources of information useful when building a DB application:

- Getting Started with Berkeley DB for C++ [http://www.oracle.com/technology/documentation/ berkeley-db/db/gsg/CXX/BerkeleyDB-Core-Cxx-GSG.pdf]

- Getting Started with Transaction Processing for C++ [http://www.oracle.com/technology/ documentation/berkeley-db/db/gsg_txn/CXX/BerkeleyDB-Core-Cxx-Txn.pdf]

- Berkeley DB Getting Started with Replicated Applications for C++ [http://www.oracle.com/ technology/documentation/berkeley-db/db/gsg_db_rep/CXX/Replication_CXX_GSG.pdf]

- Berkeley DB C API [http://www.oracle.com/technology/documentation/berkeley-db/db/ api_reference/C/BDB-C_APIReference.pdf]

- Berkeley DB STL API [http://www.oracle.com/technology/documentation/berkeley-db/db/ api_reference/STL/BDB-STL_APIReference.pdf]

- Berkeley DB TCL API [http://www.oracle.com/technology/documentation/berkeley-db/db/
  api_reference/TCL/BDB-TCL_APIReference.pdf]

- Berkeley DB Programmer's Reference Guide [http://www.oracle.com/technology/documentation/
  berkeley-db/db/programmer_reference/BDB_Prog_Reference.pdf]

Chapter 1. Introduction to Berkeley DB APIs

Welcome to the Berkeley DB API Reference Manual for C++

DB is a general-purpose embedded database engine that is capable of providing a wealth of data management services. It is designed from the ground up for high-throughput applications requiring in-process, bullet-proof management of mission-critical data. DB can gracefully scale from managing a few bytes to terabytes of data. For the most part, DB is limited only by your system's available physical resources.

This manual describes the various APIs and command line utilities available for use in the DB library.

For a general description of using DB beyond the reference material available in this manual, see the Getting Started Guides which are identified in this manual's preface.

This manual is broken into chapters, each one of which describes a series of APIs designed to work with one particular aspect of the DB library. In many cases, each such chapter is organized around a "handle", or class, which provides an interface to DB structures such as databases, environments or locks. However, in some cases, methods for multiple handles are combined together when they are used to control or interface with some isolated DB functionality. See, for example, the The DbLsn Handle chapter.

Within each chapter, methods, functions and command line utilities are organized alphabetically.

Chapter 2. The Db Handle

The Db is the handle for a single Berkeley DB database. A Berkeley DB database provides a mechanism for organizing key-data pairs of information. From the perspective of some database systems, a Berkeley DB database could be thought of as a single table within a larger database.

You create a Db handle using the Db constructor. For most database activities, you must then open the handle using the Db::open() method. When you are done with them, handles must be closed using the Db::close() method.

Alternatively, you can create a Db and then rename, remove or verify the database without performing an open. See Db::rename(), Db::remove() or Db::verify() for information on these activities.

It is possible to create databases such that they are organized within a *database environment*. Environments are optional for simple Berkeley DB applications that do not use transactions, recovery, replication or any other advanced features. For simple Berkeley DB applications, environments still offer some advantages. For example, they provide some organizational benefits on-disk (all databases are located on disk relative to the environment). Also, if you are using multiple databases, then environments allow your databases to share a common in-memory cache, which makes for more efficient usage of your hardware's resources.

See DbEnv for information on using database environments.

You specify the underlying organization of the data in the database (e.g. BTree, Hash, Queue, and Recno) when you open the database. When you create a database, you are free to specify any of the available database types. On subsequent opens, you must either specify the access method used when you first opened the database, or you can specify DB_UNKNOWN in order to have this information retrieved for you. See the Db::open() method for information on specifying database types.

Database and Related Methods

Database Operations	Description
Db::associate()	Associate a secondary index
Db::associate_foreign()	Associate a foreign index
Db::close()	Close a database
Db	Create a database handle
Db::compact()	Compact a database
Db::del()	Delete items from a database
Db::err()	Error message
Db::exists()	Return if an item appears in a database
Db::fd()	Return a file descriptor from a database
Db::get()	Get items from a database
Db::get_byteswapped()	Return if the underlying database is in host order
Db::get_dbname()	Return the file and database name
Db::get_multiple()	Return if the database handle references multiple databases
Db::get_open_flags()	Returns the flags specified to Db::open
Db::get_type()	Return the database type
Db::join()	Perform a database join on cursors
Db::key_range()	Return estimate of key location
Db::open()	Open a database
Db::put()	Store items into a database
Db::remove()	Remove a database
Db::rename()	Rename a database
Db::set_priority(), Db::get_priority()	Set/get cache page priority
Db::stat()	Database statistics
Db::stat_print()	Display database statistics
Db::sync()	Flush a database to stable storage
Db::truncate()	Empty a database
Db::upgrade()	Upgrade a database
Db::verify()	Verify/salvage a database
Db::cursor()	Create a cursor handle
Database Configuration	
Db::set_alloc()	Set local space allocation functions
Db::set_cachesize(), Db::get_cachesize()	Set/get the database cache size

Database Operations	Description
Db::set_create_dir(), Db::get_create_dir()	Set/get the directory in which a database is placed
Db::set_dup_compare()	Set a duplicate comparison function
Db::set_encrypt(), Db::get_encrypt_flags()	Set/get the database cryptographic key
Db::set_errcall()	Set error message callback
Db::set_errfile(), Db::get_errfile()	Set/get error message FILE
Db::set_error_stream()	Set C++ ostream used for error messages
Db::set_errpfx(), Db::get_errpfx()	Set/get error message prefix
Db::set_feedback()	Set feedback callback
Db::set_flags(), Db::get_flags()	Set/get general database configuration
Db::set_lorder(), Db::get_lorder()	Set/get the database byte order
Db::set_message_stream()	Set C++ ostream used for informational messages
Db::set_msgcall()	Set informational message callback
Db::set_msgfile(), Db::get_msgfile()	Set/get informational message FILE
Db::set_pagesize(), Db::get_pagesize()	Set/get the underlying database page size
Db::set_partition()	Set database partitioning
Db::set_partition_dirs(), Db::get_partition_dirs()	Set/get the directories used for database partitions
Btree/Recno Configuration	
Db::set_append_recno()	Set record append callback
Db::set_bt_compare()	Set a Btree comparison function
Db::set_bt_compress()	Set Btree compression functions
Db::set_bt_minkey(), Db::get_bt_minkey()	Set/get the minimum number of keys per Btree page
Db::set_bt_prefix()	Set a Btree prefix comparison function
Db::set_re_delim(), Db::get_re_delim()	Set/get the variable-length record delimiter
Db::set_re_len(), Db::get_re_len()	Set/get the fixed-length record length
Db::set_re_pad(), Db::get_re_pad()	Set/get the fixed-length record pad byte
Db::set_re_source(), Db::get_re_source()	Set/get the backing Recno text file
Hash Configuration	
Db::set_h_compare()	Set a Hash comparison function
Db::set_h_ffactor(), Db::get_h_ffactor()	Set/get the Hash table density
Db::set_h_hash()	Set a hashing function
Db::set_h_nelem(), Db::get_h_nelem()	Set/get the Hash table size
Queue Configuration	
Db::set_q_extentsize(), Db::get_q_extentsize()	Set/get Queue database extent size

Db::associate()

```
#include <db_cxx.h>

int
Db::associate(DbTxn *txnid, Db *secondary,
    int (*callback)(Db *secondary,
    const Dbt *key, const Dbt *data, Dbt *result), u_int32_t flags);
```

The Db::associate() function is used to declare one database a secondary index for a primary database. The Db handle that you call the associate() method from is the primary database.

After a secondary database has been "associated" with a primary database, all updates to the primary will be automatically reflected in the secondary and all reads from the secondary will return corresponding data from the primary. Note that as primary keys must be unique for secondary indices to work, the primary database must be configured without support for duplicate data items. See Secondary Indices in the *Berkeley DB Programmer's Reference Guide* for more information.

The Db::associate() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

callback

The **callback** parameter is a callback function that creates the set of secondary keys corresponding to a given primary key and data pair.

The callback parameter may be NULL if both the primary and secondary database handles were opened with the DB_RDONLY flag.

The callback takes four arguments:

- secondary

  The **secondary** parameter is the database handle for the secondary.

- key

  The **key** parameter is a Dbt referencing the primary key.

- data

  The **data** parameter is a Dbt referencing the primary data item.

- result

  The **result** parameter is a zeroed Dbt in which the callback function should fill in **data** and **size** fields that describe the secondary key or keys.

☞ Berkeley DB is not re-entrant. Callback functions should not attempt to make library calls (for example, to release locks or close open handles). Re-entering Berkeley DB is not guaranteed to work correctly, and the results are undefined.

The result Dbt can have the following flags set in its **flags** field:

- DB_DBT_APPMALLOC

  If the callback function needs to allocate memory for the **result** data field (rather than simply pointing into the primary key or datum), DB_DBT_APPMALLOC should be set in the **flags** field of the **result** Dbt, which indicates that Berkeley DB should free the memory when it is done with it.

- DB_DBT_MULTIPLE

  To return multiple secondary keys, DB_DBT_MULTIPLE should be set in the **flags** field of the **result** Dbt, which indicates Berkeley DB should treat the **size** field as the number of secondary keys (zero or more), and the **data** field as a pointer to an array of that number of Dbts describing the set of secondary keys.

  When multiple secondary keys are returned, keys may not be repeated. In other words, there must be no repeated record numbers in the array for Recno and Queue databases, and keys must not compare equally using the secondary database's comparison function for Btree and Hash databases. If keys are repeated, operations may fail and the secondary may become inconsistent with the primary.

  The DB_DBT_APPMALLOC flag may be set for any Dbt in the array of returned Dbt's to indicate that Berkeley DB should free the memory referenced by that particular Dbt's data field when it is done with it.

  The DB_DBT_APPMALLOC flag may be combined with DB_DBT_MULTIPLE in the **result** Dbt's **flag** field to indicate that Berkeley DB should free the array once it is done with all of the returned keys.

In addition, the callback can optionally return the following special value:

- DB_DONOTINDEX

  If any key/data pair in the primary yields a null secondary key and should be left out of the secondary index, the callback function may optionally return DB_DONOTINDEX. Otherwise, the callback function should return 0 in case of success or an error outside of the Berkeley DB name space in case of failure; the error code will be returned from the Berkeley DB call that initiated the callback.

  If the callback function returns DB_DONOTINDEX for any key/data pairs in the primary database, the secondary index will not contain any reference to those key/data pairs, and such operations as cursor iterations and range queries will reflect only the corresponding subset of the database. If this is not desirable, the application should ensure that the callback function is well-defined for all possible values and never returns DB_DONOTINDEX.

  Returning DB_DONOTINDEX is equivalent to setting DB_DBT_MULTIPLE on the **result** Dbt and setting the **size** field to zero.

flags

The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

- DB_CREATE

  If the secondary database is empty, walk through the primary and create an index to it in the empty secondary. This operation is potentially very expensive.

  If the secondary database has been opened in an environment configured with transactions, the entire secondary index creation is performed in the context of a single transaction.

  Care should be taken not to use a newly-populated secondary database in another thread of control until the Db::associate() call has returned successfully in the first thread.

  If transactions are not being used, care should be taken not to modify a primary database being used to populate a secondary database, in another thread of control, until the Db::associate() call has returned successfully in the first thread. If transactions are being used, Berkeley DB will perform appropriate locking and the application need not do any special operation ordering.

- DB_IMMUTABLE_KEY

  Specifies the secondary key is immutable.

  This flag can be used to optimize updates when the secondary key in a primary record will never be changed after the primary record is inserted. For immutable secondary keys, a best effort is made to avoid calling the secondary callback function when primary records are updated. This optimization may reduce the overhead of update operations significantly if the callback function is expensive.

  Be sure to specify this flag only if the secondary key in the primary record is never changed. If this rule is violated, the secondary index will become corrupted, that is, it will become out of sync with the primary.

secondary

The **secondary** parameter should be an open database handle of either a newly created and empty database that is to be used to store a secondary index, or of a database that was previously associated with the same primary and contains a secondary index. Note that it is not safe to associate as a secondary database a handle that is in use by another thread of control or has open cursors. If the handle was opened with the DB_THREAD flag it is safe to use it in multiple threads of control after the Db::associate() method has returned. Note also that either secondary keys must be unique or the secondary database must be configured with support for duplicate data items.

txnid

If the operation is part of an application-specified transaction, the **txnid** parameter is a transaction handle returned from DbEnv::txn_begin(); if the operation is part of a Berkeley DB Concurrent Data Store group, the **txnid** parameter is a handle returned from DbEnv::cdsgroup_begin(); otherwise NULL. If no transaction handle is specified, but the operation occurs in a transactional database, the operation will be implicitly transaction protected.

Errors

The `Db::associate()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DbRepHandleDeadException or DB_REP_HANDLE_DEAD

When a client synchronizes with the master, it is possible for committed transactions to be rolled back. This invalidates all the database and cursor handles opened in the replication environment. Once this occurs, an attempt to use such a handle will throw a DbRepHandleDeadException (if your application is configured to throw exceptions), or return DB_REP_HANDLE_DEAD. The application will need to discard the handle and open a new one in order to continue processing.

DbDeadlockException or DB_REP_LOCKOUT

The operation was blocked by client/master synchronization.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_REP_LOCKOUT is returned.

EINVAL

If the secondary database handle has already been associated with this or another database handle; the secondary database handle is not open; the primary database has been configured to allow duplicates; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::associate_foreign()

```
#include <db_cxx.h>

int
DB::associate_foreign(Db *secondary,,
    int (*callback)(Db *secondary,
    const Dbt *key, Dbt *data, const Dbt *foreignkey, int *changed),
    u_int32_t flags);
```

The `Db::associate_foreign()` function is used to declare one database a foreign constraint for a secondary database. The Db handle that you call the `associate_foreign()` method from is the foreign database.

After a foreign database has been "associated" with a secondary database, all keys inserted into the secondary must exist in the foreign database. Attempting to add a record with a foreign key that does not exist in the foreign database will cause the put method to fail and return `DB_FOREIGN_CONFLICT`.

Deletions in the foreign database affect the secondary in a manner defined by the flags parameter. See Foreign Indices in the *Berkeley DB Programmer's Reference Guide* for more information.

The `Db::associate_foreign()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

callback

The **callback** parameter is a callback function that nullifies the foreign key portion of a data Dbt.

The callback parameter must be NULL if either DB_FOREIGN_ABORT or DB_FOREIGN_CASCADE is set.

The callback takes four arguments:

- secondary

  The **secondary** parameter is the database handle for the secondary.

- key

  The **key** parameter is a Dbt referencing the primary key.

- data

  The **data** parameter is a Dbt referencing the primary data item to be updated.

- foreignkey

  The **foreignkey** parameter is a Dbt referencing the foreign key which is being deleted.

- changed

The **changed** parameter is a pointer to a boolean value, indicated whether **data** has changed.

☞ Berkeley DB is not re-entrant. Callback functions should not attempt to make library calls (for example, to release locks or close open handles). Re-entering Berkeley DB is not guaranteed to work correctly, and the results are undefined.

flags

The **flags** parameter must be set to one of the following values:

• DB_FOREIGN_ABORT

Abort the deletion of a key in the foreign database and return DB_FOREIGN_CONFLICT if that key exists in the secondary database. The deletion should be protected by a transaction to ensure database integrity after the aborted delete.

• DB_FOREIGN_CASCADE

The deletion of a key in the foreign database will also delete that key from the secondary database (and the corresponding entry in the secondary's primary database.)

• DB_FOREIGN_NULLIFY

The deletion of a key in the foreign database will call the nullification function passed to associate_foreign and update the secondary database with the changed data.

secondary

The **secondary** parameter should be an open database handle of a database that contains a secondary index who's keys also exist in the **foreign** database.

Errors

The Db::associate_foreign() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DbRepHandleDeadException or DB_REP_HANDLE_DEAD

When a client synchronizes with the master, it is possible for committed transactions to be rolled back. This invalidates all the database and cursor handles opened in the replication environment. Once this occurs, an attempt to use such a handle will throw a DbRepHandleDeadException (if your application is configured to throw exceptions), or return DB_REP_HANDLE_DEAD. The application will need to discard the handle and open a new one in order to continue processing.

DbDeadlockException or DB_REP_LOCKOUT

The operation was blocked by client/master synchronization.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_REP_LOCKOUT is returned.

EINVAL

If the foreign database handle is a secondary index; the foreign database handle has been configured to allow duplicates; the foreign database handle is a renumbering recno database; callback is configured and DB_FOREIGN_NULLIFY is not; DB_FOREIGN_NULLIFY is configured and callback is not.

Class

Db

See Also

Database and Related Methods

Db::close()

```
#include <db_cxx.h>

int
Db::close(u_int32_t flags);
```

The `Db::close()` method flushes any cached database information to disk, closes any open cursors, frees any allocated resources, and closes any underlying files.

Although closing a database handle will close any open cursors, it is recommended that applications explicitly close all their Dbc handles before closing the database. The reason why is that when the cursor is explicitly closed, the memory allocated for it is reclaimed; however, this will *not* happen if you close a database while cursors are still opened.

The same rule, for the same reasons, hold true for DbTxn handles. Simply make sure you close all your transaction handles before closing your database handle.

Because key/data pairs are cached in memory, applications should make a point to always either close database handles or sync their data to disk (using the Db::sync() method) before exiting, to ensure that any data cached in main memory are reflected in the underlying file system.

When called on a database that is the primary database for a secondary index, the primary database should be closed only after all secondary indices referencing it have been closed.

When multiple threads are using the Db concurrently, only a single thread may call the `Db::close()` method.

The Db handle may not be accessed again after `Db::close()` is called, regardless of its return.

The `Db::close()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or be set to the following value:

- DB_NOSYNC

  Do not flush cached information to disk. This flag is a dangerous option. It should be set only if the application is doing logging (with transactions) so that the database is recoverable after a system or application crash, or if the database is always generated from scratch after any system or application crash.

  It is important to understand that flushing cached information to disk only minimizes the window of opportunity for corrupted data. Although unlikely, it is possible for database corruption to happen if a system or application crash occurs while writing data to the database. To ensure that database corruption never occurs, applications must either: use transactions and logging with automatic recovery; use logging and application-specific recovery; or edit a copy of the database, and once all

applications using the database have successfully called `Db::close()`, atomically replace the original database with the updated copy.

Note that this flag only works when the database has been opened using an environment.

Errors

The `Db::close()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db

```
#include <db_cxx.h>

class Db {
public:
        Db(DbEnv *dbenv, u_int32_t flags);
        ~Db();

         DB *Db::get_DB();
         const DB *Db::get_const_DB() const;
         static Db *Db::get_Db(DB *db);
         static const Db *Db::get_const_Db(const DB *db);
         ...
};
```

The Db handle is the handle for a Berkeley DB database, which may or may not be part of a database environment.

Db handles are free-threaded if the DB_THREAD flag is specified to the Db::open() method when the database is opened or if the database environment in which the database is opened is free-threaded. The handle should not be closed while any other handle that refers to the database is in use; for example, database handles must not be closed while cursor handles into the database remain open, or transactions that include operations on the database have not yet been committed or aborted. Once the Db::close(), Db::remove(), Db::rename(), or Db::verify() methods are called, the handle may not be accessed again, regardless of the method's return.

The constructor creates a Db object that is the handle for a Berkeley DB database. The constructor allocates memory internally; calling the Db::close(), Db::remove(), or Db::rename() methods will free that memory.

Note that destroying the Db object is synonomous with calling Db::close(0).

Each Db object has an associated DB struct, which is used by the underlying implementation of Berkeley DB and its C-language API. The Db::get_DB() method returns a pointer to this struct. Given a const Db object, Db::get_const_DB() returns a const pointer to the same struct.

Given a DB struct, the Db::get_Db() method returns the corresponding Db object, if there is one. If the DB object was not associated with a Db (that is, it was not returned from a call to the Db::get_DB() method), then the result of Db::get_Db() is undefined. Given a const DB struct, Db::get_const_Db() returns the associated const Dbobject, if there is one.

These methods may be useful for Berkeley DB applications including both C and C++ language software. It should not be necessary to use these calls in a purely C++ application.

Parameters

dbenv

If no **dbenv** value is specified, the database is standalone; that is, it is not part of any Berkeley DB environment.

If a **dbenv** value is specified, the database is created within the specified Berkeley DB environment. The database access methods automatically make calls to the other subsystems in Berkeley DB, based on the enclosing environment. For example, if the environment has been configured to use locking, the access methods will automatically acquire the correct locks when reading and writing pages of the database.

flags

The **flags** parameter is currently unused, and must be set to 0.

- DB_CXX_NO_EXCEPTION

  The Berkeley DB C++ API supports two different error behaviors. By default, whenever an error occurs, an exception is thrown that encapsulates the error information. This generally allows for cleaner logic for transaction processing because a try block can surround a single transaction. However, if this flag is specified, exceptions are not thrown; instead, each individual function returns an error code.

  If a **dbenv** value is specified, this flag is ignored, and the error behavior of the specified environment is used instead.

Class

Db

See Also

Database and Related Methods

Db::compact()

```
#include <db_cxx.h>

int
Db::compact(DbTxn *txnid,
    Dbt *start, Dbt *stop, DB_COMPACT *c_data, u_int32_t flags, Dbt *end);
```

The `Db::compact()` method compacts Btree and Recno access method databases, and optionally returns unused Btree, Hash or Recno database pages to the underlying filesystem.

The `Db::compact()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

txnid

If the operation is part of an application-specified transaction, the **txnid** parameter is a transaction handle returned from DbEnv::txn_begin(); if the operation is part of a Berkeley DB Concurrent Data Store group, the **txnid** parameter is a handle returned from DbEnv::cdsgroup_begin(); otherwise NULL.

If a transaction handle is supplied to this method, then the operation is performed using that transaction. In this event, large sections of the tree may be locked during the course of the transaction.

If no transaction handle is specified, but the operation occurs in a transactional database, the operation will be implicitly transaction protected using multiple transactions. These transactions will be periodically committed to avoid locking large sections of the tree. Any deadlocks encountered cause the compaction operation to be retried from the point of the last transaction commit.

start

If non-NULL, the **start** parameter is the starting point for compaction in a Btree or Recno database. Compaction will start at the smallest key greater than or equal to the specified key. If NULL, compaction will start at the beginning of the database.

stop

If non-NULL, the **stop** parameter is the stopping point for compaction in a Btree or Recno database. Compaction will stop at the page with the smallest key greater than the specified key. If NULL, compaction will stop at the end of the database.

c_data

If non-NULL, the **c_data** parameter contains additional compaction configuration parameters, and returns compaction operation statistics, in a structure of type DB_COMPACT.

The following input configuration fields are available from the DB_COMPACT structure:

- int compact_fillpercent;

If non-zero, this provides the goal for filling pages, specified as a percentage between 1 and 100. Any page in a Btree or Recno databases not at or above this percentage full will be considered for compaction. The default behavior is to consider every page for compaction, regardless of its page fill percentage.

- int compact_pages;

If non-zero, the call will return after the specified number of pages have been freed, or no more pages can be freed.

- db_timeout_t compact_timeout;

If non-zero, and no **txnid** parameter was specified, this parameter identifies the lock timeout used for implicit transactions, in microseconds.

The following output statistics fields are available from the DB_COMPACT structure:

- u_int32_t compact_deadlock;

An output statistics parameter: if no **txnid** parameter was specified, the number of deadlocks which occurred.

- u_int32_t compact_pages_examine;

An output statistics parameter: the number of database pages reviewed during the compaction phase.

- u_int32_t compact_pages_free;

An output statistics parameter: the number of database pages freed during the compaction phase.

- u_int32_t compact_levels;

An output statistics parameter: the number of levels removed from the Btree or Recno database during the compaction phase.

- u_int32_t compact_pages_truncated;

An output statistics parameter: the number of database pages returned to the filesystem.

flags

The **flags** parameter must be set to 0 or one of the following values:

- DB_FREELIST_ONLY

Do no page compaction, only returning pages to the filesystem that are already free and at the end of the file. This flag must be set if the database is a Hash access method database.

- DB_FREE_SPACE

Return pages to the filesystem when possible. If this flag is not specified, pages emptied as a result of compaction will be placed on the free list for re-use, but never returned to the filesystem.

Note that only pages at the end of a file can be returned to the filesystem. Because of the one-pass nature of the compaction algorithm, any unemptied page near the end of the file inhibits returning pages to the file system. A repeated call to the Db::compact() method with a low **compact_fillpercent** may be used to return pages in this case.

end

If non-NULL, the **end** parameter will be filled in with the database key marking the end of the compaction operation in a Btree or Recno database. This is generally the first key of the page where the operation stopped.

Errors

The Db::compact() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DbDeadlockException or DB_LOCK_DEADLOCK

A transactional database environment operation was selected to resolve a deadlock.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_LOCK_DEADLOCK is returned.

DbLockNotGrantedException or DB_LOCK_NOTGRANTED

A Berkeley DB Concurrent Data Store database environment configured for lock timeouts was unable to grant a lock in the allowed time.

DbLockNotGrantedException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_LOCK_NOTGRANTED is returned.

DbRepHandleDeadException or DB_REP_HANDLE_DEAD

When a client synchronizes with the master, it is possible for committed transactions to be rolled back. This invalidates all the database and cursor handles opened in the replication environment. Once this occurs, an attempt to use such a handle will throw a DbRepHandleDeadException (if your application is configured to throw exceptions), or return DB_REP_HANDLE_DEAD. The application will need to discard the handle and open a new one in order to continue processing.

DbDeadlockException or DB_REP_LOCKOUT

The operation was blocked by client/master synchronization.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_REP_LOCKOUT is returned.

EACCES

An attempt was made to modify a read-only database.

EINVAL

An invalid flag value or parameter was specified.

Class

[Db](#)

See Also

[Database and Related Methods](#)

Db::del()

```
#include <db_cxx.h>

int
Db::del(DbTxn *txnid, Dbt *key, u_int32_t flags);
```

The `Db::del()` method removes key/data pairs from the database. The key/data pair associated with the specified **key** is discarded from the database. In the presence of duplicate key values, all records associated with the designated key will be discarded.

When called on a database that has been made into a secondary index using the Db::associate() method, the `Db::del()` method deletes the key/data pair from the primary database and all secondary indices.

The `Db::del()` method will return DB_NOTFOUND if the specified key is not in the database. The `Db::del()` method will return DB_KEYEMPTY if the database is a Queue or Recno database and the specified key exists, but was never explicitly created by the application or was later deleted. Unless otherwise specified, the `Db::del()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or one of the following values:

- DB_CONSUME

    If the database is of type DB_QUEUE then this flag may be set to force the head of the queue to move to the first non-deleted item in the queue. Normally this is only done if the deleted item is exactly at the head when deleted.

- DB_MULTIPLE

    Delete multiple data items using keys from the buffer to which the **key** parameter refers.

    To delete records in bulk by key with the btree or hash access methods, construct a bulk buffer in the **key** Dbt using DbMultipleDataBuilder. To delete records in bulk by record number, construct a bulk buffer in the **key** Dbt using DbMultipleRecnoDataBuilder with a data size of zero.

    A successful bulk delete operation is logically equivalent to a loop through each key/data pair, performing a Db::del() for each one.

    See the DBT and Bulk Operations for more information on working with bulk updates.

    The DB_MULTIPLE flag may only be used alone.

- DB_MULTIPLE_KEY

    Delete multiple data items using keys and data from the buffer to which the **key** parameter refers.

To delete records in bulk with the btree or hash access methods, construct a bulk buffer in the **key** Dbt using DbMultipleKeyDataBuilder. To delete records in bulk with the recno or hash access methods, construct a bulk buffer in the **key** Dbt using DbMultipleRecnoDataBuilder.

See the DBT and Bulk Operations for more information on working with bulk updates.

The DB_MULTIPLE_KEY flag may only be used alone.

key

The key Dbt operated on.

txnid

If the operation is part of an application-specified transaction, the **txnid** parameter is a transaction handle returned from DbEnv::txn_begin(); if the operation is part of a Berkeley DB Concurrent Data Store group, the **txnid** parameter is a handle returned from DbEnv::cdsgroup_begin(); otherwise NULL. If no transaction handle is specified, but the operation occurs in a transactional database, the operation will be implicitly transaction protected.

Errors

The Db::del() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DB_FOREIGN_CONFLICT

A foreign key constraint violation has occurred. This can be caused by one of two things:

1. An attempt was made to add a record to a constrained database, and the key used for that record does not exist in the foreign key database.

2. DB_FOREIGN_ABORT (page 10) was declared for a foreign key database, and then subsequently a record was deleted from the foreign key database without first removing it from the constrained secondary database.

DbDeadlockException or DB_LOCK_DEADLOCK

A transactional database environment operation was selected to resolve a deadlock.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_LOCK_DEADLOCK is returned.

DbLockNotGrantedException or DB_LOCK_NOTGRANTED

A Berkeley DB Concurrent Data Store database environment configured for lock timeouts was unable to grant a lock in the allowed time.

DbLockNotGrantedException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_LOCK_NOTGRANTED is returned.

DbRepHandleDeadException or DB_REP_HANDLE_DEAD

When a client synchronizes with the master, it is possible for committed transactions to be rolled back. This invalidates all the database and cursor handles opened in the replication environment. Once this occurs, an attempt to use such a handle will throw a DbRepHandleDeadException (if your application is configured to throw exceptions), or return DB_REP_HANDLE_DEAD. The application will need to discard the handle and open a new one in order to continue processing.

DbDeadlockException or DB_REP_LOCKOUT

The operation was blocked by client/master synchronization.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_REP_LOCKOUT is returned.

DB_SECONDARY_BAD

A secondary index references a nonexistent primary key.

EACCES

An attempt was made to modify a read-only database.

EINVAL

An invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::err()

```
#include <db_cxx.h>

Db::err(int error, const char *fmt, ...);

Db::errx(const char *fmt, ...);
```

The DbEnv::err(), DbEnv::errx(), Db::err() and Db::errx() methods provide error-messaging functionality for applications written using the Berkeley DB library.

The Db::err() and DbEnv::err() methods construct an error message consisting of the following elements:

- **An optional prefix string**

  If no error callback function has been set using the DbEnv::set_errcall() method, any prefix string specified using the DbEnv::set_errpfx() method, followed by two separating characters: a colon and a <space> character.

- **An optional printf-style message**

  The supplied message **fmt**, if non-NULL, in which the ANSI C X3.159-1989 (ANSI C) printf function specifies how subsequent parameters are converted for output.

- **A separator**

  Two separating characters: a colon and a <space> character.

- **A standard error string**

  The standard system or Berkeley DB library error string associated with the **error** value, as returned by the DbEnv::strerror() method.

The Db::errx() and DbEnv::errx() methods are the same as the Db::err() and DbEnv::err() methods, except they do not append the final separator characters and standard error string to the error message.

This constructed error message is then handled as follows:

- If an error callback function has been set (see Db::set_errcall() and DbEnv::set_errcall()), that function is called with two parameters: any prefix string specified (see Db::set_errpfx() and DbEnv::set_errpfx()) and the error message.

- If a C library FILE * has been set (see Db::set_errfile() and DbEnv::set_errfile()), the error message is written to that output stream.

- If a C++ ostream has been set (see DbEnv::set_error_stream() and Db::set_error_stream()), the error message is written to that stream.

- If none of these output options have been configured, the error message is written to stderr, the standard error output stream.

Parameters

error

The **error** parameter is the error value for which the DbEnv::err() and `Db::err()` methods will display an explanatory string.

fmt

The **fmt** parameter is an optional printf-style message to display.

Class

Db

See Also

Database and Related Methods

Db::exists()

```
#include <db_cxx.h>

int
Db::exists(DbTxn *txnid, Dbt *key, u_int32_t flags);
```

The `Db::exists()` method returns whether the specified key appears in the database.

The `Db::exists()` method will return DB_NOTFOUND if the specified key is not in the database. The `Db::exists()` method will return DB_KEYEMPTY if the database is a Queue or Recno database and the specified key exists, but was never explicitly created by the application or was later deleted.

Parameters

flags

The **flags** parameter must be set to zero or by bitwise inclusively **OR**'ing together one or more of the following values:

- DB_READ_COMMITTED

    Configure a transactional read operation to have degree 2 isolation (the read is not repeatable).

- DB_READ_UNCOMMITTED

    Configure a transactional read operation to have degree 1 isolation, reading modified but not yet committed data. Silently ignored if the DB_READ_UNCOMMITTED flag was not specified when the underlying database was opened.

- DB_RMW

    Acquire write locks instead of read locks when doing the read, if locking is configured. Setting this flag can eliminate deadlock during a read-modify-write cycle by acquiring the write lock during the read part of the cycle so that another thread of control acquiring a read lock for the same item, in its own read-modify-write cycle, will not result in deadlock.

    Because the `Db::exists()` method will not hold locks across Berkeley DB calls in non-transactional operations, the DB_RMW flag to the `Db::exists()` call is meaningful only in the presence of transactions.

key

The key Dbt operated on.

txnid

If the operation is part of an application-specified transaction, the **txnid** parameter is a transaction handle returned from DbEnv::txn_begin(); if the operation is part of a Berkeley DB Concurrent Data Store group, the **txnid** parameter is a handle returned from DbEnv::cdsgroup_begin(); otherwise NULL.

If no transaction handle is specified, but the operation occurs in a transactional database, the operation will be implicitly transaction protected.

Class

Db

See Also

Database and Related Methods

Db::fd()

```
#include <db_cxx.h>

int
Db::fd(int *fdp);
```

The `Db::fd()` method provides access to a file descriptor representative of the underlying database. A file descriptor referring to the same file will be returned to all processes that call Db::open() with the same **file** parameter.

This file descriptor may be safely used as a parameter to the **fcntl**(2) and **flock**(2) locking functions.

The `Db::fd()` method only supports a coarse-grained form of locking. Applications should instead use the Berkeley DB lock manager where possible.

The `Db::fd()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

fdp

The **fdp** parameter references memory into which the current file descriptor is copied.

Class

Db

See Also

Database and Related Methods

Db::get()

```
#include <db_cxx.h>

int
Db::get(DbTxn *txnid, Dbt *key, Dbt *data, u_int32_t flags);

int
Db::pget(DbTxn *txnid, Dbt *key, Dbt *pkey, Dbt *data, u_int32_t flags);
```

The Db::get() method retrieves key/data pairs from the database. The address and length of the data associated with the specified **key** are returned in the structure to which **data** refers.

In the presence of duplicate key values, Db::get() will return the first data item for the designated key. Duplicates are sorted by:

- Their sort order, if a duplicate sort function was specified.

- Any explicit cursor designated insertion.

- By insert order. This is the default behavior.

Retrieval of duplicates requires the use of cursor operations. See Dbc::get() for details.

When called on a database that has been made into a secondary index using the Db::associate() method, the Db::get() and Db::pget() methods return the key from the secondary index and the data item from the primary database. In addition, the Db::pget() method returns the key from the primary database. In databases that are not secondary indices, the Db::pget() method will always fail.

The Db::get() method will return DB_NOTFOUND if the specified key is not in the database. The Db::get() method will return DB_KEYEMPTY if the database is a Queue or Recno database and the specified key exists, but was never explicitly created by the application or was later deleted. Unless otherwise specified, the Db::get() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

data

The data Dbt operated on.

flags

The **flags** parameter must be set to 0 or one of the following values:

- DB_CONSUME

  Return the record number and data from the available record closest to the head of the queue, and delete the record. The record number will be returned in **key**, as described in Dbt. The data will be returned in the **data** parameter. A record is available if it is not deleted and is not currently locked. The underlying database must be of type Queue for DB_CONSUME to be specified.

- DB_CONSUME_WAIT

  The DB_CONSUME_WAIT flag is the same as the DB_CONSUME flag, except that if the Queue database is empty, the thread of control will wait until there is data in the queue before returning. The underlying database must be of type Queue for DB_CONSUME_WAIT to be specified.

  If lock or transaction timeouts have been specified, the `Db::get()` method with the DB_CONSUME_WAIT flag may return DB_LOCK_NOTGRANTED. This failure, by itself, does not require the enclosing transaction be aborted.

- DB_GET_BOTH

  Retrieve the key/data pair only if both the key and data match the arguments.

  When using a secondary index handle, the `DB_GET_BOTH`: flag causes:

  - the `Db::pget()` version of this method to retun the secondary key/primary key/data tuple only if both the primary and secondary keys match the arguments.

  - the `Db::get()` version of this method to result in an error.

- DB_SET_RECNO

  Retrieve the specified numbered key/data pair from a database. Upon return, both the **key** and **data** items will have been filled in.

  The **data** field of the specified **key** must be a pointer to a logical record number (that is, a **db_recno_t**). This record number determines the record to be retrieved.

  For DB_SET_RECNO to be specified, the underlying database must be of type Btree, and it must have been created with the DB_RECNUM flag.

In addition, the following flags may be set by bitwise inclusively **OR**'ing them into the **flags** parameter:

- DB_IGNORE_LEASE

  Return the data item irrespective of the state of master leases. The item will be returned under all conditions: if master leases are not configured, if the request is made to a client, if the request is made to a master with a valid lease, or if the request is made to a master without a valid lease.

- DB_MULTIPLE

  Return multiple data items in the buffer to which the **data** parameter refers.

  In the case of Btree or Hash databases, all of the data items associated with the specified key are entered into the buffer. In the case of Queue or Recno databases, all of the data items in the database, starting at, and subsequent to, the specified key, are entered into the buffer.

  The buffer to which the **data** parameter refers must be provided from user memory (see DB_DBT_USERMEM). The buffer must be at least as large as the page size of the underlying database, aligned for unsigned integer access, and be a multiple of 1024 bytes in size. If the buffer size is

insufficient, then upon return from the call the size field of the **data** parameter will have been set to an estimated buffer size, and the error DB_BUFFER_SMALL is returned. (The size is an estimate as the exact size needed may not be known until all entries are read. It is best to initially provide a relatively large buffer, but applications should be prepared to resize the buffer as necessary and repeatedly call the method.)

The DB_MULTIPLE flag may only be used alone, or with the DB_GET_BOTH and DB_SET_RECNO options. The DB_MULTIPLE flag may not be used when accessing databases made into secondary indices using the Db::associate() method.

See the DBT and Bulk Operations for more information on working with bulk get.

- DB_READ_COMMITTED

  Configure a transactional get operation to have degree 2 isolation (the read is not repeatable).

- DB_READ_UNCOMMITTED

  Configure a transactional get operation to have degree 1 isolation, reading modified but not yet committed data. Silently ignored if the DB_READ_UNCOMMITTED flag was not specified when the underlying database was opened.

- DB_RMW

  Acquire write locks instead of read locks when doing the read, if locking is configured. Setting this flag can eliminate deadlock during a read-modify-write cycle by acquiring the write lock during the read part of the cycle so that another thread of control acquiring a read lock for the same item, in its own read-modify-write cycle, will not result in deadlock.

  Because the Db::get() method will not hold locks across Berkeley DB calls in non-transactional operations, the DB_RMW flag to the Db::get() call is meaningful only in the presence of transactions.

key

The key Dbt operated on.

pkey

The **pkey** parameter is the return key from the primary database.

txnid

If the operation is part of an application-specified transaction, the **txnid** parameter is a transaction handle returned from DbEnv::txn_begin(); if the operation is part of a Berkeley DB Concurrent Data Store group, the **txnid** parameter is a handle returned from DbEnv::cdsgroup_begin(); otherwise NULL. If no transaction handle is specified, but the operation occurs in a transactional database, the operation will be implicitly transaction protected.

Errors

The `Db::get()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DbMemoryException or DB_BUFFER_SMALL

The requested item could not be returned due to undersized buffer.

DbMemoryException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_BUFFER_SMALL` is returned.

DbDeadlockException or DB_LOCK_DEADLOCK

A transactional database environment operation was selected to resolve a deadlock.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_LOCK_DEADLOCK` is returned.

DbLockNotGrantedException or DB_LOCK_NOTGRANTED

A Berkeley DB Concurrent Data Store database environment configured for lock timeouts was unable to grant a lock in the allowed time.

DbLockNotGrantedException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_LOCK_NOTGRANTED` is returned.

DbLockNotGrantedException or DB_LOCK_NOTGRANTED

The `DB_CONSUME_WAIT` flag was specified, lock or transaction timers were configured and the lock could not be granted before the wait-time expired.

DbLockNotGrantedException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_LOCK_NOTGRANTED` is returned.

DbRepHandleDeadException or DB_REP_HANDLE_DEAD

When a client synchronizes with the master, it is possible for committed transactions to be rolled back. This invalidates all the database and cursor handles opened in the replication environment. Once this occurs, an attempt to use such a handle will throw a DbRepHandleDeadException (if your application is configured to throw exceptions), or return `DB_REP_HANDLE_DEAD`. The application will need to discard the handle and open a new one in order to continue processing.

DB_REP_LEASE_EXPIRED

The operation failed because the site's replication master lease has expired.

DbDeadlockException or DB_REP_LOCKOUT

The operation was blocked by client/master synchronization.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_REP_LOCKOUT` is returned.

DB_SECONDARY_BAD

A secondary index references a nonexistent primary key.

EINVAL

If a record number of 0 was specified; the DB_THREAD flag was specified to the Db::open() method and none of the DB_DBT_MALLOC, DB_DBT_REALLOC or DB_DBT_USERMEM flags were set in the Dbt; the `Db::pget()` method was called with a Db handle that does not refer to a secondary index; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::get_bt_minkey()

```
#include <db_cxx.h>

int
Db::get_bt_minkey(u_int32_t *bt_minkeyp);
```

The `Db::get_bt_minkey()` method returns the minimum number of key/data pairs intended to be stored on any single Btree leaf page. This value can be set using the Db::set_bt_minkey() method.

The `Db::get_bt_minkey()` method may be called at any time during the life of the application.

The `Db::get_bt_minkey()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

bt_minkeyp

The `Db::get_bt_minkey()` method returns the minimum number of key/data pairs intended to be stored on any single Btree leaf page in **bt_minkeyp**.

Class

Db

See Also

Database and Related Methods, Db::set_bt_minkey()

Db::get_byteswapped()

```
#include <db_cxx.h>

int
Db::get_byteswapped(int *isswapped);
```

The `Db::get_byteswapped()` method returns whether the underlying database files were created on an architecture of the same byte order as the current one, or if they were not (that is, big-endian on a little-endian machine, or vice versa). This information may be used to determine whether application data needs to be adjusted for this architecture or not.

The `Db::get_byteswapped()` method may not be called before the Db::open() method is called.

The `Db::get_byteswapped()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

isswapped

If the underlying database files were created on an architecture of the same byte order as the current one, 0 is stored into the memory location referenced by **isswapped**. If the underlying database files were created on an architecture of a different byte order as the current one, 1 is stored into the memory location referenced by **isswapped**.

Errors

The `Db::get_byteswapped()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called before Db::open() was called; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::get_cachesize()

```
#include <db_cxx.h>

int
Db::get_cachesize(u_int32_t *gbytesp, u_int32_t *bytesp, int *ncachep);
```

The `Db::get_cachesize()` method returns the current size and composition of the cache. These values may be set using the Db::set_cachesize() method.

The `Db::get_cachesize()` method may be called at any time during the life of the application.

The `Db::get_cachesize()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

bytesp

The **bytesp** parameter references memory into which the additional bytes of memory in the cache is copied.

gbytesp

The **gbytesp** parameter references memory into which the gigabytes of memory in the cache is copied.

ncachep

The **ncachep** parameter references memory into which the number of caches is copied.

Class

Db

See Also

Database and Related Methods, Db::set_cachesize()

Db::get_create_dir()

```
#include <db_cxx.h>

int
Db::get_create_dir(const char **dirp);
```

Determine which directory a database file will be created in or was found in.

The `Db::get_create_dir()` method may be called at any time.

The `Db::get_create_dir()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

dirp

The **dirp** will be set to the directory specified in the call to Db::set_create_dir() method on this handle or to the directory that the database was found in after Db::open() has been called.

Errors

The `Db::get_create_dir()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::get_dbname()

```
#include <db_cxx.h>

int
Db::get_dbname(const char **filenamep, const char **dbnamep);
```

The `Db::get_dbname()` method returns the filename and database name used by the Db handle.

The `Db::get_dbname()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

filenamep

The **filenamep** parameter references memory into which a pointer to the current filename is copied.

dbnamep

The **dbnamep** parameter references memory into which a pointer to the current database name is copied.

Class

Db

See Also

Database and Related Methods

Db::get_encrypt_flags()

```
#include <db_cxx.h>

int
Db::get_encrypt_flags(u_int32_t *flagsp);
```

The `Db::get_encrypt_flags()` method returns the encryption flags. This flag can be set using the `Db::set_encrypt()` method.

The `Db::get_encrypt_flags()` method may be called at any time during the life of the application.

The `Db::get_encrypt_flags()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flagsp

The `Db::get_encrypt_flags()` method returns the encryption flags in **flagsp**.

Class

Db

See Also

Database and Related Methods, Db::set_encrypt()

Db::get_errfile()

```
#include <db_cxx.h>

void Db::get_errfile(FILE **errfilep);
```

The `Db::get_errfile()` method returns the `FILE *`, as set by the Db::set_errfile() method.

The `Db::get_errfile()` method may be called at any time during the life of the application.

Parameters

errfilep

The `Db::get_errfile()` method returns the FILE * in **errfilep**.

Class

Db

See Also

Database and Related Methods, Db::set_errfile()

Db::get_errpfx()

```
#include <db_cxx.h>

void Db::get_errpfx(const char **errpfxp);
```

The `Db::get_errpfx()` method returns the error prefix.

The `Db::get_errpfx()` method may be called at any time during the life of the application.

Parameters

errpfxp

The `Db::get_errpfx()` method returns a reference to the error prefix in **errpfxp**.

Class

Db

See Also

Database and Related Methods, Db::set_errpfx()

Db::get_flags()

```
#include <db_cxx.h>

int Db::get_flags(u_int32_t *flagsp);
```

The `Db::get_flags()` method returns the current database flags as set by the Db::set_flags() method.

The `Db::get_flags()` method may be called at any time during the life of the application.

The `Db::get_flags()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flagsp

The `Db::get_flags()` method returns the current flags in **flagsp**.

Class

Db

See Also

Database and Related Methods, Db::set_flags()

Db::get_h_ffactor()

```
#include <db_cxx.h>

int Db::get_h_ffactor(u_int32_t *h_ffactorp);
```

The `Db::get_h_ffactor()` method returns the hash table density as set by the Db::set_h_ffactor() method. The hash table density is the number of items that Berkeley DB tries to place in a hash bucket before splitting the hash bucket.

The `Db::get_h_ffactor()` method may be called at any time during the life of the application.

The `Db::get_h_ffactor()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

h_ffactorp

The `Db::get_h_ffactor()` method returns the hash table density in **h_ffactorp**.

Class

Db

See Also

Database and Related Methods, Db::set_h_ffactor()

DB C++ API

Db::get_h_nelem()

```
#include <db_cxx.h>

int
Db::get_h_nelem(u_int32_t *h_nelemp);
```

The `Db::get_h_nelem()` method returns the estimate of the final size of the hash table as set by the Db::set_h_nelem() method.

The `Db::get_h_nelem()` method may be called at any time during the life of the application.

The `Db::get_h_nelem()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

h_nelemp

The `Db::get_h_nelem()` method returns the estimate of the final size of the hash table in **h_nelemp**.

Class

Db

See Also

Database and Related Methods, Db::set_h_nelem()

Db::get_lorder()

```
#include <db_cxx.h>

int
Db::get_lorder(int *lorderp);
```

The `Db::get_lorder()` method returns the database byte order; a byte order of 4,321 indicates a big endian order, and a byte order of 1,234 indicates a little endian order. This value is set using the Db::set_lorder() method.

The `Db::get_lorder()` method may be called at any time during the life of the application.

The `Db::get_lorder()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

lorderp

The `Db::get_lorder()` method returns the database byte order in **lorderp**.

Class

Db

See Also

Database and Related Methods, Db::set_lorder()

Db::get_msgfile()

```
#include <db_cxx.h>

void Db::get_msgfile(FILE **msgfilep);
```

The `Db::get_msgfile()` method returns the `FILE *` used to output informational or statistical messages. This file handle is configured using the Db::set_msgfile() method.

The `Db::get_msgfile()` method may be called at any time during the life of the application.

Parameters

msgfilep

The `Db::get_msgfile()` method returns the `FILE *` in **msgfilep**.

Class

Db

See Also

Database and Related Methods, Db::set_msgfile()

Db::get_multiple()

```
#include <db_cxx.h>

int
Db::get_multiple()
```

This method returns non-zero if the Db handle references a physical file supporting multiple databases, and 0 otherwise.

In this case, the Db handle is a handle on a database whose key values are the names of the databases stored in the physical file and whose data values are opaque objects. No keys or data values may be modified or stored using the database handle.

This method may not be called before the Db::open() method is called.

Class

Db

See Also

Database and Related Methods

Db::get_open_flags()

```
#include <db_cxx.h>

int
Db::get_open_flags(u_int32_t *flagsp);
```

The `Db::get_open_flags()` method returns the current open method flags. That is, this method returns the flags that were specified when Db::open() was called.

The `Db::get_open_flags()` method may not be called before the `Db::open()` method is called.

The `Db::get_open_flags()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flagsp

The `Db::get_open_flags()` method returns the current open method flags in **flagsp**.

Class

Db

See Also

Database and Related Methods

Db::get_partition_callback()

```
#include <db_cxx.h>

int
Db::get_partition_callback(u_int32_t *partsp,
 u_int32_t (**callback_fcn) (DB *dbp, DBT *key);
```

The `Db::get_partition_callback()` method returns the partitioning information as set by the `Db::set_partition()` method.

The `Db::get_partition_callback()` method may be called at any time during the life of the application.

The `Db::get_partition_callback()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

partsp

The `Db::get_partition_callback()` method returns number of partitions in the **partsp** parameter.

callback_fcn

The **callback_fcn** parameter will be set to the partitioning function.

Class

Db

See Also

Database and Related Methods, Db::set_partition()

Db::get_partition_dirs()

```
#include <db_cxx.h>

int
Db::get_partition_dirs(const char ***dirsp);
```

Determine which directorise the database partitions files will be created in or were found in.

The `Db::get_partition_dirs()` method may be called at any time.

The `Db::get_partition_dirs()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

dirsp

The **dirsp** will be set to the array of directories specified in the call to Db::set_partition_dirs() method on this handle or to the directoreies that the database partitions were found in after Db::open() has been called.

Errors

The `Db::get_partition_dirs()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::get_partition_keys()

```
#include <db_cxx.h>

int
Db::get_partition_keys(u_int32_t *partsp, DBT *keysp);
```

The `Db::get_partition_keys()` method returns the partitioning information as set by the
Db::set_partition() method.

The `Db::get_partition_keys()` method may be called at any time during the life of the application.

The `Db::get_partition_keys()` method either returns a non-zero error value or throws an exception
that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

partsp

The `Db::get_partition_keys()` method returns number of partitions in the **partsp** parameter.

keysp

The **keysp** parameter will be set to the array of partitioning keys.

Class

Db

See Also

Database and Related Methods, Db::set_partition()

Db::get_pagesize()

```
#include <db_cxx.h>

int
Db::get_pagesize(u_int32_t *pagesizep);
```

The `Db::get_pagesize()` method returns the database's current page size, as set by the `Db::set_pagesize()` method. Note that if `Db::set_pagesize()` was not called by your application, then the default pagesize is selected based on the underlying filesystem I/O block size. If you call `Db::get_pagesize()` before you have opened the database, the value returned by this method is therefore the underlying filesystem I/O block size.

The `Db::get_pagesize()` method may be called only after the database has been opened.

The `Db::get_pagesize()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

pagesizep

The `Db::get_pagesize()` method returns the page size in **pagesizep**.

Class

Db

See Also

Database and Related Methods, Db::set_pagesize()

Db::get_priority()

```
#include <db_cxx.h>

int
Db::get_priority(DB_CACHE_PRIORITY *priorityp);
```

The `Db::get_priority()` method returns the cache priority for pages referenced by the Db handle. This priority value is set using the Db::set_priority() method.

The `Db::get_priority()` method may be called only after the database has been opened.

The `Db::get_priority()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

priorityp

The `Db::get_priority()` method returns a reference to the cache priority in **priorityp**. See Db::set_priority() for a list of possible priorities.

Class

Db

See Also

Database and Related Methods, Db::set_priority()

Db::get_q_extentsize()

```
#include <db_cxx.h>

int
Db::get_q_extentsize(u_int32_t *extentsizep);
```

The `Db::get_q_extentsize()` method returns the number of pages in an extent. This value is used only for Queue databases and is set using the Db::set_q_extentsize() method.

The `Db::get_q_extentsize()` method may be called only after the database has been opened.

The `Db::get_q_extentsize()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

extentsizep

The `Db::get_q_extentsize()` method returns the number of pages in an extent in **extentsizep**. If used on a handle that has not yet been opened, `0` is returned.

Class

Db

See Also

Database and Related Methods, Db::set_q_extentsize()

Db::get_re_delim()

```
#include <db_cxx.h>

int
Db::get_re_delim(int *delimp);
```

The `Db::get_re_delim()` method returns the delimiting byte, which is used to mark the end of a record in the backing source file for the Recno access method. This value is set using the Db::set_re_delim() method.

The `Db::get_re_delim()` method may be called only after the database has been opened.

The `Db::get_re_delim()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

delimp

The `Db::get_re_delim()` method returns the delimiting byte in **delimp**. If this method is called on a handle that has not yet been opened, then the default delimiting byte is returned. See Db::set_re_delim() for details.

Class

Db

See Also

Database and Related Methods, Db::set_re_delim()

Db::get_re_len()

```
#include <db_cxx.h>

int
Db::get_re_len(u_int32_t *re_lenp);
```

The `Db::get_re_len()` method returns the length of the records held in a Queue access method database. This value can be set using the Db::set_re_len() method.

The `Db::get_re_len()` method may be called only after the database has been opened.

The `Db::get_re_len()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

re_lenp

The `Db::get_re_len()` method returns the record length in **re_lenp**. If the record length has never been set using Db::set_re_len(), then `0` is returned.

Class

Db

See Also

Database and Related Methods, Db::set_re_len()

Db::get_re_pad()

```
#include <db_cxx.h>

int
Db::get_re_pad(int *re_padp);
```

The `Db::get_re_pad()` method returns the pad character used for short, fixed-length records used by the Queue and Recno access methods. This character is set using the Db::set_re_pad() method.

The `Db::get_re_pad()` method may be called only after the database has been opened.

The `Db::get_re_pad()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

re_padp

The `Db::get_re_pad()` method returns the pad character in **re_padp**. If used on a handle that has not yet been opened, the default pad character is returned. See the Db::set_re_pad() method description for what that default value is.

Class

Db

See Also

Database and Related Methods, Db::set_re_pad()

Db::get_re_source()

```
#include <db_cxx.h>

int
Db::get_re_source(const char **sourcep);
```

The `Db::get_re_source()` method returns the source file used by the Recno access method. This file is configured for the Recno access method using the Db::set_re_source() method.

The `Db::get_re_source()` method may be called only after the database has been opened.

The `Db::get_re_source()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

sourcep

The `Db::get_re_source()` method returns a reference to the source file in **sourcep**.

Class

Db

See Also

Database and Related Methods, Db::set_re_source()

Db::get_type()

```
#include <db_cxx.h>

int
Db::get_type(DBTYPE *type);
```

The `Db::get_type()` method returns the type of the underlying access method (and file format). The type value is one of DB_BTREE, DB_HASH, DB_RECNO, or DB_QUEUE. This value may be used to determine the type of the database after a return from Db::open() with the **type** parameter set to DB_UNKNOWN.

The `Db::get_type()` method may not be called before the Db::open() method is called.

The `Db::get_type()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

type

The **type** parameter references memory into which the type of the underlying access method is copied.

Errors

The `Db::get_type()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called before Db::open() was called; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::join()

```
#include <db_cxx.h>

int
Db::join(Dbc **curslist, Dbc **dbcp, u_int32_t flags);
```

The `Db::join()` method creates a specialized join cursor for use in performing equality or natural joins on secondary indices. For information on how to organize your data to use this functionality, see Equality join.

The `Db::join()` method is called using the Db handle of the primary database.

The join cursor supports only the Dbc::get() and Dbc::close() cursor functions:

- **Dbc::get()**

  Iterates over the values associated with the keys to which each item in **curslist** was initialized. Any data value that appears in all items specified by the **curslist** parameter is then used as a key into the **primary**, and the key/data pair found in the **primary** is returned. The **flags** parameter must be set to 0 or the following value:

  - **DB_JOIN_ITEM**

    Do not use the data value found in all the cursors as a lookup key for the **primary**, but simply return it in the key parameter instead. The data parameter is left unchanged.

  In addition, the following flag may be set by bitwise inclusively **OR**'ing it into the **flags** parameter:

  - **DB_READ_UNCOMMITTED**

    Configure a transactional join operation to have degree 1 isolation, reading modified but not yet committed data. Silently ignored if the DB_READ_UNCOMMITTED flag was not specified when the underlying database was opened.

  - **DB_RMW**

    Acquire write locks instead of read locks when doing the read, if locking is configured. Setting this flag can eliminate deadlock during a read-modify-write cycle by acquiring the write lock during the read part of the cycle so that another thread of control acquiring a read lock for the same item, in its own read-modify-write cycle, will not result in deadlock.

- **Dbc::close()**

  Close the returned cursor and release all resources. (Closing the cursors in **curslist** is the responsibility of the caller.)

The `Db::join()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

curslist

The **curslist** parameter contains a NULL terminated array of cursors. Each cursor must have been initialized to refer to the key on which the underlying database should be joined. Typically, this initialization is done by a Dbc::get() call with the DB_SET flag specified. Once the cursors have been passed as part of a **curslist**, they should not be accessed or modified until the newly created join cursor has been closed, or else inconsistent results may be returned.

Joined values are retrieved by doing a sequential iteration over the first cursor in the **curslist** parameter, and a nested iteration over each secondary cursor in the order they are specified in the **curslist** parameter. This requires database traversals to search for the current datum in all the cursors after the first. For this reason, the best join performance normally results from sorting the cursors from the one that refers to the least number of data items to the one that refers to the most. By default, Db::join() does this sort on behalf of its caller.

For the returned join cursor to be used in a transaction-protected manner, the cursors listed in **curslist** must have been created within the context of the same transaction.

dbcp

The newly created join cursor is returned in the memory location to which **dbcp** refers.

flags

The **flags** parameter must be set to 0 or the following value:

* DB_JOIN_NOSORT

  Do not sort the cursors based on the number of data items to which they refer. If the data are structured so that cursors with many data items also share many common elements, higher performance will result from listing those cursors before cursors with fewer data items; that is, a sort order other than the default. The DB_JOIN_NOSORT flag permits applications to perform join optimization prior to calling the Db::join() method.

Errors

The Db::join() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DbRepHandleDeadException or DB_REP_HANDLE_DEAD

When a client synchronizes with the master, it is possible for committed transactions to be rolled back. This invalidates all the database and cursor handles opened in the replication environment. Once this occurs, an attempt to use such a handle will throw a DbRepHandleDeadException (if your application is configured to throw exceptions), or return DB_REP_HANDLE_DEAD. The application will need to discard the handle and open a new one in order to continue processing.

DbDeadlockException or DB_REP_LOCKOUT

The operation was blocked by client/master synchronization.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_REP_LOCKOUT is returned.

DB_SECONDARY_BAD

A secondary index references a nonexistent primary key.

EINVAL

If cursor methods other than Dbc::get() or Dbc::close() were called; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::key_range()

```
#include <db_cxx.h>

int
Db::key_range(DbTxn *txnid
    Dbt *key, DB_KEY_RANGE *key_range, u_int32_t flags);
```

The `Db::key_range()` method returns an estimate of the proportion of keys that are less than, equal to, and greater than the specified key. The underlying database must be of type Btree.

The `Db::key_range()` method fills in a structure of type DB_KEY_RANGE. The following data fields are available from the DB_KEY_RANGE structure:

- **double less;**

  A value between 0 and 1, the proportion of keys less than the specified key.

- **double equal;**

  A value between 0 and 1, the proportion of keys equal to the specified key.

- **double greater;**

  A value between 0 and 1, the proportion of keys greater than the specified key.

Values are in the range of 0 to 1; for example, if the field **less** is 0.05, 5% of the keys in the database are less than the **key** parameter. The value for **equal** will be zero if there is no matching key, and will be non-zero otherwise.

The `Db::key_range()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

key

The key Dbt operated on.

key_range

The estimates are returned in the **key_range** parameter, which contains three elements of type double: **less**, **equal**, and **greater**. Values are in the range of 0 to 1; for example, if the field **less** is 0.05, 5% of the keys in the database are less than the **key** parameter. The value for **equal** will be zero if there is no matching key, and will be non-zero otherwise.

txnid

If the operation is part of an application-specified transaction, the **txnid** parameter is a transaction handle returned from DbEnv::txn_begin(); if the operation is part of a Berkeley DB Concurrent Data Store group, the **txnid** parameter is a handle returned from DbEnv::cdsgroup_begin(); otherwise NULL.

If no transaction handle is specified, but the operation occurs in a transactional database, the operation will be implicitly transaction protected. The `Db::key_range()` method does not retain the locks it acquires for the life of the transaction, so estimates may not be repeatable.

flags

The **flags** parameter is currently unused, and must be set to 0.

Errors

The `Db::key_range()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DbDeadlockException or DB_LOCK_DEADLOCK

A transactional database environment operation was selected to resolve a deadlock.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_LOCK_DEADLOCK` is returned.

DbLockNotGrantedException or DB_LOCK_NOTGRANTED

A Berkeley DB Concurrent Data Store database environment configured for lock timeouts was unable to grant a lock in the allowed time.

DbLockNotGrantedException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_LOCK_NOTGRANTED` is returned.

DbRepHandleDeadException or DB_REP_HANDLE_DEAD

When a client synchronizes with the master, it is possible for committed transactions to be rolled back. This invalidates all the database and cursor handles opened in the replication environment. Once this occurs, an attempt to use such a handle will throw a DbRepHandleDeadException (if your application is configured to throw exceptions), or return `DB_REP_HANDLE_DEAD`. The application will need to discard the handle and open a new one in order to continue processing.

DbDeadlockException or DB_REP_LOCKOUT

The operation was blocked by client/master synchronization.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_REP_LOCKOUT` is returned.

EINVAL

If the underlying database was not of type Btree; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::open()

```
#include <db_cxx.h>

int
Db::open(DbTxn *txnid, const char *file,
    const char *database, DBTYPE type, u_int32_t flags, int mode);
```

The `Db::open()` method opens the database represented by the **file** and **database**.

The currently supported Berkeley DB file formats (or *access methods*) are Btree, Hash, Queue, and Recno. The Btree format is a representation of a sorted, balanced tree structure. The Hash format is an extensible, dynamic hashing scheme. The Queue format supports fast access to fixed-length records accessed sequentially or by logical record number. The Recno format supports fixed- or variable-length records, accessed sequentially or by logical record number, and optionally backed by a flat text file.

Storage and retrieval for the Berkeley DB access methods are based on key/data pairs; see Dbt for more information.

Calling `Db::open()` is a relatively expensive operation, and maintaining a set of open databases will normally be preferable to repeatedly opening and closing the database for each new query.

The `Db::open()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success. If `Db::open()` fails, the Db::close() method must be called to discard the Db handle.

Parameters

database

The **database** parameter is optional, and allows applications to have multiple databases in a single file. Although no **database** parameter needs to be specified, it is an error to attempt to open a second database in a **file** that was not initially created using a **database** name. Further, the **database** parameter is not supported by the Queue format. Finally, when opening multiple databases in the same physical file, it is important to consider locking and memory cache issues; see Opening multiple databases in a single file for more information.

If both the **database** and **file** parameters are NULL, the database is strictly temporary and cannot be opened by any other thread of control. Thus the database can only be accessed by sharing the single database handle that created it, in circumstances where doing so is safe.

If the **database** parameter is not set to NULL, the database can be opened by other threads of control and will be replicated to client sites in any replication group, regardless of whether the **file** parameter is set to NULL.

file

The **file** parameter is used as the name of an underlying file that will be used to back the database; see File naming for more information.

In-memory databases never intended to be preserved on disk may be created by setting the **file** parameter to NULL. Whether other threads of control can access this database is driven entirely by whether the **database** parameter is set to NULL.

When using a Unicode build on Windows (the default), the **file** argument will be interpreted as a UTF-8 string, which is equivalent to ASCII for Latin characters.

flags

The **flags** parameter must be set to zero or by bitwise inclusively **OR**'ing together one or more of the following values:

- DB_AUTO_COMMIT

  Enclose the `Db::open()` call within a transaction. If the call succeeds, the open operation will be recoverable and all subsequent database modification operations based on this handle will be transactionally protected. If the call fails, no database will have been created.

- DB_CREATE

  Create the database. If the database does not already exist and the DB_CREATE flag is not specified, the `Db::open()` will fail.

- DB_EXCL

  Return an error if the database already exists. The DB_EXCL flag is only meaningful when specified with the DB_CREATE. flag.

- DB_MULTIVERSION

  Open the database with support for multiversion concurrency control. This will cause updates to the database to follow a copy-on-write protocol, which is required to support snapshot isolation. The DB_MULTIVERSION flag requires that the database be transactionally protected during its open and is not supported by the queue format.

- DB_NOMMAP

  Do not map this database into process memory (see the DbEnv::set_mp_mmapsize() method for further information).

- DB_RDONLY

  Open the database for reading only. Any attempt to modify items in the database will fail, regardless of the actual permissions of any underlying files.

- DB_READ_UNCOMMITTED

  Support transactional read operations with degree 1 isolation. Read operations on the database may request the return of modified but not yet committed data. This flag must be specified on all Db handles used to perform dirty reads or database updates, otherwise requests for dirty reads may not be honored and the read may block.

- DB_THREAD

    Cause the Db handle returned by `Db::open()` to be *free-threaded*; that is, concurrently usable by multiple threads in the address space.

- DB_TRUNCATE

    Physically truncate the underlying file, discarding all previous databases it might have held. Underlying filesystem primitives are used to implement this flag. For this reason, it is applicable only to the file and cannot be used to discard databases within a file.

    The `DB_TRUNCATE` flag cannot be lock or transaction-protected, and it is an error to specify it in a locking or transaction-protected environment.

mode

On Windows systems, the mode parameter is ignored.

On UNIX systems or in IEEE/ANSI Std 1003.1 (POSIX) environments, files created by the database open are created with mode **mode** (as described in **chmod**(2)) and modified by the process' umask value at the time of creation (see **umask**(2)). Created files are owned by the process owner; the group ownership of created files is based on the system and directory defaults, and is not further specified by Berkeley DB. System shared memory segments created by the database open are created with mode **mode**, unmodified by the process' umask value. If **mode** is 0, the database open will use a default mode of readable and writable by both owner and group.

txnid

If the operation is part of an application-specified transaction, the **txnid** parameter is a transaction handle returned from DbEnv::txn_begin(); if the operation is part of a Berkeley DB Concurrent Data Store group, the **txnid** parameter is a handle returned from DbEnv::cdsgroup_begin(); otherwise NULL. If no transaction handle is specified, but the DB_AUTO_COMMIT flag is specified, the operation will be implicitly transaction protected. Note that transactionally protected operations on a Db handle requires the Db handle itself be transactionally protected during its open. Also note that the transaction must be committed before the handle is closed; see Berkeley DB handles for more information.

type

The **type** parameter is of type DBTYPE, and must be set to one of `DB_BTREE`, `DB_HASH`, `DB_QUEUE`, `DB_RECNO`, or `DB_UNKNOWN`. If **type** is DB_UNKNOWN, the database must already exist and `Db::open()` will automatically determine its type. The Db::get_type() method may be used to determine the underlying type of databases opened using DB_UNKNOWN.

It is an error to specify the incorrect **type** for a database that already exists.

Environment Variables

If the database was opened within a database environment, the environment variable **DB_HOME** may be used as the path of the database environment home.

Db::open() is affected by any database directory specified using the DbEnv::set_data_dir() method, or by setting the "set_data_dir" string in the environment's DB_CONFIG file.

- **TMPDIR**

    If the **file** and **dbenv** parameters to Db::open() are NULL, the environment variable **TMPDIR** may be used as a directory in which to create temporary backing files

Errors

The Db::open() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DbDeadlockException or DB_LOCK_DEADLOCK

A transactional database environment operation was selected to resolve a deadlock.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_LOCK_DEADLOCK is returned.

DbLockNotGrantedException or DB_LOCK_NOTGRANTED

A Berkeley DB Concurrent Data Store database environment configured for lock timeouts was unable to grant a lock in the allowed time.

DbLockNotGrantedException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_LOCK_NOTGRANTED is returned.

ENOENT

The file or directory does not exist.

ENOENT

A nonexistent **re_source** file was specified.

DB_OLD_VERSION

The database cannot be opened without being first upgraded.

EEXIST

DB_CREATE and DB_EXCL were specified and the database exists.

EINVAL

If an unknown database type, page size, hash function, pad byte, byte order, or a flag value or parameter that is incompatible with the specified database was specified; the DB_THREAD flag was specified and fast mutexes are not available for this architecture; the DB_THREAD flag was specified to Db::open(), but was not specified to the DbEnv::open() call for the environment in which the Db handle was created; a backing flat text file was specified with either the DB_THREAD flag or the provided database environment supports transaction processing; or if an invalid flag value or parameter was specified.

DbRepHandleDeadException or DB_REP_HANDLE_DEAD

When a client synchronizes with the master, it is possible for committed transactions to be rolled back. This invalidates all the database and cursor handles opened in the replication environment. Once this occurs, an attempt to use such a handle will throw a DbRepHandleDeadException (if your application is configured to throw exceptions), or return DB_REP_HANDLE_DEAD. The application will need to discard the handle and open a new one in order to continue processing.

DbDeadlockException or DB_REP_LOCKOUT

The operation was blocked by client/master synchronization.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_REP_LOCKOUT is returned.

Class

Db

See Also

Database and Related Methods

Db::put()

```
#include <db_cxx.h>

int
Db::put(DbTxn *txnid, Dbt *key, Dbt *data, u_int32_t flags);
```

The `Db::put()` method stores key/data pairs in the database. The default behavior of the `Db::put()` function is to enter the new key/data pair, replacing any previously existing key if duplicates are disallowed, or adding a duplicate data item if duplicates are allowed. If the database supports duplicates, the `Db::put()` method adds the new data value at the end of the duplicate set. If the database supports sorted duplicates, the new data value is inserted at the correct sorted location.

Unless otherwise specified, the `Db::put()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or one of the following values:

- `DB_APPEND`

  Append the key/data pair to the end of the database. For the DB_APPEND flag to be specified, the underlying database must be a Queue or Recno database. The record number allocated to the record is returned in the specified **key**.

  There is a minor behavioral difference between the Recno and Queue access methods for the DB_APPEND flag. If a transaction enclosing a `Db::put()` operation with the DB_APPEND flag aborts, the record number may be reallocated in a subsequent `DB_APPEND` operation if you are using the Recno access method, but it will not be reallocated if you are using the Queue access method.

- `DB_NODUPDATA`

  In the case of the Btree and Hash access methods, enter the new key/data pair only if it does not already appear in the database.

  The DB_NODUPDATA flag may only be specified if the underlying database has been configured to support sorted duplicates. The DB_NODUPDATA flag may not be specified to the Queue or Recno access methods.

  The `Db::put()` method will return DB_KEYEXIST if DB_NODUPDATA is set and the key/data pair already appears in the database.

- `DB_NOOVERWRITE`

  Enter the new key/data pair only if the key does not already appear in the database. The `Db::put()` method call with the DB_NOOVERWRITE flag set will fail if the key already exists in the database, even if the database supports duplicates.

The `Db::put()` method will return DB_KEYEXIST if DB_NOOVERWRITE is set and the key already appears in the database.

This enforcement of uniqueness of keys applies only to the primary key. The behavior of insertions into secondary databases is not affected by the DB_NOOVERWRITE flag. In particular, the insertion of a record that would result in the creation of a duplicate key in a secondary database that allows duplicates would not be prevented by the use of this flag.

- DB_MULTIPLE

  Put multiple data items using keys from the buffer to which the **key** parameter refers and data values from the buffer to which the **data** parameter refers.

  To put records in bulk with the btree or hash access methods, construct bulk buffers in the **key** and **data** Dbt using DbMultipleDataBuilder. To put records in bulk with the recno or queue access methods, construct bulk buffers in the **data** Dbt as before, but construct the **key** Dbt using DbMultipleRecnoDataBuilder with a data size of zero.

  A successful bulk operation is logically equivalent to a loop through each key/data pair, performing a Db::put() for each one.

  See DBT and Bulk Operations for more information on working with bulk updates.

  The DB_MULTIPLE flag may only be used alone, or with the DB_OVERWRITE_DUP option.

- DB_MULTIPLE_KEY

  Put multiple data items using keys and data from the buffer to which the **key** parameter refers.

  To put records in bulk with the btree or hash access methods, construct a bulk buffer in the **key** Dbt using DbMultipleKeyDataBuilder. To put records in bulk with the recno or queue access methods, construct a bulk buffer in the **key** Dbt using DbMultipleRecnoDataBuilder.

  See DBT and Bulk Operations for more information on working with bulk updates.

  The DB_MULTIPLE_KEY flag may only be used alone, or with the DB_OVERWRITE_DUP option.

- DB_OVERWRITE_DUP

  Ignore duplicate records when overwriting records in a database configured for sorted duplicates.

  Normally, if a database is configured for sorted duplicates, an attempt to put a record that compares identically to a record already existing in the database will fail. Using this flag causes the put to silently proceed, without failure.

  This flag is extremely useful when performing bulk puts (using the DB_MULTIPLE or DB_MULTIPLE_KEY flags). Depending on the number of records you are writing to the database with a bulk put, you may not want the operation to fail in the event that a duplicate record is encountered. Using this flag along with the DB_MULTIPLE or DB_MULTIPLE_KEY flags allows the bulk put to complete, even if a duplicate record is encountered.

This flag is also useful if you are using a custom comparison function that compares only part of the data portion of a record. In this case, two records can compare equally when, in fact, they are not equal. This flag allows the put to complete, even if your custom comparison routine claims the two records are equal.

data

The data Dbt operated on.

key

The key Dbt operated on.

txnid

If the operation is part of an application-specified transaction, the **txnid** parameter is a transaction handle returned from DbEnv::txn_begin(); if the operation is part of a Berkeley DB Concurrent Data Store group, the **txnid** parameter is a handle returned from DbEnv::cdsgroup_begin(); otherwise NULL. If no transaction handle is specified, but the operation occurs in a transactional database, the operation will be implicitly transaction protected.

Errors

The Db::put() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DB_FOREIGN_CONFLICT

A foreign key constraint violation has occurred. This can be caused by one of two things:

1. An attempt was made to add a record to a constrained database, and the key used for that record does not exist in the foreign key database.

2. DB_FOREIGN_ABORT (page 10) was declared for a foreign key database, and then subsequently a record was deleted from the foreign key database without first removing it from the constrained secondary database.

DbDeadlockException or DB_LOCK_DEADLOCK

A transactional database environment operation was selected to resolve a deadlock.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_LOCK_DEADLOCK is returned.

DbLockNotGrantedException or DB_LOCK_NOTGRANTED

A Berkeley DB Concurrent Data Store database environment configured for lock timeouts was unable to grant a lock in the allowed time.

DbLockNotGrantedException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_LOCK_NOTGRANTED is returned.

DbRepHandleDeadException or DB_REP_HANDLE_DEAD

When a client synchronizes with the master, it is possible for committed transactions to be rolled back. This invalidates all the database and cursor handles opened in the replication environment. Once this occurs, an attempt to use such a handle will throw a DbRepHandleDeadException (if your application is configured to throw exceptions), or return DB_REP_HANDLE_DEAD. The application will need to discard the handle and open a new one in order to continue processing.

DbDeadlockException or DB_REP_LOCKOUT

The operation was blocked by client/master synchronization.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_REP_LOCKOUT is returned.

EACCES

An attempt was made to modify a read-only database.

EINVAL

If a record number of 0 was specified; an attempt was made to add a record to a fixed-length database that was too large to fit; an attempt was made to do a partial put; an attempt was made to add a record to a secondary index; or if an invalid flag value or parameter was specified.

ENOSPC

A btree exceeded the maximum btree depth (255).

Class

Db

See Also

Database and Related Methods

Db::remove()

```
#include <db_cxx.h>

int
Db::remove(const char *file, const char *database, u_int32_t flags);
```

The `Db::remove()` method removes the database specified by the **file** and **database** parameters. If no **database** is specified, the underlying file represented by **file** is removed, incidentally removing all of the databases it contained.

Applications should never remove databases with open Db handles, or in the case of removing a file, when any database in the file has an open handle. For example, some architectures do not permit the removal of files with open system handles. On these architectures, attempts to remove databases currently in use by any thread of control in the system may fail.

The `Db::remove()` method should not be called if the remove is intended to be transactionally safe; the DbEnv::dbremove() method should be used instead.

The `Db::remove()` method may not be called after calling the Db::open() method on any Db handle. If the Db::open() method has already been called on a Db handle, close the existing handle and create a new one before calling `Db::remove. ()`

The Db handle may not be accessed again after `Db::remove()` is called, regardless of its return.

The `Db::remove()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

database

The **database** parameter is the database to be removed.

file

The **file** parameter is the physical file which contains the database(s) to be removed.

Environment Variables

If the database was opened within a database environment, the environment variable `DB_HOME` may be used as the path of the database environment home.

`Db::remove()` is affected by any database directory specified using the DbEnv::set_data_dir() method, or by setting the "set_data_dir" string in the environment's DB_CONFIG file.

Errors

The `Db::remove()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after Db::open() was called; or if an invalid flag value or parameter was specified.

ENOENT

The file or directory does not exist.

Class

Db

See Also

Database and Related Methods

Db::rename()

```
#include <db_cxx.h>

int
Db::rename(const char *file,
    const char *database, const char *newname, u_int32_t flags);
```

The `Db::rename()` method renames the database specified by the **file** and **database** parameters to **newname**. If no **database** is specified, the underlying file represented by **file** is renamed, incidentally renaming all of the databases it contained.

Applications should not rename databases that are currently in use. If an underlying file is being renamed and logging is currently enabled in the database environment, no database in the file may be open when the `Db::rename()` method is called. In particular, some architectures do not permit renaming files with open handles. On these architectures, attempts to rename databases that are currently in use by any thread of control in the system may fail.

The `Db::rename()` method should not be called if the rename is intended to be transactionally safe; the DbEnv::dbrename() method should be used instead.

The `Db::rename()` method may not be called after calling the Db::open() method on any Db handle. If the Db::open() method has already been called on a Db handle, close the existing handle and create a new one before calling `Db::rename()`.

The Db handle may not be accessed again after `Db::rename()` is called, regardless of its return.

The `Db::rename()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

database

The **database** parameter is the database to be renamed.

file

The **file** parameter is the physical file which contains the database(s) to be renamed.

When using a Unicode build on Windows (the default), the **file** argument will be interpreted as a UTF-8 string, which is equivalent to ASCII for Latin characters.

flags

The **flags** parameter is currently unused, and must be set to 0.

newname

The **newname** parameter is the new name of the database or file.

Environment Variables

If the database was opened within a database environment, the environment variable `DB_HOME` may be used as the path of the database environment home.

`Db::rename()` is affected by any database directory specified using the DbEnv::set_data_dir() method, or by setting the "set_data_dir" string in the environment's DB_CONFIG file.

Errors

The `Db::rename()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after Db::open() was called; or if an invalid flag value or parameter was specified.

ENOENT

The file or directory does not exist.

Class

Db

See Also

Database and Related Methods

Db::set_alloc()

```
#include <db_cxx.h>

int
Db::set_alloc(db_malloc_fcn_type app_malloc,
    db_realloc_fcn_type app_realloc,
    db_free_fcn_type app_free);
```

Set the allocation functions used by the DbEnv and Db methods to allocate or free memory owned by the application.

There are a number of interfaces in Berkeley DB where memory is allocated by the library and then given to the application. For example, the DB_DBT_MALLOC flag, when specified in the Dbt object, will cause the Db methods to allocate and reallocate memory which then becomes the responsibility of the calling application. (See Dbt for more information.) Other examples are the Berkeley DB interfaces which return statistical information to the application: Db::stat(), DbEnv::lock_stat(), DbEnv::log_archive(), DbEnv::log_stat(), DbEnv::memp_stat(), and DbEnv::txn_stat(). There is one method in Berkeley DB where memory is allocated by the application and then given to the library: Db::associate().

On systems in which there may be multiple library versions of the standard allocation routines (notably Windows NT), transferring memory between the library and the application will fail because the Berkeley DB library allocates memory from a different heap than the application uses to free it. To avoid this problem, the DbEnv::set_alloc() and Db::set_alloc() methods can be used to pass Berkeley DB references to the application's allocation routines.

It is not an error to specify only one or two of the possible allocation function parameters to these interfaces; however, in that case the specified interfaces must be compatible with the standard library interfaces, as they will be used together. The functions specified must match the calling conventions of the ANSI C X3.159-1989 (ANSI C) library routines of the same name.

Because databases opened within Berkeley DB environments use the allocation interfaces specified to the environment, it is an error to attempt to set those interfaces in a database created within an environment.

The Db::set_alloc() method may not be called after the Db::open() method is called.

The Db::set_alloc() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Errors

The Db::set_alloc() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If called in a database environment, or called after Db::open() was called; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::set_append_recno()

```
#include <db_cxx.h>

int
Db::set_append_recno(int (*db_append_recno_fcn)(DB *dbp, Dbt *data, db_recno_t recno));
```

When using the DB_APPEND option of the Db::put() method, it may be useful to modify the stored data based on the generated key. If a callback function is specified using the Db::set_append_recno() method, it will be called after the record number has been selected, but before the data has been stored.

The Db::set_append_recno() method configures operations performed using the specified Db handle, not all operations performed on the underlying database.

The Db::set_append_recno() method may not be called after the Db::open() method is called.

The Db::set_append_recno() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

☞ Berkeley DB is not re-entrant. Callback functions should not attempt to make library calls (for example, to release locks or close open handles). Re-entering Berkeley DB is not guaranteed to work correctly, and the results are undefined.

Parameters

db_append_recno_fcn

The **db_append_recno_fcn** parameter is a function to call after the record number has been selected but before the data has been stored into the database. The function takes three parameters:

- dbp

  The **dbp** parameter is the enclosing database handle.

- data

  The **data** parameter is the data Dbt to be stored.

- recno

  The **recno** parameter is the generated record number.

The called function may modify the data Dbt. If the function needs to allocate memory for the **data** field, the **flags** field of the returned Dbt should be set to DB_DBT_APPMALLOC, which indicates that Berkeley DB should free the memory when it is done with it.

The callback function must return 0 on success and **errno** or a value outside of the Berkeley DB error name space on failure.

DB C++ API

Errors

The `Db::set_append_recno()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after Db::open() was called; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::set_bt_compare()

```
#include <db_cxx.h>

extern "C" {
    typedef int (*bt_compare_fcn_type)(DB *db, const DBT *dbt1, const DBT *dbt2);
};
int
Db::set_bt_compare(bt_compare_fcn_type bt_compare_fcn);
```

Set the Btree key comparison function. The comparison function is called whenever it is necessary to compare a key specified by the application with a key currently stored in the tree.

If no comparison function is specified, the keys are compared lexically, with shorter keys collating before longer keys.

The `Db::set_bt_compare()` method configures operations performed using the specified Db handle, not all operations performed on the underlying database.

The `Db::set_bt_compare()` method may not be called after the Db::open() method is called. If the database already exists when Db::open() is called, the information specified to `Db::set_bt_compare()` must be the same as that historically used to create the database or corruption can occur.

The `Db::set_bt_compare()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

bt_compare_fcn

The **bt_compare_fcn** function is the application-specified Btree comparison function. The comparison function takes three parameters:

- db

  The **db** parameter is the enclosing database handle.

- dbt1

  The **dbt1** parameter is the Dbt representing the application supplied key.

- dbt2

  The **dbt2** parameter is the Dbt representing the current tree's key.

The **bt_compare_fcn** function must return an integer value less than, equal to, or greater than zero if the first key parameter is considered to be respectively less than, equal to, or greater than the second key parameter. In addition, the comparison function must cause the keys in the database to be *well-ordered*. The comparison function must correctly handle any key values used by the application (possibly including zero-length keys). In addition, when Btree key prefix comparison is being performed (see Db::set_bt_prefix() for more information), the comparison routine may be passed a prefix of any

database key. The **data** and **size** fields of the Dbt are the only fields that may be used for the purposes of this comparison, and no particular alignment of the memory to which by the **data** field refers may be assumed.

Errors

The `Db::set_bt_compare()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after Db::open() was called; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::set_bt_compress()

```
#include <db_cxx.h>

extern "C" {
    typedef int (*bt_compress_fcn_type)(DB *db, const DBT *prevKey,
        const DBT *prevData, const DBT *key, const DBT *data, DBT *dest);
    typedef int (*bt_decompress_fcn_typ)(DB *db, const DBT *prevKey,
        const DBT *prevData, DBT *compressed, DBT *destKey,
        DBT *destData);
};
int
Db::set_bt_compress(bt_compress_fcn_type bt_compress_fcn,
    bt_decompress_fcn_type bt_decompress_fcn);
```

Set the Btree compression and decompression functions. The compression function is called whenever a key/data pair is added to the tree and the decompression function is called whenever data is requested from the tree.

If NULL function pointers are specified, then default compression and decompression functions are used. Berkeley DB's default compression function performs prefix compression on all keys and prefix compression on data values for duplicate keys. If using default compression, both the default compression and decompression functions must be used.

The `Db::set_bt_compress()` method configures operations performed using the specified Db handle, not all operations performed on the underlying database.

The `Db::set_bt_compress()` method may not be called after the Db::open() method is called. If the database already exists when Db::open() is called, the information specified to `Db::set_bt_compress()` must be the same as that historically used to create the database or corruption can occur.

The `Db::set_bt_compress()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

bt_compress_fcn

The **bt_compress_fcn** function is the application-specified Btree compression function. The compression function takes six parameters:

- db

  The **db** parameter is the enclosing database handle.

- prevKey

  The **prevKey** parameter is the Dbt representing the key immediately preceding the application supplied key.

- prevData

The **prevData** parameter is the Dbt representing the data associated with **prevKey**.

- key

    The **key** parameter is the Dbt representing the application supplied key.

- data

    The **data** parameter is the Dbt representing the application supplied data.

- dest

    The **dest** parameter is the Dbt representing the data stored in the tree, where the function should write the compressed data.

The **bt_compress_fcn** function must return 0 on success and a non-zero value on failure. If the compressed data cannot fit in **dest->set_data()** (the size of which is returned by **dest->get_ulen()**), the function should identify the required buffer size in **dest->set_size()** and return DB_BUFFER_SMALL.

bt_decompress_fcn

The **bt_decompress_fcn** function is the application-specified Btree decompression function. The decompression function takes six parameters:

- db

    The **db** parameter is the enclosing database handle.

- prevKey

    The **prevKey** parameter is the Dbt representing the key immediately preceding the key being decompressed.

- prevData

    The **prevData** parameter is the Dbt representing the data associated with **prevKey**.

- compressed

    The **compressed** parameter is the Dbt representing the data stored in the tree, that is, the compressed data.

- key

    The **key** parameter is the Dbt where the decompression function should store the decompressed key.

- data

    The **data** parameter is the Dbt where the decompression function should store the decompressed key.

The **bt_decompress_fcn** function must return 0 on success and a non-zero value on failure. If the decompressed data cannot fit in **key->set_data()** or **data->set_data()** (the size of which is returned by the Dbt's **get_ulen()** method), the function should identify the required buffer size using the Dbt's **set_size()** method and return DB_BUFFER_SMALL.

Errors

The Db::set_bt_compress() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after Db::open() was called; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::set_bt_minkey()

```
#include <db_cxx.h>

int
Db::set_bt_minkey(u_int32_t bt_minkey);
```

Set the minimum number of key/data pairs intended to be stored on any single Btree leaf page.

This value is used to determine if key or data items will be stored on overflow pages instead of Btree leaf pages. For more information on the specific algorithm used, see Minimum keys per page. The **bt_minkey** value specified must be at least 2; if **bt_minkey** is not explicitly set, a value of 2 is used.

The Db::set_bt_minkey() method configures a database, not only operations performed using the specified Db handle.

The Db::set_bt_minkey() method may not be called after the Db::open() method is called. If the database already exists when Db::open() is called, the information specified to Db::set_bt_minkey() will be ignored.

The Db::set_bt_minkey() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

bt_minkey

The **bt_minkey** parameter is the minimum number of key/data pairs intended to be stored on any single Btree leaf page.

Errors

The Db::set_bt_minkey() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after Db::open() was called; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::set_bt_prefix()

```
#include <db_cxx.h>

extern "C" {
    typedef size_t (*bt_prefix_fcn_type)(DB *, const DBT *, const DBT *);
};
int
Db::set_bt_prefix(bt_prefix_fcn_type bt_prefix_fcn);
```

Set the Btree prefix function. The prefix function is used to determine the amount by which keys stored on the Btree internal pages can be safely truncated without losing their uniqueness. See the Btree prefix comparison section of the Berkeley DB Reference Guide for more details about how this works. The usefulness of this is data-dependent, but can produce significantly reduced tree sizes and search times in some data sets.

If no prefix function or key comparison function is specified by the application, a default lexical comparison function is used as the prefix function. If no prefix function is specified and a key comparison function is specified, no prefix function is used. It is an error to specify a prefix function without also specifying a Btree key comparison function.

The `Db::set_bt_prefix()` method configures operations performed using the specified Db handle, not all operations performed on the underlying database.

The `Db::set_bt_prefix()` method may not be called after the Db::open() method is called. If the database already exists when Db::open() is called, the information specified to `Db::set_bt_prefix()` must be the same as that historically used to create the database or corruption can occur.

The `Db::set_bt_prefix()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

bt_prefix_fcn

The **bt_prefix_fcn** function is the application-specific Btree prefix function. The prefix function takes three parameters:

- db

  The **db** parameter is the enclosing database handle.

- dbt1

  The **dbt1** parameter is a Dbt representing a database key.

- dbt2

  The **dbt2** parameter is a Dbt representing a database key.

The **bt_prefix_fcn** function must return the number of bytes of the second key parameter that would be required by the Btree key comparison function to determine the second key parameter's ordering relationship with respect to the first key parameter. If the two keys are equal, the key length should be returned. The prefix function must correctly handle any key values used by the application (possibly including zero-length keys). The **data** and **size** fields of the Dbt are the only fields that may be used for the purposes of this determination, and no particular alignment of the memory to which the **data** field refers may be assumed.

Errors

The Db::set_bt_prefix() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after Db::open() was called; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::set_cachesize()

```
#include <db_cxx.h>

int
Db::set_cachesize(u_int32_t gbytes, u_int32_t bytes, int ncache);
```

Set the size of the shared memory buffer pool -- that is, the cache. The cache should be the size of the normal working data set of the application, with some small amount of additional memory for unusual situations. (Note: the working set is not the same as the number of pages accessed simultaneously, and is usually much larger.)

The default cache size is 256KB, and may not be specified as less than 20KB. Any cache size less than 500MB is automatically increased by 25% to account for buffer pool overhead; cache sizes larger than 500MB are used as specified. The maximum size of a single cache is 4GB on 32-bit systems and 10TB on 64-bit systems. (All sizes are in powers-of-two, that is, 256KB is 2^18 not 256,000.) For information on tuning the Berkeley DB cache size, see Selecting a cache size.

It is possible to specify caches to Berkeley DB large enough they cannot be allocated contiguously on some architectures. For example, some releases of Solaris limit the amount of memory that may be allocated contiguously by a process. If **ncache** is 0 or 1, the cache will be allocated contiguously in memory. If it is greater than 1, the cache will be split across **ncache** separate regions, where the **region size** is equal to the initial cache size divided by **ncache**.

Because databases opened within Berkeley DB environments use the cache specified to the environment, it is an error to attempt to set a cache in a database created within an environment.

The Db::set_cachesize() method may not be called after the Db::open() method is called.

The Db::set_cachesize() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

bytes

The size of the cache is set to **gbytes** gigabytes plus **bytes**.

gbytes

The size of the cache is set to **gbytes** gigabytes plus **bytes**.

ncache

The **ncache** parameter is the number of caches to create.

Errors

The Db::set_cachesize() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the specified cache size was impossibly small; the method was called after Db::open() was called; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::set_create_dir()

```
#include <db_cxx.h>

int
Db::set_create_dir(const char *dir);
```

Specify which directory a database should be created in or looked for.

The `Db::set_create_dir()` method may not be called after the Db::open() method is called.

The `Db::set_create_dir()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

dir

The **dir** will be used to create or locate the database file specified in the Db::open() method call. The directory must be one of the directories in the environment list specified by DbEnv::add_data_dir().

Errors

The `Db::set_create_dir()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::set_dup_compare()

```
#include <db_cxx.h>

extern "C" {
    typedef int (*dup_compare_fcn_type)(DB *db, const DBT *dbt1,
    const DBT *dbt2);
};
int
Db::set_dup_compare(dup_compare_fcn_type dup_compare_fcn);
```

Set the duplicate data item comparison function. The comparison function is called whenever it is necessary to compare a data item specified by the application with a data item currently stored in the database. Calling Db::set_dup_compare() implies calling Db::set_flags() with the DB_DUPSORT flag.

If no comparison function is specified, the data items are compared lexically, with shorter data items collating before longer data items.

The Db::set_dup_compare() method may not be called after the Db::open() method is called. If the database already exists when Db::open() is called, the information specified to Db::set_dup_compare() must be the same as that historically used to create the database or corruption can occur.

The Db::set_dup_compare() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

dup_compare_fcn

The **dup_compare_fcn** function is the application-specified duplicate data item comparison function. The function takes three arguments:

- db

  The **db** parameter is the enclosing database handle.

- dbt1

  The **dbt1** parameter is a Dbt representing the application supplied data item.

- dbt2

  The **dbt2** parameter is a Dbt representing the current tree's data item.

The **dup_compare_fcn** function must return an integer value less than, equal to, or greater than zero if the first data item parameter is considered to be respectively less than, equal to, or greater than the second data item parameter. In addition, the comparison function must cause the data items in the set to be *well-ordered*. The comparison function must correctly handle any data item values used by the application (possibly including zero-length data items). The **data** and **size** fields of the Dbt are the only fields that may be used for the purposes of this comparison, and no particular alignment of the memory to which the **data** field refers may be assumed.

Errors

The `Db::set_dup_compare()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::set_encrypt()

```
#include <db_cxx.h>

int
Db::set_encrypt(const char *passwd, u_int32_t flags);
```

Set the password used by the Berkeley DB library to perform encryption and decryption.

Because databases opened within Berkeley DB environments use the password specified to the environment, it is an error to attempt to set a password in a database created within an environment.

The Db::set_encrypt() method may not be called after the Db::open() method is called.

The Db::set_encrypt() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or the following value:

- DB_ENCRYPT_AES

  Use the Rijndael/AES (also known as the Advanced Encryption Standard and Federal Information Processing Standard (FIPS) 197) algorithm for encryption or decryption.

passwd

The **passwd** parameter is the password used to perform encryption and decryption.

Errors

The Db::set_encrypt() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after Db::open() was called; or if an invalid flag value or parameter was specified.

EOPNOTSUPP

Cryptography is not available in this Berkeley DB release.

Class

Db

See Also

Database and Related Methods

Db::set_errcall()

```
#include <db_cxx.h>

void Db::set_errcall(void (*db_errcall_fcn)
    (const DbEnv *dbenv, const char *errpfx, const char *msg));
```

When an error occurs in the Berkeley DB library, an exception is thrown or an error return value is returned by the interface. In some cases, however, the **errno** value may be insufficient to completely describe the cause of the error, especially during initial application debugging.

The DbEnv::set_errcall() and Db::set_errcall() methods are used to enhance the mechanism for reporting error messages to the application. In some cases, when an error occurs, Berkeley DB will call **db_errcall_fcn()** with additional error information. It is up to the **db_errcall_fcn()** function to display the error message in an appropriate manner.

Setting **db_errcall_fcn** to NULL unconfigures the callback interface.

Alternatively, you can use the DbEnv::set_error_stream() and Db::set_error_stream() methods to display the additional information via an output stream, or the Db::set_errfile() or Db::set_errfile() methods to display the additional information via a C library FILE *. You should not mix these approaches.

This error-logging enhancement does not slow performance or significantly increase application size, and may be run during normal operation as well as during application debugging.

For Db handles opened inside of Berkeley DB environments, calling the Db::set_errcall() method affects the entire environment and is equivalent to calling the DbEnv::set_errcall() method.

When used on a database that was *not* opened in an environment, the Db::set_errcall() method configures operations performed using the specified Db handle, not all operations performed on the underlying database.

The Db::set_errcall() method may be called at any time during the life of the application.

☞ Berkeley DB is not re-entrant. Callback functions should not attempt to make library calls (for example, to release locks or close open handles). Re-entering Berkeley DB is not guaranteed to work correctly, and the results are undefined.

Parameters

db_errcall_fcn

The **db_errcall_fcn** parameter is the application-specified error reporting function. The function takes three parameters:

- dbenv

  The **dbenv** parameter is the enclosing database environment.

- errpfx

The **errpfx** parameter is the prefix string (as previously set by Db::set_errpfx() or DbEnv::set_errpfx()).

- `msg`

The **msg** parameter is the error message string.

Class

Db

See Also

Database and Related Methods

Db::set_errfile()

```
#include <db_cxx.h>

void Db::set_errfile(FILE *errfile);
```

When an error occurs in the Berkeley DB library, an exception is thrown or an error return value is returned by the interface. In some cases, however, the **errno** value may be insufficient to completely describe the cause of the error, especially during initial application debugging.

The DbEnv::set_errfile() and Db::set_errfile() methods are used to enhance the mechanism for reporting error messages to the application by setting a C library FILE * to be used for displaying additional Berkeley DB error messages. In some cases, when an error occurs, Berkeley DB will output an additional error message to the specified file reference.

Alternatively, you can use the DbEnv::set_error_stream() and Db::set_error_stream() methods to display the additional messages via an output stream, or the DbEnv::set_errcall() or Db::set_errcall() methods to capture the additional error information in a way that does not use C library FILE *'s. You should not mix these approaches.

The error message will consist of the prefix string and a colon ("**:**") (if a prefix string was previously specified using Db::set_errpfx() or DbEnv::set_errpfx()), an error string, and a trailing <newline> character.

The default configuration when applications first create Db or DbEnv handles is as if the DbEnv::set_errfile() or Db::set_errfile() methods were called with the standard error output (stderr) specified as the FILE * argument. Applications wanting no output at all can turn off this default configuration by calling the DbEnv::set_errfile() or Db::set_errfile() methods with NULL as the FILE * argument. Additionally, explicitly configuring the error output channel using any of the following methods will also turn off this default output for the application:

- Db::set_errfile()

- DbEnv::set_errfile()

- DbEnv::set_errcall()

- Db::set_errcall()

- DbEnv::set_error_stream()

- Db::set_error_stream()

This error logging enhancement does not slow performance or significantly increase application size, and may be run during normal operation as well as during application debugging.

For Db handles opened inside of Berkeley DB environments, calling the Db::set_errfile() method affects the entire environment and is equivalent to calling the DbEnv::set_errfile() method.

When used on a database that was *not* opened in an environment, the `Db::set_errfile()` method configures operations performed using the specified Db handle, not all operations performed on the underlying database.

The `Db::set_errfile()` method may be called at any time during the life of the application.

Parameters

errfile

The **errfile** parameter is a C library `FILE *` to be used for displaying additional Berkeley DB error information.

Class

Db

See Also

Database and Related Methods

Db::set_error_stream()

```
#include <db_cxx.h>

void Db::set_error_stream(class ostream*);
```

When an error occurs in the Berkeley DB library, an exception is thrown or an **errno** value is returned by the interface. In some cases, however, the **errno** value may be insufficient to completely describe the cause of the error, especially during initial application debugging.

The DbEnv::set_error_stream() and Db::set_error_stream() methods are used to enhance the mechanism for reporting error messages to the application by setting the C++ ostream used for displaying additional Berkeley DB error messages. In some cases, when an error occurs, Berkeley DB will output an additional error message to the specified stream.

The error message will consist of the prefix string and a colon ("**:**") (if a prefix string was previously specified using Db::set_errpfx(), an error string, and a trailing <newline> character.

Setting **stream** to NULL unconfigures the interface.

Alternatively, you can use the DbEnv::set_errfile() or Db::set_errfile() methods to display the additional information via a C Library FILE *, or the DbEnv::set_errcall() and Db::set_errcall() methods to capture the additional error information in a way that does not use either output streams or C Library FILE *'s. You should not mix these approaches.

This error-logging enhancement does not slow performance or significantly increase application size, and may be run during normal operation as well as during application debugging.

For Db handles opened inside of Berkeley DB environments, calling the Db::set_error_stream() method affects the entire environment and is equivalent to calling the DbEnv::set_error_stream() method.

The Db::set_error_stream() method may be called at any time during the life of the application.

Parameters

stream

The **stream** parameter is the application-specified output stream to be used for additional error information.

Class

Db

See Also

Database and Related Methods

Db::set_errpfx()

```
#include <db_cxx.h>

void Db::set_errpfx(const char *errpfx);
```

Set the prefix string that appears before error messages issued by Berkeley DB.

The `Db::set_errpfx()` and DbEnv::set_errpfx() methods do not copy the memory to which the **errpfx** parameter refers; rather, they maintain a reference to it. Although this allows applications to modify the error message prefix at any time (without repeatedly calling the interfaces), it means the memory must be maintained until the handle is closed.

For Db handles opened inside of Berkeley DB environments, calling the `Db::set_errpfx()` method affects the entire environment and is equivalent to calling the DbEnv::set_errpfx() method.

The `Db::set_errpfx()` method configures operations performed using the specified Db handle, not all operations performed on the underlying database.

The `Db::set_errpfx()` method may be called at any time during the life of the application.

Parameters

errpfx

The **errpfx** parameter is the application-specified error prefix for additional error messages.

Class

Db

See Also

Database and Related Methods

Db::set_feedback()

```
#include <db_cxx.h>

int
Db::set_feedback(void (*db_feedback_fcn)(DB *dbp, int opcode, int percent));
```

Some operations performed by the Berkeley DB library can take non-trivial amounts of time. The `Db::set_feedback()` method can be used by applications to monitor progress within these operations. When an operation is likely to take a long time, Berkeley DB will call the specified callback function with progress information.

It is up to the callback function to display this information in an appropriate manner.

The `Db::set_feedback()` method may be called at any time during the life of the application.

The `Db::set_feedback()` method returns a non-zero error value on failure and 0 on success.

☞ Berkeley DB is not re-entrant. Callback functions should not attempt to make library calls (for example, to release locks or close open handles). Re-entering Berkeley DB is not guaranteed to work correctly, and the results are undefined.

Parameters

db_feedback_fcn

The **db_feedback_fcn** parameter is the application-specified feedback function called to report Berkeley DB operation progress. The callback function must take three parameters:

• dbp

   The **dbp** parameter is a reference to the enclosing database.

• opcode

   The **opcode** parameter is an operation code. The **opcode** parameter may take on any of the following values:

   • DB_UPGRADE

      The underlying database is being upgraded.

   • DB_VERIFY

      The underlying database is being verified.

• percent

   The **percent** parameter is the percent of the operation that has been completed, specified as an integer value between 0 and 100.

Class

Db

See Also

Database and Related Methods

Db::set_flags()

```
#include <db_cxx.h>

int
Db::set_flags(u_int32_t flags);
```

Configure a database. Calling `Db::set_flags()` is additive; there is no way to clear flags.

The `Db::set_flags()` method may not be called after the Db::open() method is called.

The `Db::set_flags()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

General

The following flags may be specified for any Berkeley DB access method:

- `DB_CHKSUM`

  Do checksum verification of pages read into the cache from the backing filestore. Berkeley DB uses the SHA1 Secure Hash Algorithm if encryption is configured and a general hash algorithm if it is not.

  Calling `Db::set_flags()` with the DB_CHKSUM flag only affects the specified Db handle (and any other Berkeley DB handles opened within the scope of that handle).

  If the database already exists when Db::open() is called, the DB_CHKSUM flag will be ignored.

- `DB_ENCRYPT`

  Encrypt the database using the cryptographic password specified to the DbEnv::set_encrypt() or Db::set_encrypt() methods.

  Calling `Db::set_flags()` with the DB_ENCRYPT flag only affects the specified Db handle (and any other Berkeley DB handles opened within the scope of that handle).

  If the database already exists when Db::open() is called, the DB_ENCRYPT flag must be the same as the existing database or an error will be returned.

  Encrypted databases are not portable between machines of different byte orders, that is, encrypted databases created on big-endian machines cannot be read on little-endian machines, and vice versa.

- `DB_TXN_NOT_DURABLE`

If set, Berkeley DB will not write log records for this database. This means that updates of this database exhibit the ACI (atomicity, consistency, and isolation) properties, but not D (durability); that is, database integrity will be maintained, but if the application or system fails, integrity will not persist. The database file must be verified and/or restored from backup after a failure. In order to ensure integrity after application shut down, the database handles must be closed without specifying DB_NOSYNC, or all database changes must be flushed from the database environment cache using either the DbEnv::txn_checkpoint() or DbEnv::memp_sync() methods. All database handles for a single physical file must set DB_TXN_NOT_DURABLE, including database handles for different databases in a physical file.

Calling `Db::set_flags()` with the DB_TXN_NOT_DURABLE flag only affects the specified Db handle (and any other Berkeley DB handles opened within the scope of that handle).

Btree

The following flags may be specified for the Btree access method:

- `DB_DUP`

  Permit duplicate data items in the database; that is, insertion when the key of the key/data pair being inserted already exists in the database will be successful. The ordering of duplicates in the database is determined by the order of insertion, unless the ordering is otherwise specified by use of a cursor operation or a duplicate sort function.

  The DB_DUPSORT flag is preferred to DB_DUP for performance reasons. The DB_DUP flag should only be used by applications wanting to order duplicate data items manually.

  Calling `Db::set_flags()` with the DB_DUP flag affects the database, including all threads of control accessing the database.

  If the database already exists when Db::open() is called, the DB_DUP flag must be the same as the existing database or an error will be returned.

  It is an error to specify both DB_DUP and DB_RECNUM.

- `DB_DUPSORT`

  Permit duplicate data items in the database; that is, insertion when the key of the key/data pair being inserted already exists in the database will be successful. The ordering of duplicates in the database is determined by the duplicate comparison function. If the application does not specify a comparison function using the Db::set_dup_compare() method, a default lexical comparison will be used. It is an error to specify both DB_DUPSORT and DB_RECNUM.

  Calling `Db::set_flags()` with the DB_DUPSORT flag affects the database, including all threads of control accessing the database.

  If the database already exists when Db::open() is called, the DB_DUPSORT flag must be the same as the existing database or an error will be returned.

- `DB_RECNUM`

Support retrieval from the Btree using record numbers. For more information, see the DB_SET_RECNO flag to the Db::get() and Dbc::get() methods.

Logical record numbers in Btree databases are mutable in the face of record insertion or deletion. See the DB_RENUMBER flag in the Recno access method information for further discussion.

Maintaining record counts within a Btree introduces a serious point of contention, namely the page locations where the record counts are stored. In addition, the entire database must be locked during both insertions and deletions, effectively single-threading the database for those operations. Specifying DB_RECNUM can result in serious performance degradation for some applications and data sets.

It is an error to specify both DB_DUP and DB_RECNUM.

Calling Db::set_flags() with the DB_RECNUM flag affects the database, including all threads of control accessing the database.

If the database already exists when Db::open() is called, the DB_RECNUM flag must be the same as the existing database or an error will be returned.

- DB_REVSPLITOFF

Turn off reverse splitting in the Btree. As pages are emptied in a database, the Berkeley DB Btree implementation attempts to coalesce empty pages into higher-level pages in order to keep the database as small as possible and minimize search time. This can hurt performance in applications with cyclical data demands; that is, applications where the database grows and shrinks repeatedly. For example, because Berkeley DB does page-level locking, the maximum level of concurrency in a database of two pages is far smaller than that in a database of 100 pages, so a database that has shrunk to a minimal size can cause severe deadlocking when a new cycle of data insertion begins.

Calling Db::set_flags() with the DB_REVSPLITOFF flag only affects the specified Db handle (and any other Berkeley DB handles opened within the scope of that handle).

Hash

The following flags may be specified for the Hash access method:

- DB_DUP

Permit duplicate data items in the database; that is, insertion when the key of the key/data pair being inserted already exists in the database will be successful. The ordering of duplicates in the database is determined by the order of insertion, unless the ordering is otherwise specified by use of a cursor operation.

The DB_DUPSORT flag is preferred to DB_DUP for performance reasons. The DB_DUP flag should only be used by applications wanting to order duplicate data items manually.

Calling Db::set_flags() with the DB_DUP flag affects the database, including all threads of control accessing the database.

If the database already exists when Db::open() is called, the DB_DUP flag must be the same as the existing database or an error will be returned.

- DB_DUPSORT

  Permit duplicate data items in the database; that is, insertion when the key of the key/data pair being inserted already exists in the database will be successful. The ordering of duplicates in the database is determined by the duplicate comparison function. If the application does not specify a comparison function using the Db::set_dup_compare() method, a default lexical comparison will be used.

  Calling Db::set_flags() with the DB_DUPSORT flag affects the database, including all threads of control accessing the database.

  If the database already exists when Db::open() is called, the DB_DUPSORT flag must be the same as the existing database or an error will be returned.

Queue

The following flags may be specified for the Queue access method:

- DB_INORDER

  The DB_INORDER flag modifies the operation of the DB_CONSUME or DB_CONSUME_WAIT flags to Db::get() to return key/data pairs in order. That is, they will always return the key/data item from the head of the queue.

  The default behavior of queue databases is optimized for multiple readers, and does not guarantee that record will be retrieved in the order they are added to the queue. Specifically, if a writing thread adds multiple records to an empty queue, reading threads may skip some of the initial records when the next Db::get() call returns.

  This flag modifies the Db::get() call to verify that the record being returned is in fact the head of the queue. This will increase contention and reduce concurrency when there are many reading threads.

  Calling Db::set_flags() with the DB_INORDER flag only affects the specified Db handle (and any other Berkeley DB handles opened within the scope of that handle).

Recno

The following flags may be specified for the Recno access method:

- DB_RENUMBER

  Specifying the DB_RENUMBER flag causes the logical record numbers to be mutable, and change as records are added to and deleted from the database.

  Using the Db::put() or Dbc::put() interfaces to create new records will cause the creation of multiple records if the record number is more than one greater than the largest record currently in the database. For example, creating record 28, when record 25 was previously the last record in the database, will create records 26 and 27 as well as 28. Attempts to retrieve records that were created in this manner will result in an error return of DB_KEYEMPTY.

If a created record is not at the end of the database, all records following the new record will be automatically renumbered upward by one. For example, the creation of a new record numbered 8 causes records numbered 8 and greater to be renumbered upward by one. If a cursor was positioned to record number 8 or greater before the insertion, it will be shifted upward one logical record, continuing to refer to the same record as it did before.

If a deleted record is not at the end of the database, all records following the removed record will be automatically renumbered downward by one. For example, deleting the record numbered 8 causes records numbered 9 and greater to be renumbered downward by one. If a cursor was positioned to record number 9 or greater before the removal, it will be shifted downward one logical record, continuing to refer to the same record as it did before.

If a record is deleted, all cursors that were positioned on that record prior to the removal will no longer be positioned on a valid entry. This includes cursors used to delete an item. For example, if a cursor was positioned to record number 8 before the removal of that record, subsequent calls to Dbc::get() with flags of DB_CURRENT will result in an error return of DB_KEYEMPTY until the cursor is moved to another record. A call to Dbc::get() with flags of DB_NEXT will return the new record numbered 8 - which is the record that was numbered 9 prior to the delete (if such a record existed).

For these reasons, concurrent access to a Recno database with the DB_RENUMBER flag specified may be largely meaningless, although it is supported.

Calling Db::set_flags() with the DB_RENUMBER flag affects the database, including all threads of control accessing the database.

If the database already exists when Db::open() is called, the DB_RENUMBER flag must be the same as the existing database or an error will be returned.

- DB_SNAPSHOT

  This flag specifies that any specified **re_source** file be read in its entirety when Db::open() is called. If this flag is not specified, the **re_source** file may be read lazily.

  See the Db::set_re_source() method for information on the **re_source** file.

  Calling Db::set_flags() with the DB_SNAPSHOT flag only affects the specified Db handle (and any other Berkeley DB handles opened within the scope of that handle).

Errors

The Db::set_flags() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::set_h_compare()

```
#include <db_cxx.h>

extern "C" {
    typedef int (*compare_fcn_type)(DB *db, const DBT *dbt1, const DBT *dbt2);
};
int
Db::set_h_compare(compare_fcn_type compare_fcn);
```

Set the Hash key comparison function. The comparison function is called whenever it is necessary to compare a key specified by the application with a key currently stored in the database.

If no comparison function is specified, a byte-by-byte comparison is performed.

The Db::set_h_compare() method configures operations performed using the specified Db handle, not all operations performed on the underlying database.

The Db::set_h_compare() method may not be called after the Db::open() method is called. If the database already exists when Db::open() is called, the information specified to Db::set_h_compare() must be the same as that historically used to create the database or corruption can occur.

The Db::set_h_compare() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

compare_fcn

The **compare_fcn** function is the application-specified Hash comparison function. The comparison function takes three parameters:

- db

  The **db** parameter is the enclosing database handle.

- dbt1

  The **dbt1** parameter is the Dbt representing the application supplied key.

- dbt2

  The **dbt2** parameter is the Dbt representing the current database's key.

The **compare_fcn** function must return an integer value less than, equal to, or greater than zero if the first key parameter is considered to be respectively less than, equal to, or greater than the second key parameter. The comparison function must correctly handle any key values used by the application (possibly including zero-length keys). The **data** and **size** fields of the Dbt are the only fields that may be used for the purposes of this comparison, and no particular alignment of the memory to which by the **data** field refers may be assumed.

Errors

The `Db::set_h_compare()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after Db::open() was called; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::set_h_ffactor()

```
#include <db_cxx.h>

int
Db::set_h_ffactor(u_int32_t h_ffactor);
```

Set the desired density within the hash table. If no value is specified, the fill factor will be selected dynamically as pages are filled.

The density is an approximation of the number of keys allowed to accumulate in any one bucket, determining when the hash table grows or shrinks. If you know the average sizes of the keys and data in your data set, setting the fill factor can enhance performance. A reasonable rule computing fill factor is to set it to the following:

```
            (pagesize - 32) / (average_key_size + average_data_size + 8)
```

The Db::set_h_ffactor() method configures a database, not only operations performed using the specified Db handle.

The Db::set_h_ffactor() method may not be called after the Db::open() method is called. If the database already exists when Db::open() is called, the information specified to Db::set_h_ffactor() will be ignored.

The Db::set_h_ffactor() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

h_ffactor

The **h_ffactor** parameter is the desired density within the hash table.

Errors

The Db::set_h_ffactor() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after Db::open() was called; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::set_h_hash()

```
#include <db_cxx.h>

extern "C" {
    typedef u_int32_t (*h_hash_fcn_type)
 (DB *, const void *bytes, u_int32_t length);
};
int
Db::set_h_hash(h_hash_fcn_type h_hash_fcn);
```

Set a user-defined hash function; if no hash function is specified, a default hash function is used. Because no hash function performs equally well on all possible data, the user may find that the built-in hash function performs poorly with a particular data set.

The Db::set_h_hash() method configures operations performed using the specified Db handle, not all operations performed on the underlying database.

The Db::set_h_hash() method may not be called after the Db::open() method is called. If the database already exists when Db::open() is called, the information specified to Db::set_h_hash() must be the same as that historically used to create the database or corruption can occur.

The Db::set_h_hash() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

h_hash_fcn

The **h_hash_fcn** parameter is the application-specified hash function.

Application-specified hash functions take a pointer to a byte string and a length as parameters, and return a value of type **u_int32_t**. The hash function must handle any key values used by the application (possibly including zero-length keys).

Errors

The Db::set_h_hash() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after Db::open() was called; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::set_h_nelem()

```
#include <db_cxx.h>

int
Db::set_h_nelem(u_int32_t h_nelem);
```

Set an estimate of the final size of the hash table.

In order for the estimate to be used when creating the database, the Db::set_h_ffactor() method must also be called. If the estimate or fill factor are not set or are set too low, hash tables will still expand gracefully as keys are entered, although a slight performance degradation may be noticed.

The Db::set_h_nelem() method configures a database, not only operations performed using the specified Db handle.

The Db::set_h_nelem() method may not be called after the Db::open() method is called. If the database already exists when Db::open() is called, the information specified to Db::set_h_nelem() will be ignored.

The Db::set_h_nelem() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

h_nelem

The **h_nelem** parameter is an estimate of the final size of the hash table.

Errors

The Db::set_h_nelem() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after Db::open() was called; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::set_lorder()

```
#include <db_cxx.h>

int
Db::set_lorder(int lorder);
```

Set the byte order for integers in the stored database metadata. The host byte order of the machine where the Berkeley DB library was compiled will be used if no byte order is set.

The access methods provide no guarantees about the byte ordering of the application data stored in the database, and applications are responsible for maintaining any necessary ordering.

The Db::set_lorder() method configures a database, not only operations performed using the specified Db handle.

The Db::set_lorder() method may not be called after the Db::open() method is called. If the database already exists when Db::open() is called, the information specified to Db::set_lorder() will be ignored.

If creating additional databases in a single physical file, information specified to Db::set_lorder() will be ignored and the byte order of the existing databases will be used.

The Db::set_lorder() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

lorder

The **lorder** parameter should represent the byte order as an integer; for example, big endian order is the number 4,321, and little endian order is the number 1,234.

Errors

The Db::set_lorder() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after Db::open() was called; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::set_message_stream()

```
#include <db_cxx.h>

void Db::set_message_stream(class ostream*);
```

There are interfaces in the Berkeley DB library which either directly output informational messages or statistical information, or configure the library to output such messages when performing other operations. For example, the DbEnv::set_verbose() and DbEnv::stat_print() methods.

The DbEnv::set_message_stream() and Db::set_message_stream() methods are used to display these messages for the application. In this case, the message will include a trailing <newline> character.

Setting **stream** to NULL unconfigures the interface.

Alternatively, you can use the DbEnv::set_msgfile() or Db::set_msgfile() methods to display the additional information via a C Library FILE *, or the DbEnv::set_msgcall() and Db::set_msgcall() methods to capture the additional error information in a way that does not use either output streams or C Library FILE *'s. You should not mix these approaches.

For Db handles opened inside of Berkeley DB environments, calling the Db::set_message_stream() method affects the entire environment and is equivalent to calling the DbEnv::set_message_stream() method.

The Db::set_message_stream() method configures operations performed using the specified Db handle, not all operations performed on the underlying database.

The Db::set_message_stream() method may be called at any time during the life of the application.

Parameters

stream

The **stream** parameter is the application-specified output stream to be used for additional message information.

Class

Db

See Also

Database and Related Methods

Db::set_msgcall()

```
#include <db_cxx.h>

void Db::set_msgcall(void (*db_msgcall_fcn)(const DbEnv *dbenv, char *msg));
```

There are interfaces in the Berkeley DB library which either directly output informational messages or statistical information, or configure the library to output such messages when performing other operations, for example, DbEnv::set_verbose() and DbEnv::stat_print().

The DbEnv::set_msgcall() and Db::set_msgcall() methods are used to pass these messages to the application, and Berkeley DB will call **db_msgcall_fcn** with each message. It is up to the **db_msgcall_fcn** function to display the message in an appropriate manner.

Setting **db_msgcall_fcn** to NULL unconfigures the callback interface.

Alternatively, you can use the DbEnv::set_error_stream() and Db::set_error_stream() methods to display the messages via an output stream, or the Db::set_msgfile() or Db::set_msgfile() methods to display the messages via a C library FILE *. You should not mix these approaches.

For Db handles opened inside of Berkeley DB environments, calling the Db::set_msgcall() method affects the entire environment and is equivalent to calling the DbEnv::set_msgcall() method.

The Db::set_msgcall() method configures operations performed using the specified Db handle, not all operations performed on the underlying database.

The Db::set_msgcall() method may be called at any time during the life of the application.

☞ Berkeley DB is not re-entrant. Callback functions should not attempt to make library calls (for example, to release locks or close open handles). Re-entering Berkeley DB is not guaranteed to work correctly, and the results are undefined.

Parameters

db_msgcall_fcn

The **db_msgcall_fcn** parameter is the application-specified message reporting function. The function takes two parameters:

- dbenv

  The **dbenv** parameter is the enclosing database environment.

- msg

  The **msg** parameter is the message string.

Class

Db

See Also

Database and Related Methods

Db::set_msgfile()

```
#include <db_cxx.h>

void Db::set_msgfile(FILE *msgfile);
```

There are interfaces in the Berkeley DB library which either directly output informational messages or statistical information, or configure the library to output such messages when performing other operations, for example, DbEnv::set_verbose() and DbEnv::stat_print().

The DbEnv::set_msgfile() and Db::set_msgfile() methods are used to display these messages for the application. In this case the message will include a trailing <newline> character.

Setting **msgfile** to NULL unconfigures the interface.

Alternatively, you can use the DbEnv::set_message_stream() and Db::set_message_stream() methods to display the messages via an output stream, or the DbEnv::set_msgcall() or Db::set_msgcall() methods to capture the additional error information in a way that does not use C library FILE *'s. You should not mix these approaches.

For Db handles opened inside of Berkeley DB environments, calling the Db::set_msgfile() method affects the entire environment and is equivalent to calling the DbEnv::set_msgfile() method.

The Db::set_msgfile() method configures operations performed using the specified Db handle, not all operations performed on the underlying database.

The Db::set_msgfile() method may be called at any time during the life of the application.

Parameters

msgfile

The **msgfile** parameter is a C library FILE * to be used for displaying messages.

Class

Db

See Also

Database and Related Methods

Db::set_pagesize()

```
#include <db_cxx.h>

int
Db::set_pagesize(u_int32_t pagesize);
```

Set the size of the pages used to hold items in the database, in bytes. The minimum page size is 512 bytes, the maximum page size is 64K bytes, and the page size must be a power-of-two. If the page size is not explicitly set, one is selected based on the underlying filesystem I/O block size. The automatically selected size has a lower limit of 512 bytes and an upper limit of 16K bytes.

For information on tuning the Berkeley DB page size, see Selecting a page size.

The `Db::set_pagesize()` method configures a database, not only operations performed using the specified Db handle.

The `Db::set_pagesize()` method may not be called after the Db::open() method is called. If the database already exists when Db::open() is called, the information specified to `Db::set_pagesize()` will be ignored.

If creating additional databases in a single physical file, information specified to `Db::set_pagesize()` will be ignored and the page size of the existing databases will be used.

The `Db::set_pagesize()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

The `Db::set_pagesize()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

pagesize

The **pagesize** parameter sets the database page size.

Errors

The `Db::set_pagesize()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after Db::open() was called; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::set_partition()

```
#include <db_cxx.h>

int
Db::set_partition(u_int32_t parts, DBT *kyes,
 u_int32_t (*db_partition_fcn) (Db *db, DBT *key));
```

Set up partitioning for a database. Partitioning may be used on either BTREE or HASH databases. Partitions may be specified by either a set of keys specifying a range of values in each partition or with a callback function that returns the number of the partition to put a specific key. Partition range keys may only be specified for BTREE databases.

Partitions are implimented as separate database files and can help reduce contention within a logical database. Contention can come from multiple threads of control accessing database pages simultaniously. Typically these pages are the root of a btree and the metadata page which contains allocation information in both BTREE and HASH databases. Each partition has its own metadata and root pages.

Parameters

Exactly one of the parameters **keys** and **partition_fcn** must be NULL.

parts

The **parts** parameter is the number of partitions to create. The value must be 2 or greater.

keys

The **keys** parameter is an array of DBT structures containing the keys that specify the range of key values to be stored in each partition. Each key specifies the minimum value that may be stored in the corresponding partition. The number of keys must be one less than the number of partitions specified by the **parts** parameter since the first partition will hold any key less than the first key in the array.

db_partition_fcn

The **db_partition_fcn** parameter is the application-specified partitioning function. The function returns an integer which will be used modulo the number of partitions specified by the **parts** parameter. The function will be called with two parameters:

• db

  The **db** parameter is the database handle.

• key

  The **key** parameter is the key for which a partition number should be returned.

Class

Db

See Also

Database and Related Methods

Db::set_partition_dirs()

```
#include <db_cxx.h>

int
Db::set_partition_dirs(const char **dirs);
```

Specify which directories the database extents should be created in or looked for. If the number of directories is less than the number of partitions, the directories will be used in a round robin fashion.

The `Db::set_partition_dirs()` method may not be called after the Db::open() method is called.

The `Db::set_partition_dirs()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

dirs

The **dirs** points to an array of directories that will be used to create or locate the database extent files specified in the Db::open() method call. The directories must be included in the environment list specified by DbEnv::add_data_dir().

Errors

The `Db::set_partition_dirs()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::set_priority()

```
#include <db_cxx.h>

int
Db::set_priority(DB_CACHE_PRIORITY priority);
```

Set the cache priority for pages referenced by the Db handle.

The priority of a page biases the replacement algorithm to be more or less likely to discard a page when space is needed in the buffer pool. The bias is temporary, and pages will eventually be discarded if they are not referenced again. The Db::set_priority() method is only advisory, and does not guarantee pages will be treated in a specific way.

The Db::set_priority() method may be called at any time during the life of the application.

The Db::set_priority() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

priority

The **priority** parameter must be set to one of the following values:

- DB_PRIORITY_VERY_LOW

  The lowest priority: pages are the most likely to be discarded.

- DB_PRIORITY_LOW

  The next lowest priority.

- DB_PRIORITY_DEFAULT

  The default priority.

- DB_PRIORITY_HIGH

  The next highest priority.

- DB_PRIORITY_VERY_HIGH

  The highest priority: pages are the least likely to be discarded.

Class

Db

See Also

Database and Related Methods

Db::set_q_extentsize()

```
#include <db_cxx.h>

int
Db::set_q_extentsize(u_int32_t extentsize);
```

Set the size of the extents used to hold pages in a Queue database, specified as a number of pages. Each extent is created as a separate physical file. If no extent size is set, the default behavior is to create only a single underlying database file.

For information on tuning the extent size, see Selecting a extent size.

The `Db::set_q_extentsize()` method configures a database, not only operations performed using the specified Db handle.

The `Db::set_q_extentsize()` method may not be called after the Db::open() method is called. If the database already exists when Db::open() is called, the information specified to `Db::set_q_extentsize()` will be ignored.

The `Db::set_q_extentsize()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

extentsize

The **extentsize** parameter is the number of pages in a Queue database extent.

Errors

The `Db::set_q_extentsize()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after Db::open() was called; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::set_re_delim()

```
#include <db_cxx.h>

int
Db::set_re_delim(int re_delim);
```

Set the delimiting byte used to mark the end of a record in the backing source file for the Recno access method.

This byte is used for variable length records if the **re_source** file is specified using the Db::set_re_source() method. If the **re_source** file is specified and no delimiting byte was specified, <newline> characters (that is, ASCII 0x0a) are interpreted as end-of-record markers.

The Db::set_re_delim() method configures a database, not only operations performed using the specified Db handle.

The Db::set_re_delim() method may not be called after the Db::open() method is called. If the database already exists when Db::open() is called, the information specified to Db::set_re_delim() will be ignored.

The Db::set_re_delim() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

re_delim

The **re_delim** parameter is the delimiting byte used to mark the end of a record.

Errors

The Db::set_re_delim() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after Db::open() was called; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::set_re_len()

```
#include <db_cxx.h>

int
Db::set_re_len(u_int32_t re_len);
```

For the Queue access method, specify that the records are of length **re_len**. For the Queue access method, the record length must be enough smaller than the database's page size that at least one record plus the database page's metadata information can fit on each database page.

For the Recno access method, specify that the records are fixed-length, not byte-delimited, and are of length **re_len**.

Any records added to the database that are less than **re_len** bytes long are automatically padded (see Db::set_re_pad() for more information).

Any attempt to insert records into the database that are greater than **re_len** bytes long will cause the call to fail immediately and return an error.

The Db::set_re_len() method configures a database, not only operations performed using the specified Db handle.

The Db::set_re_len() method may not be called after the Db::open() method is called. If the database already exists when Db::open() is called, the information specified to Db::set_re_len() will be ignored.

The Db::set_re_len() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

re_len

The **re_len** parameter is the length of a Queue or Recno database record, in bytes.

Errors

The Db::set_re_len() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after Db::open() was called; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::set_re_pad()

```
#include <db_cxx.h>

int
Db::set_re_pad(int re_pad);
```

Set the padding character for short, fixed-length records for the Queue and Recno access methods.

If no pad character is specified, <space> characters (that is, ASCII 0x20) are used for padding.

The Db::set_re_pad() method configures a database, not only operations performed using the specified Db handle.

The Db::set_re_pad() method may not be called after the Db::open() method is called. If the database already exists when Db::open() is called, the information specified to Db::set_re_pad() will be ignored.

The Db::set_re_pad() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

re_pad

The **re_pad** parameter is the pad character for fixed-length records for the Queue and Recno access methods.

Errors

The Db::set_re_pad() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after Db::open() was called; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::set_re_source()

```
#include <db_cxx.h>

int
Db::set_re_source(char *source);
```

Set the underlying source file for the Recno access method. The purpose of the **source** value is to provide fast access and modification to databases that are normally stored as flat text files.

The **source** parameter specifies an underlying flat text database file that is read to initialize a transient record number index. In the case of variable length records, the records are separated, as specified by Db::set_re_delim(). For example, standard UNIX byte stream files can be interpreted as a sequence of variable length records separated by <newline> characters.

In addition, when cached data would normally be written back to the underlying database file (for example, the Db::close() or Db::sync() methods are called), the in-memory copy of the database will be written back to the **source** file.

By default, the backing source file is read lazily; that is, records are not read from the file until they are requested by the application. **If multiple processes (not threads) are accessing a Recno database concurrently, and are either inserting or deleting records, the backing source file must be read in its entirety before more than a single process accesses the database, and only that process should specify the backing source file as part of the Db::open() call. See the DB_SNAPSHOT flag for more information.**

Reading and writing the backing source file specified by source cannot be transaction-protected because it involves filesystem operations that are not part of the Db transaction methodology. For this reason, if a temporary database is used to hold the records, it is possible to lose the contents of the **source** file, for example, if the system crashes at the right instant. If a file is used to hold the database, normal database recovery on that file can be used to prevent information loss, although it is still possible that the contents of **source** will be lost if the system crashes.

The **source** file must already exist (but may be zero-length) when Db::open() is called.

It is not an error to specify a read-only **source** file when creating a database, nor is it an error to modify the resulting database. However, any attempt to write the changes to the backing source file using either the Db::sync() or Db::close() methods will fail, of course. Specify the DB_NOSYNC flag to the Db::close() method to stop it from attempting to write the changes to the backing file; instead, they will be silently discarded.

For all of the previous reasons, the **source** field is generally used to specify databases that are read-only for Berkeley DB applications; and that are either generated on the fly by software tools or modified using a different mechanism — for example, a text editor.

The Db::set_re_source() method configures operations performed using the specified Db handle, not all operations performed on the underlying database.

The `Db::set_re_source()` method may not be called after the Db::open() method is called. If the database already exists when Db::open() is called, the information specified to `Db::set_re_source()` must be the same as that historically used to create the database or corruption can occur.

The `Db::set_re_source()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

source

The backing flat text database file for a Recno database.

When using a Unicode build on Windows (the default), the **source** argument will be interpreted as a UTF-8 string, which is equivalent to ASCII for Latin characters.

Errors

The `Db::set_re_source()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after Db::open() was called; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::stat()

```
#include <db_cxx.h>

int
Db::stat(void *sp, u_int32_t flags);
```

The `Db::stat()` method creates a statistical structure and copies a pointer to it into user-specified memory locations. Specifically, if **sp** is non-NULL, a pointer to the statistics for the database are copied into the memory location to which it refers.

The `Db::stat()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or one of the following values:

- `DB_FAST_STAT`

  Return only the values which do not require traversal of the database. Among other things, this flag makes it possible for applications to request key and record counts without incurring the performance penalty of traversing the entire database.

- `DB_READ_COMMITTED`

  Database items read during a transactional call will have degree 2 isolation. This ensures the stability of the data items read during the stat operation but permits that data to be modified or deleted by other transactions prior to the commit of the specified transaction.

- `DB_READ_UNCOMMITTED`

  Database items read during a transactional call will have degree 1 isolation, including modified but not yet committed data. Silently ignored if the DB_READ_UNCOMMITTED flag was not specified when the underlying database was opened.

txnid

If the operation is part of an application-specified transaction, the **txnid** parameter is a transaction handle returned from DbEnv::txn_begin(); if the operation is part of a Berkeley DB Concurrent Data Store group, the **txnid** parameter is a handle returned from DbEnv::cdsgroup_begin(); otherwise NULL. If no transaction handle is specified, but the operation occurs in a transactional database, the operation will be implicitly transaction protected.

Statistical Structure

Statistical structures are stored in allocated memory. If application-specific allocation routines have been declared (see DbEnv::set_alloc() for more information), they are used to allocate the memory; otherwise, the standard C library **malloc**(3) is used. The caller is responsible for deallocating the

memory. To deallocate the memory, free the memory reference; references inside the returned memory need not be individually freed.

If the DB_FAST_STAT flag has not been specified, the `Db::stat()` method will access some of or all the pages in the database, incurring a severe performance penalty as well as possibly flushing the underlying buffer pool.

In the presence of multiple threads or processes accessing an active database, the information returned by DB->stat may be out-of-date.

If the database was not opened read-only and the DB_FAST_STAT flag was not specified, the cached key and record numbers will be updated after the statistical information has been gathered.

The `Db::stat()` method may not be called before the Db::open() method is called.

The `Db::stat()` method returns a non-zero error value on failure and 0 on success.

Hash Statistics

In the case of a Hash database, the statistics are stored in a structure of type `DB_HASH_STAT`. The following fields will be filled in:

- `u_int32_t hash_magic;`

  Magic number that identifies the file as a Hash file. Returned if DB_FAST_STAT is set.

- `u_int32_t hash_version;`

  The version of the Hash database. Returned if DB_FAST_STAT is set.

- `u_int32_t hash_nkeys;`

  The number of unique keys in the database. If DB_FAST_STAT was specified the count will be the last saved value unless it has never been calculated, in which case it will be 0. Returned if DB_FAST_STAT is set.

- `u_int32_t hash_ndata;`

  The number of key/data pairs in the database. If DB_FAST_STAT was specified the count will be the last saved value unless it has never been calculated, in which case it will be 0. Returned if DB_FAST_STAT is set.

- `u_int32_t hash_pagecnt;`

  The number of pages in the database. Returned if DB_FAST_STAT is set.

- `u_int32_t hash_pagesize;`

  The underlying database page (and bucket) size, in bytes. Returned if DB_FAST_STAT is set.

- `u_int32_t hash_ffactor;`

The desired fill factor (number of items per bucket) specified at database-creation time. Returned if DB_FAST_STAT is set.

- `u_int32_t hash_buckets;`

  The number of hash buckets. Returned if DB_FAST_STAT is set.

- `u_int32_t hash_free;`

  The number of pages on the free list.

- `uintmax_t hash_bfree;`

  The number of bytes free on bucket pages.

- `u_int32_t hash_bigpages;`

  The number of big key/data pages.

- `uintmax_t hash_big_bfree;`

  The number of bytes free on big item pages.

- `u_int32_t hash_overflows;`

  The number of overflow pages (overflow pages are pages that contain items that did not fit in the main bucket page).

- `uintmax_t hash_ovfl_free;`

  The number of bytes free on overflow pages.

- `u_int32_t hash_dup;`

  The number of duplicate pages.

- `uintmax_t hash_dup_free;`

  The number of bytes free on duplicate pages.

Btree and Recno Statistics

In the case of a Btree or Recno database, the statistics are stored in a structure of type DB_BTREE_STAT. The following fields will be filled in:

- `u_int32_t bt_magic;`

  Magic number that identifies the file as a Btree database. Returned if DB_FAST_STAT is set.

- `u_int32_t bt_version;`

  The version of the Btree database. Returned if DB_FAST_STAT is set.

- u_int32_t bt_nkeys;

  For the Btree Access Method, the number of keys in the database. If the DB_FAST_STAT flag is not specified or the database was configured to support record numbers (see DB_RECNUM), the count will be exact. Otherwise, the count will be the last saved value unless it has never been calculated, in which case it will be 0.

  For the Recno Access Method, the number of records in the database. If the database was configured with mutable record numbers (see DB_RENUMBER), the count will be exact. Otherwise, if the DB_FAST_STAT flag is specified the count will be exact but will include deleted and implicitly created records; if the DB_FAST_STAT flag is not specified, the count will be exact and will not include deleted or implicitly created records.

  Returned if DB_FAST_STAT is set.

- u_int32_t bt_ndata;

  For the Btree Access Method, the number of key/data pairs in the database. If the DB_FAST_STAT flag is not specified, the count will be exact. Otherwise, the count will be the last saved value unless it has never been calculated, in which case it will be 0.

  For the Recno Access Method, the number of records in the database. If the database was configured with mutable record numbers (see DB_RENUMBER), the count will be exact. Otherwise, if the DB_FAST_STAT flag is specified the count will be exact but will include deleted and implicitly created records; if the DB_FAST_STAT flag is not specified, the count will be exact and will not include deleted or implicitly created records.

  Returned if DB_FAST_STAT is set.

- u_int32_t bt_pagecnt;

  The number of pages in the database. Returned if DB_FAST_STAT is set.

- u_int32_t bt_pagesize;

  The underlying database page size, in bytes. Returned if DB_FAST_STAT is set.

- u_int32_t bt_minkey;

  The minimum keys per page. Returned if DB_FAST_STAT is set.

- u_int32_t bt_re_len;

  The length of fixed-length records. Returned if DB_FAST_STAT is set.

- u_int32_t bt_re_pad;

  The padding byte value for fixed-length records. Returned if DB_FAST_STAT is set.

- u_int32_t bt_levels;

  Number of levels in the database.

- u_int32_t bt_int_pg;

  Number of database internal pages.

- u_int32_t bt_leaf_pg;

  Number of database leaf pages.

- u_int32_t bt_dup_pg;

  Number of database duplicate pages.

- u_int32_t bt_over_pg;

  Number of database overflow pages.

- u_int32_t bt_empty_pg;

  Number of empty database pages.

- u_int32_t bt_free;

  Number of pages on the free list.

- uintmax_t bt_int_pgfree;

  Number of bytes free in database internal pages.

- uintmax_t bt_leaf_pgfree;

  Number of bytes free in database leaf pages.

- uintmax_t bt_dup_pgfree;

  Number of bytes free in database duplicate pages.

- uintmax_t bt_over_pgfree;

  Number of bytes free in database overflow pages.

Queue Statistics

In the case of a Queue database, the statistics are stored in a structure of type DB_QUEUE_STAT. The following fields will be filled in:

- u_int32_t qs_magic;

  Magic number that identifies the file as a Queue file. Returned if DB_FAST_STAT is set.

- u_int32_t qs_version;

  The version of the Queue file type. Returned if DB_FAST_STAT is set.

- `u_int32_t qs_nkeys;`

  The number of records in the database. If DB_FAST_STAT was specified the count will be the last saved value unless it has never been calculated, in which case it will be 0. Returned if DB_FAST_STAT is set.

- `u_int32_t qs_ndata;`

  The number of records in the database. If DB_FAST_STAT was specified the count will be the last saved value unless it has never been calculated, in which case it will be 0. Returned if DB_FAST_STAT is set.

- `u_int32_t qs_pagesize;`

  Underlying database page size, in bytes. Returned if DB_FAST_STAT is set.

- `u_int32_t qs_extentsize;`

  Underlying database extent size, in pages. Returned if DB_FAST_STAT is set.

- `u_int32_t qs_pages;`

  Number of pages in the database.

- `u_int32_t qs_re_len;`

  The length of the records. Returned if DB_FAST_STAT is set.

- `u_int32_t qs_re_pad;`

  The padding byte value for the records. Returned if DB_FAST_STAT is set.

- `u_int32_t qs_pgfree;`

  Number of bytes free in database pages.

- `u_int32_t qs_first_recno;`

  First undeleted record in the database. Returned if DB_FAST_STAT is set.

- `u_int32_t qs_cur_recno;`

  Next available record number. Returned if DB_FAST_STAT is set.

Errors

The `Db::stat()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DbRepHandleDeadException or DB_REP_HANDLE_DEAD

When a client synchronizes with the master, it is possible for committed transactions to be rolled back. This invalidates all the database and cursor handles opened in the replication environment. Once this occurs, an attempt to use such a handle will throw a DbRepHandleDeadException (if your application is configured to throw exceptions), or return DB_REP_HANDLE_DEAD. The application will need to discard the handle and open a new one in order to continue processing.

DbDeadlockException or DB_REP_LOCKOUT

The operation was blocked by client/master synchronization.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_REP_LOCKOUT is returned.

EINVAL

An invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::stat_print()

```
#include <db_cxx.h>

int
Db::stat_print(u_int32_t flags);
```

The `Db::stat_print()` method displays the database statistical information, as described for the `Db::stat()` method. The information is printed to a specified output channel (see the DbEnv::set_msgfile() method for more information), or passed to an application callback function (see the DbEnv::set_msgcall() method for more information).

The `Db::stat_print()` method may not be called before the Db::open() method is called.

The `Db::stat_print()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

- `DB_FAST_STAT`

  Return only the values which do not require traversal of the database. Among other things, this flag makes it possible for applications to request key and record counts without incurring the performance penalty of traversing the entire database.

- `DB_STAT_ALL`

  Display all available information.

Class

Db

See Also

Database and Related Methods

Db::sync()

```
#include <db_cxx.h>

int
Db::sync(u_int32_t flags);
```

The `Db::sync()` method flushes any cached information to disk.

If the database is in memory only, the `Db::sync()` method has no effect and will always succeed.

It is important to understand that flushing cached information to disk only minimizes the window of opportunity for corrupted data. Although unlikely, it is possible for database corruption to happen if a system or application crash occurs while writing data to the database. To ensure that database corruption never occurs, applications must either: use transactions and logging with automatic recovery; use logging and application-specific recovery; or edit a copy of the database, and once all applications using the database have successfully called Db::close(), atomically replace the original database with the updated copy.

The `Db::sync()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter is currently unused, and must be set to 0.

Errors

The `Db::sync()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DbRepHandleDeadException or DB_REP_HANDLE_DEAD

When a client synchronizes with the master, it is possible for committed transactions to be rolled back. This invalidates all the database and cursor handles opened in the replication environment. Once this occurs, an attempt to use such a handle will throw a DbRepHandleDeadException (if your application is configured to throw exceptions), or return `DB_REP_HANDLE_DEAD`. The application will need to discard the handle and open a new one in order to continue processing.

DbDeadlockException or DB_REP_LOCKOUT

The operation was blocked by client/master synchronization.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_REP_LOCKOUT` is returned.

EINVAL

An invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::truncate()

```
#include <db_cxx.h>

int
Db::truncate(DbTxn *txnid, u_int32_t *countp, u_int32_t flags);
```

The `Db::truncate()` method empties the database, discarding all records it contains. The number of records discarded from the database is returned in **countp**.

When called on a database configured with secondary indices using the Db::associate() method, the `Db::truncate()` method truncates the primary database and all secondary indices. A count of the records discarded from the primary database is returned.

It is an error to call the `Db::truncate()` method on a database with open cursors.

The `Db::truncate()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

countp

The **countp** parameter references memory into which the number of records discarded from the database is copied.

flags

The **flags** parameter is currently unused, and must be set to 0.

txnid

If the operation is part of an application-specified transaction, the **txnid** parameter is a transaction handle returned from DbEnv::txn_begin(); if the operation is part of a Berkeley DB Concurrent Data Store group, the **txnid** parameter is a handle returned from DbEnv::cdsgroup_begin(); otherwise NULL. If no transaction handle is specified, but the operation occurs in a transactional database, the operation will be implicitly transaction protected.

Errors

The `Db::truncate()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DbDeadlockException or DB_LOCK_DEADLOCK

A transactional database environment operation was selected to resolve a deadlock.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_LOCK_DEADLOCK` is returned.

DbLockNotGrantedException or DB_LOCK_NOTGRANTED

A Berkeley DB Concurrent Data Store database environment configured for lock timeouts was unable to grant a lock in the allowed time.

DbLockNotGrantedException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_LOCK_NOTGRANTED` is returned.

EINVAL

If there are open cursors in the database; or if an invalid flag value or parameter was specified.

Class

Db

See Also

Database and Related Methods

Db::upgrade()

```
#include <db_cxx.h>

int
Db::upgrade(const char *file, u_int32_t flags);
```

The `Db::upgrade()` method upgrades all of the databases included in the file **file**, if necessary. If no upgrade is necessary, `Db::upgrade()` always returns success.

Database upgrades are done in place and are destructive. For example, if pages need to be allocated and no disk space is available, the database may be left corrupted. Backups should be made before databases are upgraded. See Upgrading databases for more information.

Unlike all other database operations, `Db::upgrade()` may only be done on a system with the same byte-order as the database.

The `Db::upgrade()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

file

The **file** parameter is the physical file containing the databases to be upgraded.

flags

The **flags** parameter must be set to 0 or the following value:

- `DB_DUPSORT`

  This flag is only meaningful when upgrading databases from releases before the Berkeley DB 3.1 release.

  As part of the upgrade from the Berkeley DB 3.0 release to the 3.1 release, the on-disk format of duplicate data items changed. To correctly upgrade the format requires applications to specify whether duplicate data items in the database are sorted or not. Specifying the DB_DUPSORT flag informs `Db::upgrade()` that the duplicates are sorted; otherwise they are assumed to be unsorted. Incorrectly specifying the value of this flag may lead to database corruption.

  Further, because the `Db::upgrade()` method upgrades a physical file (including all the databases it contains), it is not possible to use `Db::upgrade()` to upgrade files in which some of the databases it includes have sorted duplicate data items, and some of the databases it includes have unsorted duplicate data items. If the file does not have more than a single database, if the databases do not support duplicate data items, or if all of the databases that support duplicate data items support the same style of duplicates (either sorted or unsorted), `Db::upgrade()` will work correctly as long as the DB_DUPSORT flag is correctly specified. Otherwise, the file cannot be upgraded using `Db::upgrade;()` it must be upgraded manually by dumping and reloading the databases.

Environment Variables

If the database was opened within a database environment, the environment variable `DB_HOME` may be used as the path of the database environment home.

`Db::upgrade()` is affected by any database directory specified using the DbEnv::set_data_dir() method, or by setting the "set_data_dir" string in the environment's DB_CONFIG file.

Errors

The `Db::upgrade()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DB_OLD_VERSION

The database cannot be upgraded by this version of the Berkeley DB software.

Class

Db

See Also

Database and Related Methods

Db::verify()

```
#include <db_cxx.h>

int
Db::verify(const char *file,
    const char *database, ostream *outfile, u_int32_t flags);
```

The `Db::verify()` method verifies the integrity of all databases in the file specified by the **file** parameter, and optionally outputs the databases' key/data pairs to the file stream specified by the **outfile** parameter.

The `Db::verify()` method does not perform any locking, even in Berkeley DB environments that are configured with a locking subsystem. As such, it should only be used on files that are not being modified by another thread of control.

The `Db::verify()` method may not be called after the Db::open() method is called.

The Db handle may not be accessed again after `Db::verify()` is called, regardless of its return.

The `Db::verify()` method will return DB_VERIFY_BAD if a database is corrupted. When the DB_SALVAGE flag is specified, the DB_VERIFY_BAD return means that all key/data pairs in the file may not have been successfully output. Unless otherwise specified, the `Db::verify()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

database

The **database** parameter is the database in **file** on which the database checks for btree and duplicate sort order and for hashing are to be performed. See the DB_ORDERCHKONLY flag for more information.

The database parameter must be set to NULL except when the DB_ORDERCHKONLY flag is set.

file

The **file** parameter is the physical file in which the databases to be verified are found.

flags

The **flags** parameter must be set to 0 or the following value:

- DB_SALVAGE

  Write the key/data pairs from all databases in the file to the file stream named in the **outfile** parameter. Key values are written for Btree, Hash and Queue databases, but not for Recno databases.

  The output format is the same as that specified for the db_dump utility, and can be used as input for the db_load utility.

Because the key/data pairs are output in page order as opposed to the sort order used by db_dump, using `Db::verify()` to dump key/data pairs normally produces less than optimal loads for Btree databases.

In addition, the following flags may be set by bitwise inclusively **OR**'ing them into the **flags** parameter:

- DB_AGGRESSIVE

  Output **all** the key/data pairs in the file that can be found. By default, `Db::verify()` does not assume corruption. For example, if a key/data pair on a page is marked as deleted, it is not then written to the output file. When DB_AGGRESSIVE is specified, corruption is assumed, and any key/data pair that can be found is written. In this case, key/data pairs that are corrupted or have been deleted may appear in the output (even if the file being salvaged is in no way corrupt), and the output will almost certainly require editing before being loaded into a database.

- DB_PRINTABLE

  When using the DB_SALVAGE flag, if characters in either the key or data items are printing characters (as defined by **isprint**(3)), use printing characters to represent them. This flag permits users to use standard text editors and tools to modify the contents of databases or selectively remove data from salvager output.

  Note: different systems may have different notions about what characters are considered *printing characters*, and databases dumped in this manner may be less portable to external systems.

- DB_NOORDERCHK

  Skip the database checks for btree and duplicate sort order and for hashing.

  The `Db::verify()` method normally verifies that btree keys and duplicate items are correctly sorted, and hash keys are correctly hashed. If the file being verified contains multiple databases using differing sorting or hashing algorithms, some of them must necessarily fail database verification because only one sort order or hash function can be specified before `Db::verify()` is called. To verify files with multiple databases having differing sorting orders or hashing functions, first perform verification of the file as a whole by using the DB_NOORDERCHK flag, and then individually verify the sort order and hashing function for each database in the file using the DB_ORDERCHKONLY flag.

- DB_ORDERCHKONLY

  Perform the database checks for btree and duplicate sort order and for hashing, skipped by DB_NOORDERCHK.

  When this flag is specified, a **database** parameter should also be specified, indicating the database in the physical file which is to be checked. This flag is only safe to use on databases that have already successfully been verified using `Db::verify()` with the DB_NOORDERCHK flag set.

outfile

The **outfile** parameter is an optional file stream to which the databases' key/data pairs are written.

Environment Variables

If the database was opened within a database environment, the environment variable `DB_HOME` may be used as the path of the database environment home.

`Db::verify()` is affected by any database directory specified using the DbEnv::set_data_dir() method, or by setting the "set_data_dir" string in the environment's DB_CONFIG file.

Errors

The `Db::verify()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after Db::open() was called; or if an invalid flag value or parameter was specified.

ENOENT

The file or directory does not exist.

Class

Db

See Also

Database and Related Methods

Chapter 3. The Dbc Handle

A Dbc object is a handle for a cursor into a Berkeley DB database.

Dbc handles are not free-threaded. Cursor handles may be shared by multiple threads if access is serialized by the application.

You create a Dbc using the Db::cursor() method.

If the cursor is to be used to perform operations on behalf of a transaction, the cursor must be opened and closed within the context of that single transaction.

Once Dbc::close() has been called, the handle may not be accessed again, regardless of the method's return.

Database Cursors and Related Methods

Database Cursors and Related Methods	Description
Db::cursor()	Create a cursor handle
Dbc::close()	Close a cursor handle
Dbc::cmp()	Compare two cursors for equality.
Dbc::count()	Return count of duplicates for current key
Dbc::del()	Delete current key/data pair
Dbc::dup()	Duplicate the cursor handle
Dbc::get()	Retrieve by cursor
Dbc::put()	Store by cursor
Dbc::set_priority(), Dbc::get_priority()	Set/get the cursor's cache priority

Db::cursor()

```
#include <db_cxx.h>

int
Db::cursor(DbTxn *txnid, Dbc **cursorp, u_int32_t flags);
```

The `Db::cursor()` method returns a created database cursor.

Cursors may span threads, but only serially, that is, the application must serialize access to the cursor handle.

The `Db::cursor()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

cursorp

The **cursorp** parameter references memory into which a pointer to the allocated cursor is copied.

flags

The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

- DB_BULK

  Configure a cursor to optimize for bulk operations. Each successive operation on a cursor configured with this flag attempts to continue on the same database page as the previous operation, falling back to a search if a different page is required. This avoids searching if there is a high degree of locality between cursor operations. This flag is currently only effective with the btree access method: for other access methods it is ignored.

- DB_READ_COMMITTED

  Configure a transactional cursor to have degree 2 isolation. This ensures the stability of the current data item read by this cursor but permits data read by this cursor to be modified or deleted prior to the commit of the transaction for this cursor.

- DB_READ_UNCOMMITTED

  Configure a transactional cursor to have degree 1 isolation. Read operations performed by the cursor may return modified but not yet committed data. Silently ignored if the DB_READ_UNCOMMITTED flag was not specified when the underlying database was opened.

- DB_WRITECURSOR

  Specify that the cursor will be used to update the database. The underlying database environment must have been opened using the DB_INIT_CDB flag.

- DB_TXN_SNAPSHOT

    Configure a transactional cursor to operate with read-only snapshot isolation. For databases with the DB_MULTIVERSION flag set, data values will be read as they are when the cursor is opened, without taking read locks.

    This flag implicitly begins a transaction that is committed when the cursor is closed.

    This flag is silently ignored if DB_MULTIVERSION is not set on the underlying database or if a transaction is supplied in the **txnid** parameter.

txnid

To transaction-protect cursor operations, cursors must be opened and closed within the context of a transaction. The **txnid** parameter specifies the transaction context in which the cursor may be used.

Cursor operations are not automatically transaction-protected, even if the DB_AUTO_COMMIT flag is specified to the DbEnv::set_flags() or Db::open() methods. If cursor operations are to be transaction-protected, the **txnid** parameter must be a transaction handle returned from DbEnv::txn_begin(); otherwise, NULL.

Errors

The `Db::cursor()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DbRepHandleDeadException or DB_REP_HANDLE_DEAD

When a client synchronizes with the master, it is possible for committed transactions to be rolled back. This invalidates all the database and cursor handles opened in the replication environment. Once this occurs, an attempt to use such a handle will throw a DbRepHandleDeadException (if your application is configured to throw exceptions), or return `DB_REP_HANDLE_DEAD`. The application will need to discard the handle and open a new one in order to continue processing.

DbDeadlockException or DB_REP_LOCKOUT

The operation was blocked by client/master synchronization.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_REP_LOCKOUT` is returned.

EINVAL

An invalid flag value or parameter was specified.

Class

Db

See Also

Database Cursors and Related Methods

Dbc::close()

```
#include <db_cxx.h>

int
Dbc::close(void);
```

The Dbc::close() method discards the cursor.

It is possible for the Dbc::close() method to return DB_LOCK_DEADLOCK, signaling that any enclosing transaction should be aborted. If the application is already intending to abort the transaction, this error should be ignored, and the application should proceed.

After Dbc::close() has been called, regardless of its return, the cursor handle may not be used again.

The Dbc::close() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Errors

The Dbc::close() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DbDeadlockException or DB_LOCK_DEADLOCK

A transactional database environment operation was selected to resolve a deadlock.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_LOCK_DEADLOCK is returned.

DbLockNotGrantedException or DB_LOCK_NOTGRANTED

A Berkeley DB Concurrent Data Store database environment configured for lock timeouts was unable to grant a lock in the allowed time.

DbLockNotGrantedException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_LOCK_NOTGRANTED is returned.

EINVAL

If the cursor is already closed; or if an invalid flag value or parameter was specified.

Class

Dbc

See Also

Database Cursors and Related Methods

Dbc::cmp()

```
#include <db_cxx.h>

int
Dbc::cmp(Dbc *other_cursor, int *result, u_int32_t flags);
```

The `Dbc::cmp()` method compares two cursors for equality. Two cursors are equal if and only if they are positioned on the same item in the same database.

The `Dbc::cmp()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

other_cursor

The **other_cursor** parameter references another cursor handle that will be used as the comparator.

result

If the call is successful and both cursors are positioned on the same item, result is set to zero. If the call is successful but the cursors are not positioned on the same item, result is set to a non-zero value. If the call is unsuccessful, the value of result should be ignored.

flags

The **flags** parameter is currently unused, and must be set to 0.

Errors

The `Dbc::cmp()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

- If either of the cursors are already closed.

- If the cursors have been opened against different databases.

- If either of the cursors have not been positioned.

- If the other_dbc parameter is NULL.

- If the result parameter is NULL.

Class

Dbc

See Also

Database Cursors and Related Methods

Dbc::count()

```
#include <db_cxx.h>

int
Dbc::count(db_recno_t *countp, u_int32_t flags);
```

The `Dbc::count()` method returns a count of the number of data items for the key to which the cursor refers.

The `Dbc::count()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

countp

The **countp** parameter references memory into which the count of the number of duplicate data items is copied.

flags

The **flags** parameter is currently unused, and must be set to 0.

Errors

The `Dbc::count()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DbRepHandleDeadException or DB_REP_HANDLE_DEAD

When a client synchronizes with the master, it is possible for committed transactions to be rolled back. This invalidates all the database and cursor handles opened in the replication environment. Once this occurs, an attempt to use such a handle will throw a DbRepHandleDeadException (if your application is configured to throw exceptions), or return DB_REP_HANDLE_DEAD. The application will need to discard the handle and open a new one in order to continue processing.

DbDeadlockException or DB_REP_LOCKOUT

The operation was blocked by client/master synchronization.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_REP_LOCKOUT is returned.

EINVAL

If the cursor has not been initialized; or if an invalid flag value or parameter was specified.

Class

Dbc

See Also

Database Cursors and Related Methods

Dbc::del()

```
#include <db_cxx.h>

int
Dbc::del(u_int32_t flags);
```

The `Dbc::del()` method deletes the key/data pair to which the cursor refers.

When called on a cursor opened on a database that has been made into a secondary index using the Db::associate() method, the Db::del() method deletes the key/data pair from the primary database and all secondary indices.

The cursor position is unchanged after a delete, and subsequent calls to cursor functions expecting the cursor to refer to an existing key will fail.

The `Dbc::del()` method will return DB_KEYEMPTY if the element has already been deleted. The `Dbc::del()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or one of the following values:

• DB_CONSUME

  If the database is of type DB_QUEUE then this flag may be set to force the head of the queue to move to the first non-deleted item in the queue. Normally this is only done if the deleted item is exactly at the head when deleted.

Errors

The `Dbc::del()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DB_FOREIGN_CONFLICT

A foreign key constraint violation has occurred. This can be caused by one of two things:

1. An attempt was made to add a record to a constrained database, and the key used for that record does not exist in the foreign key database.

2. DB_FOREIGN_ABORT (page 10) was declared for a foreign key database, and then subsequently a record was deleted from the foreign key database without first removing it from the constrained secondary database.

DbDeadlockException or DB_LOCK_DEADLOCK

A transactional database environment operation was selected to resolve a deadlock.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_LOCK_DEADLOCK` is returned.

DbLockNotGrantedException or DB_LOCK_NOTGRANTED

A Berkeley DB Concurrent Data Store database environment configured for lock timeouts was unable to grant a lock in the allowed time.

DbLockNotGrantedException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_LOCK_NOTGRANTED` is returned.

DbRepHandleDeadException or DB_REP_HANDLE_DEAD

When a client synchronizes with the master, it is possible for committed transactions to be rolled back. This invalidates all the database and cursor handles opened in the replication environment. Once this occurs, an attempt to use such a handle will throw a DbRepHandleDeadException (if your application is configured to throw exceptions), or return `DB_REP_HANDLE_DEAD`. The application will need to discard the handle and open a new one in order to continue processing.

DbDeadlockException or DB_REP_LOCKOUT

The operation was blocked by client/master synchronization.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_REP_LOCKOUT` is returned.

DB_SECONDARY_BAD

A secondary index references a nonexistent primary key.

EACCES

An attempt was made to modify a read-only database.

EINVAL

If the cursor has not been initialized; or if an invalid flag value or parameter was specified.

EPERM

Write attempted on read-only cursor when the DB_INIT_CDB flag was specified to DbEnv::open().

Class

Dbc

See Also

Database Cursors and Related Methods

Dbc::dup()

```
#include <db_cxx.h>

int
Dbc::dup(Dbc **cursorp, u_int32_t flags);
```

The `Dbc::dup()` method creates a new cursor that uses the same transaction and locker ID as the original cursor. This is useful when an application is using locking and requires two or more cursors in the same thread of control.

The `Dbc::dup()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

cursorp

The `Dbc::dup()` method returns the newly created cursor in **cursorp**.

flags

The **flags** parameter must be set to 0 or the following flag:

• `DB_POSITION`

   The newly created cursor is initialized to refer to the same position in the database as the original cursor (if any) and hold the same locks (if any). If the DB_POSITION flag is not specified, or the original cursor does not hold a database position and locks, the created cursor is uninitialized and will behave like a cursor newly created using the Db::cursor() method.

Errors

The `Dbc::dup()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DbRepHandleDeadException or DB_REP_HANDLE_DEAD

When a client synchronizes with the master, it is possible for committed transactions to be rolled back. This invalidates all the database and cursor handles opened in the replication environment. Once this occurs, an attempt to use such a handle will throw a DbRepHandleDeadException (if your application is configured to throw exceptions), or return `DB_REP_HANDLE_DEAD`. The application will need to discard the handle and open a new one in order to continue processing.

DbDeadlockException or DB_REP_LOCKOUT

The operation was blocked by client/master synchronization.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_REP_LOCKOUT` is returned.

EINVAL

An invalid flag value or parameter was specified.

Class

Dbc

See Also

Database Cursors and Related Methods

Dbc::get()

```
#include <db_cxx.h>

int
Dbc::get(Dbt *key, Dbt *data, u_int32_t flags);

int
Dbc::pget(Dbt *key, Dbt *pkey, Dbt *data, u_int32_t flags);
```

The `Dbc::get()` method retrieves key/data pairs from the database. The address and length of the key are returned in the object to which **key** refers (except for the case of the DB_SET flag, in which the **key** object is unchanged), and the address and length of the data are returned in the object to which **data** refers.

When called on a cursor opened on a database that has been made into a secondary index using the Db::associate() method, the `Dbc::get()` and `Dbc::pget()` methods return the key from the secondary index and the data item from the primary database. In addition, the `Dbc::pget()` method returns the key from the primary database. In databases that are not secondary indices, the `Dbc::pget()` method will always fail.

Modifications to the database during a sequential scan will be reflected in the scan; that is, records inserted behind a cursor will not be returned while records inserted in front of a cursor will be returned.

In Queue and Recno databases, missing entries (that is, entries that were never explicitly created or that were created and then deleted) will be skipped during a sequential scan.

Unless otherwise specified, the `Dbc::get()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

If `Dbc::get()` fails for any reason, the state of the cursor will be unchanged.

Parameters

data

The data Dbt operated on.

flags

The **flags** parameter must be set to one of the following values:

- DB_CURRENT

  Return the key/data pair to which the cursor refers.

  The `Dbc::get()` method will return DB_KEYEMPTY if DB_CURRENT is set and the cursor key/data pair was deleted.

- DB_FIRST

The cursor is set to refer to the first key/data pair of the database, and that pair is returned. If the first key has duplicate values, the first data item in the set of duplicates is returned.

If the database is a Queue or Recno database, `Dbc::get()` using the DB_FIRST flag will ignore any keys that exist but were never explicitly created by the application, or were created and later deleted.

The `Dbc::get()` method will return DB_NOTFOUND if DB_FIRST is set and the database is empty.

- DB_GET_BOTH

Move the cursor to the specified key/data pair of the database. The cursor is positioned to a key/data pair if both the key and data match the values provided on the key and data parameters.

In all other ways, this flag is identical to the DB_SET flag.

When used with `Dbc::pget()` on a secondary index handle, both the secondary and primary keys must be matched by the secondary and primary key item in the database. It is an error to use the DB_GET_BOTH flag with the `Dbc::get()` version of this method and a cursor that has been opened on a secondary index handle.

- DB_GET_BOTH_RANGE

Move the cursor to the specified key/data pair of the database. The key parameter must be an exact match with a key in the database. The data item retrieved is the item in a duplicate set that is the smallest value which is greater than or equal to the value provided by the data parameter (as determined by the comparison function). If this flag is specified on a database configured without sorted duplicate support, the behavior is identical to the DB_GET_BOTH flag. Returns the datum associated with the given key/data pair.

In all other ways, this flag is identical to the DB_GET_BOTH flag.

- DB_GET_RECNO

Return the record number associated with the cursor. The record number will be returned in **data**, as described in Dbt. The **key** parameter is ignored.

For DB_GET_RECNO to be specified, the underlying database must be of type Btree, and it must have been created with the DB_RECNUM flag.

When called on a cursor opened on a database that has been made into a secondary index, the `Dbc::get()` and `Dbc::pget()` methods return the record number of the primary database in **data**. In addition, the `Dbc::pget()` method returns the record number of the secondary index in **pkey**. If either underlying database is not of type Btree or is not created with the DB_RECNUM flag, the out-of-band record number of 0 is returned.

- DB_JOIN_ITEM

Do not use the data value found in all of the cursors as a lookup key for the primary database, but simply return it in the key parameter instead. The data parameter is left unchanged.

For DB_JOIN_ITEM to be specified, the underlying cursor must have been returned from the Db::join() method.

- DB_LAST

  The cursor is set to refer to the last key/data pair of the database, and that pair is returned. If the last key has duplicate values, the last data item in the set of duplicates is returned.

  If the database is a Queue or Recno database, Dbc::get() using the DB_LAST flag will ignore any keys that exist but were never explicitly created by the application, or were created and later deleted.

  The Dbc::get() method will return DB_NOTFOUND if DB_LAST is set and the database is empty.

- DB_NEXT

  If the cursor is not yet initialized, DB_NEXT is identical to DB_FIRST. Otherwise, the cursor is moved to the next key/data pair of the database, and that pair is returned. In the presence of duplicate key values, the value of the key may not change.

  If the database is a Queue or Recno database, Dbc::get() using the DB_NEXT flag will skip any keys that exist but were never explicitly created by the application, or those that were created and later deleted.

  The Dbc::get() method will return DB_NOTFOUND if DB_NEXT is set and the cursor is already on the last record in the database.

- DB_NEXT_DUP

  If the next key/data pair of the database is a duplicate data record for the current key/data pair, the cursor is moved to the next key/data pair of the database, and that pair is returned.

  The Dbc::get() method will return DB_NOTFOUND if DB_NEXT_DUP is set and the next key/data pair of the database is not a duplicate data record for the current key/data pair.

- DB_NEXT_NODUP

  If the cursor is not yet initialized, DB_NEXT_NODUP is identical to DB_FIRST. Otherwise, the cursor is moved to the next non-duplicate key of the database, and that key/data pair is returned.

  If the database is a Queue or Recno database, Dbc::get() using the DB_NEXT_NODUP flag will ignore any keys that exist but were never explicitly created by the application, or those that were created and later deleted.

  The Dbc::get() method will return DB_NOTFOUND if DB_NEXT_NODUP is set and no non-duplicate key/data pairs exist after the cursor position in the database.

- DB_PREV

If the cursor is not yet initialized, DB_PREV is identical to DB_LAST. Otherwise, the cursor is moved to the previous key/data pair of the database, and that pair is returned. In the presence of duplicate key values, the value of the key may not change.

If the database is a Queue or Recno database, `Dbc::get()` using the DB_PREV flag will skip any keys that exist but were never explicitly created by the application, or those that were created and later deleted.

The `Dbc::get()` method will return DB_NOTFOUND if DB_PREV is set and the cursor is already on the first record in the database.

- DB_PREV_DUP

  If the previous key/data pair of the database is a duplicate data record for the current key/data pair, the cursor is moved to the previous key/data pair of the database, and that pair is returned.

  The `Dbc::get()` method will return DB_NOTFOUND if DB_PREV_DUP is set and the previous key/data pair of the database is not a duplicate data record for the current key/data pair.

- DB_PREV_NODUP

  If the cursor is not yet initialized, DB_PREV_NODUP is identical to DB_LAST. Otherwise, the cursor is moved to the previous non-duplicate key of the database, and that key/data pair is returned.

  If the database is a Queue or Recno database, `Dbc::get()` using the DB_PREV_NODUP flag will ignore any keys that exist but were never explicitly created by the application, or those that were created and later deleted.

  The `Dbc::get()` method will return DB_NOTFOUND if DB_PREV_NODUP is set and no non-duplicate key/data pairs exist before the cursor position in the database.

- DB_SET

  Move the cursor to the specified key/data pair of the database, and return the datum associated with the given key.

  The `Dbc::get()` method will return DB_NOTFOUND if DB_SET is set and no matching keys are found. The `Dbc::get()` method will return DB_KEYEMPTY if DB_SET is set and the database is a Queue or Recno database, and the specified key exists, but was never explicitly created by the application or was later deleted. In the presence of duplicate key values, `Dbc::get()` will return the first data item for the given key.

- DB_SET_RANGE

  Move the cursor to the specified key/data pair of the database. In the case of the Btree access method, the key is returned as well as the data item and the returned key/data pair is the smallest key greater than or equal to the specified key (as determined by the Btree comparison function), permitting partial key matches and range searches.

  In all other ways the behavior of this flag is the same as the DB_SET flag.

- DB_SET_RECNO

  Move the cursor to the specific numbered record of the database, and return the associated key/data pair. The **data** field of the specified **key** must be a pointer to a memory location from which a db_recno_t may be read, as described in Dbt. This memory location will be read to determine the record to be retrieved.

  For DB_SET_RECNO to be specified, the underlying database must be of type Btree, and it must have been created with the DB_RECNUM flag.

In addition, the following flags may be set by bitwise inclusively **OR**'ing them into the **flags** parameter:

- DB_IGNORE_LEASE

  This flag is relevant only when using a replicated environment.

  Return the data item irrespective of the state of master leases. The item will be returned under all conditions: if master leases are not configured, if the request is made to a client, if the request is made to a master with a valid lease, or if the request is made to a master without a valid lease.

- DB_READ_COMMITTED

  Configure a transactional get operation to have degree 2 isolation (the read is not repeatable).

- DB_READ_UNCOMMITTED

  Database items read during a transactional call will have degree 1 isolation, including modified but not yet committed data. Silently ignored if the DB_READ_UNCOMMITTED flag was not specified when the underlying database was opened.

- DB_MULTIPLE

  Return multiple data items in the **data** parameter.

  In the case of Btree or Hash databases, duplicate data items for the current key, starting at the current cursor position, are entered into the buffer. Subsequent calls with both the DB_NEXT_DUP and DB_MULTIPLE flags specified will return additional duplicate data items associated with the current key or DB_NOTFOUND if there are no additional duplicate data items to return. Subsequent calls with both the DB_NEXT and DB_MULTIPLE flags specified will return additional duplicate data items associated with the current key or if there are no additional duplicate data items will return the next key and its data items or DB_NOTFOUND if there are no additional keys in the database.

  In the case of Queue or Recno databases, data items starting at the current cursor position are entered into the buffer. The record number of the first record will be returned in the **key** parameter. The record number of each subsequent returned record must be calculated from this value. Subsequent calls with the DB_MULTIPLE flag specified will return additional data items or DB_NOTFOUND if there are no additional data items to return.

  The buffer to which the **data** parameter refers must be provided from user memory (see DB_DBT_USERMEM). The buffer must be at least as large as the page size of the underlying database, aligned for unsigned integer access, and be a multiple of 1024 bytes in size. If the buffer size is

insufficient, then upon return from the call the size field of the **data** parameter will have been set to an estimated buffer size, and the error DB_BUFFER_SMALL is returned. (The size is an estimate as the exact size needed may not be known until all entries are read. It is best to initially provide a relatively large buffer, but applications should be prepared to resize the buffer as necessary and repeatedly call the method.)

The multiple data items can be iterated over using the DbMultipleDataIterator class.

The DB_MULTIPLE flag may only be used with the DB_CURRENT, DB_FIRST, DB_GET_BOTH, DB_GET_BOTH_RANGE, DB_NEXT, DB_NEXT_DUP, DB_NEXT_NODUP, DB_SET, DB_SET_RANGE, and DB_SET_RECNO options. The DB_MULTIPLE flag may not be used when accessing databases made into secondary indices using the Db::associate() method.

- DB_MULTIPLE_KEY

  Return multiple key and data pairs in the **data** parameter.

  Key and data pairs, starting at the current cursor position, are entered into the buffer. Subsequent calls with both the DB_NEXT and DB_MULTIPLE_KEY flags specified will return additional key and data pairs or DB_NOTFOUND if there are no additional key and data items to return.

  In the case of Btree or Hash databases, the multiple key and data pairs can be iterated over using the DbMultipleKeyDataIterator class.

  In the case of Queue or Recno databases, the multiple record number and data pairs can be iterated over using the DbMultipleRecnoDataIterator class.

  The buffer to which the **data** parameter refers must be provided from user memory (see DB_DBT_USERMEM). The buffer must be at least as large as the page size of the underlying database, aligned for unsigned integer access, and be a multiple of 1024 bytes in size. If the buffer size is insufficient, then upon return from the call the size field of the **data** parameter will have been set to an estimated buffer size, and the error DB_BUFFER_SMALL is returned. (The size is an estimate as the exact size needed may not be known until all entries are read. It is best to initially provide a relatively large buffer, but applications should be prepared to resize the buffer as necessary and repeatedly call the method.)

  The DB_MULTIPLE_KEY flag may only be used with the DB_CURRENT, DB_FIRST, DB_GET_BOTH, DB_GET_BOTH_RANGE, DB_NEXT, DB_NEXT_DUP, DB_NEXT_NODUP, DB_SET, DB_SET_RANGE, and DB_SET_RECNO options. The DB_MULTIPLE_KEY flag may not be used when accessing databases made into secondary indices using the Db::associate() method.

- DB_RMW

  Acquire write locks instead of read locks when doing the read, if locking is configured. Setting this flag can eliminate deadlock during a read-modify-write cycle by acquiring the write lock during the read part of the cycle so that another thread of control acquiring a read lock for the same item, in its own read-modify-write cycle, will not result in deadlock.

key

The key Dbt operated on.

pkey

The secondary index key Dbt operated on.

Errors

The `Dbc::get()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DbMemoryException or DB_BUFFER_SMALL

The requested item could not be returned due to undersized buffer.

DbMemoryException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_BUFFER_SMALL` is returned.

DbDeadlockException or DB_LOCK_DEADLOCK

A transactional database environment operation was selected to resolve a deadlock.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_LOCK_DEADLOCK` is returned.

DbLockNotGrantedException or DB_LOCK_NOTGRANTED

A Berkeley DB Concurrent Data Store database environment configured for lock timeouts was unable to grant a lock in the allowed time.

DbLockNotGrantedException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_LOCK_NOTGRANTED` is returned.

DbRepHandleDeadException or DB_REP_HANDLE_DEAD

When a client synchronizes with the master, it is possible for committed transactions to be rolled back. This invalidates all the database and cursor handles opened in the replication environment. Once this occurs, an attempt to use such a handle will throw a DbRepHandleDeadException (if your application is configured to throw exceptions), or return `DB_REP_HANDLE_DEAD`. The application will need to discard the handle and open a new one in order to continue processing.

DB_REP_LEASE_EXPIRED

The operation failed because the site's replication master lease has expired.

DbDeadlockException or DB_REP_LOCKOUT

The operation was blocked by client/master synchronization.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_REP_LOCKOUT` is returned.

DB_SECONDARY_BAD

A secondary index references a nonexistent primary key.

EINVAL

If the DB_CURRENT, DB_NEXT_DUP or DB_PREV_DUP flags were specified and the cursor has not been initialized; the `Dbc::pget()` method was called with a cursor that does not refer to a secondary index; or if an invalid flag value or parameter was specified.

Class

Dbc

See Also

Database Cursors and Related Methods

Dbc::get_priority()

```
#include <db_cxx.h>

int
Dbc::get_priority(DB_CACHE_PRIORITY *priorityp);
```

The `Dbc::get_priority()` method returns the cache priority for pages referenced by the Dbc handle.

The `Dbc::get_priority()` method may be called at any time during the life of the application.

The `Dbc::get_priority()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

priorityp

The `Dbc::get_priority()` method returns a reference to the cache priority for pages referenced by the Dbc handle in **priorityp**.

Class

Dbc

See Also

Database Cursors and Related Methods

Dbc::put()

```
#include <db_cxx.h>

int
Dbc::put(Dbt *key, Dbt *data, u_int32_t flags);
```

The `Dbc::put()` method stores key/data pairs into the database.

Unless otherwise specified, the `Dbc::put()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

If `Dbc::put()` fails for any reason, the state of the cursor will be unchanged. If `Dbc::put()` succeeds and an item is inserted into the database, the cursor is always positioned to refer to the newly inserted item.

Parameters

data

The data Dbt operated on.

flags

The **flags** parameter must be set to one of the following values:

- DB_AFTER

  In the case of the Btree and Hash access methods, insert the data element as a duplicate element of the key to which the cursor refers. The new element appears immediately after the current cursor position. It is an error to specify DB_AFTER if the underlying Btree or Hash database is not configured for unsorted duplicate data items. The **key** parameter is ignored.

  In the case of the Recno access method, it is an error to specify DB_AFTER if the underlying Recno database was not created with the DB_RENUMBER flag. If the DB_RENUMBER flag was specified, a new key is created, all records after the inserted item are automatically renumbered, and the key of the new record is returned in the structure to which the **key** parameter refers. The initial value of the **key** parameter is ignored. See Db::open() for more information.

  The DB_AFTER flag may not be specified to the Queue access method.

  The `Dbc::put()` method will return DB_NOTFOUND if the current cursor record has already been deleted and the underlying access method is Hash.

- DB_BEFORE

  In the case of the Btree and Hash access methods, insert the data element as a duplicate element of the key to which the cursor refers. The new element appears immediately before the current cursor position. It is an error to specify DB_AFTER if the underlying Btree or Hash database is not configured for unsorted duplicate data items. The **key** parameter is ignored.

In the case of the Recno access method, it is an error to specify DB_BEFORE if the underlying Recno database was not created with the DB_RENUMBER flag. If the DB_RENUMBER flag was specified, a new key is created, the current record and all records after it are automatically renumbered, and the key of the new record is returned in the structure to which the **key** parameter refers. The initial value of the **key** parameter is ignored. See Db::open() for more information.

The DB_BEFORE flag may not be specified to the Queue access method.

The Dbc::put() method will return DB_NOTFOUND if the current cursor record has already been deleted and the underlying access method is Hash.

- DB_CURRENT

    Overwrite the data of the key/data pair to which the cursor refers with the specified data item. The **key** parameter is ignored.

    The Dbc::put() method will return DB_NOTFOUND if the current cursor record has already been deleted.

- DB_KEYFIRST

    Insert the specified key/data pair into the database.

    If the underlying database supports duplicate data items, and if the key already exists in the database and a duplicate sort function has been specified, the inserted data item is added in its sorted location. If the key already exists in the database and no duplicate sort function has been specified, the inserted data item is added as the first of the data items for that key.

- DB_KEYLAST

    Insert the specified key/data pair into the database.

    If the underlying database supports duplicate data items, and if the key already exists in the database and a duplicate sort function has been specified, the inserted data item is added in its sorted location. If the key already exists in the database, and no duplicate sort function has been specified, the inserted data item is added as the last of the data items for that key.

- DB_NODUPDATA

    In the case of the Btree and Hash access methods, insert the specified key/data pair into the database, unless a key/data pair comparing equally to it already exists in the database. If a matching key/data pair already exists in the database, DB_KEYEXIST is returned. The DB_NODUPDATA flag may only be specified if the underlying database has been configured to support sorted duplicate data items.

    The DB_NODUPDATA flag may not be specified to the Queue or Recno access methods.

key

The key Dbt operated on.

Errors

The `Dbc::put()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DB_KEYEXIST

An attempt was made to insert a duplicate key into a database not configured for duplicate data.

DB_FOREIGN_CONFLICT

A foreign key constraint violation has occurred. This can be caused by one of two things:

1. An attempt was made to add a record to a constrained database, and the key used for that record does not exist in the foreign key database.

2. DB_FOREIGN_ABORT (page 10) was declared for a foreign key database, and then subsequently a record was deleted from the foreign key database without first removing it from the constrained secondary database.

DbDeadlockException or DB_LOCK_DEADLOCK

A transactional database environment operation was selected to resolve a deadlock.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_LOCK_DEADLOCK` is returned.

DbLockNotGrantedException or DB_LOCK_NOTGRANTED

A Berkeley DB Concurrent Data Store database environment configured for lock timeouts was unable to grant a lock in the allowed time.

DbLockNotGrantedException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_LOCK_NOTGRANTED` is returned.

DbRepHandleDeadException or DB_REP_HANDLE_DEAD

When a client synchronizes with the master, it is possible for committed transactions to be rolled back. This invalidates all the database and cursor handles opened in the replication environment. Once this occurs, an attempt to use such a handle will throw a DbRepHandleDeadException (if your application is configured to throw exceptions), or return `DB_REP_HANDLE_DEAD`. The application will need to discard the handle and open a new one in order to continue processing.

DbDeadlockException or DB_REP_LOCKOUT

The operation was blocked by client/master synchronization.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_REP_LOCKOUT` is returned.

EACCES

An attempt was made to modify a read-only database.

EINVAL

If the DB_AFTER, DB_BEFORE or DB_CURRENT flags were specified and the cursor has not been initialized; the DB_AFTER or DB_BEFORE flags were specified and a duplicate sort function has been specified; the DB_CURRENT flag was specified, a duplicate sort function has been specified, and the data item of the referenced key/data pair does not compare equally to the **data** parameter; the DB_AFTER or DB_BEFORE flags were specified, and the underlying access method is Queue; an attempt was made to add a record to a fixed-length database that was too large to fit; an attempt was made to add a record to a secondary index; or if an invalid flag value or parameter was specified.

EPERM

Write attempted on read-only cursor when the DB_INIT_CDB flag was specified to DbEnv::open().

Class

Dbc

See Also

Database Cursors and Related Methods

Dbc::set_priority()

```
#include <db_cxx.h>

int
Dbc::set_priority(DB_CACHE_PRIORITY priority);
```

Set the cache priority for pages referenced by the Dbc handle.

The priority of a page biases the replacement algorithm to be more or less likely to discard a page when space is needed in the buffer pool. The bias is temporary, and pages will eventually be discarded if they are not referenced again. The Dbc::set_priority() method is only advisory, and does not guarantee pages will be treated in a specific way.

The Dbc::set_priority() method may be called at any time during the life of the application.

The Dbc::set_priority() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

priority

The **priority** parameter must be set to one of the following values:

- DB_PRIORITY_VERY_LOW

  The lowest priority: pages are the most likely to be discarded.

- DB_PRIORITY_LOW

  The next lowest priority.

- DB_PRIORITY_DEFAULT

  The default priority.

- DB_PRIORITY_HIGH

  The next highest priority.

- DB_PRIORITY_VERY_HIGH

  The highest priority: pages are the least likely to be discarded.

Class

Dbc

See Also

Database Cursors and Related Methods

Chapter 4. The Dbt Handle

```
#include <db_cxx.h>

class Dbt {
public:
        Dbt(void *data, size_t size);
        Dbt();
        Dbt(const Dbt &);
        Dbt &operator = (const Dbt &);
        ~Dbt();

        void *get_data() const;
        void set_data(void *);

        u_int32_t get_size() const;
        void set_size(u_int32_t);

        u_int32_t get_ulen() const;
        void set_ulen(u_int32_t);

        u_int32_t get_dlen() const;
        void set_dlen(u_int32_t);

        u_int32_t get_doff() const;
        void set_doff(u_int32_t);

        u_int32_t get_flags() const;
        void set_flags(u_int32_t);

        DBT *Dbt::get_DBT();
        const DBT *Dbt::get_const_DBT() const;
        static Dbt *Dbt::get_Dbt(DBT *dbt);
        static const Dbt *Dbt::get_const_Dbt(const DBT *dbt);
};
```

The Dbt class is used to encode key and data items in a Berkeley DB database.

Storage and retrieval for the Db access methods are based on key/data pairs. Both key and data items are represented by Dbt objects. Key and data byte strings may refer to strings of zero length up to strings of essentially unlimited length. See Database limits for more information.

In the case when the **flags** structure element is set to 0, when the application is providing Berkeley DB a key or data item to store into the database, Berkeley DB expects the **data** object to point to a byte string of **size** bytes. When returning a key/data item to the application, Berkeley DB will store into the **data** object a pointer to a byte string of **size** bytes, and the memory to which the pointer refers will be allocated and managed by Berkeley DB. Note that using the default flags for returned Dbts is only compatible with single threaded usage of Berkeley DB.

Access to `Dbt` objects is not re-entrant. In particular, if multiple threads simultaneously access the same Dbt object using Db API calls, the results are undefined, and may result in a crash. One easy way to avoid problems is to use `Dbt` objects that are constructed as stack variables.

Each `Dbt` object has an associated `DBT` struct, which is used by the underlying implementation of Berkeley DB and its C-language API. The `Dbt::get_DBT()` method returns a pointer to this struct. Given a const `Dbt` object, `Dbt::get_const_DBT()` returns a const pointer to the same struct.

Given a `DBT` struct, the `Dbt::get_Dbt()` method returns the corresponding `Dbt` object, if there is one. If the `DBT` object was not associated with a `Dbt` (that is, it was not returned from a call to `Dbt::get_DBT()`), then the result of `Dbt::get_Dbt()` is undefined. Given a const `DBT` struct, `Dbt::get_const_Dbt()` returns the associated const `Dbt` object, if there is one.

These methods may be useful for Berkeley DB applications including both C and C++ language software. It should not be necessary to use these calls in a purely C++ application.

- `Dbt::set_data(void *data)`

  Set the data array.

  The `data` parameter is an array of bytes to be used to set the content for the `Dbt`.

- `Dbt::get_data()`

  Return the data array.

- `Dbt::set_size(u_int32_t size)`

  Sets the byte size of the data array, in bytes.

- `Dbt::get_size()`

  Return the data array size.

- `Dbt::set_ulen(u_int32_t value)`

  Set the byte size of the user-specified buffer.

  Note that applications can determine the length of a record by setting the `ulen` field to 0 and checking the return value in the **size** field. See the `DB_DBT_USERMEM` flag for more information.

- `Dbt::get_ulen()`

  Return the length in bytes of the user-specified buffer.

  Note that applications can determine the length of a record by setting the `ulen` field to 0 and checking the return value in the **size** field. See the `DB_DBT_USERMEM` flag for more information.

- `Dbt::set_dlen(u_int32_t dlen)`

Set the length of the partial record being read or written by the application, in bytes. See the `DB_DBT_PARTIAL` flag for more information.

- `Dbt::get_dlen()`

  Return the length of the partial record, in bytes.

- `Dbt::set_doff(u_int32_t value)`

  Sets the offset of the partial record being read or written by the application, in bytes. See the `DB_DBT_PARTIAL` flag for more information.

- `Dbt::get_doff()`

  Return the offset of the partial record, in bytes.

- `Dbt::set_flags(u_int32_t flags)`

  Set the object flag value.

  The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

  - `DB_DBT_MALLOC`

    When this flag is set, Berkeley DB will allocate memory for the returned key or data item (using **malloc**(3), or the user-specified malloc function), and return a pointer to it in the **data** field of the key or data `DBT` structure. Because any allocated memory becomes the responsibility of the calling application, the caller must determine whether memory was allocated using the returned value of the **data** field.

    It is an error to specify more than one of `DB_DBT_MALLOC`, `DB_DBT_REALLOC`, and `DB_DBT_USERMEM`.

  - `DB_DBT_REALLOC`

    When this flag is set Berkeley DB will allocate memory for the returned key or data item (using **realloc**(3), or the user-specified realloc function), and return a pointer to it in the **data** field of the key or data DBT structure. Because any allocated memory becomes the responsibility of the calling application, the caller must determine whether memory was allocated using the returned value of the **data** field.

    It is an error to specify more than one of `DB_DBT_MALLOC`, `DB_DBT_REALLOC`, and `DB_DBT_USERMEM`.

  - `DB_DBT_USERMEM`

    The *data* field of the key or data structure must refer to memory that is at least *ulen* bytes in length. If the length of the requested item is less than or equal to that number of bytes, the item is copied into the memory to which the *data* field refers. Otherwise, the *size* field is set to the length needed for the requested item, and the error `DB_BUFFER_SMALL` is returned.

    It is an error to specify more than one of `DB_DBT_MALLOC`, `DB_DBT_REALLOC`, and `DB_DBT_USERMEM`.

If `DB_DBT_MALLOC` or `DB_DBT_REALLOC` is specified, Berkeley DB allocates a properly sized byte array to contain the data. This can be convenient if you know little about the nature of the data, specifically the size of data in the database. However, if your application makes repeated calls to retrieve keys or data, you may notice increased garbage collection due to this allocation. If you know the maximum size of data you are retrieving, you might decrease the memory burden and speed your application by allocating your own byte array and using `DB_DBT_USERMEM`. Even if you don't know the maximum size, you can use this option and reallocate your array whenever your retrieval API call returns an `DB_BUFFER_SMALL` error or throws an exception encapsulating an `DB_BUFFER_SMALL`.

- `DB_DBT_PARTIAL`

  Do partial retrieval or storage of an item. If the calling application is doing a get, the **dlen** bytes starting **doff** bytes from the beginning of the retrieved data record are returned as if they comprised the entire record. If any or all of the specified bytes do not exist in the record, the get is successful, and any existing bytes are returned.

  For example, if the data portion of a retrieved record was 100 bytes, and a partial retrieval was done using a DBT having a **dlen** field of 20 and a **doff** field of 85, the get call would succeed, the **data** field would refer to the last 15 bytes of the record, and the **size** field would be set to 15.

  If the calling application is doing a put, the **dlen** bytes starting **doff** bytes from the beginning of the specified key's data record are replaced by the data specified by the **data** and **size** structure elements. If **dlen** is smaller than **size** the record will grow; if **dlen** is larger than **size** the record will shrink. If the specified bytes do not exist, the record will be extended using nul bytes as necessary, and the put call will succeed.

  It is an error to attempt a partial put using the Db::put() method in a database that supports duplicate records. Partial puts in databases supporting duplicate records must be done using a Dbc::put() method.

  It is an error to attempt a partial put with differing **dlen** and **size** values in Queue or Recno databases with fixed-length records.

  For example, if the data portion of a retrieved record was 100 bytes, and a partial put was done using a DBT having a **dlen** field of 20, a **doff** field of 85, and a **size** field of 30, the resulting record would be 115 bytes in length, where the last 30 bytes would be those specified by the put call.

- `DB_DBT_APPMALLOC`

  After an application-supplied callback routine passed to either Db::associate() or Db::set_append_recno() is executed, the **data** field of a DBT may refer to memory allocated with **malloc**(3) or **realloc**(3). In that case, the callback sets the `DB_DBT_APPMALLOC` flag in the DBT so that Berkeley DB will call **free**(3) to deallocate the memory when it is no longer required.

- `DB_DBT_MULTIPLE`

  Set in a secondary key creation callback routine passed to Db::associate() to indicate that multiple secondary keys should be associated with the given primary key/data pair. If set, the **size** field indicates the number of secondary keys and the **data** field refers to an array of that number of DBT structures.

The `DB_DBT_APPMALLOC` flag may be set on any of the DBT structures to indicate that their **data** field needs to be freed.

DBT and Bulk Operations

DBT and Bulk Operations	Description
DbMultipleIterator	Base class for bulk get retrieval
DbMultipleDataIterator	Bulk retrieval iterator for data items
DbMultipleKeyDataIterator	Bulk retrieval iterator for key/data pairs
DbMultipleRecnoDataIterator	Bulk retrieval iterator for record number / data item pairs
DbMultipleBuilder	Base class for bulk buffer building
DbMultipleDataBuilder	Bulk buffer builder for data items
DbMultipleKeyDataBuilder	Bulk buffer builder for key/data pairs
DbMultipleRecnoDataBuilder	Bulk buffer builder for record number / data pairs

DbMultipleIterator

```
#include <db_cxx.h>

class DbMultipleIterator
{ };
```

The `DbMultipleIterator` class is a shared package-private base class for the three types of bulk-return Iterator; it should never be instantiated directly, but it handles the functionality shared by its subclasses.

Class

DbMultipleIterator

See Also

DBT and Bulk Operations

DbMultipleDataIterator

```
#include <db_cxx.h>

class DbMultipleDataIterator
{
public:
    DbMultipleDataIterator(const Dbt &dbt);

    bool next(Dbt &data);
};
```

If either of the DB_MULTIPLE or DB_MULTIPLE_KEY flags were specified to the Db::get() or Dbc::get() methods, the data Dbt returned by those interfaces will refer to a buffer that is filled with data. Access to that data is through the classes.

The DbMultipleDataIterator class is used to iterate through data returned using the DB_MULTIPLE flag from a database belonging to any access method.

The constructor takes the The Dbt Handle returned by the call to Db::get() or Dbc::get() that used the DB_MULTIPLE flag.

☞   All instances of the bulk retrieval classes may be used only once, and to traverse the bulk retrieval buffer in the forward direction only. However, they are nondestructive, so multiple iterators can be instantiated and used on the same returned data Dbt.

Parameters are:

• dbt

    The **dbt** parameter is a data The Dbt Handle returned by the call to Db::get() or Dbc::get() that used the DB_MULTIPLE flag.

DbMultipleDataIterator.next()

The DbMultipleDataIterator.next() method returns the next data item in the original bulk retrieval buffer.

The DbMultipleDataIterator.next() method method returns false if no more data are available, and true otherwise.

Parameters are:

• data

    The **data** parameter is a The Dbt Handle that will be filled in with a reference to a buffer, a size, and an offset that together yield the next data item in the original bulk retrieval buffer.

Class

DbMultipleIterator

See Also

DBT and Bulk Operations

DbMultipleKeyDataIterator

```
#include <db_cxx.h>

class DbMultipleKeyDataIterator
{
public:
    DbMultipleKeyDataIterator(const Dbt &dbt);

    bool next(Dbt &key, Dbt &data);
};
```

If either of the DB_MULTIPLE or DB_MULTIPLE_KEY flags were specified to the Db::get() or Dbc::get() methods, the data Dbt returned by those interfaces will refer to a buffer that is filled with data. Access to that data is through these classes.

The DbMultipleDataIterator class is used to iterate through data returned using the DB_MULTIPLE_KEY flag from a database belonging to Btree or Hash access methods.

The constructor takes the The Dbt Handle returned by the call to Db::get() or Dbc::get() that used the DB_MULTIPLE_KEY flag.

☞ All instances of the bulk retrieval classes may be used only once, and to traverse the bulk retrieval buffer in the forward direction only. However, they are nondestructive, so multiple iterators can be instantiated and used on the same returned data Dbt.

Parameters are:

• dbt

    The **dbt** parameter is a data The Dbt Handle returned by the call to Db::get() or Dbc::get() that used the DB_MULTIPLE_KEY flag.

DbMultipleKeyDataIterator.next()

The DbMultipleKeyDataIterator.next() method returns the next data item in the original bulk retrieval buffer.

The DbMultipleKeyDataIterator.next() method method returns false if no more data are available, and true otherwise.

Parameters are:

• key

    The **key** parameter is a The Dbt Handle that will be filled in with a reference to a buffer, a size, and an offset that together yield the next data item in the original bulk retrieval buffer.

• data

    The **data** parameter is a The Dbt Handle that will be filled in with a reference to a buffer, a size, and an offset that together yield the next data item in the original bulk retrieval buffer.

Class

DbMultipleIterator

See Also

DBT and Bulk Operations

DbMultipleRecnoDataIterator

```
#include <db_cxx.h>

class DbMultipleRecnoDataIterator
{
public:
    DbMultipleRecnoDataIterator(const Dbt &dbt);

    bool next(db_recno_t &key, Dbt &data);
};
```

If either of the DB_MULTIPLE or DB_MULTIPLE_KEY flags were specified to the Db::get() or Dbc::get() methods, the data Dbt returned by those interfaces will refer to a buffer that is filled with data. Access to that data is through these classes.

The DbMultipleDataIterator class is used to iterate through data returned using the DB_MULTIPLE_KEY flag from a database belonging to Queue or Recno access methods.

The constructor takes the The Dbt Handle returned by the call to Db::get() or Dbc::get() that used the DB_MULTIPLE_KEY flag.

☞   All instances of the bulk retrieval classes may be used only once, and to traverse the bulk retrieval buffer in the forward direction only. However, they are nondestructive, so multiple iterators can be instantiated and used on the same returned data Dbt.

Parameters are:

• dbt

   The **dbt** parameter is a data The Dbt Handle returned by the call to Db::get() or Dbc::get() that used the DB_MULTIPLE_KEY flag.

DbMultipleRecnoDataIterator.next()

The DbMultipleRecnoDataIterator.next() method returns the next data item in the original bulk retrieval buffer.

The DbMultipleRecnoDataIterator.next() method method returns false if no more data are available, and true otherwise.

Parameters are:

• key

   The **key** parameter is a The Dbt Handle that will be filled in with a reference to a buffer, a size, and an offset that together yield the next data item in the original bulk retrieval buffer.

• data

   The **data** parameter is a The Dbt Handle that will be filled in with a reference to a buffer, a size, and an offset that together yield the next data item in the original bulk retrieval buffer.

Class

DbMultipleIterator

See Also

DBT and Bulk Operations

DbMultipleBuilder

```
#include <db_cxx.h>

class DbMultipleBuilder
{ };
```

The `DbMultipleBuilder` class is a shared package-private base class for the three types of bulk buffer builders; it should never be instantiated directly, but it handles the functionality shared by its subclasses.

Class

Dbt

See Also

DBT and Bulk Operations

DbMultipleDataBuilder

```
#include <db_cxx.h>

class DbMultipleDataBuilder
{
public:
    DbMultipleDataBuilder(Dbt &dbt);

    bool append(void *dbuf, size_t dlen);
    bool reserve(void *&ddest, size_t dlen);
};
```

This class builds a bulk buffer for use when the DB_MULTIPLE flag is specified to either the Db::put() or Db::del() methods. The buffer in the Dbt passed to the constructor is filled by calls to DbMultipleDataBuilder.append() or DbMultipleDataBuilder.reserve().

The constructor takes a The Dbt Handle that must be configured to contain a buffer managed by the application, with the ulen field set to the size of the buffer.

☞ All instances of the bulk retrieval classes may be used only once, and to build the bulk buffer in the forward direction only.

Parameters are:

• dbt

   The **dbt** parameter is a The Dbt Handle that must already be configured to contain a buffer managed by the application, with the ulen field set to the size of the buffer, which must be a multiple of 4.

DbMultipleDataBuilder.append()

The DbMultipleDataBuilder.append() method copies a data item to the end of the buffer.

The DbMultipleDataBuilder.append() method returns false if the data does not fit in the buffer and true otherwise.

Parameters are:

• dbuf

   A pointer to the data to be copied into the bulk buffer.

• dlen

   The number of bytes to be copied.

DbMultipleDataBuilder.reserve()

The `DbMultipleDataBuilder.reserve()` method reserves space for the next data item in the bulk buffer. Unlike the `append()`, no data is actually copied into the bulk buffer by `reserve()`: copying the data is the responsibility of the application.

The `DbMultipleDataBuilder.reserve()` method returns `false` if the data does not fit in the buffer and `true` otherwise.

Parameters are:

- `ddest`

  Set to a pointer to the position in the bulk buffer reserved for the data item, if enough space is available.

- `dlen`

  The number of bytes to reserve.

Class

Dbt

See Also

DBT and Bulk Operations

DbMultipleKeyDataBuilder

```
#include <db_cxx.h>

class DbMultipleKeyDataBuilder
{
public:
    DbMultipleKeyDataBuilder(Dbt &dbt);

    bool append(void *kbuf, size_t klen, void *dbuf, size_t dlen);
    bool reserve(void *&kdest, size_t klen, void *&ddest, size_t dlen);
};
```

This class builds a bulk buffer for use when the DB_MULTIPLE_KEY flag is specified to either the Db::put() or Db::del() methods with the btree or hash access methods. The buffer in the Dbt passed to the constructor is filled by calls to DbMultipleKeyDataBuilder.append() or DbMultipleKeyDataBuilder.reserve().

The constructor takes a The Dbt Handle that must be configured to contain a buffer managed by the application, with the ulen field set to the size of the buffer.

☞ All instances of the bulk retrieval classes may be used only once, and to build the bulk buffer in the forward direction only.

Parameters are:

• dbt

  The **dbt** parameter is a The Dbt Handle that must already be configured to contain a buffer managed by the application, with the ulen field set to the size of the buffer, which must be a multiple of 4.

DbMultipleKeyDataBuilder.append()

The DbMultipleKeyDataBuilder.append() method copies a key/data pair to the end of the buffer.

The DbMultipleKeyDataBuilder.append() method returns false if the key/data pair does not fit in the buffer and true otherwise.

Parameters are:

• kbuf

  A pointer to the key to be copied into the bulk buffer.

• klen

  The number of bytes of the key to be copied.

• dbuf

  A pointer to the data item to be copied into the bulk buffer.

- `dlen`

   The number of bytes of the data item to be copied.

DbMultipleKeyDataBuilder.reserve()

The `DbMultipleKeyDataBuilder.reserve()` method reserves space for the next key/data pair in the bulk buffer. Unlike the `append()`, no data is actually copied into the bulk buffer by `reserve()`: copying the data is the responsibility of the application.

The `DbMultipleKeyDataBuilder.reserve()` method returns `false` if the data does not fit in the buffer and `true` otherwise.

Parameters are:

- `kdest`

   Set to a pointer to the position in the bulk buffer reserved for the key, if enough space is available.

- `klen`

   The number of bytes to reserve for the key.

- `ddest`

   Set to a pointer to the position in the bulk buffer reserved for the data item, if enough space is available.

- `dlen`

   The number of bytes to reserve for the data item.

Class

DbMultipleBuilder

See Also

DBT and Bulk Operations

DbMultipleRecnoDataBuilder

```
#include <db_cxx.h>

class DbMultipleRecnoDataBuilder
{
public:
    DbMultipleRecnoDataBuilder(Dbt &dbt);

    bool append(db_recno_t recno, void *dbuf, size_t dlen);
    bool reserve(db_recno_t recno, void *&ddest, size_t dlen);
};
```

This class builds a bulk buffer for use when the DB_MULTIPLE_KEY flag is specified to either the Db::put() or Db::del() methods with the recno or queue access methods, or for the key when the DB_MULTIPLE flag is used. The buffer in the Dbt passed to the constructor is filled by calls to DbMultipleRecnoDataBuilder.append() or DbMultipleRecnoDataBuilder.reserve().

The constructor takes a The Dbt Handle that must be configured to contain a buffer managed by the application, with the ulen field set to the size of the buffer.

☞   All instances of the bulk retrieval classes may be used only once, and to build the bulk buffer in
         the forward direction only.

Parameters are:

• dbt

    The **dbt** parameter is a The Dbt Handle that must already be configured to contain a buffer managed
    by the application, with the ulen field set to the size of the buffer, which must be a multiple of 4.

bool append(db_recno_t recno, void *dbuf, size_t dlen);

DbMultipleRecnoDataBuilder.append()

The DbMultipleRecnoDataBuilder.append() method copies a record number / data pair to the end of the buffer.

The DbMultipleRecnoDataBuilder.append() method returns false if the record number / data pair does not fit in the buffer and true otherwise.

Parameters are:

• recno

    The record number to append.

• dbuf

    A pointer to the data item to be copied into the bulk buffer.

- `dlen`

    The number of bytes of the data item to be copied.

DbMultipleRecnoDataBuilder.reserve()

The `DbMultipleRecnoDataBuilder.reserve()` method reserves space for the next record number / data pair in the bulk buffer. The record number is appended, but unlike the `append()`, the data is not copied into the bulk buffer by `reserve()`: copying the data is the responsibility of the application.

The `DbMultipleRecnoDataBuilder.reserve()` method returns `false` if the record does not fit in the buffer and `true` otherwise.

Parameters are:

- `recno`

    The record number to append.

- `ddest`

    Set to a pointer to the position in the bulk buffer reserved for the data item, if enough space is available.

- `dlen`

    The number of bytes to reserve for the data item.

Class

DbMultipleBuilder

See Also

DBT and Bulk Operations

Chapter 5. The DbEnv Handle

The `DbEnv` object is the handle for a Berkeley DB environment — a collection including support for some or all of caching, locking, logging and transaction subsystems, as well as databases and log files. Methods of the `DbEnv` handle are used to configure the environment as well as to operate on subsystems and databases in the environment.

`DbEnv` handles are opened using the DbEnv::open() method.

When you are done using your environment, close it using the DbEnv::close() method. Before closing your environment, make sure all open database handles are closed first. See the Db::close() method for more information.

Database Environments and Related Methods

Database Environment Operations	Description
Db::get_env()	Return the Db's underlying DbEnv handle
DbEnv::close()	Close an environment
DbEnv	Create an environment handle
DbEnv::dbremove()	Remove a database
DbEnv::dbrename()	Rename a database
DbEnv::err()	Error message
DbEnv::failchk()	Check for thread failure
DbEnv::fileid_reset()	Reset database file IDs
DbEnv::get_home()	Return environment's home directory
DbEnv::get_open_flags()	Return flags with which the environment was opened
DbEnv::lsn_reset()	Reset database file LSNs
DbEnv::open()	Open an environment
DbEnv::remove()	Remove an environment
DbEnv::stat_print()	Environment statistics
DbEnv::strerror()	Error strings
DbEnv::version()	Return version information
Environment Configuration	
DbEnv::add_data_dir()	add an environment data directory
DbEnv::set_alloc()	Set local space allocation functions
DbEnv::set_app_dispatch()	Configure application recovery callback
DbEnv::set_data_dir(), DbEnv::get_data_dirs()	Set/get the environment data directory
DbEnv::set_create_dir(), DbEnv::get_create_dir()	add an environment data directory
DbEnv::set_encrypt(), DbEnv::get_encrypt_flags()	Set/get the environment cryptographic key
DbEnv::set_event_notify()	Set event notification callback
DbEnv::set_errcall()	Set error message callbacks
DbEnv::set_errfile(), DbEnv::get_errfile()	Set/get error message FILE
DbEnv::set_error_stream()	Set C++ ostream used for error messages
DbEnv::set_errpfx(), DbEnv::get_errpfx()	Set/get error message prefix
DbEnv::set_feedback()	Set feedback callback
DbEnv::set_flags(), DbEnv::get_flags()	Environment configuration
DbEnv::set_intermediate_dir_mode(), DbEnv::get_intermediate_dir_mode()	Set/get intermediate directory creation mode

Database Environment Operations	Description
DbEnv::set_isalive()	Set thread is-alive callback
DbEnv::set_message_stream()	Set C++ ostream used for informational messages
DbEnv::set_msgcall()	Set informational message callback
DbEnv::set_msgfile(), DbEnv::get_msgfile()	Set/get informational message FILE
DbEnv::set_shm_key(), DbEnv::get_shm_key()	Set/get system memory shared segment ID
DbEnv::set_thread_count(), DbEnv::get_thread_count()	Set/get approximate thread count
DbEnv::set_thread_id()	Set thread of control ID function
DbEnv::set_thread_id_string()	Set thread of control ID format function
DbEnv::set_timeout(), DbEnv::get_timeout()	Set/get lock and transaction timeout
DbEnv::set_tmp_dir(), DbEnv::get_tmp_dir()	Set/get the environment temporary file directory
DbEnv::set_verbose(), DbEnv::get_verbose()	Set/get verbose messages
DbEnv::set_cachesize(), DbEnv::get_cachesize()	Set/get the environment cache size

Db::get_env()

```
#include <db_cxx.h>

DbEnv *
Db::get_env();
```

The `Db::get_env()` method returns the handle for the database environment underlying the database.

The `Db::get_env()` method may be called at any time during the life of the application.

Class

Db

See Also

Database and Related Methods

DbEnv::add_data_dir()

```
#include <db_cxx.h>

int
DbEnv::add_data_dir(const char *dir);
```

Add the path of a directory to be used as the location of the access method database files. Paths specified to the Db::open() function will be searched relative to this path. Paths set using this method are additive, and specifying more than one will result in each specified directory being searched for database files.

If no database directories are specified, database files must be named either by absolute paths or relative to the environment home directory. See Berkeley DB File Naming for more information.

The database environment's data directories may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "add_data_dir", one or more whitespace characters, and the directory name.

The DbEnv::add_data_dir() method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The DbEnv::add_data_dir() method may not be called after the DbEnv::open() method is called. If the database environment already exists when DbEnv::open() is called, the information specified to DbEnv::add_data_dir() must be consistent with the existing environment or corruption can occur.

The DbEnv::add_data_dir() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

dir

The **dir** parameter is a directory to be used as a location for database files.

When using a Unicode build on Windows (the default), this argument will be interpreted as a UTF-8 string, which is equivalent to ASCII for Latin characters.

Errors

The DbEnv::add_data_dir() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after DbEnv::open() was called; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::close()

```
#include <db_cxx.h>

DbEnv::close(u_int32_t flags);
```

The `DbEnv::close()` method closes the Berkeley DB environment, freeing any allocated resources and closing any underlying subsystems.

The DbEnv handle should not be closed while any other handle that refers to it is not yet closed; for example, database environment handles must not be closed while database handles remain open, or transactions in the environment have not yet been committed or aborted. Specifically, this includes the Db, Dbc, DbTxn, DbLogc and DbMpoolFile handles.

Where the environment was initialized with the DB_INIT_LOCK flag, calling `DbEnv::close()` does not release any locks still held by the closing process, providing functionality for long-lived locks. Processes that want to have all their locks released can do so by issuing the appropriate DbEnv::lock_vec() call.

Where the environment was initialized with the DB_INIT_MPOOL flag, calling `DbEnv::close()` implies calls to DbMpoolFile::close() for any remaining open files in the memory pool that were returned to this process by calls to DbMpoolFile::open(). It does not imply a call to DbMpoolFile::sync() for those files.

Where the environment was initialized with the DB_INIT_TXN flag, calling `DbEnv::close()` aborts any unresolved transactions. Applications should not depend on this behavior for transactions involving Berkeley DB databases; all such transactions should be explicitly resolved. The problem with depending on this semantic is that aborting an unresolved transaction involving database operations requires a database handle. Because the database handles should have been closed before calling `DbEnv::close()`, it will not be possible to abort the transaction, and recovery will have to be run on the Berkeley DB environment before further operations are done.

Where log cursors were created using the DbEnv::log_cursor() method, calling `DbEnv::close()` does not imply closing those cursors.

In multithreaded applications, only a single thread may call the `DbEnv::close()` method.

After `DbEnv::close()` has been called, regardless of its return, the Berkeley DB environment handle may not be accessed again.

The `DbEnv::close()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter is currently unused, and must be set to 0.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv

```
#include <db_cxx.h>

class DbEnv {
public:
        DbEnv(u_int32 flags);
        ~DbEnv();

        DB_ENV *DbEnv::get_DB_ENV();
        const DB_ENV *DbEnv::get_const_DB_ENV() const;
        static DbEnv *DbEnv::get_DbEnv(DB_ENV *dbenv);
        static const DbEnv *DbEnv::get_const_DbEnv(const DB_ENV *dbenv);

        ...
};
```

The `DbEnv` object is the handle for a Berkeley DB environment — a collection including support for some or all of caching, locking, logging and transaction subsystems, as well as databases and log files. Methods of the `DbEnv` handle are used to configure the environment as well as to operate on subsystems and databases in the environment.

`DbEnv` handles are free-threaded if the DB_THREAD flag is specified to the DbEnv::open() method when the environment is opened. The `DbEnv` handle should not be closed while any other handle remains open that is using it as a reference (for example, Db or DbTxn). Once either the DbEnv::close() or DbEnv::remove() methods are called, the handle may not be accessed again, regardless of the method's return.

The constructor creates the `DbEnv` object. The constructor allocates memory internally; calling the DbEnv::close() or DbEnv::remove() methods will free that memory.

Before the handle may be used, you must open it using the DbEnv::open() method.

The **flags** parameter must be set to 0.

- `DB_CXX_NO_EXCEPTIONS`

  The Berkeley DB C++ API supports two different error behaviors. By default, whenever an error occurs, an exception is thrown that encapsulates the error information. This generally allows for cleaner logic for transaction processing because a try block can surround a single transaction. However, if `DB_CXX_NO_EXCEPTIONS` is specified, exceptions are not thrown; instead, each individual function returns an error code.

Each `DbEnv` object has an associated `DB_ENV` structure, which is used by the underlying implementation of Berkeley DB and its C-language API. The `DbEnv::get_DB_ENV()` method returns a pointer to this struct. Given a `const DbEnv` object, `DbEnv::get_const_DB_ENV()` returns a const pointer to the same struct.

Given a `DB_ENV` struct, the `DbEnv::get_DbEnv()` method returns the corresponding `DbEnv` object, if there is one. If the `DB_ENV` struct was not associated with a `DbEnv` (that is, it was not returned from a call to `DbEnv::get_DB_ENV()`), then the result of `DbEnv::get_DbEnv()` is undefined. Given a `const DB_ENV` struct, `DbEnv::get_const_Db_Env()` returns the associated `const DbEnv` object, if there is one.

These methods may be useful for Berkeley DB applications including both C and C++ language software. It should not be necessary to use these calls in a purely C++ application.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::dbremove()

```
#include <db_cxx.h>

int
DbEnv::dbremove(DbTxn *txnid,
    const char *file, const char *database, u_int32_t flags);
```

The `DbEnv::dbremove()` method removes the database specified by the **file** and **database** parameters. If no **database** is specified, the underlying file represented by **file** is removed, incidentally removing all of the databases it contained.

Applications should never remove databases with open Db handles, or in the case of removing a file, when any database in the file has an open handle.

The `DbEnv::dbremove()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

`DbEnv::dbremove()` is affected by any database directory specified using the DbEnv::set_data_dir() method, or by setting the `set_data_dir` string in the environment's DB_CONFIG file.

Parameters

database

The **database** parameter is the database to be removed.

file

The **file** parameter is the physical file which contains the database(s) to be removed.

txnid

If the operation is part of an application-specified transaction, the **txnid** parameter is a transaction handle returned from DbEnv::txn_begin(); if the operation is part of a Berkeley DB Concurrent Data Store group, the **txnid** parameter is a handle returned from DbEnv::cdsgroup_begin(); otherwise NULL. If no transaction handle is specified, but the DB_AUTO_COMMIT flag is specified to either this method or the environment handle, the operation will be implicitly transaction protected.

flags

The **flags** parameter must be set to 0 or the following value:

• DB_AUTO_COMMIT

  Enclose the `DbEnv::dbremove()` call within a transaction. If the call succeeds, changes made by the operation will be recoverable. If the call fails, the operation will have made no changes.

Environment Variables

The environment variable DB_HOME may be used as the path of the database environment home.

Errors

The `DbEnv::dbremove()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DbDeadlockException or DB_LOCK_DEADLOCK

A transactional database environment operation was selected to resolve a deadlock.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_LOCK_DEADLOCK` is returned.

DbLockNotGrantedException or DB_LOCK_NOTGRANTED

A Berkeley DB Concurrent Data Store database environment configured for lock timeouts was unable to grant a lock in the allowed time.

DbLockNotGrantedException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `DB_LOCK_NOTGRANTED` is returned.

EINVAL

If the method was called before DbEnv::open() was called; or if an invalid flag value or parameter was specified.

ENOENT

The file or directory does not exist.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::dbrename()

```
#include <db_cxx.h>

int
DbEnv::dbrename(DbTxn *txnid, const char *file,
    const char *database, const char *newname, u_int32_t flags);
```

The `DbEnv::dbrename()` method renames the database specified by the **file** and **database** parameters to **newname**. If no **database** is specified, the underlying file represented by **file** is renamed using the value supplied to **newname**, incidentally renaming all of the databases it contained.

Applications should not rename databases that are currently in use. If an underlying file is being renamed and logging is currently enabled in the database environment, no database in the file may be open when the `DbEnv::dbrename()` method is called.

The `DbEnv::dbrename()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

`DbEnv::dbrename()` is affected by any database directory specified using the DbEnv::set_data_dir() method, or by setting the `set_data_dir` string in the environment's DB_CONFIG file.

Parameters

database

The **database** parameter is the database to be renamed.

file

The **file** parameter is the physical file which contains the database(s) to be renamed.

When using a Unicode build on Windows (the default), the **file** argument will be interpreted as a UTF-8 string, which is equivalent to ASCII for Latin characters.

flags

The **flags** parameter must be set to 0 or the following value:

- `DB_AUTO_COMMIT`

  Enclose the `DbEnv::dbrename()` call within a transaction. If the call succeeds, changes made by the operation will be recoverable. If the call fails, the operation will have made no changes.

newname

The **newname** parameter is the new name of the database or file.

txnid

If the operation is part of an application-specified transaction, the **txnid** parameter is a transaction handle returned from DbEnv::txn_begin(); if the operation is part of a Berkeley DB Concurrent Data Store group, the **txnid** parameter is a handle returned from DbEnv::cdsgroup_begin(); otherwise NULL. If no transaction handle is specified, but the DB_AUTO_COMMIT flag is specified to either this method or the environment handle, the operation will be implicitly transaction protected.

Environment Variables

The environment variable DB_HOME may be used as the path of the database environment home.

Errors

The DbEnv::dbrename() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DbDeadlockException or DB_LOCK_DEADLOCK

A transactional database environment operation was selected to resolve a deadlock.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_LOCK_DEADLOCK is returned.

DbLockNotGrantedException or DB_LOCK_NOTGRANTED

A Berkeley DB Concurrent Data Store database environment configured for lock timeouts was unable to grant a lock in the allowed time.

DbLockNotGrantedException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_LOCK_NOTGRANTED is returned.

EINVAL

If the method was called before DbEnv::open() was called; or if an invalid flag value or parameter was specified.

ENOENT

The file or directory does not exist.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::err()

```
#include <db_cxx.h>

DbEnv::err(int error, const char *fmt, ...);

DbEnv::errx(const char *fmt, ...);
```

The `DbEnv::err()`, `DbEnv::errx,()`, Db::err() and `Db::errx()` methods provide error-messaging functionality for applications written using the Berkeley DB library.

The Db::err() and DbEnv::err() methods constructs an error message consisting of the following elements:

- **An optional prefix string**

  If no error callback function has been set using the DbEnv::set_errcall() method, any prefix string specified using the DbEnv::set_errpfx() method, followed by two separating characters: a colon and a <space> character.

- **An optional printf-style message**

  The supplied message **fmt**, if non-NULL, in which the ANSI C X3.159-1989 (ANSI C) printf function specifies how subsequent parameters are converted for output.

- **A separator**

  Two separating characters: a colon and a <space> character.

- **A standard error string**

  The standard system or Berkeley DB library error string associated with the **error** value, as returned by the DbEnv::strerror() method.

This constructed error message is then handled as follows:

- If an error callback function has been set (see Db::set_errcall() and DbEnv::set_errcall()), that function is called with two parameters: any prefix string specified (see Db::set_errpfx() and DbEnv::set_errpfx()) and the error message.

- If a C library FILE * has been set (see Db::set_errfile() and DbEnv::set_errfile()), the error message is written to that output stream.

- If a C++ ostream has been set (see DbEnv::set_error_stream() and Db::set_error_stream()), the error message is written to that stream.

- If none of these output options have been configured, the error message is written to stderr, the standard error output stream.

Parameters

error

The **error** parameter is the error value for which the `DbEnv::err()` and Db::err() methods will display a explanatory string.

fmt

The **fmt** parameter is an optional printf-style message to display.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::failchk()

```
#include <db_cxx.h>

int
DbEnv::failchk(u_int32_t flags);
```

The `DbEnv::failchk()` method checks for threads of control (either a true thread or a process) that have exited while manipulating Berkeley DB library data structures, while holding a logical database lock, or with an unresolved transaction (that is, a transaction that was never aborted or committed). For more information, see Architecting Data Store and Concurrent Data Store applications, and Architecting Transactional Data Store applications, both in the *Berkeley DB Programmer's Reference Guide*.

The `DbEnv::failchk()` method is based on the `thread_id` and `is_alive` functions specified to the DbEnv::set_thread_id() and DbEnv::set_isalive() methods. Applications calling the `DbEnv::failchk()` method must have already called the DbEnv::set_isalive() method, on the same DbEnv, and must have configured their database environment using the DbEnv::set_thread_count() method.

If `DbEnv::failchk()` determines a thread of control exited while holding database read locks, it will release those locks. If `DbEnv::failchk()` determines a thread of control exited with an unresolved transaction, the transaction will be aborted. In either of these cases, `DbEnv::failchk()` will return 0 and the application may continue to use the database environment.

In either of these cases, the `DbEnv::failchk()` method will also report the process and thread IDs associated with any released locks or aborted transactions. The information is printed to a specified output channel (see the DbEnv::set_msgfile() method for more information), or passed to an application callback function (see the DbEnv::set_msgcall() method for more information).

If `DbEnv::failchk()` determines a thread of control has exited such that database environment recovery is required, it will return DB_RUNRECOVERY. In this case, the application should not continue to use the database environment. For a further description as to the actions the application should take when this failure occurs, see Handling failure in Data Store and Concurrent Data Store applications, and Handling failure in Transactional Data Store applications, both in the *Berkeley DB Programmer's Reference Guide*.

In multiprocess applications, it is recommended that the DbEnv handle used to invoke the `DbEnv::failchk()` method not be shared and therefore not *free-threaded*.

The `DbEnv::failchk()` method may not be called by the application before the DbEnv::open() method is called.

The `DbEnv::failchk()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter is currently unused, and must be set to 0.

Errors

The `DbEnv::failchk()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::fileid_reset()

```
#include <db_cxx.h>

int
DbEnv::fileid_reset(const char *file, u_int32_t flags);
```

The `DbEnv::fileid_reset()` method allows database files to be copied, and then the copy used in the same database environment as the original.

All databases contain an ID string used to identify the database in the database environment cache. If a physical database file is copied, and used in the same environment as another file with the same ID strings, corruption can occur. The `DbEnv::fileid_reset()` method creates new ID strings for all of the databases in the physical file.

The `DbEnv::fileid_reset()` method modifies the physical file, in-place. Applications should not reset IDs in files that are currently in use.

The `DbEnv::fileid_reset()` method may be called at any time during the life of the application.

The `DbEnv::fileid_reset()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

file

The name of the physical file in which new file IDs are to be created.

flags

The **flags** parameter must be set to 0 or the following value:

- `DB_ENCRYPT`

    The file contains encrypted databases.

Errors

The `DbEnv::fileid_reset()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::get_create_dir()

```
#include <db_cxx.h>

int
DbEnv::get_create_dir(const char **dirp);
```

The `DbEnv::get_create_dir()` method returns a pointer to the name of the directory to create databases in.

The `DbEnv::get_create_dir()` method may be called at any time during the life of the application.

The `DbEnv::get_create_dir()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

dirp

The `DbEnv::get_create_dir()` method returns a ponter to the name of the directory in **dirp**.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::get_data_dirs()

```
#include <db_cxx.h>

int
DbEnv::get_data_dirs(const char ***dirpp);
```

The `DbEnv::get_data_dirs()` method returns the NULL-terminated array of directories.

The `DbEnv::get_data_dirs()` method may be called at any time during the life of the application.

The `DbEnv::get_data_dirs()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

dirpp

The `DbEnv::get_data_dirs()` method returns a reference to the NULL-terminated array of directories in **dirpp**.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::get_encrypt_flags()

```
#include <db_cxx.h>

int
DbEnv::get_encrypt_flags(u_int32_t *flagsp);
```

The `DbEnv::get_encrypt_flags()` method returns the encryption flags.

The `DbEnv::get_encrypt_flags()` method may be called at any time during the life of the application.

The `DbEnv::get_encrypt_flags()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flagsp

The `DbEnv::get_encrypt_flags()` method returns the encryption flags in **flagsp**.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::get_errfile()

```
#include <db_cxx.h>

void
DbEnv::get_errfile(FILE **errfilep);
```

The `DbEnv::get_errfile()` method returns the FILE * used for displaying additional Berkeley DB error messages. This C library is set using the DbEnv::set_errfile() method.

The `DbEnv::get_errfile()` method may be called at any time during the life of the application.

Parameters

errfilep

The `DbEnv::get_errfile()` method returns the FILE * in **errfilep**.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::get_errpfx()

```
#include <db_cxx.h>

void
DbEnv::get_errpfx(const char **errpfxp);
```

The `DbEnv::get_errpfx()` method returns the error prefix that appears before error messages issued by Berkeley DB. This error prefix is set using the DbEnv::set_errpfx() method.

The `DbEnv::get_errpfx()` method may be called at any time during the life of the application.

Parameters

errpfxp

The `DbEnv::get_errpfx()` method returns a reference to the error prefix in **errpfxp**.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::get_flags()

```
#include <db_cxx.h>

int
DbEnv::get_flags(u_int32_t *flagsp)
```

The `DbEnv::get_flags()` method returns the configuration flags set for a DbEnv handle. These flags are set using the DbEnv::set_flags() method.

The `DbEnv::get_flags()` method may be called at any time during the life of the application.

The `DbEnv::get_flags()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flagsp

The `DbEnv::get_flags()` method returns the configuration flags in **flagsp**.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::get_home()

```
#include <db_cxx.h>

int
DbEnv::get_home(const char **homep);
```

The `DbEnv::get_home()` method returns the database environment home directory. This directory is normally identified when the DbEnv::open() method is called.

The `DbEnv::get_home()` method may be called at any time during the life of the application.

The `DbEnv::get_home()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::get_intermediate_dir_mode()

```
#include <db_cxx.h>

int
DbEnv::get_intermediate_dir_mode(u_int32_t *modep);
```

The `DbEnv::get_intermediate_dir_mode()` method returns the intermediate directory permissions.

Intermediate directories are directories needed for recovery. Normally, Berkeley DB does not create these directories and will do so only if the DbEnv::set_intermediate_dir_mode() method is called.

The `DbEnv::get_intermediate_dir_mode()` method may be called at any time during the life of the application.

The `DbEnv::get_intermediate_dir_mode()` method returns a non-zero error value on failure and 0 on success.

Parameters

modep

The `DbEnv::get_intermediate_dir_mode()` method returns a reference to the intermediate directory permissions in **modep**.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::get_msgfile()

```
#include <db_cxx.h>

void
DbEnv::get_msgfile(FILE **msgfilep);
```

The `DbEnv::get_msgfile()` method returns the `FILE *` used for displaying messages. This is set using the DbEnv::set_msgfile() method.

The `DbEnv::get_msgfile()` method may be called at any time during the life of the application.

Parameters

msgfilep

The `DbEnv::get_msgfile()` method returns the FILE * in **msgfilep**.

Class

DbEnv

See Also

Database Environments and Related Methods, DbEnv::set_msgfile()

DbEnv::get_open_flags()

```
#include <db_cxx.h>

int
DbEnv::get_open_flags(u_int32_t *flagsp);
```

The `DbEnv::get_open_flags()` method returns the open method flags originally used to create the database environment.

The `DbEnv::get_open_flags()` method may not be called before the `DbEnv::open()` method is called.

The `DbEnv::get_open_flags()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flagsp

The `DbEnv::get_open_flags()` method returns the open method flags originally used to create the database environment in **flagsp**.

Class

DbEnv

See Also

Database Environments and Related Methods, DbEnv::open()

DbEnv::get_shm_key()

```
#include <db_cxx.h>

int
DbEnv::get_shm_key(long *shm_keyp);
```

The `DbEnv::get_shm_key()` method returns the base segment ID. This is used for Berkeley DB environment shared memory regions created in system memory on VxWorks or systems supporting X/Open-style shared memory interfaces. It may be specified using the DbEnv::set_shm_key() method.

The `DbEnv::get_shm_key()` method may be called at any time during the life of the application.

The `DbEnv::get_shm_key()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

shm_keyp

The `DbEnv::get_shm_key()` method returns the base segment ID in **shm_keyp**.

Class

DbEnv

See Also

Database Environments and Related Methods, DbEnv::set_shm_key()

DbEnv::get_thread_count()

```
#include <db_cxx.h>

int
DbEnv::get_thread_count(u_int32_t *countp);
```

The `DbEnv::get_thread_count()` method returns the thread count as set by the DbEnv::set_thread_count() method.

The `DbEnv::get_thread_count()` method may be called at any time during the life of the application.

The `DbEnv::get_thread_count()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

countp

The `DbEnv::get_thread_count()` method returns the thread count in **countp**.

Class

DbEnv

See Also

Database Environments and Related Methods, DbEnv::set_thread_count()

DbEnv::get_timeout()

```
#include <db_cxx.h>

int
DbEnv::get_timeout(db_timeout_t *timeoutp, u_int32_t flag);
```

The `DbEnv::get_timeout()` method returns a value, in microseconds, representing either lock or transaction timeouts. These values are set using the DbEnv::set_timeout() method.

The `DbEnv::get_timeout()` method may be called at any time during the life of the application.

The `DbEnv::get_timeout()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flag

The **flags** parameter must be set to one of the following values:

- `DB_SET_LOCK_TIMEOUT`

  Return the timeout value for locks in this database environment.

- `DB_SET_TXN_TIMEOUT`

  Return the timeout value for transactions in this database environment.

timeoutp

The **timeoutp** parameter references memory into which the timeout value of the specified **flag** parameter is copied.

Class

DbEnv

See Also

Database Environments and Related Methods, DbEnv::set_timeout()

DbEnv::get_tmp_dir()

```
#include <db_cxx.h>

int
DbEnv::get_tmp_dir(const char **dirp);
```

The `DbEnv::get_tmp_dir()` method returns the database environment temporary file directory.

The `DbEnv::get_tmp_dir()` method may be called at any time during the life of the application.

The `DbEnv::get_tmp_dir()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

dirp

The `DbEnv::get_tmp_dir()` method returns a reference to the database environment temporary file directory in **dirp**.

Class

DbEnv

See Also

Database Environments and Related Methods, DbEnv::set_tmp_dir()

DbEnv::get_verbose()

```
#include <db_cxx.h>

int
DbEnv::get_verbose(u_int32_t which, int *onoffp);
```

The `DbEnv::get_verbose()` method returns whether the specified **which** parameter is currently set or not. These parameters are set using the DbEnv::set_verbose() method.

The `DbEnv::get_verbose()` method may be called at any time during the life of the application.

The `DbEnv::get_verbose()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

which

The **which** parameter is the message value for which configuration is being checked. Must be set to one of the following values:

- `DB_VERB_DEADLOCK`

  Display additional information when doing deadlock detection.

- `DB_VERB_FILEOPS`

  Display additional information when performing filesystem operations such as open, close or rename. May not be available on all platforms.

- `DB_VERB_FILEOPS_ALL`

  Display additional information when performing all filesystem operations, including read and write. May not be available on all platforms.

- `DB_VERB_RECOVERY`

  Display additional information when performing recovery.

- `DB_VERB_REGISTER`

  Display additional information concerning support for the DB_REGISTER flag to the DbEnv::open() method.

- `DB_VERB_REPLICATION`

  Display all detailed information about replication. This includes the information displayed by all of the other DB_VERB_REP_* and DB_VERB_REPMGR_* values.

- `DB_VERB_REP_ELECT`

Display detailed information about replication elections.

- DB_VERB_REP_LEASE

Display detailed information about replication master leases.

- DB_VERB_REP_MISC

Display detailed information about general replication processing not covered by the other DB_VERB_REP_* values.

- DB_VERB_REP_MSGS

Display detailed information about replication message processing.

- DB_VERB_REP_SYNC

Display detailed information about replication client synchronization.

- DB_VERB_REPMGR_CONNFAIL

Display detailed information about Replication Manager connection failures.

- DB_VERB_REPMGR_MISC

Display detailed information about general Replication Manager processing.

- DB_VERB_WAITSFOR

Display the waits-for table when doing deadlock detection.

onoffp

The **onoffp** parameter references memory into which the configuration of the specified **which** parameter is copied.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::lsn_reset()

```
#include <db_cxx.h>

int
DbEnv::lsn_reset(const char *file, u_int32_t flags);
```

The `DbEnv::lsn_reset()` method allows database files to be moved from one transactional database environment to another.

Database pages in transactional database environments contain references to the environment's log files (that is, log sequence numbers, or LSNs). Copying or moving a database file from one database environment to another, and then modifying it, can result in data corruption if the LSNs are not first cleared.

Note that LSNs should be reset before moving or copying the database file into a new database environment, rather than moving or copying the database file and then resetting the LSNs. Berkeley DB has consistency checks that may be triggered if an application calls `DbEnv::lsn_reset()` on a database in a new environment when the database LSNs still reflect the old environment.

The `DbEnv::lsn_reset()` method modifies the physical file, in-place. Applications should not reset LSNs in files that are currently in use.

The `DbEnv::lsn_reset()` method may be called at any time during the life of the application.

The `DbEnv::lsn_reset()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

file

The name of the physical file in which the LSNs are to be cleared.

flags

The **flags** parameter must be set to 0 or the following value:

• `DB_ENCRYPT`

   The file contains encrypted databases.

Errors

The `DbEnv::lsn_reset()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::open()

```
#include <db_cxx.h>

int
DbEnv::open(const char *db_home, u_int32_t flags, int mode);
```

The `DbEnv::open()` method opens a Berkeley DB environment. It provides a structure for creating a consistent environment for processes using one or more of the features of Berkeley DB.

The `DbEnv::open()` method method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success. If `DbEnv::open()` fails, the DbEnv::close() method must be called to discard the DbEnv handle.

Parameters

db_home

The **db_home** parameter is the database environment's home directory. For more information on **db_home**, and filename resolution in general, see Berkeley DB File Naming. The environment variable **DB_HOME** may be used as the path of the database home, as described in Berkeley DB File Naming.

When using a Unicode build on Windows (the default), the **db_home** argument will be interpreted as a UTF-8 string, which is equivalent to ASCII for Latin characters.

flags

The **flags** parameter specifies the subsystems that are initialized and how the application's environment affects Berkeley DB file naming, among other things. The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the values described in this section.

Because there are a large number of flags that can be specified, they have been grouped together by functionality. The first group of flags indicates which of the Berkeley DB subsystems should be initialized.

The choice of subsystems initialized for a Berkeley DB database environment is specified by the thread of control initially creating the environment. Any subsequent thread of control joining the environment will automatically be configured to use the same subsystems as were created in the environment (unless the thread of control requests a subsystem not available in the environment, which will fail). Applications joining an environment, able to adapt to whatever subsystems have been configured in the environment, should open the environment without specifying any subsystem flags. Applications joining an environment, requiring specific subsystems from their environments, should open the environment specifying those specific subsystem flags.

- `DB_INIT_CDB`

  Initialize locking for the Berkeley DB Concurrent Data Store product. In this mode, Berkeley DB provides multiple reader/single writer access. The only other subsystem that should be specified with the `DB_INIT_CDB` flag is `DB_INIT_MPOOL`.

- `DB_INIT_LOCK`

Initialize the locking subsystem. This subsystem should be used when multiple processes or threads are going to be reading and writing a Berkeley DB database, so that they do not interfere with each other. If all threads are accessing the database(s) read-only, locking is unnecessary. When the DB_INIT_LOCK flag is specified, it is usually necessary to run a deadlock detector, as well. See db_deadlock and DbEnv::lock_detect() for more information.

- DB_INIT_LOG

Initialize the logging subsystem. This subsystem should be used when recovery from application or system failure is necessary. If the log region is being created and log files are already present, the log files are reviewed; subsequent log writes are appended to the end of the log, rather than overwriting current log entries.

- DB_INIT_MPOOL

Initialize the shared memory buffer pool subsystem. This subsystem should be used whenever an application is using any Berkeley DB access method.

- DB_INIT_REP

Initialize the replication subsystem. This subsystem should be used whenever an application plans on using replication. The DB_INIT_REP flag requires the DB_INIT_TXN and DB_INIT_LOCK flags also be configured.

- DB_INIT_TXN

Initialize the transaction subsystem. This subsystem should be used when recovery and atomicity of multiple operations are important. The DB_INIT_TXN flag implies the DB_INIT_LOG flag.

The second group of flags govern what recovery, if any, is performed when the environment is initialized:

- DB_RECOVER

Run normal recovery on this environment before opening it for normal use. If this flag is set, the DB_CREATE and DB_INIT_TXN flags must also be set, because the regions will be removed and re-created, and transactions are required for application recovery.

- DB_RECOVER_FATAL

Run catastrophic recovery on this environment before opening it for normal use. If this flag is set, the DB_CREATE and DB_INIT_TXN flags must also be set, because the regions will be removed and re-created, and transactions are required for application recovery.

A standard part of the recovery process is to remove the existing Berkeley DB environment and create a new one in which to perform recovery. If the thread of control performing recovery does not specify the correct region initialization information (for example, the correct memory pool cache size), the result can be an application running in an environment with incorrect cache and other subsystem sizes. For this reason, the thread of control performing recovery should specify correct configuration information before calling the DbEnv::open() method; or it should remove the environment after recovery is completed, leaving creation of the correctly sized environment to a subsequent call to the DbEnv::open() method.

All Berkeley DB recovery processing must be single-threaded; that is, only a single thread of control may perform recovery or access a Berkeley DB environment while recovery is being performed. Because it is not an error to specify DB_RECOVER for an environment for which no recovery is required, it is reasonable programming practice for the thread of control responsible for performing recovery and creating the environment to always specify the DB_CREATE and DB_RECOVER flags during startup.

The third group of flags govern file-naming extensions in the environment:

• DB_USE_ENVIRON

  The Berkeley DB process' environment may be permitted to specify information to be used when naming files; see Berkeley DB File Naming. Because permitting users to specify which files are used can create security problems, environment information will be used in file naming for all users only if the DB_USE_ENVIRON flag is set.

• DB_USE_ENVIRON_ROOT

  The Berkeley DB process' environment may be permitted to specify information to be used when naming files; see Berkeley DB File Naming. Because permitting users to specify which files are used can create security problems, if the DB_USE_ENVIRON_ROOT flag is set, environment information will be used in file naming only for users with appropriate permissions (for example, users with a user-ID of 0 on UNIX systems).

Finally, there are a few additional unrelated flags:

• DB_CREATE

  Cause Berkeley DB subsystems to create any underlying files, as necessary.

• DB_LOCKDOWN

  Lock shared Berkeley DB environment files and memory-mapped databases into memory.

• DB_FAILCHK

  Internally call the DbEnv::failchk() method as part of opening the environment. When DB_FAILCHK is specified, a check is made to ensure all DbEnv::failchk() prerequisites are meet.

  If the DB_FAILCHK flag is used in conjunction with the DB_REGISTER flag, then a check will be made to see if the environment needs recovery. If recovery is needed, a call will be made to the DbEnv::failchk() method to release any database reads locks held by the thread of control that exited and, if needed, to abort the unresolved transaction. If DbEnv::failchk() determines environment recovery is still required, the recovery actions for DB_REGISTER will be followed.

  If the DB_FAILCHK flag is not used in conjunction with the DB_REGISTER flag, then make an internal call to DbEnv::failchk() as the last step of opening the environment. If DbEnv::failchk() determines database environment recovery is required, DB_RUNRECOVERY will be returned.

• DB_PRIVATE

Allocate region memory from the heap instead of from memory backed by the filesystem or system shared memory.

This flag implies the environment will only be accessed by a single process (although that process may be multithreaded). This flag has two effects on the Berkeley DB environment. First, all underlying data structures are allocated from per-process memory instead of from shared memory that is accessible to more than a single process. Second, mutexes are only configured to work between threads.

This flag should not be specified if more than a single process is accessing the environment because it is likely to cause database corruption and unpredictable behavior. For example, if both a server application and Berkeley DB utilities (for example, db_archive, db_checkpoint or db_stat) are expected to access the environment, the DB_PRIVATE flag should not be specified.

See Shared Memory Regions for more information.

- DB_REGISTER

  Check to see if recovery needs to be performed before opening the database environment. (For this check to be accurate, all processes using the environment must specify DB_REGISTER when opening the environment.) If recovery needs to be performed for any reason (including the initial use of the DB_REGISTER flag), and DB_RECOVER is also specified, recovery will be performed and the open will proceed normally. If recovery needs to be performed and DB_RECOVER is not specified, DB_RUNRECOVERY will be returned. If recovery does not need to be performed, the DB_RECOVER flag will be ignored. See Architecting Transactional Data Store applications for more information.

- DB_SYSTEM_MEM

  Allocate region memory from system shared memory instead of from heap memory or memory backed by the filesystem.

  See Shared Memory Regions for more information.

- DB_THREAD

  Cause the DbEnv handle returned by DbEnv::open() to be *free-threaded*; that is, concurrently usable by multiple threads in the address space. The DB_THREAD flag should be specified if the DbEnv handle will be concurrently used by more than one thread in the process, or if any Db handles opened in the scope of the DbEnv handle will be concurrently used by more than one thread in the process.

  This flag is required when using the Replication Manager.

mode

On Windows systems, the mode parameter is ignored.

On UNIX systems or in IEEE/ANSI Std 1003.1 (POSIX) environments, files created by Berkeley DB are created with mode **mode** (as described in **chmod**(2)) and modified by the process' umask value at the time of creation (see **umask**(2)). Created files are owned by the process owner; the group ownership of created files is based on the system and directory defaults, and is not further specified by Berkeley DB. System shared memory segments created by Berkeley DB are created with mode **mode**, unmodified

by the process' umask value. If **mode** is 0, Berkeley DB will use a default mode of readable and writable by both owner and group.

Errors

The `DbEnv::open()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DB_RUNRECOVERY

Either the `DB_REGISTER` flag was specified, a failure occurred, and no recovery flag was specified, or the `DB_FAILCHK` flag was specified and recovery was deemed necessary.

DB_VERSION_MISMATCH

The version of the Berkeley DB library doesn't match the version that created the database environment.

EAGAIN

The shared memory region was locked and (repeatedly) unavailable.

EINVAL

If the `DB_THREAD` flag was specified and fast mutexes are not available for this architecture; The `DB_HOME` or `TMPDIR` environment variables were set, but empty; An incorrectly formatted **NAME VALUE** entry or line was found; or if an invalid flag value or parameter was specified.

ENOSPC

HP-UX only: Due to the constraints of the PA-RISC memory architecture, HP-UX does not allow a process to map a file into its address space multiple times. For this reason, each Berkeley DB environment may be opened only once by a process on HP-UX; that is, calls to `DbEnv::open()` will fail if the specified Berkeley DB environment has been opened and not subsequently closed.

ENOENT

The file or directory does not exist.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::remove()

```
#include <db_cxx.h>

int
DbEnv::remove(const char *db_home, u_int32_t flags);
```

The DbEnv::remove() method destroys a Berkeley DB environment if it is not currently in use. The environment regions, including any backing files, are removed. Any log or database files and the environment directory are not removed.

If there are processes that have called DbEnv::open() without calling DbEnv::close() (that is, there are processes currently using the environment), DbEnv::remove() will fail without further action unless the DB_FORCE flag is set, in which case DbEnv::remove() will attempt to remove the environment, regardless of any processes still using it.

The result of attempting to forcibly destroy the environment when it is in use is unspecified. Processes using an environment often maintain open file descriptors for shared regions within it. On UNIX systems, the environment removal will usually succeed, and processes that have already joined the region will continue to run in that region without change. However, processes attempting to join the environment will either fail or create new regions. On other systems in which the **unlink**(2) system call will fail if any process has an open file descriptor for the file (for example Windows/NT), the region removal will fail.

Calling DbEnv::remove() should not be necessary for most applications because the Berkeley DB environment is cleaned up as part of normal database recovery procedures. However, applications may want to call DbEnv::remove() as part of application shut down to free up system resources. For example, if the DB_SYSTEM_MEM flag was specified to DbEnv::open(), it may be useful to call DbEnv::remove() in order to release system shared memory segments that have been allocated. Or, on architectures in which mutexes require allocation of underlying system resources, it may be useful to call DbEnv::remove() in order to release those resources. Alternatively, if recovery is not required because no database state is maintained across failures, and no system resources need to be released, it is possible to clean up an environment by simply removing all the Berkeley DB files in the database environment's directories.

In multithreaded applications, only a single thread may call the DbEnv::remove() method.

A DbEnv handle that has already been used to open an environment should not be used to call the DbEnv::remove() method; a new DbEnv handle should be created for that purpose.

After DbEnv::remove() has been called, regardless of its return, the Berkeley DB environment handle may not be accessed again.

The DbEnv::remove() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

db_home

The **db_home** parameter names the database environment to be removed.

When using a Unicode build on Windows (the default), the **db_home** argument will be interpreted as a UTF-8 string, which is equivalent to ASCII for Latin characters.

flags

The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

- DB_FORCE

  If set, the environment is removed, regardless of any processes that may still using it, and no locks are acquired during this process. (Generally, this flag is specified only when applications were unable to shut down cleanly, and there is a risk that an application may have died holding a Berkeley DB lock.)

- DB_USE_ENVIRON

  The Berkeley DB process' environment may be permitted to specify information to be used when naming files; see Berkeley DB File Naming. Because permitting users to specify which files are used can create security problems, environment information will be used in file naming for all users only if the DB_USE_ENVIRON flag is set.

- DB_USE_ENVIRON_ROOT

  The Berkeley DB process' environment may be permitted to specify information to be used when naming files; see Berkeley DB File Naming. Because permitting users to specify which files are used can create security problems, if the DB_USE_ENVIRON_ROOT flag is set, environment information will be used in file naming only for users with appropriate permissions (for example, users with a user-ID of 0 on UNIX systems).

Errors

The DbEnv::remove() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EBUSY

The shared memory region was in use and the force flag was not set.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_alloc()

```
#include <db_cxx.h>

extern "C" {
    typedef void *(*db_malloc_fcn_type)(size_t);
    typedef void *(*db_realloc_fcn_type)(void *, size_t);
    typedef void *(*db_free_fcn_type)(void *);
};

int
DbEnv::set_alloc(db_malloc_fcn_type app_malloc,
    db_realloc_fcn_type app_realloc,
    db_free_fcn_type app_free);
```

Set the allocation functions used by the DbEnv and Db methods to allocate or free memory owned by the application.

There are a number of interfaces in Berkeley DB where memory is allocated by the library and then given to the application. For example, the DB_DBT_MALLOC flag, when specified in the Dbt object, will cause the Db methods to allocate and reallocate memory which then becomes the responsibility of the calling application. Other examples are the Berkeley DB interfaces which return statistical information to the application: Db::stat(), DbEnv::lock_stat(), DbEnv::log_archive(), DbEnv::log_stat(), DbEnv::memp_stat(), and DbEnv::txn_stat(). There is one method in Berkeley DB where memory is allocated by the application and then given to the library: the callback specified to Db::associate().

On systems in which there may be multiple library versions of the standard allocation routines (notably Windows NT), transferring memory between the library and the application will fail because the Berkeley DB library allocates memory from a different heap than the application uses to free it. To avoid this problem, the DbEnv::set_alloc() and Db::set_alloc() methods can be used to pass Berkeley DB references to the application's allocation routines.

It is not an error to specify only one or two of the possible allocation function parameters to these interfaces; however, in that case the specified interfaces must be compatible with the standard library interfaces, as they will be used together. The functions specified must match the calling conventions of the ANSI C X3.159-1989 (ANSI C) library routines of the same name.

The DbEnv::set_alloc() method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The DbEnv::set_alloc() method may not be called after the DbEnv::open() method is called.

The DbEnv::set_alloc() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

app_malloc

The **app_malloc** parameter is the application-specified malloc function.

app_realloc

The **app_realloc** parameter is the application-specified realloc function.

app_free

The **app_free** parameter is the application-specified free function.

Errors

The `DbEnv::set_alloc()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after DbEnv::open() was called; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_app_dispatch()

```
#include <db_cxx.h>

int
DbEnv::set_app_dispatch(int (*tx_recover)(DbEnv *dbenv,
    Dbt *log_rec, DbLsn *lsn, db_recops op));
```

Declare a function to be called during transaction abort and recovery to process application-specific log records.

The DbEnv::set_app_dispatch() method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The DbEnv::set_app_dispatch() method may not be called after the DbEnv::open() method is called. If the database environment already exists when DbEnv::open() is called, the information specified to DbEnv::set_app_dispatch() must be consistent with the existing environment or corruption can occur.

The DbEnv::set_app_dispatch() method returns a non-zero error value on failure and 0 on success.

Parameters

tx_recover

The **tx_recover** parameter is the application's abort and recovery function. The function takes four parameters:

- dbenv

  The **dbenv** parameter is the enclosing database environment handle.

- log_rec

  The **log_rec** parameter is a log record.

- lsn

  The **lsn** parameter is a log sequence number.

- op

  The **op** parameter is one of the following values:

  - DB_TXN_BACKWARD_ROLL

    The log is being read backward to determine which transactions have been committed and to abort those operations that were not; undo the operation described by the log record.

  - DB_TXN_FORWARD_ROLL

    The log is being played forward; redo the operation described by the log record.

- `DB_TXN_ABORT`

  The log is being read backward during a transaction abort; undo the operation described by the log record.

- `DB_TXN_APPLY`

  The log is being applied on a replica site; redo the operation described by the log record.

- `DB_TXN_PRINT`

  The log is being printed for debugging purposes; print the contents of this log record in the desired format.

The DB_TXN_FORWARD_ROLL and DB_TXN_APPLY operations frequently imply the same actions, redoing changes that appear in the log record, although if a recovery function is to be used on a replication client where reads may be taking place concurrently with the processing of incoming messages, DB_TXN_APPLY operations should also perform appropriate locking. The macro DB_REDO(op) checks that the operation is one of DB_TXN_FORWARD_ROLL or DB_TXN_APPLY, and should be used in the recovery code to refer to the conditions under which operations should be redone. Similarly, the macro DB_UNDO(op) checks if the operation is one of DB_TXN_BACKWARD_ROLL or DB_TXN_ABORT.

The function must return 0 on success and either **errno** or a value outside of the Berkeley DB error name space on failure.

Errors

The `DbEnv::set_app_dispatch()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after DbEnv::open() was called; or if an invalid flag value or parameter was specified.

Class

DbEnv, DbTxn

See Also

Transaction Subsystem and Related Methods

DbEnv::set_data_dir()

```
#include <db_cxx.h>

int
DbEnv::set_data_dir(const char *dir);
```

☞ This interface has been deprecated. You should use DbEnv::add_data_dir() and DbEnv::set_create_dir() instead.

Set the path of a directory to be used as the location of the access method database files. Paths specified to the Db::open() function will be searched relative to this path. Paths set using this method are additive, and specifying more than one will result in each specified directory being searched for database files. If any directories are specified, database files will always be created in the first path specified.

If no database directories are specified, database files must be named either by absolute paths or relative to the environment home directory. See Berkeley DB File Naming for more information.

The database environment's data directories may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_data_dir", one or more whitespace characters, and the directory name.

The DbEnv::set_data_dir() method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The DbEnv::set_data_dir() method may not be called after the DbEnv::open() method is called. If the database environment already exists when DbEnv::open() is called, the information specified to DbEnv::set_data_dir() must be consistent with the existing environment or corruption can occur.

The DbEnv::set_data_dir() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

dir

The **dir** parameter is a directory to be used as a location for database files.

When using a Unicode build on Windows (the default), this argument will be interpreted as a UTF-8 string, which is equivalent to ASCII for Latin characters.

Errors

The DbEnv::set_data_dir() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after DbEnv::open() was called; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_create_dir()

```
#include <db_cxx.h>

int
DbEnv::set_create_dir(const char *dir);
```

Sets the path of a directory to be used as the location to create the access method database files. When the Db::open() function is used to create a file it will be created relative to this path.

If no database directories are specified, database files will be created either by absolute paths or relative to the environment home directory. See Berkeley DB File Naming for more information.

The database environment's create directory may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_create_dir", one or more whitespace characters, and the directory name.

The DbEnv::set_create_dir() method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The DbEnv::set_create_dir() method may be called at any time.

The DbEnv::set_create_dir() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

dir

The **dir** parameter is a directory to be used to create database files. This directory must be one of the directories specified via a call to DbEnv::add_data_dir()

When using a Unicode build on Windows (the default), this argument will be interpreted as a UTF-8 string, which is equivalent to ASCII for Latin characters.

Errors

The DbEnv::set_create_dir() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after DbEnv::open() was called; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_encrypt()

```
#include <db_cxx.h>

int
DbEnv::set_encrypt(const char *passwd, u_int32_t flags);
```

Set the password used by the Berkeley DB library to perform encryption and decryption.

The `DbEnv::set_encrypt()` method configures a database environment, not only operations performed using the specified DbEnv handle.

The `DbEnv::set_encrypt()` method may not be called after the DbEnv::open() method is called. If the database environment already exists when DbEnv::open() is called, the information specified to `DbEnv::set_encrypt()` must be consistent with the existing environment or an error will be returned.

The `DbEnv::set_encrypt()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or the following value:

- `DB_ENCRYPT_AES`

  Use the Rijndael/AES (also known as the Advanced Encryption Standard and Federal Information Processing Standard (FIPS) 197) algorithm for encryption or decryption.

passwd

The **passwd** parameter is the password used to perform encryption and decryption.

Errors

The `DbEnv::set_encrypt()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after DbEnv::open() was called; or if an invalid flag value or parameter was specified.

EOPNOTSUPP

Cryptography is not available in this Berkeley DB release.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_event_notify()

```
#include <db_cxx.h>

int
DbEnv::set_event_notify(
    void (*db_event_fcn)(DB_ENV *dbenv, u_int32_t event,
    void *event_info));
```

The `DbEnv::set_event_notify()` method configures a callback function which is called to notify the process of specific Berkeley DB events.

☞ Berkeley DB is not re-entrant. Callback functions should not attempt to make library calls (for example, to release locks or close open handles). Re-entering Berkeley DB is not guaranteed to work correctly, and the results are undefined.

The `DbEnv::set_event_notify()` method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The `DbEnv::set_event_notify()` method may be called at any time during the life of the application.

The `DbEnv::set_event_notify()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

db_event_fcn

The **db_event_fcn** parameter is the application's event notification function. The function takes three parameters:

• dbenv

    The **dbenv** parameter is the enclosing database environment handle.

• event

    The **event** parameter is one of the following values:

    • DB_EVENT_PANIC

        Errors can occur in the Berkeley DB library where the only solution is to shut down the application and run recovery (for example, if Berkeley DB is unable to allocate heap memory). In such cases, the Berkeley DB methods will return DB_RUNRECOVERY. It is often easier to simply exit the application when such errors occur rather than gracefully return up the stack.

        When **event** is set to DB_EVENT_PANIC, the database environment has failed. All threads of control in the database environment should exit the environment, and recovery should be run.

    • DB_EVENT_REG_ALIVE

Recovery is needed in an environment where the DB_REGISTER flag was specified on the DbEnv::open() method and there is a process attached to the environment. The callback function is triggered once for each process attached.

The **event_info** parameter points to a pid_t value containing the process identifier (pid) of the process the Berkeley DB library detects is attached to the environment.

- DB_EVENT_REG_PANIC

Recovery is needed in an environment where the DB_REGISTER flag was specified on the DbEnv::open() method. All threads of control in the database environment should exit the environment.

This event is different than the DB_EVENT_PANIC event because it can only be triggered when DB_REGISTER was specified. It can be used to distinguish between the case when a process dies in the environment and recovery is initiated versus the case when an error happened (for example, if Berkeley DB is unable to allocate heap memory)

- DB_EVENT_REP_CLIENT

The local site is now a replication client.

- DB_EVENT_REP_ELECTED

The local replication site has just won an election. An application using the Base replication API should arrange for a call to the DbEnv::rep_start() method after receiving this event, to reconfigure the local environment as a replication master.

Replication Manager applications may safely ignore this event. The Replication Manager calls DbEnv::rep_start() automatically on behalf of the application when appropriate (resulting in firing of the DB_EVENT_REP_MASTER event).

- DB_EVENT_REP_MASTER

The local site is now the master site of its replication group. It is the application's responsibility to begin acting as the master environment.

- DB_EVENT_REP_NEWMASTER

The replication group of which this site is a member has just established a new master; the local site is not the new master. The **event_info** parameter points to an integer containing the environment ID of the new master.

- DB_EVENT_REP_PERM_FAILED

The replication manager did not receive enough acknowledgements (based on the acknowledgement policy configured with DbEnv::repmgr_set_ack_policy()) to ensure a transaction's durability within the replication group. The transaction will be flushed to the master's local disk storage for durability.

The DB_EVENT_REP_PERM_FAILED event is provided only to applications configured for the replication manager.

- DB_EVENT_REP_STARTUPDONE

    The client has completed startup synchronization and is now processing live log records received from the master.

- DB_EVENT_WRITE_FAILED

    A Berkeley DB write to stable storage failed.

- event_info

    The **event_info** parameter may reference memory which contains additional information describing an event. By default, **event_info** is NULL; specific events may pass non-NULL values, in which case the event will also describe the memory's structure.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_errcall()

```
#include <db_cxx.h>

void DbEnv::set_errcall(void (*db_errcall_fcn)
    (const Dbenv *dbenv, const char *errpfx, const char *msg));
```

When an error occurs in the Berkeley DB library, an exception is thrown or an error return value is returned by the interface. In some cases, however, the **errno** value may be insufficient to completely describe the cause of the error, especially during initial application debugging.

The DbEnv::set_errcall() and DbEnv::set_errcall() methods are used to enhance the mechanism for reporting error messages to the application. In some cases, when an error occurs, Berkeley DB will call **db_errcall_fcn** with additional error information. It is up to the **db_errcall_fcn** function to display the error message in an appropriate manner.

Setting **db_errcall_fcn** to NULL unconfigures the callback interface.

Alternatively, you can use the DbEnv::set_error_stream() and Db::set_error_stream() methods to display the additional information via an output stream, or the Db::set_errfile() or Db::set_errfile() methods to display the additional information via a C library FILE *. You should not mix these approaches.

This error-logging enhancement does not slow performance or significantly increase application size, and may be run during normal operation as well as during application debugging.

The DbEnv::set_errcall() method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The DbEnv::set_errcall() method may be called at any time during the life of the application.

☞ Berkeley DB is not re-entrant. Callback functions should not attempt to make library calls (for example, to release locks or close open handles). Re-entering Berkeley DB is not guaranteed to work correctly, and the results are undefined.

Parameters

db_errcall_fcn

The **db_errcall_fcn** parameter is the application-specified error reporting function. The function takes three parameters:

• dbenv

  The **dbenv** parameter is the enclosing database environment.

• errpfx

  The **errpfx** parameter is the prefix string (as previously set by Db::set_errpfx() or DbEnv::set_errpfx()).

• msg

  The **msg** parameter is the error message string.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_errfile()

```
#include <db_cxx.h>

void
DbEnv::set_errfile(FILE *errfile);
```

When an error occurs in the Berkeley DB library, an exception is thrown or an error return value is returned by the interface. In some cases, however, the return value may be insufficient to completely describe the cause of the error especially during initial application debugging.

The DbEnv::set_errfile() and Db::set_errfile() methods are used to enhance the mechanism for reporting error messages to the application by setting a C library FILE * to be used for displaying additional Berkeley DB error messages. In some cases, when an error occurs, Berkeley DB will output an additional error message to the specified file reference.

Alternatively, you can use the DbEnv::set_error_stream() and Db::set_error_stream() methods to display the additional messages via an output stream, or the DbEnv::set_errcall() or Db::set_errcall() methods to capture the additional error information in a way that does not use C library FILE *s. You should not mix these approaches.

The error message will consist of the prefix string and a colon ("**:**") (if a prefix string was previously specified using Db::set_errpfx() or DbEnv::set_errpfx()), an error string, and a trailing <newline> character.

The default configuration when applications first create Db or DbEnv handles is as if the Db::set_errfile() or DbEnv::set_errfile() methods were called with the standard error output (stderr) specified as the FILE * argument. Applications wanting no output at all can turn off this default configuration by calling the Db::set_errfile() or DbEnv::set_errfile() methods with NULL as the FILE * argument. Additionally, explicitly configuring the error output channel using any of the following methods will also turn off this default output for the application:

- DbEnv::set_errfile()

- Db::set_errfile()

- DbEnv::set_errcall()

- Db::set_errcall()

- DbEnv::set_error_stream()

- Db::set_error_stream()

This error logging enhancement does not slow performance or significantly increase application size, and may be run during normal operation as well as during application debugging.

The DbEnv::set_errfile() method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The DbEnv::set_errfile() method may be called at any time during the life of the application.

Parameters

errfile

The **errfile** parameter is a C library FILE * to be used for displaying additional Berkeley DB error information.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_error_stream()

```
#include <db_cxx.h>

void DbEnv::set_error_stream(class ostream*);
```

When an error occurs in the Berkeley DB library, an exception is thrown or an **errno** value is returned by the interface. In some cases, however, the **errno** value may be insufficient to completely describe the cause of the error, especially during initial application debugging.

The `DbEnv::set_error_stream()` and Db::set_error_stream() methods are used to enhance the mechanism for reporting error messages to the application by setting the C++ ostream used for displaying additional Berkeley DB error messages. In some cases, when an error occurs, Berkeley DB will output an additional error message to the specified stream.

The error message will consist of the prefix string and a colon ("**:**") (if a prefix string was previously specified using Db::set_errpfx(), an error string, and a trailing <newline> character.

Setting **stream** to NULL unconfigures the interface.

Alternatively, you can use the DbEnv::set_errfile() or Db::set_errfile() methods to display the additional information via a C Library `FILE *`, or the DbEnv::set_errcall() and Db::set_errcall() methods to capture the additional error information in a way that does not use either output streams or C Library `FILE *`'s. You should not mix these approaches.

This error-logging enhancement does not slow performance or significantly increase application size, and may be run during normal operation as well as during application debugging.

The `DbEnv::set_error_stream()` method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The `DbEnv::set_error_stream()` method may be called at any time during the life of the application.

Parameters

stream

The **stream** parameter is the application-specified output stream to be used for additional error information.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_errpfx()

```
#include <db_cxx.h>

void
DbEnv::set_errpfx(const char *errpfx);
```

Set the prefix string that appears before error messages issued by Berkeley DB.

The Db::set_errpfx() and DbEnv::set_errpfx() methods do not copy the memory to which the **errpfx** parameter refers; rather, they maintain a reference to it. Although this allows applications to modify the error message prefix at any time (without repeatedly calling the interfaces), it means the memory must be maintained until the handle is closed.

The DbEnv::set_errpfx() method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The DbEnv::set_errpfx() method may be called at any time during the life of the application.

Parameters

errpfx

The **errpfx** parameter is the application-specified error prefix for additional error messages.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_feedback()

```
#include <db_cxx.h>

int
DbEnv::set_feedback(void (*db_feedback_fcn)(DbEnv *dbenv, int opcode,
    int percent));
```

Some operations performed by the Berkeley DB library can take non-trivial amounts of time. The
DbEnv::set_feedback() method can be used by applications to monitor progress within these operations.
When an operation is likely to take a long time, Berkeley DB will call the specified callback function
with progress information.

☞ Berkeley DB is not re-entrant. Callback functions should not attempt to make library calls (for
example, to release locks or close open handles). Re-entering Berkeley DB is not guaranteed to
work correctly, and the results are undefined.

It is up to the callback function to display this information in an appropriate manner.

The DbEnv::set_feedback() method configures operations performed using the specified DbEnv handle.

The DbEnv::set_feedback() method may be called at any time during the life of the application.

The DbEnv::set_feedback() method returns a non-zero error value on failure and 0 on success.

Parameters

db_feedback_fcn

The **db_feedback_fcn** parameter is the application-specified feedback function called to report Berkeley
DB operation progress. The callback function must take three parameters:

• dbenv

   The **dbenv** parameter is a reference to the enclosing database environment.

• opcode

   The **opcode** parameter is an operation code. The **opcode** parameter may take on any of the following
   values:

   • DB_RECOVER

      The environment is being recovered.

• percent

   The **percent** parameter is the percent of the operation that has been completed, specified as an
   integer value between 0 and 100.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_flags()

```
#include <db_cxx.h>

int
DbEnv::set_flags(u_int32_t flags, int onoff);
```

Configure a database environment.

The database environment's flag values may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_flags", one or more whitespace characters, and the method flag parameter as a string; for example, "set_flags DB_TXN_NOSYNC". Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The `DbEnv::set_flags()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set by bitwise inclusively **OR**'ing together one or more of the following values:

- DB_AUTO_COMMIT

  If set, Db handle operations for which no explicit transaction handle was specified, and which modify databases in the database environment, will be automatically enclosed within a transaction.

  Calling `DbEnv::set_flags()` with this flag only affects the specified DbEnv handle (and any other Berkeley DB handles opened within the scope of that handle). For consistent behavior across the environment, all DbEnv handles opened in the environment must either set this flag or the flag should be specified in the DB_CONFIG configuration file.

  This flag may be used to configure Berkeley DB at any time during the life of the application.

- DB_CDB_ALLDB

  If set, Berkeley DB Concurrent Data Store applications will perform locking on an environment-wide basis rather than on a per-database basis.

  Calling `DbEnv::set_flags()` with the DB_CDB_ALLDB flag only affects the specified DbEnv handle (and any other Berkeley DB handles opened within the scope of that handle). For consistent behavior across the environment, all DbEnv handles opened in the environment must either set the DB_CDB_ALLDB flag or the flag should be specified in the DB_CONFIG configuration file.

  The DB_CDB_ALLDB flag may be used to configure Berkeley DB only before the DbEnv::open() method is called.

- DB_DIRECT_DB

Turn off system buffering of Berkeley DB database files to avoid double caching.

Calling `DbEnv::set_flags()` with the DB_DIRECT_DB flag only affects the specified DbEnv handle (and any other Berkeley DB handles opened within the scope of that handle). For consistent behavior across the environment, all DbEnv handles opened in the environment must either set the DB_DIRECT_DB flag or the flag should be specified in the DB_CONFIG configuration file.

The DB_DIRECT_DB flag may be used to configure Berkeley DB at any time during the life of the application.

- DB_DSYNC_DB

Configure Berkeley DB to flush database writes to the backing disk before returning from the write system call, rather than flushing database writes explicitly in a separate system call, as necessary. This is only available on some systems (for example, systems supporting the IEEE/ANSI Std 1003.1 (POSIX) standard O_DSYNC flag, or systems supporting the Windows FILE_FLAG_WRITE_THROUGH flag). This flag may result in inaccurate file modification times and other file-level information for Berkeley DB database files. This flag will almost certainly result in a performance decrease on most systems. This flag is only applicable to certain filesysystems (for example, the Veritas VxFS filesystem), where the filesystem's support for trickling writes back to stable storage behaves badly (or more likely, has been misconfigured).

Calling `DbEnv::set_flags()` with the DB_DSYNC_DB flag only affects the specified DbEnv handle (and any other Berkeley DB handles opened within the scope of that handle). For consistent behavior across the environment, all DbEnv handles opened in the environment must either set the DB_DSYNC_DB flag or the flag should be specified in the DB_CONFIG configuration file.

The DB_DSYNC_DB flag may be used to configure Berkeley DB at any time during the life of the application.

- DB_MULTIVERSION

If set, all databases in the environment will be opened as if DB_MULTIVERSION is passed to the Db::open() method. This flag will be ignored for queue databases for which DB_MULTIVERSION is not supported.

Calling `DbEnv::set_flags()` with the DB_MULTIVERSION flag only affects the specified DbEnv handle (and any other Berkeley DB handles opened within the scope of that handle). For consistent behavior across the environment, all DbEnv handles opened in the environment must either set the DB_MULTIVERSION flag or the flag should be specified in the DB_CONFIG configuration file.

The DB_MULTIVERSION flag may be used to configure Berkeley DB at any time during the life of the application.

- DB_NOLOCKING

If set, Berkeley DB will grant all requested mutual exclusion mutexes and database locks without regard for their actual availability. This functionality should never be used for purposes other than debugging.

Calling `DbEnv::set_flags()` with the DB_NOLOCKING flag only affects the specified DbEnv handle (and any other Berkeley DB handles opened within the scope of that handle).

The DB_NOLOCKING flag may be used to configure Berkeley DB at any time during the life of the application.

• `DB_NOMMAP`

If set, Berkeley DB will copy read-only database files into the local cache instead of potentially mapping them into process memory (see the description of the DbEnv::set_mp_mmapsize() method for further information).

Calling `DbEnv::set_flags()` with the DB_NOMMAP flag only affects the specified DbEnv handle (and any other Berkeley DB handles opened within the scope of that handle). For consistent behavior across the environment, all DbEnv handles opened in the environment must either set the DB_NOMMAP flag or the flag should be specified in the DB_CONFIG configuration file.

The DB_NOMMAP flag may be used to configure Berkeley DB at any time during the life of the application.

• `DB_NOPANIC`

If set, Berkeley DB will ignore any panic state in the database environment. (Database environments in a panic state normally refuse all attempts to call Berkeley DB functions, returning DB_RUNRECOVERY.) This functionality should never be used for purposes other than debugging.

Calling `DbEnv::set_flags()` with the DB_NOPANIC flag only affects the specified DbEnv handle (and any other Berkeley DB handles opened within the scope of that handle).

The DB_NOPANIC flag may be used to configure Berkeley DB at any time during the life of the application.

• `DB_OVERWRITE`

Overwrite files stored in encrypted formats before deleting them. Berkeley DB overwrites files using alternating 0xff, 0x00 and 0xff byte patterns. For file overwriting to be effective, the underlying file must be stored on a fixed-block filesystem. Systems with journaling or logging filesystems will require operating system support and probably modification of the Berkeley DB sources.

Calling `DbEnv::set_flags()` with the DB_OVERWRITE flag only affects the specified DbEnv handle (and any other Berkeley DB handles opened within the scope of that handle).

The DB_OVERWRITE flag may be used to configure Berkeley DB at any time during the life of the application.

• `DB_PANIC_ENVIRONMENT`

If set, Berkeley DB will set the panic state for the database environment. (Database environments in a panic state normally refuse all attempts to call Berkeley DB functions, returning DB_RUNRECOVERY.) This flag may not be specified using the environment's DB_CONFIG file.

Calling `DbEnv::set_flags()` with the DB_PANIC_ENVIRONMENT flag affects the database environment, including all threads of control accessing the database environment.

The DB_PANIC_ENVIRONMENT flag may be used to configure Berkeley DB only after the DbEnv::open() method is called.

- `DB_REGION_INIT`

  In some applications, the expense of page-faulting the underlying shared memory regions can affect performance. (For example, if the page-fault occurs while holding a lock, other lock requests can convoy, and overall throughput may decrease.) If set, Berkeley DB will page-fault shared regions into memory when initially creating or joining a Berkeley DB environment. In addition, Berkeley DB will write the shared regions when creating an environment, forcing the underlying virtual memory and filesystems to instantiate both the necessary memory and the necessary disk space. This can also avoid out-of-disk space failures later on.

  Calling `DbEnv::set_flags()` with the DB_REGION_INIT flag only affects the specified DbEnv handle (and any other Berkeley DB handles opened within the scope of that handle). For consistent behavior across the environment, all DbEnv handles opened in the environment must either set the DB_REGION_INIT flag or the flag should be specified in the DB_CONFIG configuration file.

  The DB_REGION_INIT flag may be used to configure Berkeley DB at any time during the life of the application.

- `DB_TIME_NOTGRANTED`

  If set, database calls timing out based on lock or transaction timeout values will return DB_LOCK_NOTGRANTED instead of DB_LOCK_DEADLOCK. This allows applications to distinguish between operations which have deadlocked and operations which have exceeded their time limits.

  Calling `DbEnv::set_flags()` with the DB_TIME_NOTGRANTED flag only affects the specified DbEnv handle (and any other Berkeley DB handles opened within the scope of that handle). For consistent behavior across the environment, all DbEnv handles opened in the environment must either set the DB_TIME_NOTGRANTED flag or the flag should be specified in the DB_CONFIG configuration file.

  The `DB_TIME_NOTGRANTED` flag may be used to configure Berkeley DB at any time during the life of the application.

  Note that the DbEnv::lock_get() and DbEnv::lock_vec() methods are unaffected by this flag.

- `DB_TXN_NOSYNC`

  If set, Berkeley DB will not write or synchronously flush the log on transaction commit. This means that transactions exhibit the ACI (atomicity, consistency, and isolation) properties, but not D (durability); that is, database integrity will be maintained, but if the application or system fails, it is possible some number of the most recently committed transactions may be undone during recovery. The number of transactions at risk is governed by how many log updates can fit into the log buffer, how often the operating system flushes dirty buffers to disk, and how often the log is checkpointed.

  Calling `DbEnv::set_flags()` with the DB_TXN_NOSYNC flag only affects the specified DbEnv handle (and any other Berkeley DB handles opened within the scope of that handle). For consistent behavior

across the environment, all DbEnv handles opened in the environment must either set the DB_TXN_NOSYNC flag or the flag should be specified in the DB_CONFIG configuration file.

The DB_TXN_NOSYNC flag may be used to configure Berkeley DB at any time during the life of the application.

- DB_TXN_NOWAIT

  If set and a lock is unavailable for any Berkeley DB operation performed in the context of a transaction, cause the operation to return DB_LOCK_DEADLOCK (or DB_LOCK_NOTGRANTED if configured using the DB_TIME_NOTGRANTED flag).

  Calling DbEnv::set_flags() with the DB_TXN_NOWAIT flag only affects the specified DbEnv handle (and any other Berkeley DB handles opened within the scope of that handle). For consistent behavior across the environment, all DbEnv handles opened in the environment must either set the DB_TXN_NOWAIT flag or the flag should be specified in the DB_CONFIG configuration file.

  The DB_TXN_NOWAIT flag may be used to configure Berkeley DB at any time during the life of the application.

- DB_TXN_SNAPSHOT

  If set, all transactions in the environment will be started as if DB_TXN_SNAPSHOT were passed to the DbEnv::txn_begin() method, and all non-transactional cursors will be opened as if DB_TXN_SNAPSHOT were passed to the Db::cursor() method.

  Calling DbEnv::set_flags() with the DB_TXN_SNAPSHOT flag only affects the specified DbEnv handle (and any other Berkeley DB handles opened within the scope of that handle). For consistent behavior across the environment, all DbEnv handles opened in the environment must either set the DB_TXN_SNAPSHOT flag or the flag should be specified in the DB_CONFIG configuration file.

  The DB_TXN_SNAPSHOT flag may be used to configure Berkeley DB at any time during the life of the application.

- DB_TXN_WRITE_NOSYNC

  If set, Berkeley DB will write, but will not synchronously flush, the log on transaction commit. This means that transactions exhibit the ACI (atomicity, consistency, and isolation) properties, but not D (durability); that is, database integrity will be maintained, but if the system fails, it is possible some number of the most recently committed transactions may be undone during recovery. The number of transactions at risk is governed by how often the system flushes dirty buffers to disk and how often the log is checkpointed.

  Calling DbEnv::set_flags() with the DB_TXN_WRITE_NOSYNC flag only affects the specified DbEnv handle (and any other Berkeley DB handles opened within the scope of that handle). For consistent behavior across the environment, all DbEnv handles opened in the environment must either set the DB_TXN_WRITE_NOSYNC flag or the flag should be specified in the DB_CONFIG configuration file.

  The DB_TXN_WRITE_NOSYNC flag may be used to configure Berkeley DB at any time during the life of the application.

- DB_YIELDCPU

  If set, Berkeley DB will yield the processor immediately after each page or mutex acquisition. This functionality should never be used for purposes other than stress testing.

  Calling `DbEnv::set_flags()` with the DB_YIELDCPU flag only affects the specified DbEnv handle (and any other Berkeley DB handles opened within the scope of that handle). For consistent behavior across the environment, all DbEnv handles opened in the environment must either set the DB_YIELDCPU flag or the flag should be specified in the DB_CONFIG configuration file.

  The DB_YIELDCPU flag may be used to configure Berkeley DB at any time during the life of the application.

onoff

If the **onoff** parameter is zero, the specified flags are cleared; otherwise they are set.

Errors

The `DbEnv::set_flags()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_intermediate_dir_mode()

```
#include <db_cxx.h>

int
DbEnv::set_intermediate_dir_mode(u_int32_t mode);
```

By default, Berkeley DB does not create intermediate directories needed for recovery, that is, if the file **/a/b/c/mydatabase** is being recovered, and the directory path **b/c** does not exist, recovery will fail. This default behavior is because Berkeley DB does not know what permissions are appropriate for intermediate directory creation, and creating the directory might result in a security problem.

The `DbEnv::set_intermediate_dir_mode()` method causes Berkeley DB to create any intermediate directories needed during recovery, using the specified permissions.

On UNIX systems or in IEEE/ANSI Std 1003.1 (POSIX) environments, created directories are owned by the process owner; the group ownership of created directories is based on the system and directory defaults, and is not further specified by Berkeley DB.

The database environment's intermediate directory permissions may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_intermediate_dir_mode", one or more whitespace characters, and the directory permissions. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The `DbEnv::set_intermediate_dir_mode()` method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The `DbEnv::set_intermediate_dir_mode()` method may not be called after the DbEnv::open() method is called.

The `DbEnv::set_intermediate_dir_mode()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

mode

The **mode** parameter specifies the directory permissions.

Directory permissions are interpreted as a string of nine characters, using the character set **r** (read), **w** (write), **x** (execute or search), and **-** (none). The first character is the read permissions for the directory owner (set to either **r** or **-**). The second character is the write permissions for the directory owner (set to either **w** or **-**). The third character is the execute permissions for the directory owner (set to either **x** or **-**).

Similarly, the second set of three characters are the read, write and execute/search permissions for the directory group, and the third set of three characters are the read, write and execute/search permissions for all others. For example, the string **rwx------** would configure read, write and execute/search access for the owner only. The string **rwxrwx---** would configure read, write and

execute/search access for both the owner and the group. The string **rwxr-----** would configure read, write and execute/search access for the directory owner and read-only access for the directory group.

Errors

The `DbEnv::set_intermediate_dir_mode()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after DbEnv::open() was called; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_isalive()

```
#include <db_cxx.h>

int
DbEnv::set_isalive(int (*is_alive)(DbEnv *dbenv, pid_t pid,
    db_threadid_t tid, u_int32_t flags));
```

Declare a function that returns if a thread of control (either a true thread or a process) is still running. The `DbEnv::set_isalive()` method supports the DbEnv::failchk() method. For more information, see Architecting Data Store and Concurrent Data Store applications, and Architecting Transactional Data Store applications, both in the *Berkeley DB Programmer's Reference Guide*.

The `DbEnv::set_isalive()` method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The `DbEnv::set_isalive()` method may be called at any time during the life of the application.

The `DbEnv::set_isalive()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

is_alive

The **is_alive** parameter is a function which returns non-zero if the thread of control, identified by the **pid** and **tid** arguments, is still running. The function takes four arguments:

- **dbenv**

  The **dbenv** parameter is the enclosing database environment handle, allowing application access to the application-private fields of that object.

- **pid**

  The **pid** parameter is a process ID returned by the function specified to the DbEnv::set_thread_id() method.

- **tid**

  The **tid** parameter is a thread ID returned by the function specified to the DbEnv::set_thread_id() method.

- **flags**

  The **flags** parameter must be set to 0 or the following value:

  - `DB_MUTEX_PROCESS_ONLY`

    Return only if the process is alive, the thread ID should be ignored.

Errors

The `DbEnv::set_isalive()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_message_stream()

```
#include <db_cxx.h>

void DbEnv::set_message_stream(class ostream*);
```

There are interfaces in the Berkeley DB library which either directly output informational messages or statistical information, or configure the library to output such messages when performing other operations. For example, the DbEnv::set_verbose() and DbEnv::stat_print() methods.

The `DbEnv::set_message_stream()` and Db::set_message_stream() methods are used to display these messages for the application. In this case, the message will include a trailing `<newline>` character.

Setting **stream** to NULL unconfigures the interface.

Alternatively, you can use the DbEnv::set_msgfile() or Db::set_msgfile() methods to display the additional information via a C Library `FILE *`, or the DbEnv::set_msgcall() and Db::set_msgcall() methods to capture the additional error information in a way that does not use either output streams or C Library `FILE *`'s. You should not mix these approaches.

The `DbEnv::set_message_stream()` method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The `DbEnv::set_message_stream()` method may be called at any time during the life of the application.

Parameters

stream

The **stream** parameter is the application-specified output stream to be used for additional message information.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_msgcall()

```
#include <db_cxx.h>

void DbEnv::set_msgcall(void (*db_msgcall_fcn)(const DbEnv *dbenv,
    const char *msg));
```

There are interfaces in the Berkeley DB library which either directly output informational messages or statistical information, or configure the library to output such messages when performing other operations, for example, DbEnv::set_verbose() and DbEnv::stat_print().

The DbEnv::set_msgcall() and Db::set_msgcall() methods are used to pass these messages to the application, and Berkeley DB will call **db_msgcall_fcn** with each message. It is up to the **db_msgcall_fcn** function to display the message in an appropriate manner.

Setting **db_msgcall_fcn** to NULL unconfigures the callback interface.

Alternatively, you can use the DbEnv::set_error_stream() and Db::set_error_stream() methods to display the messages via an output stream, or the Db::set_msgfile() or Db::set_msgfile() methods to display the messages via a C library FILE *. You should not mix these approaches.

The DbEnv::set_msgcall() method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The DbEnv::set_msgcall() method may be called at any time during the life of the application.

☞ Berkeley DB is not re-entrant. Callback functions should not attempt to make library calls (for example, to release locks or close open handles). Re-entering Berkeley DB is not guaranteed to work correctly, and the results are undefined.

Parameters

db_msgcall_fcn

The **db_msgcall_fcn** parameter is the application-specified message reporting function. The function takes two parameters:

- dbenv

  The **dbenv** parameter is the enclosing database environment.

- msg

  The **msg** parameter is the message string.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_msgfile()

```
#include <db_cxx.h>

void
DbEnv::set_msgfile(FILE *msgfile);
```

There are interfaces in the Berkeley DB library which either directly output informational messages or statistical information, or configure the library to output such messages when performing other operations, for example, DbEnv::set_verbose() and DbEnv::stat_print().

The DbEnv::set_msgfile() and Db::set_msgfile() methods are used to display these messages for the application. In this case the message will include a trailing <newline> character.

Setting **msgfile** to NULL unconfigures the interface.

Alternatively, you can use the DbEnv::set_message_stream() and Db::set_message_stream() methods to display the messages via an output stream, or the DbEnv::set_msgcall() or Db::set_msgcall() methods to capture the additional error information in a way that does not use C library FILE *'s. You should not mix these approaches.

The DbEnv::set_msgfile() method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The DbEnv::set_msgfile() method may be called at any time during the life of the application.

Parameters

msgfile

The **msgfile** parameter is a C library FILE * to be used for displaying messages.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_shm_key()

```
#include <db_cxx.h>

int
DbEnv::set_shm_key(long shm_key);
```

Specify a base segment ID for Berkeley DB environment shared memory regions created in system memory on VxWorks or systems supporting X/Open-style shared memory interfaces; for example, UNIX systems supporting **shmget**(2) and related System V IPC interfaces.

This base segment ID will be used when Berkeley DB shared memory regions are first created. It will be incremented a small integer value each time a new shared memory region is created; that is, if the base ID is 35, the first shared memory region created will have a segment ID of 35, and the next one will have a segment ID between 36 and 40 or so. A Berkeley DB environment always creates a master shared memory region; an additional shared memory region for each of the subsystems supported by the environment (Locking, Logging, Memory Pool and Transaction); plus an additional shared memory region for each additional memory pool cache that is supported. Already existing regions with the same segment IDs will be removed. See Shared Memory Regions for more information.

The intent behind this method is two-fold: without it, applications have no way to ensure that two Berkeley DB applications don't attempt to use the same segment IDs when creating different Berkeley DB environments. In addition, by using the same segment IDs each time the environment is created, previously created segments will be removed, and the set of segments on the system will not grow without bound.

The database environment's base segment ID may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_shm_key", one or more whitespace characters, and the ID. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The DbEnv::set_shm_key() method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The DbEnv::set_shm_key() method may not be called after the DbEnv::open() method is called. If the database environment already exists when DbEnv::open() is called, the information specified to DbEnv::set_shm_key() must be consistent with the existing environment or corruption can occur.

The DbEnv::set_shm_key() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

shm_key

The **shm_key** parameter is the base segment ID for the database environment.

Errors

The `DbEnv::set_shm_key()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after DbEnv::open() was called; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_thread_count()

```
#include <db_cxx.h>

int
DbEnv::set_thread_count(u_int32_t count);
```

Declare an approximate number of threads in the database environment. The `DbEnv::set_thread_count()` method must be called prior to opening the database environment if the DbEnv::failchk() method will be used. The `DbEnv::set_thread_count()` method does not set the maximum number of threads but is used to determine memory sizing and the thread control block reclamation policy.

If a process has not configured an **is_alive** function from the DbEnv::set_isalive() method, and then attempts to join a database environment configured for failure checking with the DbEnv::failchk(), DbEnv::set_thread_id(), DbEnv::set_isalive() and `DbEnv::set_thread_count()` methods, the program may be unable to allocate a thread control block and fail to join the environment. **This is true of the standalone Berkeley DB utility programs.** To avoid problems when using the standalone Berkeley DB utility programs with environments configured for failure checking, incorporate the utility's functionality directly in the application, or call the DbEnv::failchk() method before running the utility.

The database environment's thread count may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_thread_count", one or more whitespace characters, and the thread count. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The `DbEnv::set_thread_count()` method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The `DbEnv::set_thread_count()` method may not be called after the DbEnv::open() method is called.

The `DbEnv::set_thread_count()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

count

The **count** parameter is an approximate thread count for the database environment.

Errors

The `DbEnv::set_thread_count()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after DbEnv::open() was called; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_thread_id()

```
#include <db_cxx.h>

int
DbEnv::set_thread_id(char *(*thread_id_string)(DbEnv *dbenv,
    pid_t pid, db_threadid_t tid, char *buf));
```

Declare a function that returns a unique identifier pair for the current thread of control. The
DbEnv::set_thread_id() method supports the DbEnv::failchk() method. For more information, see
Architecting Data Store and Concurrent Data Store applications, and Architecting Transactional Data
Store applications, both in the *Berkeley DB Programmer's Reference Guide*.

The DbEnv::set_thread_id() method configures operations performed using the specified DbEnv handle,
not all operations performed on the underlying database environment.

The DbEnv::set_thread_id() method may be called at any time during the life of the application.

The DbEnv::set_thread_id() method either returns a non-zero error value or throws an exception that
encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

thread_id

The **thread_id** parameter is a function which returns a unique identifier pair for a thread of control
in a Berkeley DB application. The function takes three arguments:

- **dbenv**

  The **dbenv** parameter is the enclosing database environment handle, allowing application access to
  the application-private fields of that object.

- **pid**

  The **pid** points to a memory location of type pid_t, or NULL. The process ID of the current thread
  of control may be returned in this memory location, if it is not NULL.

- **tid**

  The **tid** points to a memory location of type db_threadid_t, or NULL. The thread ID of the current
  thread of control may be returned in this memory location, if it is not NULL.

Errors

The DbEnv::set_thread_id() method may fail and throw a DbException exception, encapsulating one
of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Assigning Thread IDs

The standard system library calls to return process and thread IDs are often sufficient for this purpose (for example, `getpid()` and `pthread_self()` on POSIX systems or GetCurrentThreadID on Windows systems). However, if the Berkeley DB application dynamically creates processes or threads, some care may be necessary in assigning unique IDs. In most threading systems, process and thread IDs are available for re-use as soon as the process or thread exits. If a new process or thread is created between the time of process or thread exit, and the DbEnv::failchk() method is run, it may be possible for DbEnv::failchk() to not detect that a thread of control exited without properly releasing all Berkeley DB resources.

It may be possible to handle this problem by inhibiting process or thread creation between thread of control exit and calling the DbEnv::failchk() method. Alternatively, the **thread_id** function must be constructed to not re-use **pid/tid** pairs. For example, in a single process application, the returned process ID might be used as an incremental counter, with the returned thread ID set to the actual thread ID. Obviously, the **is_alive** function specified to the DbEnv::set_isalive() method must be compatible with any **thread_id** function specified to `DbEnv::set_thread_id()`.

The db_threadid_t type is configured to be the same type as a standard thread identifier, in Berkeley DB configurations where this type is known (for example, systems supporting pthread_t or thread_t, or DWORD on Windows). If the Berkeley DB configuration process is unable to determine the type of a standard thread identifier, the db_thread_t type is set to uintmax_t (or the largest available unsigned integral type, on systems lacking the uintmax_t type). Applications running on systems lacking a detectable standard thread type, and which are also using thread APIs where a thread identifier is not an integral value and so will not fit into the configured db_threadid_t type, must either translate between the db_threadid_t type and the thread identifier (mapping the thread identifier to a unique identifier of the appropriate size), or modify the Berkeley DB sources to use an appropriate db_threadid_t type. Note: we do not currently know of any systems where this is necessary. If your application has to solve this problem, please contact our support group and let us know.

If no **thread_id** function is specified by the application, the Berkeley DB library will identify threads of control by using the `taskIdSelf()` call on VxWorks, the `getpid()` and `GetCurrentThreadID()` calls on Windows, the `getpid()` and `pthread_self()` calls when the Berkeley DB library has been configured for POSIX pthreads or Solaris LWP threads, the `getpid()` and `thr_self()` calls when the Berkeley DB library has been configured for UI threads, and otherwise `getpid()`.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_thread_id_string()

```
#include <db_cxx>

int
DbEnv::set_thread_id(void (*thread_id)
                        (DbEnv *dbenv, pid_t *pid, db_threadid_t *tid));
```

Declare a function that formats a process ID and thread ID identifier pair for display into a caller-supplied buffer. The function must return a reference to the caller-specified buffer. The `DbEnv::set_thread_id_string()` method supports the DbEnv::set_thread_id() method.

The `DbEnv::set_thread_id_string()` method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The `DbEnv::set_thread_id_string()` method may be called at any time during the life of the application.

The `DbEnv::set_thread_id_string()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

thread_id_string

The **thread_id_string** parameter is a function which returns a buffer in which is an identifier pair formatted for display. The function takes four arguments:

- dbenv

  The **dbenv** parameter is the enclosing database environment handle, allowing application access to the application-private fields of that object.

- pid

  The **pid** argument is a process ID.

- tid

  The **tid** argument is a thread ID.

- buf

  The **buf** argument is character array of at least DB_THREADID_STRLEN bytes in length, into which the identifier pair should be formatted.

If no **thread_id_string** function is specified, the default routine displays the identifier pair as "pid/tid", that is, the process ID represented as an unsigned integer value, a slash ('/') character, then the thread ID represented as an unsigned integer value.

Errors

The `DbEnv::set_thread_id_string()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_timeout()

```
#include <db_cxx.h>

int
DbEnv::set_timeout(db_timeout_t timeout, u_int32_t flags);
```

The `DbEnv::set_timeout()` method sets timeout values for locks or transactions in the database environment, and the wait time for a process to exit the environment when DB_REGISTER recovery is needed.

DB_SET_LOCK_TIMEOUT and DB_SET_TXN_TIMEOUT timeouts are checked whenever a thread of control blocks on a lock or when deadlock detection is performed. In the case of DB_SET_LOCK_TIMEOUT, the lock is one requested explicitly through the Lock subsystem interfaces. In the case of DB_SET_TXN_TIMEOUT, the lock is one requested on behalf of a transaction. In either case, it may be a lock requested by the database access methods underlying the application. These timeouts are only checked when the lock request first blocks or when deadlock detection is performed, the accuracy of the timeout depends on how often deadlock detection is performed.

Lock and transaction timeout values specified for the database environment may be overridden on a per-lock or per-transaction basis. See DbEnv::lock_vec() and DbTxn::set_timeout() for more information.

The `DbEnv::set_timeout()` method may be called at any time during the life of the application.

The `DbEnv::set_timeout()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to one of the following values:

- DB_SET_LOCK_TIMEOUT

  Set the timeout value for locks in this database environment.

  The database environment's lock timeout value may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_lock_timeout", one or more whitespace characters, and the lock timeout value. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

  This flag configures a database environment, not only operations performed using the specified DbEnv handle.

- DB_SET_REG_TIMEOUT

  Set the timeout value on how long to wait for processes to exit the environment before recovery is started when the DbEnv::open() method was called with the DB_REGISTER flag and recovery must be performed.

This wait timeout value may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_reg_timeout", one or more whitespace characters, and the wait timeout value. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

This flag configures operations performed using the specified DbEnv handle.

- DB_SET_TXN_TIMEOUT

  Set the timeout value for transactions in this database environment.

  The database environment's transaction timeout value may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_txn_timeout", one or more whitespace characters, and the transaction timeout value. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

  This flag configures a database environment, not only operations performed using the specified DbEnv handle.

timeout

The **timeout** parameter is the timeout value. It must be specified as an unsigned 32-bit number of microseconds, limiting the maximum timeout to roughly 71 minutes.

Errors

The DbEnv::set_timeout() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_tmp_dir()

```
#include <db_cxx.h>

int
DbEnv::set_tmp_dir(const char *dir);
```

Specify the path of a directory to be used as the location of temporary files. The files created to back in-memory access method databases will be created relative to this path. These temporary files can be quite large, depending on the size of the database.

If no directories are specified, the following alternatives are checked in the specified order. The first existing directory path is used for all temporary files.

1. The value of the environment variable **TMPDIR**.

2. The value of the environment variable **TEMP**.

3. The value of the environment variable **TMP**.

4. The value of the environment variable **TempFolder**.

5. The value returned by the **GetTempPath** interface.

6. The directory **/var/tmp**.

7. The directory **/usr/tmp**.

8. The directory **/temp**.

9. The directory **/tmp**.

10. The directory **C:/temp**.

11. The directory **C:/tmp**.

☞   Environment variables are only checked if one of the DB_USE_ENVIRON or DB_USE_ENVIRON_ROOT flags were specified.

☞   The GetTempPath interface is only checked on Win/32 platforms.

The database environment's temporary file directory may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_tmp_dir", one or more whitespace characters, and the directory name. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The DbEnv::set_tmp_dir() method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The DbEnv::set_tmp_dir() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

dir

The **dir** parameter is the directory to be used to store temporary files.

When using a Unicode build on Windows (the default), the this argument will be interpreted as a UTF-8 string, which is equivalent to ASCII for Latin characters.

Errors

The `DbEnv::set_tmp_dir()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after DbEnv::open() was called; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_verbose()

```
#include <db_cxx.h>

int
DbEnv::set_verbose(u_int32_t which, int onoff);
```

The `DbEnv::set_verbose()` method turns specific additional informational and debugging messages in the Berkeley DB message output on and off. To see the additional messages, verbose messages must also be configured for the application. For more information on verbose messages, see the DbEnv::set_msgfile() method.

The database environment's messages may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_verbose", one or more whitespace characters, and the method **which** parameter as a string; for example, "set_verbose DB_VERB_RECOVERY". Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The `DbEnv::set_verbose()` method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The `DbEnv::set_verbose()` method may be called at any time during the life of the application.

The `DbEnv::set_verbose()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

onoff

If the **onoff** parameter is set to non-zero, the additional messages are output.

which

The **which** parameter must be set to one of the following values:

- `DB_VERB_DEADLOCK`

  Display additional information when doing deadlock detection.

- `DB_VERB_FILEOPS`

  Display additional information when performing filesystem operations such as open, close or rename. May not be available on all platforms.

- `DB_VERB_FILEOPS_ALL`

  Display additional information when performing all filesystem operations, including read and write. May not be available on all platforms.

- `DB_VERB_RECOVERY`

Display additional information when performing recovery.

- DB_VERB_REGISTER

Display additional information concerning support for the DB_REGISTER flag to the DbEnv::open() method.

- DB_VERB_REPLICATION

Display all detailed information about replication. This includes the information displayed by all of the other DB_VERB_REP_* and DB_VERB_REPMGR_* values.

- DB_VERB_REP_ELECT

Display detailed information about replication elections.

- DB_VERB_REP_LEASE

Display detailed information about replication master leases.

- DB_VERB_REP_MISC

Display detailed information about general replication processing not covered by the other DB_VERB_REP_* values.

- DB_VERB_REP_MSGS

Display detailed information about replication message processing.

- DB_VERB_REP_SYNC

Display detailed information about replication client synchronization.

- DB_VERB_REPMGR_CONNFAIL

Display detailed information about Replication Manager connection failures.

- DB_VERB_REPMGR_MISC

Display detailed information about general Replication Manager processing.

- DB_VERB_WAITSFOR

Display the waits-for table when doing deadlock detection.

Errors

The DbEnv::set_verbose() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

[DbEnv](#)

See Also

[Database Environments and Related Methods](#)

DbEnv::stat_print()

```
#include <db_cxx.h>

int
DbEnv::stat_print(u_int32_t flags);
```

The `DbEnv::stat_print()` method displays the default statistical information. The information is printed to a specified output channel (see the DbEnv::set_msgfile() method for more information), or passed to an application callback function (see the DbEnv::set_msgcall() method for more information).

The `DbEnv::stat_print()` method may not be called before the DbEnv::open() method is called.

The `DbEnv::stat_print()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

- `DB_STAT_ALL`

  Display all available information.

- `DB_STAT_CLEAR`

  Reset statistics after displaying their values.

- `DB_STAT_SUBSYSTEM`

  Display information for all configured subsystems.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::strerror()

```
#include <db_cxx.h>

static char *
DbEnv::strerror(int error);
```

The `DbEnv::strerror()` method returns an error message string corresponding to the error number **error** parameter.

This function is a superset of the ANSI C X3.159-1989 (ANSI C) **strerror**(3) function. If the error number **error** is greater than or equal to 0, then the string returned by the system function **strerror**(3) is returned. If the error number is less than 0, an error string appropriate to the corresponding Berkeley DB library error is returned. See Error returns to applications for more information.

Parameters

error

The **error** parameter is the error number for which an error message string is wanted.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::version()

```
#include <db_cxx.h>

static char *
DbEnv::version(int *major, int *minor, int *patch);
```

The `DbEnv::version()` method returns a pointer to a string, suitable for display, containing Berkeley DB version information.

Parameters

major

If **major** is non-NULL, the major version of the Berkeley DB release is copied to the memory to which it refers.

minor

If **minor** is non-NULL, the minor version of the Berkeley DB release is copied to the memory to which it refers.

patch

If **patch** is non-NULL, the patch version of the Berkeley DB release is copied to the memory to which it refers.

Class

DbEnv

See Also

Database Environments and Related Methods

Chapter 6. The DbException Class

```
#include <db_cxx.h>
class DbException {
public:
    int get_errno() const;
    virtual const char *what() const;
    DbEnv *get_env() const;
};
```

This information describes the `DbException` class and how it is used by the various Berkeley DB classes.

Most methods in the Berkeley DB classes return an `int`, but also throw an exception. This allows for two different error behaviors. By default, the Berkeley DB C++ API is configured to throw an exception whenever a serious error occurs. This generally allows for cleaner logic for transaction processing because a try block can surround a single transaction. Alternatively, Berkeley DB can be configured to not throw exceptions, and instead have the individual function return an error code, by setting the DB_CXX_NO_EXCEPTIONS for the Db and DbEnv constructors.

A `DbException` object contains an informational string, an `errno`, and a reference to the environment from which the exception was thrown. The `errno` can be obtained by using the `DbException::get_errno()` method, and can be used, in standard cases, to determine the type of the exception. The informational string can be obtained by using the `DbException::what()`. And, the environment can be obtained using the `DbException::get_env()` method.

We expect in the future that this class will inherit from the standard class exception, but certain language implementation bugs currently prevent this on some platforms.

Some methods may return non-zero values without issuing an exception. This occurs in situations that are not normally considered an error, but when some informational status is returned. For example, the Db::get() method returns DB_NOTFOUND when a requested key does not appear in the database.

DB C++ Exceptions

DB C++ Exceptions	Description
DbDeadlockException	Exception class for deadlocks
DbLockNotGrantedException	Exception class for lock request failures
DbMemoryException	Exception class for insufficient memory
DbRepHandleDeadException	Exception class for database and cursor handles that are invalidated in a replicated application.
DbRunRecoveryException	Exception class for failures requiring recovery

DbDeadlockException

```
#include <db_cxx.h>

class DbDeadlockException : public DbException { ... };
```

This information describes the DbDeadlockException class and how it is used by the various Berkeley DB classes.

A DbDeadlockException is thrown when multiple threads competing for a lock are deadlocked, when a lock request has timed out (and DB_TIME_NOTGRANTED has not been set in the environment), or when a lock request would need to block and the transaction has been configured to not wait for locks. One of the threads' transactions is selected for termination, and a DbDeadlockException is thrown to that thread.

The DbException errno value is set to DB_LOCK_DEADLOCK.

DbLockNotGrantedException

```
#include <db_cxx.h>

class DbLockNotGrantedException : public DbException {
public:
    db_lockop_t get_op() const;
    db_lockmode_t get_mode() const;
    const Dbt* get_obj() const;
    DbLock *get_lock() const;
    int get_index() const;
};
```

This information describes the `DbLockNotGrantedException` class and how it is used by the various Berkeley DB classes.

A `DbLockNotGrantedException` is thrown when lock or transaction timeouts have been configured, a database operation has timed out, and the DB_TIME_NOTGRANTED configuration flag has been specified.

Additionally `DbLockNotGrantedException` is thrown when a Berkeley DB Concurrent Data Store database environment configured for lock timeouts was unable to grant a lock in the allowed time.

Finally, `DbLockNotGrantedException` is thrown when a lock requested using the DbEnv::lock_get() or DbEnv::lock_vec() methods, where the DB_LOCK_NOWAIT flag or lock timers were configured, could not be granted before the wait-time expired.

The DbException errno value is set to `DB_LOCK_NOTGRANTED`.

The following getter methods are available on this class:

- `get_op()`

    Returns `DB_LOCK_GET` when DbEnv::lock_get() was called, and returns the **op** for the failed `DB_LOCKREQ` when DbEnv::lock_vec() was called. If this exception is raised due to a database operation, `DB_LOCK_GET` is returned.

- `get_mode()`

    Returns the **mode** parameter when DbEnv::lock_get() was called, and returns the **mode** for the failed `DB_LOCKREQ` when DbEnv::lock_vec() was called. If this exception is raised due to a database operation, `DB_LOCK_NG` is returned.

- `get_obj()`

    Returns the **object** parameter when DbEnv::lock_get() was called, and returns the **object** for the failed `DB_LOCKREQ` when DbEnv::lock_vec() was called. The Dbt pointer may or may not refer valid memory, depending on whether the Dbt used in the call to the failed DbEnv::lock_get() or DbEnv::lock_vec() method is still in scope and has not been deleted.

- `get_lock()`

Returns NULL when DbEnv::lock_get() was called, and returns the **lock** in the failed DB_LOCKREQ when DbEnv::lock_vec() was called. If this exception is raised due to a database operation, NULL is returned.

- get_index()

  Returns -1 when DbEnv::lock_get() was called, and returns the index of the failed DB_LOCKREQ when DbEnv::lock_vec() was called. If this exception is raised due to a database operation, 0 is returned.

DbMemoryException

```
#include <db_cxx.h>

class DbMemoryException : public DbException {
public:
    Dbt *get_dbt() const;
};
```

This information describes the `DbMemoryException` class and how it is used by the various Berkeley DB classes.

A `DbMemoryException` is thrown when there is insufficient memory to complete an operation, and there is the possibility of recovering. An example is during a Db::get() or Dbc::get() operation with the Dbt flags set to DB_DBT_USERMEM.

The DbException errno value is set to `DB_BUFFER_SMALL` or `ENOMEM`.

The `get_dbt()` method returns the Dbt with insufficient memory to complete the operation, causing the `DbMemoryException` to be thrown. The Dbt pointer may or may not refer to valid memory, depending on whether the Dbt used in the call to the failed Berkeley DB method is still in scope and has not been deleted.

DbRepHandleDeadException

```
#include <db_cxx.h>

class DbRepHandleDeadException : public DbException {
};
```

This information describes the DbRepHandleDead class and how it is used by the various Berkeley DB classes.

A DbRepHandleDeadException is seen only for replicated applications. When a client synchronizes with the master, it is possible for committed transactions to be rolled back. This invalidates all the database and cursor handles opened in the replication environment.

This exception is therefore thrown when the application attempts to access a database or cursor handle that has been invalidated due to a transaction roll back.

When this exception is seen, the application must abandon the attempted operation, discard the handle, and then open a new one before proceeding with the abandoned operation.

DbRunRecoveryException

```
#include <db_cxx.h>

class DbRunRecoveryException : public DbException { ... };
```

This information describes the `DbRunRecoveryException` class and how it is used by the various Berkeley DB classes.

Errors can occur in the Berkeley DB library where the only solution is to shut down the application and run recovery (for example, if Berkeley DB is unable to allocate heap memory). When a fatal error occurs in Berkeley DB, methods will throw a `DbRunRecoveryException`, at which point all subsequent Berkeley DB calls will also fail in the same way. When this occurs, recovery should be performed.

The DbException errno value is set to `DB_RUNRECOVERY`.

Chapter 7. The DbLock Handle

```
#include <db_cxx.h>

class DbLock {
public:
        DbLock();
        DbLock(const DbLock &);
        DbLock &operator = (const DbLock &);
        ~DbLock();
};
```

The locking interfaces for the Berkeley DB database environment are methods of the DbEnv handle. The DbLock object is the handle for a single lock, and has no methods of its own.

Locking Subsystem and Related Methods

Locking Subsystem and Related Methods	Description
DbDeadlockException	Exception class for deadlocks
DbLockNotGrantedException	Exception class for lock request failures
DbEnv::lock_detect()	Perform deadlock detection
DbEnv::lock_get()	Acquire a lock
DbEnv::lock_id()	Acquire a locker ID
DbEnv::lock_id_free()	Release a locker ID
DbEnv::lock_put()	Release a lock
DbEnv::lock_stat()	Return lock subsystem statistics
DbEnv::lock_stat_print()	Print lock subsystem statistics
DbEnv::lock_vec()	Acquire/release locks
DbEnv::cdsgroup_begin()	Get a locker ID in Berkeley DB Concurrent Data Store
Locking Subsystem Configuration	
DbEnv::set_timeout(), DbEnv::get_timeout()	Set/get lock and transaction timeout
DbEnv::set_lk_conflicts(), DbEnv::get_lk_conflicts()	Set/get lock conflicts matrix
DbEnv::set_lk_detect(), DbEnv::get_lk_detect()	Set/get automatic deadlock detection
DbEnv::set_lk_max_lockers(), DbEnv::get_lk_max_lockers()	Set/get maximum number of lockers
DbEnv::set_lk_max_locks(), DbEnv::get_lk_max_locks()	Set/get maximum number of locks
DbEnv::set_lk_max_objects(), DbEnv::get_lk_max_objects()	Set/get maximum number of lock objects
DbEnv::set_lk_partitions(), DbEnv::get_lk_partitions()	Set/get number of lock partitions

DbEnv::get_lk_conflicts()

```
#include <db_cxx.h>

int
DbEnv::get_lk_conflicts(const u_int8_t **lk_conflictsp, int *lk_modesp);
```

The `DbEnv::get_lk_conflicts()` method returns the current conflicts array. You can specify a conflicts array using DbEnv::set_lk_conflicts()

The `DbEnv::get_lk_conflicts()` method may be called at any time during the life of the application.

The `DbEnv::get_lk_conflicts()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

lk_conflictsp

The **lk_conflictsp** parameter references memory into which a pointer to the current conflicts array is copied.

lk_modesp

The **lk_modesp** parameter references memory into which the size of the current conflicts array is copied.

Errors

The `DbEnv::get_lk_conflicts()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

The method was called on an environment which had been opened without being configured for locking.

Class

DbEnv, DbLock

See Also

Locking Subsystem and Related Methods, DbEnv::set_lk_conflicts()

DbEnv::get_lk_detect()

```
#include <db_cxx.h>

int
DbEnv::get_lk_detect(u_int32_t *lk_detectp);
```

The `DbEnv::get_lk_detect()` method returns the deadlock detector configuration. You can manage this using the DbEnv::set_lk_detect() method.

The `DbEnv::get_lk_detect()` method may be called at any time during the life of the application.

The `DbEnv::get_lk_detect()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

lk_detectp

The `DbEnv::get_lk_detect()` method returns the deadlock detector configuration in **lk_detectp**.

Errors

The `DbEnv::get_lk_detect()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

The method was called on an environment which had been opened without being configured for locking.

Class

DbEnv, DbLock

See Also

Locking Subsystem and Related Methods, DbEnv::set_lk_detect()

DbEnv::get_lk_max_lockers()

```
#include <db_cxx.h>

int
DbEnv::get_lk_max_lockers(u_int32_t *, lk_maxp);
```

The `DbEnv::get_lk_max_lockers()` method returns the maximum number of potential lockers. You can configure this using the DbEnv::set_lk_max_lockers() method.

The `DbEnv::get_lk_max_lockers()` method may be called at any time during the life of the application.

The `DbEnv::get_lk_max_lockers()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

lk_maxp

The `DbEnv::get_lk_max_lockers()` method returns the maximum number of lockers in **lk_maxp**.

Errors

The `DbEnv::get_lk_max_lockers()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

The method was called on an environment which had been opened without being configured for locking.

Class

DbEnv, DbLock

See Also

Locking Subsystem and Related Methods, DbEnv::set_lk_max_lockers()

DbEnv::get_lk_max_locks()

```
#include <db_cxx.h>

int
DbEnv::get_lk_max_locks(u_int32_t *lk_maxp);
```

The `DbEnv::get_lk_max_locks()` method returns the maximum number of potential locks. You can configure this using the DbEnv::set_lk_max_locks() method.

The `DbEnv::get_lk_max_locks()` method may be called at any time during the life of the application.

The `DbEnv::get_lk_max_locks()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

lk_maxp

The `DbEnv::get_lk_max_locks()` method returns the maximum number of locks in **lk_maxp**.

Errors

The `DbEnv::get_lk_max_locks()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

The method was called on an environment which had been opened without being configured for locking.

Class

DbEnv, DbLock

See Also

Locking Subsystem and Related Methods, DbEnv::set_lk_max_locks()

DbEnv::get_lk_max_objects()

```
#include <db_cxx.h>

int
DbEnv::get_lk_max_objects(u_int32_t *lk_maxp);
```

The `DbEnv::get_lk_max_objects()` method returns the maximum number of locked objects. You can configure this using the DbEnv::set_lk_max_objects() method.

The `DbEnv::get_lk_max_objects()` method may be called at any time during the life of the application.

The `DbEnv::get_lk_max_objects()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

lk_maxp

The `DbEnv::get_lk_max_objects()` method returns the maximum number of potentially locked objects in **lk_maxp**.

Errors

The `DbEnv::get_lk_max_objects()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

The method was called on an environment which had been opened without being configured for locking.

Class

DbEnv, DbLock

See Also

Locking Subsystem and Related Methods, DbEnv::set_lk_max_objects()

DbEnv::get_lk_partitions()

```
#include <db_cxx.h>

int
DbEnv::get_lk_partitions(u_int32_t *lk_partitions);
```

The `DbEnv::get_lk_partitions()` method returns the number of lock table partitions used in the Berkeley DB environment. You can configure this using the DbEnv::set_lk_partitions() method.

The `DbEnv::get_lk_partitions()` method may be called at any time during the life of the application.

The `DbEnv::get_lk_partitions()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

lk_partitions

The `DbEnv::get_lk_partitions()` method returns the number of partitions in **lk_partitions**.

Errors

The `DbEnv::get_lk_partitions()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

The method was called on an environment which had been opened without being configured for locking.

Class

DbEnv, DbLock

See Also

Locking Subsystem and Related Methods, DbEnv::set_lk_partitions()

DbEnv::set_lk_conflicts()

```
#include <db_cxx.h>

int
DbEnv::set_lk_conflicts(u_int8_t *conflicts, int nmodes);
```

Set the locking conflicts matrix.

If `DbEnv::set_lk_conflicts()` is never called, a standard conflicts array is used; see Standard Lock Modes for more information.

The `DbEnv::set_lk_conflicts()` method configures a database environment, not only operations performed using the specified DbEnv handle.

The `DbEnv::set_lk_conflicts()` method may not be called after the DbEnv::open() method is called. If the database environment already exists when DbEnv::open() is called, the information specified to `DbEnv::set_lk_conflicts()` will be ignored.

The `DbEnv::set_lk_conflicts()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

conflicts

The **conflicts** parameter is the new locking conflicts matrix. The **conflicts** parameter is an **nmodes** by **nmodes** array. A non-0 value for the array element indicates that requested_mode and held_mode conflict:

```
        conflicts[requested_mode][held_mode]
```

The *not-granted* mode must be represented by 0.

nmodes

The **nmodes** parameter is the size of the lock conflicts matrix.

Errors

The `DbEnv::set_lk_conflicts()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after DbEnv::open() was called; or if an invalid flag value or parameter was specified.

ENOMEM

The conflicts array could not be copied.

Class

DbEnv, DbLock

See Also

Locking Subsystem and Related Methods

DbEnv::set_lk_detect()

```
#include <db_cxx.h>

int
DbEnv::set_lk_detect(u_int32_t detect);
```

Set if the deadlock detector is to be run whenever a lock conflict occurs, and specify what lock request(s) should be rejected. As transactions acquire locks on behalf of a single locker ID, rejecting a lock request associated with a transaction normally requires the transaction be aborted.

The database environment's deadlock detector configuration may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_lk_detect", one or more whitespace characters, and the method **detect** parameter as a string; for example, "set_lk_detect DB_LOCK_OLDEST". Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The `DbEnv::set_lk_detect()` method configures a database environment, not only operations performed using the specified DbEnv handle.

The `DbEnv::set_lk_detect()` method may be called at any time during the life of the application.

The `DbEnv::set_lk_detect()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

detect

The **detect** parameter configures the deadlock detector. The specified value must be one of the following list:

- DB_LOCK_DEFAULT

  Use whatever lock policy was specified when the database environment was created. If no lock policy has yet been specified, set the lock policy to DB_LOCK_RANDOM.

- DB_LOCK_EXPIRE

  Reject lock requests which have timed out. No other deadlock detection is performed.

- DB_LOCK_MAXLOCKS

  Reject the lock request for the locker ID with the most locks.

- DB_LOCK_MAXWRITE

  Reject the lock request for the locker ID with the most write locks.

- DB_LOCK_MINLOCKS

Reject the lock request for the locker ID with the fewest locks.

- DB_LOCK_MINWRITE

  Reject the lock request for the locker ID with the fewest write locks.

- DB_LOCK_OLDEST

  Reject the lock request for the locker ID with the oldest lock.

- DB_LOCK_RANDOM

  Reject the lock request for a random locker ID.

- DB_LOCK_YOUNGEST

  Reject the lock request for the locker ID with the youngest lock.

Errors

The DbEnv::set_lk_detect() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv, DbLock

See Also

Locking Subsystem and Related Methods

DbEnv::set_lk_max_lockers()

```
#include <db_cxx.h>

int
DbEnv::set_lk_max_lockers(u_int32_t max);
```

Set the maximum number of locking entities supported by the Berkeley DB environment. This value is used by DbEnv::open() to estimate how much space to allocate for various lock-table data structures. The default value is 1000 lockers. For specific information on configuring the size of the lock subsystem, see Configuring locking: sizing the system.

The database environment's maximum number of lockers may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_lk_max_lockers", one or more whitespace characters, and the number of lockers. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The `DbEnv::set_lk_max_lockers()` method configures a database environment, not only operations performed using the specified DbEnv handle.

The `DbEnv::set_lk_max_lockers()` method may not be called after the DbEnv::open() method is called. If the database environment already exists when DbEnv::open() is called, the information specified to `DbEnv::set_lk_max_lockers()` will be ignored.

The `DbEnv::set_lk_max_lockers()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

max

The **max** parameter is the maximum number simultaneous locking entities supported by the Berkeley DB environment.

Errors

The `DbEnv::set_lk_max_lockers()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after DbEnv::open() was called; or if an invalid flag value or parameter was specified.

Class

DbEnv, DbLock

See Also

Locking Subsystem and Related Methods

DbEnv::set_lk_max_locks()

```
#include <db_cxx.h>

int
DbEnv::set_lk_max_locks(u_int32_t max);
```

Set the maximum number of locks supported by the Berkeley DB environment. This value is used by DbEnv::open() to estimate how much space to allocate for various lock-table data structures. The default value is 1000 locks. For specific information on configuring the size of the lock subsystem, see Configuring locking: sizing the system.

The database environment's maximum number of locks may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_lk_max_locks", one or more whitespace characters, and the number of locks. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The DbEnv::set_lk_max_locks() method configures a database environment, not only operations performed using the specified DbEnv handle.

The DbEnv::set_lk_max_locks() method may not be called after the DbEnv::open() method is called. If the database environment already exists when DbEnv::open() is called, the information specified to DbEnv::set_lk_max_locks() will be ignored.

The DbEnv::set_lk_max_locks() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

max

The **max** parameter is the maximum number of locks supported by the Berkeley DB environment.

Errors

The DbEnv::set_lk_max_locks() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after DbEnv::open() was called; or if an invalid flag value or parameter was specified.

Class

DbEnv, DbLock

See Also

Locking Subsystem and Related Methods

DbEnv::set_lk_max_objects()

```
#include <db_cxx.h>

int
DbEnv::set_lk_max_objects(u_int32_t max);
```

Set the maximum number of locked objects supported by the Berkeley DB environment. This value is used by DbEnv::open() to estimate how much space to allocate for various lock-table data structures. The default value is 1000 objects. For specific information on configuring the size of the lock subsystem, see Configuring locking: sizing the system.

The database environment's maximum number of objects may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_lk_max_objects", one or more whitespace characters, and the number of objects. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The `DbEnv::set_lk_max_objects()` method configures a database environment, not only operations performed using the specified DbEnv handle.

The `DbEnv::set_lk_max_objects()` method may not be called after the DbEnv::open() method is called. If the database environment already exists when DbEnv::open() is called, the information specified to `DbEnv::set_lk_max_objects()` will be ignored.

The `DbEnv::set_lk_max_objects()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

max

The **max** parameter is the maximum number of locked objects supported by the Berkeley DB environment.

Errors

The `DbEnv::set_lk_max_objects()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after DbEnv::open() was called; or if an invalid flag value or parameter was specified.

Class

DbEnv, DbLock

See Also

Locking Subsystem and Related Methods

DbEnv::set_lk_partitions()

```
#include <db_cxx.h>

int
DbEnv::set_lk_partitions(u_int32_t partitions);
```

Set the number of lock table partitions in the Berkeley DB environment. The default value is 10 times the number of CPUs on the system if there is more than one CPU. Increasing the number of partitions can provide for greater throughput on a system with multiple CPUs and more than one thread contending for the lock manager. On single processor systems more than one partition may increase the overhead of the lock manager. Systems often report threading contexts as CPUs. If your system does this, set the number of partitions to 1 to get optimal performance.

The database environment's number of partitions may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_lk_partitions", one or more whitespace characters, and the number of partitions. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The `DbEnv::set_lk_partitions()` method configures a database environment, not only operations performed using the specified DbEnv handle.

The `DbEnv::set_lk_partitions()` method may not be called after the DbEnv::open() method is called. If the database environment already exists when DbEnv::open() is called, the information specified to `DbEnv::set_lk_partitions()` will be ignored.

The `DbEnv::set_lk_partitions()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

partitions

The **partitions** parameter is the number of partitions to be configured in the Berkeley DB environment.

Errors

The `DbEnv::set_lk_partitions()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after DbEnv::open() was called; or if an invalid flag value or parameter was specified.

Class

DbEnv, DbLock

See Also

Locking Subsystem and Related Methods

DbEnv::lock_detect()

```
#include <db_cxx.h>

int
DbEnv::lock_detect(u_int32_t flags, u_int32_t atype, int *rejected);
```

The `DbEnv::lock_detect()` method runs one iteration of the deadlock detector. The deadlock detector traverses the lock table and marks one of the participating lock requesters for rejection in each deadlock it finds.

The `DbEnv::lock_detect()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

rejected

If the **rejected** parameter is non-NULL, the memory location to which it refers will be set to the number of lock requests that were rejected.

atype

The **atype** parameter specifies which lock request(s) to reject. It must be set to one of the following list:

- `DB_LOCK_DEFAULT`

  Use whatever lock policy was specified when the database environment was created. If no lock policy has yet been specified, set the lock policy to DB_LOCK_RANDOM.

- `DB_LOCK_EXPIRE`

  Reject lock requests which have timed out. No other deadlock detection is performed.

- `DB_LOCK_MAXLOCKS`

  Reject the lock request for the locker ID with the most locks.

- `DB_LOCK_MAXWRITE`

  Reject the lock request for the locker ID with the most write locks.

- `DB_LOCK_MINLOCKS`

  Reject the lock request for the locker ID with the fewest locks.

- `DB_LOCK_MINWRITE`

  Reject the lock request for the locker ID with the fewest write locks.

- `DB_LOCK_OLDEST`

Reject the lock request for the locker ID with the oldest lock.

- DB_LOCK_RANDOM

  Reject the lock request for a random locker ID.

- DB_LOCK_YOUNGEST

  Reject the lock request for the locker ID with the youngest lock.

flags

The **flags** parameter is currently unused, and must be set to 0.

Errors

The `DbEnv::lock_detect()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv, DbLock

See Also

Locking Subsystem and Related Methods

DbEnv::lock_get()

```
#include <db_cxx.h>

int
DbEnv::lock_get(u_int32_t locker, u_int32_t flags,
    const Dbt *object, const db_lockmode_t lock_mode, DbLock *lock);
```

The `DbEnv::lock_get()` method acquires a lock from the lock table, returning information about it in the **lock** parameter.

The `DbEnv::lock_get()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

locker

The **locker** parameter is an unsigned 32-bit integer quantity. It represents the entity requesting the lock.

flags

The **flags** parameter must be set to 0 or the following value:

• DB_LOCK_NOWAIT

   If a lock cannot be granted because the requested lock conflicts with an existing lock, return DB_LOCK_NOTGRANTED immediately instead of waiting for the lock to become available.

object

The **object** parameter is an untyped byte string that specifies the object to be locked. Applications using the locking subsystem directly while also doing locking via the Berkeley DB access methods must take care not to inadvertently lock objects that happen to be equal to the unique file IDs used to lock files. See Access method locking conventions in the *Berkeley DB Programmer's Reference Guide* for more information.

lock_mode

The **lock_mode** parameter is used as an index into the environment's lock conflict matrix. When using the default lock conflict matrix, **lock_mode** must be set to one of the following values:

• DB_LOCK_READ

   read (shared)

• DB_LOCK_WRITE

   write (exclusive)

- DB_LOCK_IWRITE

  intention to write (shared)

- DB_LOCK_IREAD

  intention to read (shared)

- DB_LOCK_IWR

  intention to read and write (shared)

See DbEnv::set_lk_conflicts() and Standard Lock Modes for more information on the lock conflict matrix.

lock

The DbEnv::lock_get() method returns the lock information in **lock**.

Errors

The DbEnv::lock_get() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DbDeadlockException or DB_LOCK_DEADLOCK

A transactional database environment operation was selected to resolve a deadlock.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_LOCK_DEADLOCK is returned.

DbLockNotGrantedException or DB_LOCK_NOTGRANTED

A Berkeley DB Concurrent Data Store database environment configured for lock timeouts was unable to grant a lock in the allowed time.

DbLockNotGrantedException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_LOCK_NOTGRANTED is returned.

DbLockNotGrantedException or DB_LOCK_NOTGRANTED

The DB_LOCK_NOWAIT flag or lock timers were configured and the lock could not be granted before the wait-time expired.

DbLockNotGrantedException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_LOCK_NOTGRANTED is returned.

EINVAL

An invalid flag value or parameter was specified.

EINVAL

The method was called on an environment which had been opened without being configured for locking.

ENOMEM

The maximum number of locks has been reached.

Class

DbEnv, DbLock

See Also

Locking Subsystem and Related Methods

DbEnv::lock_id()

```
#include <db_cxx.h>

int
DbEnv::lock_id(u_int32_t *idp);
```

The `DbEnv::lock_id()` method copies a locker ID, which is guaranteed to be unique in the environment's lock table, into the memory location to which **idp** refers.

The DbEnv::lock_id_free() method should be called to return the locker ID to the Berkeley DB library when it is no longer needed.

The `DbEnv::lock_id()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

idp

The **idp** parameter references memory into which the allocated locker ID is copied.

Class

DbEnv, DbLock

See Also

Locking Subsystem and Related Methods

DbEnv::lock_id_free()

```
#include <db_cxx.h>

int
DbEnv::lock_id_free(u_int32_t id);
```

The `DbEnv::lock_id_free()` method frees a locker ID allocated by the DbEnv::lock_id() method.

The `DbEnv::lock_id_free()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

id

The **id** parameter is the locker id to be freed.

Errors

The `DbEnv::lock_id_free()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the locker ID is invalid or locks are still held by this locker ID; or if an invalid flag value or parameter was specified.

Class

DbEnv, DbLock

See Also

Locking Subsystem and Related Methods

DbEnv::lock_put()

```
#include <db_cxx.h>

int
DbEnv::lock_put(DbLock *lock);
```

The `DbEnv::lock_put()` method releases **lock**.

The `DbEnv::lock_put()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

lock

The **lock** parameter is the lock to be released.

Errors

The `DbEnv::lock_put()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv, DbLock

See Also

Locking Subsystem and Related Methods

DbEnv::lock_stat()

```
#include <db_cxx.h>

int
DbEnv::lock_stat(DB_LOCK_STAT **statp, u_int32_t flags);
```

The `DbEnv::lock_stat()` method returns the locking subsystem statistics.

The `DbEnv::lock_stat()` method creates a statistical structure of type DB_LOCK_STAT and copies a pointer to it into a user-specified memory location.

Statistical structures are stored in allocated memory. If application-specific allocation routines have been declared (see DbEnv::set_alloc() for more information), they are used to allocate the memory; otherwise, the standard C library **malloc**(3) is used. The caller is responsible for deallocating the memory. To deallocate the memory, free the memory reference; references inside the returned memory need not be individually freed.

The following DB_LOCK_STAT fields will be filled in:

- **u_int32_t st_id;**

  The last allocated locker ID.

- **u_int32_t st_cur_maxid;**

  The current maximum unused locker ID.

- **int st_nmodes;**

  The number of lock modes.

- **u_int32_t st_maxlocks;**

  The maximum number of locks possible.

- **u_int32_t st_maxlockers;**

  The maximum number of lockers possible.

- **u_int32_t st_maxobjects;**

  The maximum number of lock objects possible.

- **u_int32_t st_partitions;**

  The number of lock table partitions.

- **u_int32_t st_nlocks;**

  The number of current locks.

- **u_int32_t st_maxnlocks;**

  The maximum number of locks at any one time. Note that if there is more than one partition, this is the sum of the maximum across all partitions.

- **u_int32_t st_maxhlocks;**

  The maximum number of locks in any hash bucket at any one time.

- **uintmax_t st_locksteals;**

  The maximum number of locks stolen by an empty partition.

- **uintmax_t st_maxlsteals;**

  The maximum number of lock steals for any one partition.

- **u_int32_t st_nlockers;**

  The number of current lockers.

- **u_int32_t st_maxnlockers;**

  The maximum number of lockers at any one time.

- **u_int32_t st_nobjects;**

  The number of current lock objects.

- **u_int32_t st_maxnobjects;**

  The maximum number of lock objects at any one time. Note that if there is more than one partition this is the sum of the maximum across all partitions.

- **u_int32_t st_maxhobjects;**

  The maximum number of objects in any hash bucket at any one time.

- **uintmax_t st_objectsteals;**

  The maximum number of objects stolen by an empty partition.

- **uintmax_t st_maxosteals;**

  The maximum number of object steals for any one partition.

- **uintmax_t st_nrequests;**

  The total number of locks requested.

- **uintmax_t st_nreleases;**

  The total number of locks released.

- **uintmax_t st_nupgrade;**

  The total number of locks upgraded.

- **uintmax_t st_ndowngrade;**

  The total number of locks downgraded.

- **uintmax_t st_lock_wait;**

  The number of lock requests not immediately available due to conflicts, for which the thread of control waited.

- **uintmax_t st_lock_nowait;**

  The number of lock requests not immediately available due to conflicts, for which the thread of control did not wait.

- **uintmax_t st_ndeadlocks;**

  The number of deadlocks.

- **db_timeout_t st_locktimeout;**

  Lock timeout value.

- **uintmax_t st_nlocktimeouts;**

  The number of lock requests that have timed out.

- **u_int32_t st_txntimeout;**

  Transaction timeout value.

- **uintmax_t st_ntxntimeouts;**

  The number of transactions that have timed out. This value is also a component of **st_ndeadlocks**, the total number of deadlocks detected.

- **uintmax_t st_objs_wait;**

  The number of requests to allocate or deallocate an object for which the thread of control waited.

- **uintmax_t st_objs_nowait;**

  The number of requests to allocate or deallocate an object for which the thread of control did not wait.

- **uintmax_t st_lockers_wait;**

  The number of requests to allocate or deallocate a locker for which the thread of control waited.

- **uintmax_t st_lockers_nowait;**

The number of requests to allocate or deallocate a locker for which the thread of control did not wait.

- **u_int32_t st_hash_len;**

Maximum length of a lock hash bucket.

- **roff_t st_regsize;**

The size of the lock region, in bytes.

- **uintmax_t st_part_wait;**

The number of times that a thread of control was forced to wait before obtaining the lock partition mutex.

- **uintmax_t st_part_nowait;**

The number of times that a thread of control was able to obtain the lock partition mutex without waiting.

- **uintmax_t st_part_max_wait;**

The maximum number of times that a thread of control was forced to wait before obtaining any one lock partition mutex.

- **uintmax_t st_part_max_nowait;**

The number of times that a thread of control was able to obtain any one lock partition mutex without waiting.

- **uintmax_t st_region_wait;**

The number of times that a thread of control was forced to wait before obtaining the lock region mutex.

- **uintmax_t st_region_nowait;**

The number of times that a thread of control was able to obtain the lock region mutex without waiting.

The `DbEnv::lock_stat()` method may not be called before the DbEnv::open() method is called.

The `DbEnv::lock_stat()` method returns a non-zero error value on failure and 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or the following value:

- `DB_STAT_CLEAR`

Reset statistics after returning their values.

statp

The **statp** parameter references memory into which a pointer to the allocated statistics structure is copied.

Errors

The `DbEnv::lock_stat()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv, DbLock

See Also

Locking Subsystem and Related Methods

DbEnv::lock_stat_print()

```
#include <db_cxx.h>

int
DbEnv::lock_stat_print(u_int32_t flags);
```

The `DbEnv::lock_stat_print()` method displays the locking subsystem statistical information, as described for the `DbEnv::lock_stat()` method. The information is printed to a specified output channel (see the DbEnv::set_msgfile() method for more information), or passed to an application callback function (see the DbEnv::set_msgcall() method for more information).

The `DbEnv::lock_stat_print()` method may not be called before the DbEnv::open() method is called.

The `DbEnv::lock_stat_print()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

- `DB_STAT_ALL`

  Display all available information.

- `DB_STAT_CLEAR`

  Reset statistics after displaying their values.

- `DB_STAT_LOCK_CONF`

  Display the lock conflict matrix.

- `DB_STAT_LOCK_LOCKERS`

  Display the lockers within hash chains.

- `DB_STAT_LOCK_OBJECTS`

  Display the lock objects within hash chains.

- `DB_STAT_LOCK_PARAMS`

  Display the locking subsystem parameters.

Class

DbEnv, DbLock

See Also

Locking Subsystem and Related Methods

DbEnv::lock_vec()

```
#include <db_cxx.h>

int
DbEnv::lock_vec(u_int32_t locker, u_int32_t flags,
    DB_LOCKREQ list[], int nlist, DB_LOCKREQ **elistp);
```

The `DbEnv::lock_vec()` method atomically obtains and releases one or more locks from the lock table. The `DbEnv::lock_vec()` method is intended to support acquisition or trading of multiple locks under one lock table semaphore, as is needed for lock coupling or in multigranularity locking for lock escalation.

If any of the requested locks cannot be acquired, or any of the locks to be released cannot be released, the operations before the failing operation are guaranteed to have completed successfully, and `DbEnv::lock_vec()` returns a non-zero value. In addition, if **elistp** is not NULL, it is set to point to the DB_LOCKREQ entry that was being processed when the error occurred.

Unless otherwise specified, the `DbEnv::lock_vec()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

elistp

If an error occurs, and the **elistp** parameter is non-NULL, it is set to point to the DB_LOCKREQ entry that was being processed when the error occurred.

flags

The **flags** parameter must be set to 0 or the following value:

- `DB_LOCK_NOWAIT`

  If a lock cannot be granted because the requested lock conflicts with an existing lock, return DB_LOCK_NOTGRANTED immediately instead of waiting for the lock to become available. In this case, if non-NULL, **elistp** identifies the request that was not granted.

locker

The **locker** parameter is an unsigned 32-bit integer quantity. It represents the entity requesting or releasing the lock.

list

The **list** array provided to `DbEnv::lock_vec()` is typedef'd as DB_LOCKREQ.

To ensure compatibility with future releases of Berkeley DB, all fields of the DB_LOCKREQ structure that are not explicitly set should be initialized to 0 before the first time the structure is used. Do this by declaring the structure external or static, or by calling **memset**(3).

A DB_LOCKREQ structure has at least the following fields:

- `lockop_t op;`

  The operation to be performed, which must be set to one of the following values:

  - `DB_LOCK_GET`

    Get the lock defined by the values of the **mode** and **obj** structure fields, for the specified **locker**. Upon return from `DbEnv::lock_vec()`, if the **lock** field is non-NULL, a reference to the acquired lock is stored there. (This reference is invalidated by any call to `DbEnv::lock_vec()` or DbEnv::lock_put() that releases the lock.)

  - `DB_LOCK_GET_TIMEOUT`

    Identical to DB_LOCK_GET except that the value in the **timeout** structure field overrides any previously specified timeout value for this lock. A value of 0 turns off any previously specified timeout.

  - `DB_LOCK_PUT`

    The lock to which the **lock** structure field refers is released. The **locker** parameter, and **mode** and **obj** fields are ignored.

  - `DB_LOCK_PUT_ALL`

    All locks held by the specified **locker** are released. The **lock**, **mode**, and **obj** structure fields are ignored. Locks acquired in operations performed by the current call to `DbEnv::lock_vec()` which appear before the DB_LOCK_PUT_ALL operation are released; those acquired in operations appearing after the DB_LOCK_PUT_ALL operation are not released.

  - `DB_LOCK_PUT_OBJ`

    All locks held on **obj** are released. The **locker** parameter and the **lock** and **mode** structure fields are ignored. Locks acquired in operations performed by the current call to `DbEnv::lock_vec()` that appear before the DB_LOCK_PUT_OBJ operation are released; those acquired in operations appearing after the DB_LOCK_PUT_OBJ operation are not released.

  - `DB_LOCK_TIMEOUT`

    Cause the specified **locker** to timeout immediately. If the database environment has not configured automatic deadlock detection, the transaction will timeout the next time deadlock detection is performed. As transactions acquire locks on behalf of a single locker ID, timing out the locker ID associated with a transaction will time out the transaction itself.

- `DB_LOCK lock;`

  A lock reference.

- `const lockmode_t mode;`

  The lock mode, used as an index into the environment's lock conflict matrix. When using the default lock conflict matrix, **mode** must be set to one of the following values:

- DB_LOCK_READ

  read (shared)

- DB_LOCK_WRITE

  write (exclusive)

- DB_LOCK_IWRITE

  intention to write (shared)

- DB_LOCK_IREAD

  intention to read (shared)

- DB_LOCK_IWR

  intention to read and write (shared)

  See DbEnv::set_lk_conflicts() and Standard Lock Modes for more information on the lock conflict matrix.

- **const DBT obj;**

  An untyped byte string that specifies the object to be locked or released. Applications using the locking subsystem directly while also doing locking via the Berkeley DB access methods must take care not to inadvertently lock objects that happen to be equal to the unique file IDs used to lock files. See Access method locking conventions in the *Berkeley DB Programmer's Reference Guide* for more information.

- **u_int32_t timeout;**

  The lock timeout value.

nlist

The **nlist** parameter specifies the number of elements in the **list** array.

Errors

The DbEnv::lock_vec() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DbDeadlockException or DB_LOCK_DEADLOCK

A transactional database environment operation was selected to resolve a deadlock.

DbDeadlockException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_LOCK_DEADLOCK is returned.

DbLockNotGrantedException or DB_LOCK_NOTGRANTED

A Berkeley DB Concurrent Data Store database environment configured for lock timeouts was unable to grant a lock in the allowed time.

DbLockNotGrantedException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_LOCK_NOTGRANTED is returned.

DbLockNotGrantedException or DB_LOCK_NOTGRANTED

The DB_LOCK_NOWAIT flag or lock timers were configured and the lock could not be granted before the wait-time expired.

DbLockNotGrantedException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, DB_LOCK_NOTGRANTED is returned.

EINVAL

An invalid flag value or parameter was specified.

ENOMEM

The maximum number of locks has been reached.

Class

DbEnv, DbLock

See Also

Locking Subsystem and Related Methods

Chapter 8. The DbLsn Handle

```
#include <db_cxx.h>

class DbLsn : public DB_LSN { ... };
```

The DbLsn object is a *log sequence number* which specifies a unique location in a log file. A DbLsn consists of two unsigned 32-bit integers -- one specifies the log file number, and the other specifies an offset in the log file.

Logging Subsystem and Related Methods

Logging Subsystem and Related Methods	Description
DbEnv::log_archive()	List log and database files
DbEnv::log_file()	Map Log Sequence Numbers to log files
DbEnv::log_flush()	Flush log records
DbEnv::log_printf()	Append informational message to the log
DbEnv::log_put()	Write a log record
DbEnv::log_stat()	Return log subsystem statistics
DbEnv::log_stat_print()	Print log subsystem statistics
DbEnv::log_compare()	Compare two Log Sequence Numbers
Logging Subsystem Cursors	
DbEnv::log_cursor()	Create a log cursor handle
The DbLogc Handle	A log cursor handle
DbLogc::close()	Close a log cursor
DbLogc::get()	Retrieve a log record
Logging Subsystem Configuration	
DbEnv::log_set_config(), DbEnv::log_get_config()	Configure the logging subsystem
DbEnv::set_lg_bsize(), DbEnv::get_lg_bsize()	Set/get log buffer size
DbEnv::set_lg_dir(), DbEnv::get_lg_dir()	Set/get the environment logging directory
DbEnv::set_lg_filemode(), DbEnv::get_lg_filemode()	Set/get log file mode
DbEnv::set_lg_max(), DbEnv::get_lg_max()	Set/get log file size
DbEnv::set_lg_regionmax(), DbEnv::get_lg_regionmax()	Set/get logging region size

DbEnv::get_lg_bsize()

```
#include <db_cxx.h>

int
DbEnv::get_lg_bsize(u_int32_t *lg_bsizep);
```

The `DbEnv::get_lg_bsize()` method returns the size of the log buffer, in bytes. You can manage this value using the DbEnv::set_lg_bsize() method.

The `DbEnv::get_lg_bsize()` method may be called at any time during the life of the application.

The `DbEnv::get_lg_bsize()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

lg_bsizep

The `DbEnv::get_lg_bsize()` method returns the size of the log buffer, in bytes in **lg_bsizep**.

Class

DbEnv, DbLogc, DbLsn

See Also

Logging Subsystem and Related Methods, DbEnv::set_lg_bsize()

DbEnv::get_lg_dir()

```
#include <db_cxx.h>

int
DbEnv::get_lg_dir(const char **dirp);
```

The `DbEnv::get_lg_dir()` method returns the log directory, which is the location for logging files. You can manage this value using the DbEnv::set_lg_dir() method.

The `DbEnv::get_lg_dir()` method may be called at any time during the life of the application.

The `DbEnv::get_lg_dir()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

dirp

The `DbEnv::get_lg_dir()` method returns a reference to the log directory in **dirp**.

Class

DbEnv, DbLogc, DbLsn

See Also

Logging Subsystem and Related Methods, DbEnv::set_lg_dir()

DbEnv::get_lg_filemode()

```
#include <db_cxx.h>

int
DbEnv::get_lg_filemode(int *lg_modep);
```

The `DbEnv::set_lg_filemode()` method returns the log file mode. You can manage this value using the DbEnv::set_lg_filemode() method.

The `DbEnv::set_lg_filemode()` method may be called at any time during the life of the application.

The `DbEnv::set_lg_filemode()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

lg_modep

The `DbEnv::set_lg_filemode()` method returns the log file mode in **lg_modep**.

Class

DbEnv, DbLogc, DbLsn

See Also

Logging Subsystem and Related Methods, DbEnv::set_lg_filemode()

DbEnv::get_lg_max()

```
#include <db_cxx.h>

int
DbEnv::get_lg_max(u_int32_t *lg_maxp);
```

The `DbEnv::get_lg_max()` method returns the maximum log file size. You can manage this value using the DbEnv::set_lg_max() method.

The `DbEnv::get_lg_max()` method may be called at any time during the life of the application.

The `DbEnv::get_lg_max()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

lg_maxp

The `DbEnv::get_lg_max()` method returns the maximum log file size in **lg_maxp**.

Class

DbEnv, DbLogc, DbLsn

See Also

Logging Subsystem and Related Methods, DbEnv::set_lg_max()

DbEnv::get_lg_regionmax()

```
#include <db_cxx.h>

int
DbEnv::get_lg_regionmax(u_int32_t *lg_regionmaxp);
```

The `DbEnv::get_lg_regionmax()` method returns the size of the underlying logging subsystem region. You can manage this value using the DbEnv::set_lg_regionmax() method.

The `DbEnv::get_lg_regionmax()` method may be called at any time during the life of the application.

The `DbEnv::get_lg_regionmax()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

lg_regionmaxp

The `DbEnv::get_lg_regionmax()` method returns the size of the underlying logging subsystem region in **lg_regionmaxp**.

Class

DbEnv, DbLogc, DbLsn

See Also

Logging Subsystem and Related Methods, DbEnv::set_lg_regionmax()

DbEnv::log_archive()

```
#include <db_cxx.h>

int
DbEnv::log_archive(char *(*listp)[], u_int32_t flags);
```

The `DbEnv::log_archive()` method returns an array of log or database filenames.

By default, `DbEnv::log_archive()` returns the names of all of the log files that are no longer in use (for example, that are no longer involved in active transactions), and that may safely be archived for catastrophic recovery and then removed from the system. If there are no filenames to return, the memory location to which **listp** refers will be set to NULL.

Arrays of log filenames are stored in allocated memory. If application-specific allocation routines have been declared (see DbEnv::set_alloc() for more information), they are used to allocate the memory; otherwise, the standard C library **malloc**(3) is used. The caller is responsible for deallocating the memory. To deallocate the memory, free the memory reference; references inside the returned memory need not be individually freed.

Log cursor handles (returned by the DbEnv::log_cursor() method) may have open file descriptors for log files in the database environment. Also, the Berkeley DB interfaces to the database environment logging subsystem (for example, DbEnv::log_put() and DbTxn::abort()) may allocate log cursors and have open file descriptors for log files as well. On operating systems where filesystem related system calls (for example, rename and unlink on Windows/NT) can fail if a process has an open file descriptor for the affected file, attempting to move or remove the log files listed by `DbEnv::log_archive()` may fail. All Berkeley DB internal use of log cursors operates on active log files only and furthermore, is short-lived in nature. So, an application seeing such a failure should be restructured to close any open log cursors it may have, and otherwise to retry the operation until it succeeds. (Although the latter is not likely to be necessary; it is hard to imagine a reason to move or rename a log file in which transactions are being logged or aborted.)

See db_archive for more information on database archival procedures.

The `DbEnv::log_archive()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

- `DB_ARCH_ABS`

  All pathnames are returned as absolute pathnames, instead of relative to the database home directory.

- `DB_ARCH_DATA`

Return the database files that need to be archived in order to recover the database from catastrophic failure. If any of the database files have not been accessed during the lifetime of the current log files, `DbEnv::log_archive()` will not include them in this list. It is also possible that some of the files referred to by the log have since been deleted from the system.

The `DB_ARCH_DATA` and `DB_ARCH_LOG` flags are mutually exclusive.

- `DB_ARCH_LOG`

  Return all the log filenames, regardless of whether or not they are in use.

  The `DB_ARCH_DATA` and `DB_ARCH_LOG flags` are mutually exclusive.

- `DB_ARCH_REMOVE`

  Remove log files that are no longer needed; no filenames are returned. Automatic log file removal is likely to make catastrophic recovery impossible.

  The `DB_ARCH_REMOVE` flag may not be specified with any other flag.

listp

The **listp** parameter references memory into which the allocated array of log or database filenames is copied. If there are no filenames to return, the memory location to which **listp** refers will be set to NULL.

Errors

The `DbEnv::log_archive()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv, DbLogc, DbLsn

See Also

Logging Subsystem and Related Methods

DbEnv::log_cursor()

```
#include <db_cxx.h>

int
DbEnv::log_cursor(DbLogc **cursorp, u_int32_t flags);
```

The `DbEnv::log_cursor()` method returns a created log cursor.

The `DbEnv::log_cursor()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

cursorp

The **cursorp** parameter references memory into which a pointer to the created log cursor is copied.

flags

The **flags** parameter is currently unused, and must be set to 0.

Errors

The `DbEnv::log_cursor()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv, DbLogc, DbLsn

See Also

Logging Subsystem and Related Methods

DbEnv::log_file()

```
#include <db_cxx.h>

int
DbEnv::log_file(const DbLsn *lsn, char *namep, size_t len);
```

The `DbEnv::log_file()` method maps `DbLsn` structures to filenames, returning the name of the file containing the record named by **lsn**.

This mapping of `DbLsn` structures to files is needed for database administration. For example, a transaction manager typically records the earliest DbLsn needed for restart, and the database administrator may want to archive log files to tape when they contain only DbLsn entries before the earliest one needed for restart.

The `DbEnv::log_file()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

lsn

The **lsn** parameter is the `DbLsn` structure for which a filename is wanted.

namep

The **namep** parameter references memory into which the name of the file containing the record named by **lsn** is copied.

len

The **len** parameter is the length of the **namep** buffer in bytes. If **namep** is too short to hold the filename, `DbEnv::log_file()` will fail. (Log filenames are always 14 characters long.)

Errors

The `DbEnv::log_file()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If supplied buffer was too small to hold the log filename; or if an invalid flag value or parameter was specified.

Class

DbEnv, DbLogc, DbLsn

See Also

Logging Subsystem and Related Methods

DbEnv::log_flush()

```
#include <db_cxx.h>

int
DbEnv::log_flush(const DbLsn *lsn);
```

The `DbEnv::log_flush()` method writes log records to disk.

The `DbEnv::log_flush()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

lsn

All log records with DbLsn values less than or equal to the **lsn** parameter are written to disk. If **lsn** is NULL, all records in the log are flushed.

Errors

The `DbEnv::log_flush()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv, DbLogc, DbLsn

See Also

Logging Subsystem and Related Methods

DbEnv::log_get_config()

```
#include <db_cxx.h>

int
DbEnv::log_get_config(u_int32_t which, int *onoffp)
```

The `DbEnv::log_get_config()` method returns whether the specified **which** parameter is currently set or not. You can manage this value using the DbEnv::log_set_config() method.

The `DbEnv::log_get_config()` method may be called at any time during the life of the application.

The `DbEnv::log_get_config()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

which

The **which** parameter is the message value for which configuration is being checked. Must be set to one of the following values:

- `DB_LOG_DIRECT`

  System buffering is turned off for Berkeley DB log files to avoid double caching.

- `DB_LOG_DSYNC`

  Berkeley DB is configured to flush log writes to the backing disk before returning from the write system call, rather than flushing log writes explicitly in a separate system call, as necessary.

- `DB_LOG_AUTO_REMOVE`

  Berkeley DB automatically removes log files that are no longer needed.

- `DB_LOG_IN_MEMORY`

  Transaction logs are maintained in memory rather than on disk. This means that transactions exhibit the ACI (atomicity, consistency, and isolation) properties, but not D (durability).

- `DB_LOG_ZERO`

  All pages of a log file are zeroed when that log file is created.

onoffp

The **onoffp** parameter references memory into which the configuration of the specified **which** parameter is copied.

If the returned **onoff** value is zero, the parameter is off; otherwise, on.

Class

DbEnv

See Also

Logging Subsystem and Related Methods, DbEnv::log_set_config()

DbEnv::log_printf()

```
#include <db_cxx.h>

int
DbEnv::log_printf(DB_TXN *txnid, const char *fmt, ...);
```

The `DbEnv::log_printf()` method appends an informational message to the Berkeley DB database environment log files.

The `DbEnv::log_printf()` method allows applications to include information in the database environment log files, for later review using the db_printlog utility. This method is intended for debugging and performance tuning.

The `DbEnv::log_printf()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

txnid

If the logged message refers to an application-specified transaction, the **txnid** parameter is a transaction handle returned from DbEnv::txn_begin(); otherwise NULL.

fmt

A format string that specifies how subsequent arguments (or arguments accessed via the variable-length argument facilities of stdarg(3)) are converted for output. The format string may contain any formatting directives supported by the underlying C library vsnprintf(3) function.

Errors

The `DbEnv::log_printf()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv, DbLogc, DbLsn

See Also

Logging Subsystem and Related Methods

DbEnv::log_put()

```
#include <db_cxx.h>

int
DbEnv::log_put(DbLsn *lsn, const Dbt *data, u_int32_t flags);
```

The `DbEnv::log_put()` method appends records to the log. The DbLsn of the put record is returned in the **lsn** parameter.

The `DbEnv::log_put()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

data

The **data** parameter is the record to write to the log.

The caller is responsible for providing any necessary structure to **data**. (For example, in a write-ahead logging protocol, the application must understand what part of **data** is an operation code, what part is redo information, and what part is undo information. In addition, most transaction managers will store in **data** the DbLsn of the previous log record for the same transaction, to support chaining back through the transaction's log records during undo.)

flags

The **flags** parameter must be set to 0 or the following value:

- DB_FLUSH

  The log is forced to disk after this record is written, guaranteeing that all records with DbLsn values less than or equal to the one being "put" are on disk before `DbEnv::log_put()` returns.

lsn

The **lsn** parameter references memory into which the DbLsn of the put record is copied.

Errors

The `DbEnv::log_put()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the record to be logged is larger than the maximum log record; or if an invalid flag value or parameter was specified.

Class

DbEnv, DbLogc, DbLsn

See Also

Logging Subsystem and Related Methods

DbEnv::log_set_config()

```
#include <db_cxx.h>

int
DbEnv::log_set_config(u_int32_t flags, int onoff);
```

The `DbEnv::log_set_config()` method configures the Berkeley DB logging subsystem.

The `DbEnv::log_set_config()` method configures a database environment, not only operations performed using the specified DbEnv handle.

The `DbEnv::log_set_config()` method may be called at any time during the life of the application.

The `DbEnv::log_set_config()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set by bitwise inclusively **OR**'ing together one or more of the following values:

- `DB_LOG_DIRECT`

  Turn off system buffering of Berkeley DB log files to avoid double caching.

  Calling `DbEnv::log_set_config()` with the DB_LOG_DIRECT flag only affects the specified DbEnv handle (and any other Berkeley DB handles opened within the scope of that handle). For consistent behavior across the environment, all DbEnv handles opened in the environment must either set the DB_LOG_DIRECT flag or the flag should be specified in the DB_CONFIG configuration file.

  The `DB_LOG_DIRECT` flag may be used to configure Berkeley DB at any time during the life of the application.

- `DB_LOG_DSYNC`

  Configure Berkeley DB to flush log writes to the backing disk before returning from the write system call, rather than flushing log writes explicitly in a separate system call, as necessary. This is only available on some systems (for example, systems supporting the IEEE/ANSI Std 1003.1 (POSIX) standard O_DSYNC flag, or systems supporting the Windows FILE_FLAG_WRITE_THROUGH flag). This flag may result in inaccurate file modification times and other file-level information for Berkeley DB log files. This flag may offer a performance increase on some systems and a performance decrease on others.

  Calling `DbEnv::log_set_config()` with the DB_LOG_DSYNC flag only affects the specified DbEnv handle (and any other Berkeley DB handles opened within the scope of that handle). For consistent behavior across the environment, all DbEnv handles opened in the environment must either set the DB_LOG_DSYNC flag or the flag should be specified in the DB_CONFIG configuration file.

The DB_LOG_DSYNC flag may be used to configure Berkeley DB at any time during the life of the application.

- DB_LOG_AUTO_REMOVE

  If set, Berkeley DB will automatically remove log files that are no longer needed.

  Automatic log file removal is likely to make catastrophic recovery impossible.

  Replication applications will rarely want to configure automatic log file removal as it increases the likelihood a master will be unable to satisfy a client's request for a recent log record.

  Calling DbEnv::log_set_config() with the DB_LOG_AUTO_REMOVE flag affects the database environment, including all threads of control accessing the database environment.

  The DB_LOG_AUTO_REMOVE flag may be used to configure Berkeley DB at any time during the life of the application.

- DB_LOG_IN_MEMORY

  If set, maintain transaction logs in memory rather than on disk. This means that transactions exhibit the ACI (atomicity, consistency, and isolation) properties, but not D (durability); that is, database integrity will be maintained, but if the application or system fails, integrity will not persist. All database files must be verified and/or restored from a replication group master or archival backup after application or system failure.

  When in-memory logs are configured and no more log buffer space is available, Berkeley DB methods may return an additional error value, DB_LOG_BUFFER_FULL. When choosing log buffer and file sizes for in-memory logs, applications should ensure the in-memory log buffer size is large enough that no transaction will ever span the entire buffer, and avoid a state where the in-memory buffer is full and no space can be freed because a transaction that started in the first log "file" is still active.

  Calling DbEnv::log_set_config() with the DB_LOG_IN_MEMORY flag affects the database environment, including all threads of control accessing the database environment.

  The DB_LOG_IN_MEMORY flag may be used to configure Berkeley DB only before the DbEnv::open() method is called.

- DB_LOG_ZERO

  If set, zero all pages of a log file when that log file is created. This has shown to provide greater transaction throughput in some environments. The log file will be zeroed by the thread which needs to re-create the new log file. Other threads may not write to the log file while this is happening.

  Calling DbEnv::log_set_config() with the DB_LOG_ZERO flag affects only the current environment handle.

  The DB_LOG_ZERO flag may be used to configure Berkeley DB at any time.

onoff

If the **onoff** parameter is zero, the specified flags are cleared; otherwise they are set.

Errors

The `DbEnv::log_set_config()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv

See Also

Logging Subsystem and Related Methods

DbEnv::log_stat()

```
#include <db_cxx.h>

int
DbEnv::log_stat(DB_LOG_STAT **statp, u_int32_t flags);
```

The `DbEnv::log_stat()` method returns the logging subsystem statistics.

The `DbEnv::log_stat()` method creates a statistical structure of type `DB_LOG_STAT` and copies a pointer to it into a user-specified memory location.

Statistical structures are stored in allocated memory. If application-specific allocation routines have been declared (see DbEnv::set_alloc() for more information), they are used to allocate the memory; otherwise, the standard C library **malloc**(3) is used. The caller is responsible for deallocating the memory. To deallocate the memory, free the memory reference; references inside the returned memory need not be individually freed.

The following `DB_LOG_STAT` fields will be filled in:

- **u_int32_t st_magic;**

  The magic number that identifies a file as a log file.

- **u_int32_t st_version;**

  The version of the log file type.

- **int st_mode;**

  The mode of any created log files.

- **u_int32_t st_lg_bsize;**

  The in-memory log record cache size.

- **u_int32_t st_lg_size;**

  The log file size.

- **uintmax_t st_record;**

  The number of records written to this log.

- **u_int32_t st_w_mbytes;**

  The number of megabytes written to this log.

- **u_int32_t st_w_bytes;**

  The number of bytes over and above **st_w_mbytes** written to this log.

- **u_int32_t st_wc_mbytes;**

  The number of megabytes written to this log since the last checkpoint.

- **u_int32_t st_wc_bytes;**

  The number of bytes over and above **st_wc_mbytes** written to this log since the last checkpoint.

- **uintmax_t st_wcount;**

  The number of times the log has been written to disk.

- **uintmax_t st_wcount_fill;**

  The number of times the log has been written to disk because the in-memory log record cache filled up.

- **uintmax_t st_rcount;**

  The number of times the log has been read from disk.

- **uintmax_t st_scount;**

  The number of times the log has been flushed to disk.

- **u_int32_t st_cur_file;**

  The current log file number.

- **u_int32_t st_cur_offset;**

  The byte offset in the current log file.

- **u_int32_t st_disk_file;**

  The log file number of the last record known to be on disk.

- **u_int32_t st_disk_offset;**

  The byte offset of the last record known to be on disk.

- **u_int32_t st_maxcommitperflush;**

  The maximum number of commits contained in a single log flush.

- **u_int32_t st_mincommitperflush;**

  The minimum number of commits contained in a single log flush that contained a commit.

- **roff_t st_regsize;**

  The size of the log region, in bytes.

- **uintmax_t st_region_wait;**

  The number of times that a thread of control was forced to wait before obtaining the log region mutex.

- **uintmax_t st_region_nowait;**

  The number of times that a thread of control was able to obtain the log region mutex without waiting.

The `DbEnv::log_stat()` method may not be called before the DbEnv::open() method is called.

The `DbEnv::log_stat()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or the following value:

- `DB_STAT_CLEAR`

  Reset statistics after returning their values.

statp

The **statp** parameter references memory into which a pointer to the allocated statistics structure is copied.

Errors

The `DbEnv::log_stat()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv, DbLogc, DbLsn

See Also

Logging Subsystem and Related Methods

DbEnv::log_stat_print()

```
#include <db_cxx.h>

int
DbEnv::log_stat_print(u_int32_t flags);
```

The `DbEnv::log_stat_print()` method displays the logging subsystem statistical information, as described for the `DbEnv::log_stat()` method. The information is printed to a specified output channel (see the DbEnv::set_msgfile() method for more information), or passed to an application callback function (see the DbEnv::set_msgcall() method for more information).

The `DbEnv::log_stat_print()` method may not be called before the DbEnv::open() method is called.

The `DbEnv::log_stat_print()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

- `DB_STAT_ALL`

  Display all available information.

- `DB_STAT_CLEAR`

  Reset statistics after displaying their values.

Class

DbEnv, DbLogc, DbLsn

See Also

Logging Subsystem and Related Methods

DbEnv::set_lg_bsize()

```
#include <db_cxx.h>

int
DbEnv::set_lg_bsize(u_int32_t lg_bsize);
```

Set the size of the in-memory log buffer, in bytes.

When the logging subsystem is configured for on-disk logging, the default size of the in-memory log buffer is approximately 32KB. Log information is stored in-memory until the storage space fills up or transaction commit forces the information to be flushed to stable storage. In the presence of long-running transactions or transactions producing large amounts of data, larger buffer sizes can increase throughput.

When the logging subsystem is configured for in-memory logging, the default size of the in-memory log buffer is 1MB. Log information is stored in-memory until the storage space fills up or transaction abort or commit frees up the memory for new transactions. In the presence of long-running transactions or transactions producing large amounts of data, the buffer size must be sufficient to hold all log information that can accumulate during the longest running transaction. When choosing log buffer and file sizes for in-memory logs, applications should ensure the in-memory log buffer size is large enough that no transaction will ever span the entire buffer, and avoid a state where the in-memory buffer is full and no space can be freed because a transaction that started in the first log "file" is still active.

The database environment's log buffer size may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_lg_bsize", one or more whitespace characters, and the size in bytes. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The `DbEnv::set_lg_bsize()` method configures a database environment, not only operations performed using the specified DbEnv handle.

The `DbEnv::set_lg_bsize()` method may not be called after the DbEnv::open() method is called. If the database environment already exists when DbEnv::open() is called, the information specified to `DbEnv::set_lg_bsize()` will be ignored.

The `DbEnv::set_lg_bsize()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

lg_bsize

The **lg_bsize** parameter is the size of the in-memory log buffer, in bytes.

Errors

The `DbEnv::set_lg_bsize()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv, DbLogc, DbLsn

See Also

Logging Subsystem and Related Methods

DbEnv::set_lg_dir()

```
#include <db_cxx.h>

int
DbEnv::set_lg_dir(const char *dir);
```

The path of a directory to be used as the location of logging files. Log files created by the Log Manager subsystem will be created in this directory.

If no logging directory is specified, log files are created in the environment home directory. See Berkeley DB File Naming for more information.

For the greatest degree of recoverability from system or application failure, database files and log files should be located on separate physical devices.

The database environment's logging directory may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_lg_dir", one or more whitespace characters, and the directory name. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The `DbEnv::set_lg_dir()` method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The `DbEnv::set_lg_dir()` method may not be called after the DbEnv::open() method is called. If the database environment already exists when DbEnv::open() is called, the information specified to `DbEnv::set_lg_dir()` must be consistent with the existing environment or corruption can occur.

The `DbEnv::set_lg_dir()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

dir

The **dir** parameter is the directory used to store the logging files.

When using a Unicode build on Windows (the default), the **dir** argument will be interpreted as a UTF-8 string, which is equivalent to ASCII for Latin characters.

Errors

The `DbEnv::set_lg_dir()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after DbEnv::open() was called; or if an invalid flag value or parameter was specified.

Class

DbEnv, DbLogc, DbLsn

See Also

Logging Subsystem and Related Methods

DbEnv::set_lg_filemode()

```
#include <db_cxx.h>

int
DbEnv::set_lg_filemode(int lg_filemode);
```

Set the absolute file mode for created log files. This method is **only** useful for the rare Berkeley DB application that does not control its umask value.

Normally, if Berkeley DB applications set their umask appropriately, all processes in the application suite will have read permission on the log files created by any process in the application suite. However, if the Berkeley DB application is a library, a process using the library might set its umask to a value preventing other processes in the application suite from reading the log files it creates. In this rare case, the DbEnv::set_lg_filemode() method can be used to set the mode of created log files to an absolute value.

The database environment's log file mode may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_lg_filemode", one or more whitespace characters, and the absolute mode of created log files. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The DbEnv::set_lg_filemode() method configures a database environment, not only operations performed using the specified DbEnv handle.

The DbEnv::set_lg_filemode() method may be called at any time during the life of the application.

The DbEnv::set_lg_filemode() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

lg_filemode

The **lg_filemode** parameter is the absolute mode of the created log file.

Class

DbEnv, DbLogc, DbLsn

See Also

Logging Subsystem and Related Methods

DbEnv::set_lg_max()

```
#include <db_cxx.h>

int
DbEnv::set_lg_max(u_int32_t lg_max);
```

Set the maximum size of a single file in the log, in bytes. Because DbLsn file offsets are unsigned four-byte values, the set value may not be larger than the maximum unsigned four-byte value.

When the logging subsystem is configured for on-disk logging, the default size of a log file is 10MB.

When the logging subsystem is configured for in-memory logging, the default size of a log file is 256KB. In addition, the configured log buffer size must be larger than the log file size. (The logging subsystem divides memory configured for in-memory log records into "files", as database environments configured for in-memory log records may exchange log records with other members of a replication group, and those members may be configured to store log records on-disk.) When choosing log buffer and file sizes for in-memory logs, applications should ensure the in-memory log buffer size is large enough that no transaction will ever span the entire buffer, and avoid a state where the in-memory buffer is full and no space can be freed because a transaction that started in the first log "file" is still active.

See Log File Limits for more information.

The database environment's log file size may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_lg_max", one or more whitespace characters, and the size in bytes. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The DbEnv::set_lg_max() method configures a database environment, not only operations performed using the specified DbEnv handle.

The DbEnv::set_lg_max() method may be called at any time during the life of the application.

If no size is specified by the application, the size last specified for the database region will be used, or if no database region previously existed, the default will be used.

The DbEnv::set_lg_max() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

lg_max

The **lg_max** parameter is the size of a single log file, in bytes.

Errors

The DbEnv::set_lg_max() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the size of the log file is less than four times the size of the in-memory log buffer; the specified log file size was too large; or if an invalid flag value or parameter was specified.

Class

DbEnv, DbLogc, DbLsn

See Also

Logging Subsystem and Related Methods

DbEnv::set_lg_regionmax()

```
#include <db_cxx.h>

int
DbEnv::set_lg_regionmax(u_int32_t lg_regionmax);
```

Set the size of the underlying logging area of the Berkeley DB environment, in bytes. By default, or if the value is set to 0, the minimum region size is used, approximately 128KB. The log region is used to store filenames, and so may need to be increased in size if a large number of files will be opened and registered with the specified Berkeley DB environment's log manager.

The database environment's log region size may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_lg_regionmax", one or more whitespace characters, and the size in bytes. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The DbEnv::set_lg_regionmax() method configures a database environment, not only operations performed using the specified DbEnv handle.

The DbEnv::set_lg_regionmax() method may not be called after the DbEnv::open() method is called. If the database environment already exists when DbEnv::open() is called, the information specified to DbEnv::set_lg_regionmax() will be ignored.

The DbEnv::set_lg_regionmax() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

lg_regionmax

The **lg_regionmax** parameter is the size of the logging area in the Berkeley DB environment, in bytes.

Errors

The DbEnv::set_lg_regionmax() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after DbEnv::open() was called; or if an invalid flag value or parameter was specified.

Class

DbEnv, DbLogc, DbLsn

See Also

Logging Subsystem and Related Methods

The DbLogc Handle

```
#include <db_cxx.h>

int
DbEnv::log_cursor(DbLogc **cursorp, u_int32_t flags);
```

The DbLogc object is the handle for a cursor into the log files, supporting sequential access to the records stored in log files. The handle is not free-threaded. Once the DbLogc::close() method is called, the handle may not be accessed again, regardless of that method's return.

For more information, see the DbLsn handle.

DbLogc::close()

```
#include <db_cxx.h>

int
DbLogc::close(u_int32_t flags);
```

The DbLogc::close() method discards the log cursor. After DbLogc::close() has been called, regardless of its return, the cursor handle may not be used again.

The DbLogc::close() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter is currently unused, and must be set to 0.

Errors

The DbLogc::close() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the cursor is already closed; or if an invalid flag value or parameter was specified.

Class

DbEnv, DbLogc, DbLsn

See Also

Logging Subsystem and Related Methods

DbLogc::get()

```
#include <db_cxx.h>

int
DbLogc::get(DbLsn *lsn, Dbt *data, u_int32_t flags);
```

The `DbLogc::get()` method returns records from the log.

Unless otherwise specified, the `DbLogc::get()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

data

The data field of the **data** structure is set to the record retrieved, and the size field indicates the number of bytes in the record. See Dbt for a description of other fields in the **data** structure. The DB_DBT_MALLOC, DB_DBT_REALLOC and DB_DBT_USERMEM flags may be specified for any Dbt used for data retrieval.

flags

The **flags** parameter must be set to one of the following values:

- DB_CURRENT

  Return the log record to which the log currently refers.

- DB_FIRST

  The first record from any of the log files found in the log directory is returned in the **data** parameter. The **lsn** parameter is overwritten with the DbLsn of the record returned.

  The `DbLogc::get()` method will return DB_NOTFOUND if DB_FIRST is set and the log is empty.

- DB_LAST

  The last record in the log is returned in the **data** parameter. The **lsn** parameter is overwritten with the DbLsn of the record returned.

  The `DbLogc::get()` method will return DB_NOTFOUND if DB_LAST is set and the log is empty.

- DB_NEXT

  The current log position is advanced to the next record in the log, and that record is returned in the **data** parameter. The **lsn** parameter is overwritten with the DbLsn of the record returned.

  If the cursor has not been initialized via DB_FIRST, DB_LAST, DB_SET, DB_NEXT, or DB_PREV, `DbLogc::get()` will return the first record in the log.

The `DbLogc::get()` method will return DB_NOTFOUND if DB_NEXT is set and the last log record has already been returned or the log is empty.

- DB_PREV

  The current log position is advanced to the previous record in the log, and that record is returned in the **data** parameter. The **lsn** parameter is overwritten with the DbLsn of the record returned.

  If the cursor has not been initialized via DB_FIRST, DB_LAST, DB_SET, DB_NEXT, or DB_PREV, `DbLogc::get()` will return the last record in the log.

  The `DbLogc::get()` method will return DB_NOTFOUND if DB_PREV is set and the first log record has already been returned or the log is empty.

- DB_SET

  Retrieve the record specified by the **lsn** parameter.

lsn

When the **flag** parameter is set to DB_CURRENT, DB_FIRST, DB_LAST, DB_NEXT or DB_PREV, the **lsn** parameter is overwritten with the DbLsn value of the record retrieved. When **flag** is set to DB_SET, the **lsn** parameter is the DbLsn value of the record to be retrieved.

Errors

The `DbLogc::get()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the DB_CURRENT flag was set and the log cursor has not yet been initialized; the DB_CURRENT, DB_NEXT, or DB_PREV flags were set and the log was opened with the DB_THREAD flag set; the DB_SET flag was set and the specified log sequence number does not appear in the log; or if an invalid flag value or parameter was specified.

Class

DbEnv, DbLogc, DbLsn

See Also

Logging Subsystem and Related Methods

DbEnv::log_compare()

```
#include <db_cxx.h>

static int
DbEnv::log_compare(const DbLsn *lsn0, const DbLsn *lsn1);
```

The `DbEnv::log_compare()` method allows the caller to compare two `DbLsn` structures, returning 0 if they are equal, 1 if **lsn0** is greater than **lsn1**, and -1 if **lsn0** is less than **lsn1**.

Parameters

lsn0

The **lsn0** parameter is one of the `DbLsn` structures to be compared.

lsn1

The **lsn1** parameter is one of the `DbLsn` structures to be compared.

Class

DbEnv, DbLogc, DbLsn

See Also

Logging Subsystem and Related Methods

Chapter 9. The DbMpoolFile Handle

```
#include <db_cxx.h>

class DbMpoolFile {
public:
        DB_MPOOLFILE *DbMpoolFile::get_DB_MPOOLFILE();
        const DB_MPOOLFILE *DbMpoolFile::get_const_DB_MPOOLFILE() const; ...
};
```

The memory pool interfaces for the Berkeley DB database environment are methods of the DbEnv handle. The DbEnv memory pool methods and the DB_MPOOLFILE class provide general-purpose, page-oriented buffer management of files. Although designed to work with the other Dbclasses, they are also useful for more general purposes. The memory pools are referred to in this document as simply *the cache*.

The cache may be shared between processes. The cache is usually filled by pages from one or more files. Pages in the cache are replaced in LRU (least-recently-used) order, with each new page replacing the page that has been unused the longest. Pages retrieved from the cache using DbMpoolFile::get() are *pinned* in the cache until they are returned to the control of the cache using the DbMpoolFile::put() method.

The DbMpoolFile object is the handle for a file in the cache. The handle is not free-threaded. Once the DbMpoolFile::close() method is called, the handle may not be accessed again, regardless of that method's return.

Memory Pools and Related Methods

Memory Pools and Related Methods	Description
DbMemoryException	Exception class for insufficient memory
Db::get_mpf()	Return the DbMpoolFile for a Db
DbEnv::memp_stat()	Return cache statistics
DbEnv::memp_stat_print()	Print cache statistics
DbEnv::memp_sync()	Flush all pages from the cache
DbEnv::memp_trickle()	Flush some pages from the cache
Memory Pool Configuration	
DbEnv::memp_register()	Register a custom file type
DbEnv::set_cache_max(), DbEnv::get_cache_max()	Set/get the maximum cache size
DbEnv::set_cachesize(), DbEnv::get_cachesize()	Set/get the environment cache size
DbEnv::set_mp_max_openfd(), DbEnv::get_mp_max_openfd()	Set/get the maximum number of open file descriptors
DbEnv::set_mp_max_write(), DbEnv::get_mp_max_write()	Set/get the maximum number of sequential disk writes
DbEnv::set_mp_mmapsize(), DbEnv::get_mp_mmapsize()	Set/get maximum file size to memory map when opened read-only
Memory Pool Files	
DbEnv::memp_fcreate()	Create a memory pool file handle
DbMpoolFile::close()	Close a file in the cache
DbMpoolFile::get()	Get page from a file in the cache
DbMpoolFile::open()	Open a file in the cache
DbMpoolFile::put()	Return a page to the cache
DbMpoolFile::sync()	Flush pages from a file from the cache
Memory Pool File Configuration	
DbMpoolFile::set_clear_len(), DbMpoolFile::get_clear_len()	Set/get number of bytes to clear when creating a new page
DbMpoolFile::set_fileid(), DbMpoolFile::get_fileid()	Set/get file unique identifier
DbMpoolFile::set_flags(), DbMpoolFile::get_flags()	Set/get file options
DbMpoolFile::set_ftype(), DbMpoolFile::get_ftype()	Set/get file type
DbMpoolFile::set_lsn_offset(), DbMpoolFile::get_lsn_offset()	Set/get file log-sequence-number offset
DbMpoolFile::set_maxsize(), DbMpoolFile::get_maxsize()	Set/get maximum file size

Memory Pools and Related Methods	Description
DbMpoolFile::set_pgcookie(), DbMpoolFile::get_pgcookie()	Set/get file cookie for pgin/pgout
DbMpoolFile::set_priority(), DbMpoolFile::get_priority()	Set/get cache file priority

Db::get_mpf()

```
#include <db_cxx.h>

DbMpoolFile *
Db::get_mpf();
```

The `Db::get_mpf()` method returns the handle for the cache file underlying the database.

The `Db::get_mpf()` method should be used with caution on a replication client site. This method exposes an internal structure that may not be valid after a client site synchronizes with its master site.

The `Db::get_mpf()` method may be called at any time during the life of the application.

Class

Db

See Also

Memory Pools and Related Methods

DbEnv::get_cache_max()

```
#include <db_cxx.h>

int
DbEnv::get_cache_max(u_int32_t *gbytesp, u_int32_t *bytesp);
```

The `DbEnv::get_cache_max()` method returns the maximum size of the cache as set using the
DbEnv::set_cache_max() method.

The `DbEnv::get_cache_max()` method may be called at any time during the life of the application.

The `DbEnv::get_cache_max()` method either returns a non-zero error value or throws an exception that
encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

bytesp

The **bytesp** parameter references memory into which the additional bytes of memory in the cache is
copied.

gbytesp

The **gbytesp** parameter references memory into which the gigabytes of memory in the cache is copied.

Class

DbEnv

See Also

Database Environments and Related Methods, DbEnv::set_cache_max()

DbEnv::get_cachesize()

```
#include <db_cxx.h>

int
DbEnv::get_cachesize(u_int32_t *gbytesp, u_int32_t *bytesp, int *ncachep);
```

The `DbEnv::get_cachesize()` method returns the current size and composition of the cache, as set using the DbEnv::set_cachesize() method.

The `DbEnv::get_cachesize()` method may be called at any time during the life of the application.

The `DbEnv::get_cachesize()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

bytesp

The **bytesp** parameter references memory into which the additional bytes of memory in the cache is copied.

gbytesp

The **gbytesp** parameter references memory into which the gigabytes of memory in the cache is copied.

ncachep

The **ncachep** parameter references memory into which the number of caches is copied.

Class

DbEnv

See Also

Memory Pools and Related Methods, Database Environments and Related Methods, DbEnv::set_cachesize()

DbEnv::get_mp_max_openfd()

```
#include <db_cxx.h>

int
DbEnv::get_mp_max_openfd(int *maxopenfdp);
```

Returns the maximum number of file descriptors the library will open concurrently when flushing dirty pages from the cache. This value is set by the DbEnv::set_mp_max_openfd() method.

The `DbEnv::get_mp_max_openfd()` method may be called at any time during the life of the application.

The `DbEnv::get_mp_max_openfd()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

maxopenfdp

The `DbEnv::get_mp_max_openfd()` method returns the maximum number of file descriptors open in **maxopenfdp**.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods, DbEnv::set_mp_max_openfd()

DbEnv::get_mp_max_write()

```
#include <db_cxx.h>

int
DbEnv::get_mp_max_write(int *maxwritep, db_timeout_t *maxwrite_sleepp);
```

The `DbEnv::get_mp_max_write()` method returns the current maximum number of sequential write operations and microseconds to pause that the library can schedule when flushing dirty pages from the cache. These values are set by the DbEnv::set_mp_max_write() method.

The `DbEnv::get_mp_max_write()` method may be called at any time during the life of the application.

The `DbEnv::get_mp_max_write()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

maxwritep

The **maxwritep** parameter references memory into which the maximum number of sequential write operations is copied.

maxwrite_sleepp

The **maxwrite_sleepp** parameter references memory into which the microseconds to pause before scheduling further write operations is copied.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods, DbEnv::set_mp_max_write()

DbEnv::get_mp_mmapsize()

```
#include <db_cxx.h>

int
DbEnv::get_mp_mmapsize(size_t *mp_mmapsizep);
```

The `DbEnv::get_mp_mmapsize()` method returns the the maximum file size, in bytes, for a file to be mapped into the process address space. This value can be managed using the DbEnv::set_mp_mmapsize() method.

The `DbEnv::get_mp_mmapsize()` method may be called at any time during the life of the application.

The `DbEnv::get_mp_mmapsize()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

mp_mmapsizep

The `DbEnv::get_mp_mmapsize()` method returns the maximum file map size in **mp_mmapsizep**.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods, DbEnv::set_mp_mmapsize()

DbEnv::memp_fcreate()

```
#include <db_cxx.h>

int
DbEnv::memp_fcreate(DbMpoolFile **dbmfp, u_int32_t flags);
```

The `DbEnv::memp_fcreate()` method creates a DbMpoolFile structure that is the handle for a Berkeley DB cache (that is, a shared memory buffer pool file). A pointer to this structure is returned in the memory to which **dbmfp** refers. Calling the DbMpoolFile::close() method will discard the returned handle.

The `DbEnv::memp_fcreate()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

dbmfp

The `DbEnv::memp_fcreate()` method returns a pointer to a mpool structure in **dbmfp**.

flags

The **flags** parameter is currently unused, and must be set to 0.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods

DbEnv::memp_register()

```
#include <db_cxx.h>

extern "C" {
    typedef int (*pgin_fcn_type)(DB_ENV *dbenv,
 db_pgno_t pgno, void *pgaddr, DBT *pgcookie);
    typedef int (*pgout_fcn_type)(DB_ENV *dbenv,
 db_pgno_t pgno, void *pgaddr, DBT *pgcookie);
};
int
DbEnv::memp_register(int ftype,
    pgin_fcn_type pgin_fcn, pgout_fcn_type pgout_fcn);
```

The `DbEnv::memp_register()` method registers page-in and page-out functions for files of type **ftype** in the cache.

If the **pgin_fcn** function is non-NULL, it is called each time a page is read into the cache from a file of type **ftype**, or a page is created for a file of type **ftype** (see the DB_MPOOL_CREATE flag for the DbMpoolFile::get() method).

If the **pgout_fcn** function is non-NULL, it is called each time a page is written to a file of type **ftype**.

The purpose of the `DbEnv::memp_register()` function is to support processing when pages are entered into, or flushed from, the cache. For example, this functionality might be used to do byte-endian conversion as pages are read from, or written to, the underlying file.

A file type must be specified to make it possible for unrelated threads or processes that are sharing a cache, to evict each other's pages from the cache. During initialization, applications should call `DbEnv::memp_register()` for each type of file requiring input or output processing that will be sharing the underlying cache. (No registry is necessary for the standard Berkeley DB access method types because Db::open() registers them separately.)

If a thread or process does not call `DbEnv::memp_register()` for a file type, it is impossible for it to evict pages for any file requiring input or output processing from the cache. For this reason, `DbEnv::memp_register()` should always be called by each application sharing a cache for each type of file included in the cache, regardless of whether or not the application itself uses files of that type.

The `DbEnv::memp_register()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

ftype

The **ftype** parameter specifies the type of file for which the page-in and page-out functions will be called.

The **ftype** value for a file must be a non-zero positive number less than 128 (0 and negative numbers are reserved for internal use by the Berkeley DB library).

pgin_fcn, pgout_fcn

The page-in and page-out functions.

The **pgin_fcn** and **pgout_fcn** functions are called with a reference to the current database environment, the page number being read or written, a pointer to the page being read or written, and any parameter **pgcookie** that was specified to the DbMpoolFile::set_pgcookie() method.

The **pgin_fcn** and **pgout_fcn** functions should return 0 on success, and a non-zero value on failure, in which case the shared Berkeley DB library function calling it will also fail, returning that non-zero value. The non-zero value should be selected from values outside of the Berkeley DB library namespace.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods

DbEnv::memp_stat()

```
#include <db_cxx.h>

int
DbEnv::memp_stat(DB_MPOOL_STAT **gsp,
    DB_MPOOL_FSTAT *(*fsp)[], u_int32_t flags);
```

The `DbEnv::memp_stat()` method returns the memory pool (that is, the buffer cache) subsystem statistics.

The `DbEnv::memp_stat()` method creates statistical structures of type `DB_MPOOL_STAT` and `DB_MPOOL_FSTAT`, and copy pointers to them into user-specified memory locations. The cache statistics are stored in the `DB_MPOOL_STAT` structure and the per-file cache statistics are stored the `DB_MPOOL_FSTAT` structure.

Statistical structures are stored in allocated memory. If application-specific allocation routines have been declared (see DbEnv::set_alloc() for more information), they are used to allocate the memory; otherwise, the standard C library **malloc**(3) is used. The caller is responsible for deallocating the memory. To deallocate the memory, free the memory reference; references inside the returned memory need not be individually freed.

If **gsp** is non-NULL, the global statistics for the cache **mp** are copied into the memory location to which it refers. The following `DB_MPOOL_STAT` fields will be filled in:

- **u_int32_t st_gbytes;**

  Gigabytes of cache (total cache size is st_gbytes + st_bytes).

- **u_int32_t st_bytes;**

  Bytes of cache (total cache size is st_gbytes + st_bytes).

- **u_int32_t st_ncache;**

  Number of caches.

- **u_int32_t st_max_ncache;**

  Maximum number of caches, as configured with the DbEnv::set_cache_max() method.

- **roff_t st_regsize;**

  Individual cache size, in bytes.

- **size_t st_mmapsize;**

  Maximum memory-mapped file size.

- **int st_maxopenfd;**

  Maximum open file descriptors.

- **int st_maxwrite;**

Maximum sequential buffer writes.

- **db_timeout_t st_maxwrite_sleep;**

  Microseconds to pause after writing maximum sequential buffers.

- **u_int32_t st_map;**

  Requested pages mapped into the process' address space (there is no available information about whether or not this request caused disk I/O, although examining the application page fault rate may be helpful).

- **uintmax_t st_cache_hit;**

  Requested pages found in the cache.

- **uintmax_t st_cache_miss;**

  Requested pages not found in the cache.

- **uintmax_t st_page_create;**

  Pages created in the cache.

- **uintmax_t st_page_in;**

  Pages read into the cache.

- **uintmax_t st_page_out;**

  Pages written from the cache to the backing file.

- **uintmax_t st_ro_evict;**

  Clean pages forced from the cache.

- **uintmax_t st_rw_evict;**

  Dirty pages forced from the cache.

- **uintmax_t st_page_trickle;**

  Dirty pages written using the DbEnv::memp_trickle() method.

- **u_int32_t st_pages;**

  Pages in the cache.

- **uintmax_t st_page_clean;**

  Clean pages currently in the cache.

- **uintmax_t st_page_dirty;**

Dirty pages currently in the cache.

- **uintmax_t st_hash_buckets;**

  Number of hash buckets in buffer hash table.

- **uintmax_t st_hash_searches;**

  Total number of buffer hash table lookups.

- **uintmax_t st_hash_longest;**

  Longest chain ever encountered in buffer hash table lookups.

- **uintmax_t st_hash_examined;**

  Total number of hash elements traversed during hash table lookups.

- **uintmax_t st_hash_nowait;**

  Number of times that a thread of control was able to obtain a hash bucket lock without waiting.

- **uintmax_t st_hash_wait;**

  Number of times that a thread of control was forced to wait before obtaining a hash bucket lock.

- **uintmax_t st_hash_max_nowait;**

  The number of times a thread of control was able to obtain the hash bucket lock without waiting on the bucket which had the maximum number of times that a thread of control needed to wait.

- **uintmax_t st_hash_max_wait;**

  Maximum number of times any hash bucket lock was waited for by a thread of control.

- **uintmax_t st_region_wait;**

  Number of times that a thread of control was forced to wait before obtaining a cache region mutex.

- **uintmax_t st_region_nowait;**

  Number of times that a thread of control was able to obtain a cache region mutex without waiting.

- **uintmax_t st_mvcc_frozen;**

  Number of buffers frozen.

- **uintmax_t st_mvcc_thawed;**

  Number of buffers thawed.

- **uintmax_t st_mvcc_freed;**

Number of frozen buffers freed.

- **uintmax_t st_alloc;**

  Number of page allocations.

- **uintmax_t st_alloc_buckets;**

  Number of hash buckets checked during allocation.

- **uintmax_t st_alloc_max_buckets;**

  Maximum number of hash buckets checked during an allocation.

- **uintmax_t st_alloc_pages;**

  Number of pages checked during allocation.

- **uintmax_t st_alloc_max_pages;**

  Maximum number of pages checked during an allocation.

- **uintmax_t st_io_wait;**

  Number of operations blocked waiting for I/O to complete.

- **uintmax_t st_sync_interrupted;**

  Number of mpool sync operations interrupted.

If **fsp** is non-NULL, a pointer to a NULL-terminated variable length array of statistics for individual files, in the cache **mp**, is copied into the memory location to which it refers. If no individual files currently exist in the cache, **fsp** will be set to NULL.

The per-file statistics are stored in structures of type DB_MPOOL_FSTAT. The following DB_MPOOL_FSTAT fields will be filled in for each file in the cache; that is, each element of the array:

- **char * file_name;**

  The name of the file.

- **size_t st_pagesize;**

  Page size in bytes.

- **uintmax_t st_cache_hit;**

  Requested pages found in the cache.

- **uintmax_t st_cache_miss;**

  Requested pages not found in the cache.

- **u_int32_t st_map;**

  Requested pages mapped into the process' address space.

- **uintmax_t st_page_create;**

  Pages created in the cache.

- **uintmax_t st_page_in;**

  Pages read into the cache.

- **uintmax_t st_page_out;**

  Pages written from the cache to the backing file.

The `DbEnv::memp_stat()` method may not be called before the DbEnv::open() method is called.

The `DbEnv::memp_stat()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or the following value:

- `DB_STAT_CLEAR`

  Reset statistics after returning their values.

fsp

The **fsp** parameter references memory into which a pointer to the allocated per-file statistics structures is copied.

gsp

The **gsp** parameter references memory into which a pointer to the allocated global statistics structure is copied.

Errors

The `DbEnv::memp_stat()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods

DbEnv::memp_stat_print()

```
#include <db_cxx.h>

int
DbEnv::memp_stat_print(u_int32_t flags);
```

The `DbEnv::memp_stat_print()` method displays cache subsystem statistical information, as described for the DbEnv::memp_stat() method. The information is printed to a specified output channel (see the DbEnv::set_msgfile() method for more information), or passed to an application callback function (see the DbEnv::set_msgcall() method for more information).

The `DbEnv::memp_stat_print()` method may not be called before the DbEnv::open() method is called.

The `DbEnv::memp_stat_print()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

- `DB_STAT_ALL`

  Display all available information.

- `DB_STAT_CLEAR`

  Reset statistics after displaying their values.

- `DB_STAT_MEMP_HASH`

  Display the buffers with hash chains.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods

DbEnv::memp_sync()

```
#include <db_cxx.h>

int
DbEnv::memp_sync(DbLsn *lsn);
```

The `DbEnv::memp_sync()` method flushes modified pages in the cache to their backing files.

Pages in the cache that cannot be immediately written back to disk (for example, pages that are currently in use by another thread of control) are waited for and written to disk as soon as it is possible to do so.

The `DbEnv::memp_sync()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

lsn

The purpose of the **lsn** parameter is to enable a transaction manager to ensure, as part of a checkpoint, that all pages modified by a certain time have been written to disk.

All modified pages with a a log sequence number (DbLsn) less than the **lsn** parameter are written to disk. If **lsn** is NULL, all modified pages in the cache are written to disk.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods

DbEnv::memp_trickle()

```
#include <db_cxx.h>

int
DbEnv::memp_trickle(int percent, int *nwrotep);
```

The `DbEnv::memp_trickle()` method ensures that a specified percent of the pages in the cache are clean, by writing dirty pages to their backing files.

The purpose of the `DbEnv::memp_trickle()` function is to enable a memory pool manager to ensure that a page is always available for reading in new information without having to wait for a write.

The `DbEnv::memp_trickle()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

nwrotep

The **nwrotep** parameter references memory into which the number of pages written to reach the specified percentage is copied.

percent

The **percent** parameter is the percent of the pages in the cache that should be clean.

Errors

The `DbEnv::memp_trickle()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors: following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods

DbEnv::set_cache_max()

```
#include <db_cxx.h>

int
DbEnv::set_cache_max(u_int32_t gbytes, u_int32_t bytes);
```

The `DbEnv::set_cache_max()` method sets the maximum cache size, in bytes. The specified size is rounded to the nearest multiple of the cache region size, which is the initial cache size divided by the number of regions specified to the DbEnv::set_cachesize() method. If no value is specified, it defaults to the initial cache size.

The database environment's maximum cache size may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_cache_max", one or more whitespace characters, and the size in bytes. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The `DbEnv::set_cache_max()` method configures a database environment, not only operations performed using the specified DbEnv handle.

The `DbEnv::set_cache_max()` method may be called at any time during the life of the application.

The `DbEnv::set_cache_max()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

bytes

The **bytes** parameter specifies the number of bytes which, when added to the **gbytes** parameter, specifies the maximum size of the cache.

gbytes

The **gbytes** parameter specifies the number of bytes which, when added to the **bytes** parameter, specifies the maximum size of the cache.

Errors

The `DbEnv::set_cache_max()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_cachesize()

```
#include <db_cxx.h>

int
DbEnv::set_cachesize(u_int32_t gbytes, u_int32_t bytes, int ncache);
```

Set the size of the shared memory buffer pool — that is, the cache. The cache should be the size of the normal working data set of the application, with some small amount of additional memory for unusual situations. (Note: the working set is not the same as the number of pages accessed simultaneously, and is usually much larger.)

The default cache size is 256KB, and may not be specified as less than 20KB. Any cache size less than 500MB is automatically increased by 25% to account for cache overhead; cache sizes larger than 500MB are used as specified. The maximum size of a single cache is 4GB on 32-bit systems and 10TB on 64-bit systems. (All sizes are in powers-of-two, that is, 256KB is 2^18 not 256,000.) For information on tuning the Berkeley DB cache size, see Selecting a cache size.

It is possible to specify caches to Berkeley DB large enough they cannot be allocated contiguously on some architectures. For example, some releases of Solaris limit the amount of memory that may be allocated contiguously by a process. If **ncache** is 0 or 1, the cache will be allocated contiguously in memory. If it is greater than 1, the cache will be split across **ncache** separate regions, where the **region size** is equal to the initial cache size divided by **ncache**.

The cache may be resized by calling `DbEnv::set_cachesize()` after the environment is open. The supplied size will be rounded to the nearest multiple of the region size and may not be larger than the maximum size configured with DbEnv::set_cache_max(). The **ncache** parameter is ignored when resizing the cache.

The database environment's initial cache size may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_cachesize", one or more whitespace characters, and the initial cache size specified in three parts: the gigabytes of cache, the additional bytes of cache, and the number of caches, also separated by whitespace characters. For example, "set_cachesize 2 524288000 3" would create a 2.5GB logical cache, split between three physical caches. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The `DbEnv::set_cachesize()` method configures a database environment, not only operations performed using the specified DbEnv handle.

The `DbEnv::set_cachesize()` method may be called at any time during the life of the application.

The `DbEnv::set_cachesize()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

bytes

The size of the cache is set to **gbytes** gigabytes plus **bytes**.

gbytes

The size of the cache is set to **gbytes** gigabytes plus **bytes**.

ncache

The **ncache** parameter is the number of caches to create.

Errors

The `DbEnv::set_cachesize()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the specified cache size was impossibly small; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Database Environments and Related Methods

DbEnv::set_mp_max_openfd()

```
#include <db_cxx.h>

int
DbEnv::set_mp_max_openfd(int maxopenfd);
```

The `DbEnv::set_mp_max_openfd()` method limits the number of file descriptors the library will open concurrently when flushing dirty pages from the cache.

The database environment's limit on open file descriptors to flush dirty pages may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_max_openfd", one or more whitespace characters, and the number of open file descriptors. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The DbEnv::set_mp_max_openfd() method configures a database environment, not only operations performed using the specified DbEnv handle.

The `DbEnv::set_mp_max_openfd()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

maxopenfd

The maximum number of file descriptors that may be concurrently opened by the library when flushing dirty pages from the cache.

Errors

The `DbEnv::set_mp_max_openfd()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods

DbEnv::set_mp_max_write()

```
#include <db_cxx.h>

int
DbEnv::set_mp_max_write(int maxwrite, db_timeout_t maxwrite_sleep);
```

The `DbEnv::set_mp_max_write()` method limits the number of sequential write operations scheduled by the library when flushing dirty pages from the cache.

The database environment's maximum number of sequential write operations may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_mp_max_write", one or more whitespace characters, and the maximum number of sequential writes and the number of microseconds to sleep, also separated by whitespace characters. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The `DbEnv::set_mp_max_write()` method configures a database environment, not only operations performed using the specified DbEnv handle.

The `DbEnv::set_mp_max_write()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

maxwrite

The maximum number of sequential write operations scheduled by the library when flushing dirty pages from the cache, or 0 if there is no limitation on the number of sequential write operations.

maxwrite_sleep

The number of microseconds the thread of control should pause before scheduling further write operations. It must be specified as an unsigned 32-bit number of microseconds, limiting the maximum pause to roughly 71 minutes.

Errors

The `DbEnv::set_mp_max_write()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods

DbEnv::set_mp_mmapsize()

```
#include <db_cxx.h>

int
DbEnv::set_mp_mmapsize(size_t mp_mmapsize);
```

Files that are opened read-only in the cache (and that satisfy a few other criteria) are, by default, mapped into the process address space instead of being copied into the local cache. This can result in better-than-usual performance because available virtual memory is normally much larger than the local cache, and page faults are faster than page copying on many systems. However, it can cause resource starvation in the presence of limited virtual memory, and it can result in immense process sizes in the presence of large databases.

The `DbEnv::set_mp_mmapsize()` method sets the maximum file size, in bytes, for a file to be mapped into the process address space. If no value is specified, it defaults to 10MB.

The database environment's maximum mapped file size may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_mp_mmapsize", one or more whitespace characters, and the size in bytes. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The `DbEnv::set_mp_mmapsize()` method configures a database environment, not only operations performed using the specified DbEnv handle.

The `DbEnv::set_mp_mmapsize()` method may be called at any time during the life of the application.

The `DbEnv::set_mp_mmapsize()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

mp_mmapsize

The **mp_mmapsize** parameter is the maximum file size, in bytes, for a file to be mapped into the process address space.

Errors

The `DbEnv::set_mp_mmapsize()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods

DbMpoolFile::close()

```
#include <db_cxx.h>

int
DbMpoolFile::close(u_int32_t flags);
```

The `DbMpoolFile::close()` method closes the source file indicated by the DbMpoolFile structure. Calling `DbMpoolFile::close()` does not imply a call to DbMpoolFile::sync(); that is, no pages are written to the source file as as a result of calling `DbMpoolFile::close.()`.

If the DbMpoolFile was temporary, any underlying files created for this DbMpoolFile will be removed.

After `DbMpoolFile::close()` has been called, regardless of its return, the DbMpoolFile handle may not be accessed again.

The `DbMpoolFile::close()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter is currently unused, and must be set to 0.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods

DbMpoolFile::get()

```
#include <db_cxx.h>

int
DbMpoolFile::get(db_pgno_t *pgnoaddr,
    DbTxn *txnid, u_int32_t flags, void **pagep);
```

The `DbMpoolFile::get()` method returns pages from the cache.

All pages returned by `DbMpoolFile::get()` will be retained (that is, *latched*) in the cache until a subsequent call to DbMpoolFile::put(). There is no deadlock detection amoung latches so care must be taken in the application if the DB_MPOOL_DIRTY or DB_MPOOL_EDIT flags are used as these get exlusive latches on the pages.

The returned page is **size_t** type aligned.

Fully or partially created pages have all their bytes set to a nul byte, unless the DbMpoolFile::set_clear_len() method was called to specify other behavior before the file was opened.

The `DbMpoolFile::get()` method will return `DB_PAGE_NOTFOUND` if the requested page does not exist and DB_MPOOL_CREATE was not set. Unless otherwise specified, the `DbMpoolFile::get()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

- DB_MPOOL_CREATE

  If the specified page does not exist, create it. In this case, the pgin method, if specified, is called.

- DB_MPOOL_DIRTY

  The page will be modified and must be written to the source file before being evicted from the cache. For files open with the DB_MULTIVERSION flag set, a new copy of the page will be made if this is the first time the specified transaction is modifying it. A page fetched with the `DB_MPOOL_DIRTY` flag will be **excluively latched** until a subsequent call to DbMpoolFile::put().

- DB_MPOOL_EDIT

  The page will be modified and must be written to the source file before being evicted from the cache. No copy of the page will be made, regardless of the DB_MULTIVERSION setting. This flag is only intended for use in situations where a transaction handle is not available, such as during aborts or recovery. A page fetched with the `DB_MPOOL_EDIT` flag will be **excluively latched** until a subsequent call to DbMpoolFile::put().

- DB_MPOOL_LAST

  Return the last page of the source file, and copy its page number into the memory location to which **pgnoaddr** refers.

- DB_MPOOL_NEW

  Create a new page in the file, and copy its page number into the memory location to which **pgnoaddr** refers. In this case, the pgin_fcn callback, if specified on DbEnv::memp_register(), is **not** called.

The DB_MPOOL_CREATE, DB_MPOOL_LAST, and DB_MPOOL_NEW flags are mutually exclusive.

pagep

The **pagep** parameter references memory into which a pointer to the returned page is copied.

pgnoaddr

If the **flags** parameter is set to DB_MPOOL_LAST or DB_MPOOL_NEW, the page number of the created page is copied into the memory location to which the **pgnoaddr** parameter refers. Otherwise, the **pgnoaddr** parameter is the page to create or retrieve.

☞   Page numbers begin at 0; that is, the first page in the file is page number 0, not page number 1.

txnid

If the operation is part of an application-specified transaction, the **txnid** parameter is a transaction handle returned from DbEnv::txn_begin(); otherwise NULL. A transaction is required if the file is open for multiversion concurrency control by passing DB_MULTIVERSION to DbMpoolFile::open() and the DB_MPOOL_DIRTY, DB_MPOOL_CREATE or DB_MPOOL_NEW flags were specified. Otherwise it is ignored.

Errors

The DbMpoolFile::get() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EACCES

The DB_MPOOL_DIRTY or DB_MPOOL_EDIT flag was set and the source file was not opened for writing.

EAGAIN

The page reference count has overflowed. (This should never happen unless there is a bug in the application.)

EINVAL

If the DB_MPOOL_NEW flag was set, and the source file was not opened for writing; more than one of DB_MPOOL_CREATE, DB_MPOOL_LAST, and DB_MPOOL_NEW was set; or if an invalid flag value or parameter was specified.

DB_LOCK_DEADLOCK

For transactions configured with DB_TXN_SNAPSHOT, the page has been modified since the transaction began.

ENOMEM

The cache is full, and no more pages will fit in the cache.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods

DbMpoolFile::open()

```
#include <db_cxx.h>

int
DbMpoolFile::open(const char *file, u_int32_t flags, int mode, size_t pagesize);
```

The `DbMpoolFile::open()` method opens a file in the in-memory cache.

The `DbMpoolFile::open()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

file

The **file** parameter is the name of the file to be opened. If **file** is NULL, a private temporary file is created that cannot be shared with any other process (although it may be shared with other threads of control in the same process).

When using a Unicode build on Windows (the default), the **file** argument will be interpreted as a UTF-8 string, which is equivalent to ASCII for Latin characters.

flags

The **flags** parameter must be set to zero or by bitwise inclusively **OR**'ing together one or more of the following values:

- DB_CREATE

  Create any underlying files, as necessary. If the database do not already exist and the DB_CREATE flag is not specified, the call will fail.

- DB_DIRECT

  If set and supported by the system, turn off system buffering of the file to avoid double caching.

- DB_MULTIVERSION

  Open the file with support for multiversion concurrency control. Calls to DbMpoolFile::get() with dirty pages will cause copies to be made in the cache.

- DB_NOMMAP

  Always copy this file into the local cache instead of potentially mapping it into process memory (see the DbEnv::set_mp_mmapsize() method for further information).

- DB_ODDFILESIZE

Attempts to open files which are not a multiple of the page size in length will fail, by default. If the DB_ODDFILESIZE flag is set, any partial page at the end of the file will be ignored and the open will proceed.

- DB_RDONLY

Open any underlying files for reading only. Any attempt to modify the file using the memory pool (cache) functions will fail, regardless of the actual permissions of the file.

mode

On Windows systems, the mode parameter is ignored.

On UNIX systems or in IEEE/ANSI Std 1003.1 (POSIX) environments, files created by DbMpoolFile::open() are created with mode **mode** (as described in **chmod**(2)) and modified by the process' umask value at the time of creation (see **umask**(2)). Created files are owned by the process owner; the group ownership of created files is based on the system and directory defaults, and is not further specified by Berkeley DB. System shared memory segments created by DbMpoolFile::open() are created with mode **mode**, unmodified by the process' umask value. If **mode** is 0, DbMpoolFile::open() will use a default mode of readable and writable by both owner and group.

pagesize

The **pagesize** parameter is the size, in bytes, of the unit of transfer between the application and the cache, although it is not necessarily the unit of transfer between the cache and the underlying filesystem.

Errors

The DbMpoolFile::open() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the file has already been entered into the cache, and the **pagesize** value is not the same as when the file was entered into the cache, or the length of the file is not zero or a multiple of the **pagesize**; the DB_RDONLY flag was specified for an in-memory cache; or if an invalid flag value or parameter was specified.

ENOMEM

The maximum number of open files has been reached.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods

DbMpoolFile::put()

```
#include <db_cxx.h>

int
DbMpoolFile::put(void *pgaddr, DB_CACHE_PRIORITY priority, u_int32_t flags);
```

The `DbMpoolFile::put()` method returns a reference to a page in the cache, setting the priority of the page as specified by the **priority** parameter.

The `DbMpoolFile::put()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter is currently unused, and must be set to 0.

pgaddr

The **pgaddr** parameter is the address of the page to be returned to the cache. The **pgaddr** parameter must be a value previously returned by the DbMpoolFile::get() method.

priority

Set the page's **priority** as follows:

- `DB_PRIORITY_UNCHANGED`

  The priority is unchanged.

- `DB_PRIORITY_VERY_LOW`

  The lowest priority: pages are the most likely to be discarded.

- `DB_PRIORITY_LOW`

  The next lowest priority.

- `DB_PRIORITY_DEFAULT`

  The default priority.

- `DB_PRIORITY_HIGH`

  The next highest priority.

- `DB_PRIORITY_VERY_HIGH`

  The highest priority: pages are the least likely to be discarded.

Errors

The `DbMpoolFile::put()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods

DbMpoolFile::sync()

```
#include <db_cxx.h>

int
DbMpoolFile::sync();
```

The `DbMpoolFile::sync()` method writes all modified pages associated with the DbMpoolFile back to the source file. If any of the modified pages are *pinned* (that is, currently in use), `DbMpoolFile::sync()` will ignore them.

The `DbMpoolFile::sync()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods

DbMpoolFile::get_clear_len()

```
#include <db_cxx.h>

int
DbMpoolFile::get_clear_len(u_int32_t *lenp);
```

The `DbMpoolFile::get_clear_len()` method returns the bytes to be cleared.

The `DbMpoolFile::get_clear_len()` method may be called at any time during the life of the application.

The `DbMpoolFile::get_clear_len()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

lenp

The `DbMpoolFile::get_clear_len()` method returns the bytes to be cleared in **lenp**.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods

DbMpoolFile::get_fileid()

```
#include <db_cxx.h>

int DbMpoolFile::get_fileid(u_int8_t *fileid);
```

The `DbMpoolFile::get_fileid()` method copies the file's identifier into the memory location referenced by **fileid**. The fileid specifies a unique identifier for the file, which is used so that the cache functions (that is, the shared memory buffer pool functions) are able to uniquely identify files. This is necessary for multiple processes wanting to share a file to correctly identify the file in the cache.

The `DbMpoolFile::get_fileid()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods, DbMpoolFile::set_fileid()

DbMpoolFile::get_flags()

```
#include <db_cxx.h>

int
DbMpoolFile::get_flags(u_int32_t *flagsp);
```

The `DbMpoolFile::get_flags()` method returns the flags used to configure a file in the cache.

The `DbMpoolFile::get_flags()` method may be called at any time during the life of the application.

The `DbMpoolFile::get_flags()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flagsp

The `DbMpoolFile::get_flags()` method returns the flags in **flagsp**.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods, DbMpoolFile::set_flags()

DbMpoolFile::get_ftype()

```
#include <db_cxx.h>

int
DbMpoolFile::get_ftype(int *ftypep);
```

The `DbMpoolFile::get_ftype()` method returns the file type. The file type is used for the purposes of file processing, and will be the same as is set using the DbEnv::memp_register() method.

The `DbMpoolFile::get_ftype()` method may be called at any time during the life of the application.

The `DbMpoolFile::get_ftype()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

ftypep

The `DbMpoolFile::get_ftype()` method returns the file type in **ftypep**.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods, DbMpoolFile::set_ftype()

DbMpoolFile::get_lsn_offset()

```
#include <db_cxx.h>

int
DbMpoolFile::get_lsn_offset(int32_t *lsn_offsetp);
```

The `DbMpoolFile::get_lsn_offset()` method returns the log sequence number byte offset configured for a file's pages using the DbMpoolFile::set_lsn_offset() method.

The `DbMpoolFile::get_lsn_offset()` method may be called at any time during the life of the application.

The `DbMpoolFile::get_lsn_offset()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

lsn_offsetp

The `DbMpoolFile::get_lsn_offset()` method returns the log sequence number byte offset in **lsn_offsetp**.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods, DbMpoolFile::set_lsn_offset()

DbMpoolFile::get_maxsize()

```
#include <db_cxx.h>

int
DbMpoolFile::get_maxsize(u_int32_t *gbytesp, u_int32_t *bytesp);
```

Returns the maximum size configured for the file, as configured using the DbMpoolFile::set_maxsize() method.

The `DbMpoolFile::get_maxsize()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

The `DbMpoolFile::get_maxsize()` method may be called at any time during the life of the application.

Parameters

bytesp

The **bytesp** parameter references memory into which the additional bytes of memory in the maximum file size is copied.

gbytesp

The **gbytesp** parameter references memory into which the gigabytes of memory in the maximum file size is copied.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods, DbMpoolFile::set_maxsize()

DbMpoolFile::get_pgcookie()

```
#include <db_cxx.h>

int
DbMpoolFile::get_pgcookie(DBT *dbt);
```

The `DbMpoolFile::get_pgcookie()` method returns the byte string provided to the functions registered to do input or output processing of the file's pages as they are read from or written to, the backing filesystem store. This byte string is configured using the DbMpoolFile::set_pgcookie() method.

The `DbMpoolFile::get_pgcookie()` method may be called at any time during the life of the application.

The `DbMpoolFile::get_pgcookie()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

dbt

The `DbMpoolFile::get_pgcookie()` method returns a reference to the byte string in **dbt**.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods, DbMpoolFile::set_pgcookie()

DbMpoolFile::get_priority()

```
#include <db_cxx.h>

int
DbMpoolFile::get_priority(DB_CACHE_PRIORITY *priorityp);
```

The `DbMpoolFile::get_priority()` method returns the cache priority for the file referenced by the DbMpoolFile handle. The priority of a page biases the replacement algorithm to be more or less likely to discard a page when space is needed in the cache. This value is set using the DbMpoolFile::set_priority() method.

The `DbMpoolFile::get_priority()` method may be called at any time during the life of the application.

The `DbMpoolFile::get_priority()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

priorityp

The `DbMpoolFile::get_priority()` method returns a reference to the cache priority for the file referenced by the DbMpoolFile handle in **priorityp**.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods, DbMpoolFile::set_priority()

DbMpoolFile::set_clear_len()

```
#include <db_cxx.h>

int DbMpoolFile::set_clear_len(u_int32_t len);
```

The `DbMpoolFile::set_clear_len()` method sets the number of initial bytes in a page that should be set to nul when the page is created as a result of the DB_MPOOL_CREATE or DB_MPOOL_NEW flags specified to DbMpoolFile::get(). If no clear length is specified, the entire page is cleared when it is created.

The `DbMpoolFile::set_clear_len()` method configures a file in the cache, not only operations performed using the specified DbMpoolFile handle.

The `DbMpoolFile::set_clear_len()` method may not be called after the DbMpoolFile::open() method is called. If the file is already open in the cache when DbMpoolFile::open() is called, the information specified to `DbMpoolFile::set_clear_len()` must be consistent with the existing file or an error will be returned.

The `DbMpoolFile::set_clear_len()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

len

The **len** parameter is the number of initial bytes in a page that should be set to nul when the page is created. A value of 0 results in the entire page being set to nul bytes.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods

DbMpoolFile::set_fileid()

```
#include <db_cxx.h>

int
DbMpoolFile::set_fileid(u_int8_t *fileid);
```

The `DbMpoolFile::set_fileid()` method specifies a unique identifier for the file. (The shared memory buffer pool functions must be able to uniquely identify files in order that multiple processes wanting to share a file will correctly identify it in the cache.)

On most UNIX/POSIX systems, the **fileid** field will not need to be set, and the memory pool functions will use the file's device and inode numbers for this purpose. On Windows systems, the memory pool functions use the values returned by `GetFileInformationByHandle()` by default — these values are known to be constant between processes and over reboot in the case of NTFS (in which they are the NTFS MFT indices).

On other filesystems (for example, FAT or NFS), these default values are not necessarily unique between processes or across system reboots. **Applications wanting to maintain a shared cache between processes or across system reboots, in which the cache contains pages from files stored on such filesystems, must specify a unique file identifier using the `DbMpoolFile::set_fileid()` method, and each process opening the file must provide the same unique identifier.**

This call should not be necessary for most applications. Specifically, it is not necessary if the cache is not shared between processes and is reinstantiated after each system reboot, if the application is using the Berkeley DB access methods instead of calling the pool functions explicitly, or if the files in the cache are stored on filesystems in which the default values as described previously are invariant between process and across system reboots.

The `DbMpoolFile::set_fileid()` method configures a file in the cache, not only operations performed using the specified DbMpoolFile handle.

The `DbMpoolFile::set_fileid()` method may not be called after the DbMpoolFile::open() method is called.

The `DbMpoolFile::set_fileid()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

fileid

The **fileid** parameter is the unique identifier for the file. Unique file identifiers must be a `DB_FILE_ID_LEN` length array of bytes.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods

DbMpoolFile::set_flags()

```
#include <db_cxx.h>

int
DbMpoolFile::set_flags(u_int32_t flags, bool onoff);
```

Configure a file in the cache.

To set the flags for a particular database, call the `DbMpoolFile::set_flags()` method using the DbMpoolFile handle stored in the **mpf** field of the Db handle.

The `DbMpoolFile::set_flags()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set by bitwise inclusively **OR**'ing together one or more of the following values:

- `DB_MPOOL_NOFILE`

  If set, no backing temporary file will be opened for the specified in-memory database, even if it expands to fill the entire cache. Attempts to create new database pages after the cache has been filled will fail.

  The `DB_MPOOL_NOFILE` flag configures a file in the cache, not only operations performed using the specified DbMpoolFile handle.

  The `DB_MPOOL_NOFILE` flag may be used to configure Berkeley DB at any time during the life of the application.

- `DB_MPOOL_UNLINK`

  If set, remove the file when the last reference to it is closed.

  The `DB_MPOOL_ULINK` flag configures a file in the cache, not only operations performed using the specified DbMpoolFile handle.

  The `DB_MPOOL_ULINK` flag may be used to configure Berkeley DB at any time during the life of the application.

onoff

If **onoff** is zero, the specified flags are cleared; otherwise they are set.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods

DbMpoolFile::set_ftype()

```
#include <db_cxx.h>

int
DbMpoolFile::set_flags(int ftype);
```

The `DbMpoolFile::set_ftype()` method specifies a file type for the purposes of input or output processing of the file's pages as they are read from or written to, the backing filesystem store.

The `DbMpoolFile::set_ftype()` method configures a file in the cache, not only operations performed using the specified DbMpoolFile handle.

The `DbMpoolFile::set_ftype()` method may not be called after the DbMpoolFile::open() method is called. If the file is already open in the cache when DbMpoolFile::open() is called, the information specified to `DbMpoolFile::set_ftype()` will replace the existing information.

The `DbMpoolFile::set_ftype()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

ftype

The **ftype** parameter sets the file's type for the purposes of input and output processing. The **ftype** must be the same as a **ftype** parameter previously specified to the DbEnv::memp_register() method.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods

DbMpoolFile::set_lsn_offset()

```
#include <db_cxx.h>

int DbMpoolFile::set_lsn_offset(int32_t lsn_offset);
```

The `DbMpoolFile::set_lsn_offset()` method specifies the zero-based byte offset of a log sequence number (DbLsn) on the file's pages, for the purposes of page-flushing as part of transaction checkpoint. (See the DbEnv::memp_sync() documentation for more information.)

The `DbMpoolFile::set_lsn_offset()` method configures a file in the cache, not only operations performed using the specified DbMpoolFile handle.

The `DbMpoolFile::set_lsn_offset()` method may not be called after the DbMpoolFile::open() method is called. If the file is already open in the cache when DbMpoolFile::open() is called, the information specified to `DbMpoolFile::set_lsn_offset()` must be consistent with the existing file or an error will be returned.

The `DbMpoolFile::set_lsn_offset()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

lsn_offset

The **lsn_offset** parameter is the zero-based byte offset of the log sequence number on the file's pages.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods

DbMpoolFile::set_maxsize()

```
#include <db_cxx.h>

int
DbMpoolFile::set_maxsize(u_int32_t gbytes, u_int32_t bytes);
```

Set the maximum size for the file to be **gbytes** gigabytes plus **bytes**. Attempts to allocate new pages in the file after the limit has been reached will fail.

To set the maximum file size for a particular database, call the `DbMpoolFile::set_maxsize()` method using the DbMpoolFile handle stored in the **mpf** field of the Db handle. Attempts to insert new items into the database after the limit has been reached may fail.

The `DbMpoolFile::set_maxsize()` method configures a file in the cache, not only operations performed using the specified DbMpoolFile handle.

The `DbMpoolFile::set_maxsize()` method may be called at any time during the life of the application.

The `DbMpoolFile::set_maxsize()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

bytes

The maximum size of the file is set to **gbytes** gigabytes plus **bytes**.

gbytes

The maximum size of the file is set to **gbytes** gigabytes plus **bytes**.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods

DbMpoolFile::set_pgcookie()

```
#include <db_cxx.h>

int
DbMpoolFile::set_pgcookie(DBT *pgcookie);
```

The `DbMpoolFile::set_pgcookie()` method specifies a byte string that is provided to the functions registered to do input or output processing of the file's pages as they are read from or written to, the backing filesystem store. (See the DbEnv::memp_register() documentation for more information.)

The `DbMpoolFile::set_pgcookie()` method configures a file in the cache, not only operations performed using the specified DbMpoolFile handle.

The `DbMpoolFile::set_pgcookie()` method may not be called after the DbMpoolFile::open() method is called. If the file is already open in the cache when DbMpoolFile::open() is called, the information specified to `DbMpoolFile::set_pgcookie()` will replace the existing information.

The `DbMpoolFile::set_pgcookie()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

pgcookie

The **pgcookie** parameter is a byte string provided to the functions registered to do input or output processing of the file's pages.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods

DbMpoolFile::set_priority()

```
#include <db_cxx.h>

int
DbMpoolFile::set_priority(DB_CACHE_PRIORITY priority);
```

Set the cache priority for pages referenced by the DbMpoolFile handle.

The priority of a page biases the replacement algorithm to be more or less likely to discard a page when space is needed in the cache. The bias is temporary, and pages will eventually be discarded if they are not referenced again. The DbMpoolFile::set_priority() method is only advisory, and does not guarantee pages will be treated in a specific way.

To set the priority for the pages belonging to a particular database, call the DbMpoolFile::set_priority() method using the DbMpoolFile handle returned by the Db::get_mpf() method.

The DbMpoolFile::set_priority() method configures a file in the cache, not only operations performed using the specified DbMpoolFile handle.

The DbMpoolFile::set_priority() method may be called at any time during the life of the application.

The DbMpoolFile::set_priority() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

priority

The **priority** parameter must be set to one of the following values:

- DB_PRIORITY_VERY_LOW

  The lowest priority: pages are the most likely to be discarded.

- DB_PRIORITY_LOW

  The next lowest priority.

- DB_PRIORITY_DEFAULT

  The default priority.

- DB_PRIORITY_HIGH

  The next highest priority.

- DB_PRIORITY_VERY_HIGH

  The highest priority: pages are the least likely to be discarded.

Class

DbEnv, DbMpoolFile

See Also

Memory Pools and Related Methods

Chapter 10. Mutex Methods

This chapter describes methods that can be used to manage mutexes within DB. Many of the methods described here are used to configure DB's internal mutex system. However, a series of APIs are available for use as a general-purpose, cross platform mutex management system. These methods can be used independently of DB's main purpose, which is as a high-end data management engine.

Mutex Methods

Mutexes and Related Methods	Description
DbEnv::mutex_alloc()	Allocate a mutex
DbEnv::mutex_free()	Free a mutex
DbEnv::mutex_lock()	Lock a mutex
DbEnv::mutex_stat()	Mutex statistics
DbEnv::mutex_stat_print()	Print mutex statistics
DbEnv::mutex_unlock()	Unlock a mutex
Mutex Configuration	
DbEnv::mutex_set_align(), DbEnv::mutex_get_align()	Configure mutex alignment
DbEnv::mutex_set_increment(), DbEnv::mutex_get_increment()	Configure number of additional mutexes
DbEnv::mutex_set_max(), DbEnv::mutex_get_max()	Configure total number of mutexes
DbEnv::mutex_set_tas_spins(), DbEnv::mutex_get_tas_spins()	Configure test-and-set mutex spin count

DbEnv::mutex_alloc()

```
#include <db_cxx.h>

int
DbEnv::mutex_alloc(u_int32_t flags, db_mutex_t *mutexp);
```

The `DbEnv::mutex_alloc()` method allocates a mutex and returns a reference to it into the memory specified by **mutexp**.

The `DbEnv::mutex_alloc()` method may not be called before the DbEnv::open() method is called.

The `DbEnv::mutex_alloc()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

- `DB_MUTEX_PROCESS_ONLY`

  The mutex is associated with a single process. The DbEnv::failchk() method will release mutexes held by any process which has exited.

- `DB_MUTEX_SELF_BLOCK`

  The mutex must be self-blocking. That is, if a thread of control locks the mutex and then attempts to lock the mutex again, the thread of control will block until another thread of control releases the original lock on the mutex, allowing the original thread of control to lock the mutex the second time. Attempting to re-acquire a mutex for which the `DB_MUTEX_SELF_BLOCK` flag was not specified will result in undefined behavior.

mutexp

The **mutexp** parameter references memory into which the mutex reference is copied.

Errors

The `DbEnv::mutex_alloc()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv

See Also

Mutex Methods

DbEnv::mutex_free()

```
#include <db_cxx.h>

int
DbEnv::mutex_free(db_mutex_t mutex);
```

The `DbEnv::mutex_free()` method discards a mutex allocated by DbEnv::mutex_alloc().

The `DbEnv::mutex_free()` method may not be called before the DbEnv::open() method is called.

The `DbEnv::mutex_free()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

mutex

The **mutex** parameter is a mutex previously allocated by DbEnv::mutex_alloc().

Errors

The `DbEnv::mutex_free()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv

See Also

Mutex Methods

DbEnv::mutex_get_align()

```
#include <db_cxx.h>

int
DbEnv::mutex_get_align(u_int32_t *alignp);
```

The `DbEnv::mutex_get_align()` method returns the mutex alignment, in bytes.

The `DbEnv::mutex_get_align()` method may be called at any time during the life of the application.

The `DbEnv::mutex_get_align()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

alignp

The `DbEnv::mutex_get_align()` method returns the mutex alignment, in bytes in **alignp**.

Class

DbEnv

See Also

Mutex Methods

DbEnv::mutex_get_increment()

```
#include <db_cxx.h>

int
DbEnv::mutex_get_increment(u_int32_t *incrementp);
```

The `DbEnv::mutex_get_increment()` method returns the number of additional mutexes to allocate.

The `DbEnv::mutex_get_increment()` method may be called at any time during the life of the application.

The `DbEnv::mutex_get_increment()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

incrementp

The `DbEnv::mutex_get_increment()` method returns the number of additional mutexes to allocate in incrementp.

Class

DbEnv

See Also

Mutex Methods

DbEnv::mutex_get_max()

```
#include <db_cxx.h>

int
DbEnv::mutex_get_max(u_int32_t *maxp);
```

The `DbEnv::mutex_get_max()` method returns the total number of mutexes allocated.

The `DbEnv::mutex_get_max()` method may be called at any time during the life of the application.

The `DbEnv::mutex_get_max()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

maxp

The `DbEnv::mutex_get_max()` method returns the total number of mutexes allocated in **maxp**.

Class

[DbEnv](#)

See Also

[Mutex Methods](#)

DbEnv::mutex_get_tas_spins()

```
#include <db_cxx.h>

int
DbEnv::mutex_get_tas_spins(u_int32_t *, tas_spinsp);
```

The `DbEnv::mutex_get_tas_spins()` method returns the test-and-set spin count. This value may be configured using the DbEnv::mutex_set_tas_spins() method.

The `DbEnv::mutex_get_tas_spins()` method may be called at any time during the life of the application.

The `DbEnv::mutex_get_tas_spins()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

tas_spinsp

The `DbEnv::mutex_get_tas_spins()` method returns the test-and-set spin count in **tas_spinsp**.

Class

DbEnv

See Also

Mutex Methods

DbEnv::mutex_lock()

```
#include <db_cxx.h>

int
DbEnv::mutex_lock(db_mutex_t mutex);
```

The `DbEnv::mutex_lock()` method locks the mutex allocated by DbEnv::mutex_alloc(). The thread of control calling `DbEnv::mutex_lock()` will block until the lock is available.

The `DbEnv::mutex_lock()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

mutex

The **mutex** parameter is a mutex previously allocated by DbEnv::mutex_alloc().

Errors

The `DbEnv::mutex_lock()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv

See Also

Mutex Methods

DbEnv::mutex_set_align()

```
#include <db_cxx.h>

int
DbEnv::mutex_set_align(u_int32_t align);
```

Set the mutex alignment, in bytes.

It is sometimes advantageous to align mutexes on specific byte boundaries in order to minimize cache line collisions. The `DbEnv::mutex_set_align()` method specifies an alignment for mutexes allocated by Berkeley DB.

The database environment's mutex alignment may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "mutex_set_align", one or more whitespace characters, and the mutex alignment in bytes. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The `DbEnv::mutex_set_align()` method configures a database environment, not only operations performed using the specified DbEnv handle.

The `DbEnv::mutex_set_align()` method may not be called after the DbEnv::open() method is called. If the database environment already exists when DbEnv::open() is called, the information specified to `DbEnv::mutex_set_align()` will be ignored.

The `DbEnv::mutex_set_align()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

align

The **align** parameter is the mutex alignment, in bytes. The mutex alignment must be a power-of-two.

Errors

The `DbEnv::mutex_set_align()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after DbEnv::open() was called; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Mutex Methods

DbEnv::mutex_set_increment()

```
#include <db_cxx.h>

int
DbEnv::mutex_set_increment(u_int32_t increment);
```

Configure the number of additional mutexes to allocate.

If an application will allocate mutexes for its own use, the `DbEnv::mutex_set_increment()` method is used to add a number of mutexes to the default allocation.

Calling the `DbEnv::mutex_set_increment()` method discards any value previously set using the `DbEnv::mutex_set_max()` method.

The database environment's number of additional mutexes may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "mutex_set_increment", one or more whitespace characters, and the number of additional mutexes. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The `DbEnv::mutex_set_increment()` method configures a database environment, not only operations performed using the specified DbEnv handle.

The `DbEnv::mutex_set_increment()` method may not be called after the DbEnv::open() method is called. If the database environment already exists when DbEnv::open() is called, the information specified to `DbEnv::mutex_set_increment()` will be ignored.

The `DbEnv::mutex_set_increment()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

increment

The **increment** parameter is the number of additional mutexes to allocate.

Errors

The `DbEnv::mutex_set_increment()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after DbEnv::open() was called; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Mutex Methods

DbEnv::mutex_set_max()

```
#include <db_cxx.h>

int
DbEnv::mutex_set_max(u_int32_t max);
```

Configure the total number of mutexes to allocate.

Berkeley DB allocates a default number of mutexes based on the initial configuration of the database environment. That default calculation may be too small if the application has an unusual need for mutexes (for example, if the application opens an unexpectedly large number of databases) or too large (if the application is trying to minimize its memory footprint). The DbEnv::mutex_set_max() method is used to specify an absolute number of mutexes to allocate.

Calling the DbEnv::mutex_set_max() method discards any value previously set using the DbEnv::mutex_set_increment() method.

The database environment's total number of mutexes may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "mutex_set_max", one or more whitespace characters, and the total number of mutexes. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The DbEnv::mutex_set_max() method configures a database environment, not only operations performed using the specified DbEnv handle.

The DbEnv::mutex_set_max() method may not be called after the DbEnv::open() method is called. If the database environment already exists when DbEnv::open() is called, the information specified to DbEnv::mutex_set_max() will be ignored.

The DbEnv::mutex_set_max() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

max

The **max** parameter is the absolute number of mutexes to allocate.

Errors

The DbEnv::mutex_set_max() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv

See Also

Mutex Methods

DbEnv::mutex_set_tas_spins()

```
#include <db_cxx.h>

int
DbEnv::mutex_set_tas_spins(u_int32_t tas_spins);
```

Specify that test-and-set mutexes should spin **tas_spins** times without blocking. The value defaults to 1 on uniprocessor systems and to 50 times the number of processors on multiprocessor systems.

The database environment's test-and-set spin count may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_tas_spins", one or more whitespace characters, and the number of spins. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The DbEnv::mutex_set_tas_spins() method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The DbEnv::mutex_set_tas_spins() method may be called at any time during the life of the application.

The DbEnv::mutex_set_tas_spins() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

tas_spins

The **tas_spins** parameter is the number of spins test-and-set mutexes should execute before blocking.

Errors

The DbEnv::mutex_set_tas_spins() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv

See Also

Mutex Methods

DbEnv::mutex_stat()

```
#include <db_cxx.h>

int
DbEnv::mutex_stat(DB_MUTEX_STAT **statp, u_int32_t flags);
```

The `DbEnv::mutex_stat()` method returns the mutex subsystem statistics.

The `DbEnv::mutex_stat()` method creates a statistical structure of type `DB_MUTEX_STAT` and copies a pointer to it into a user-specified memory location.

Statistical structures are stored in allocated memory. If application-specific allocation routines have been declared (see DbEnv::set_alloc() for more information), they are used to allocate the memory; otherwise, the standard C library **malloc**(3) is used. The caller is responsible for deallocating the memory. To deallocate the memory, free the memory reference; references inside the returned memory need not be individually freed.

The following `DB_MUTEX_STAT` fields will be filled in:

- **u_int32_t st_mutex_align;**

  The mutex alignment, in bytes.

- **u_int32_t st_mutex_tas_spins;**

  The number of times test-and-set mutexes will spin without blocking.

- **int st_mutex_cnt;**

  The total number of mutexes configured.

- **u_int32_t st_mutex_free;**

  The number of mutexes currently available.

- **u_int32_t st_mutex_inuse;**

  The number of mutexes currently in use.

- **u_int32_t st_mutex_inuse_max;**

  The maximum number of mutexes ever in use.

- **roff_t st_regsize;**

  The size of the mutex region, in bytes.

- **uintmax_t st_region_wait;**

  The number of times that a thread of control was forced to wait before obtaining the mutex region mutex.

- **uintmax_t st_region_nowait;**

    The number of times that a thread of control was able to obtain the mutex region mutex without waiting.

The `DbEnv::mutex_stat()` method may not be called before the DbEnv::open() method is called.

The `DbEnv::mutex_stat()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or the following value:

- `DB_STAT_CLEAR`

    Reset statistics after returning their values.

statp

The **statp** parameter references memory into which a pointer to the allocated statistics structure is copied.

Errors

The `DbEnv::mutex_stat()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv

See Also

Mutex Methods

DbEnv::mutex_stat_print()

```
#include <db_cxx.h>

int
DbEnv::mutex_stat_print(u_int32_t flags);
```

The `DbEnv::mutex_stat_print()` method displays the mutex subsystem statistical information, as described for the `DbEnv::mutex_stat()` method. The information is printed to a specified output channel (see the DbEnv::set_msgfile() method for more information), or passed to an application callback function (see the DbEnv::set_msgcall() method for more information).

The `DbEnv::mutex_stat_print()` method may not be called before the DbEnv::open() method is called.

The `DbEnv::mutex_stat_print()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

- `DB_STAT_ALL`

  Display all available information.

- `DB_STAT_CLEAR`

  Reset statistics after displaying their values.

Class

DbEnv

See Also

Mutex Methods

DbEnv::mutex_unlock()

```
#include <db_cxx.h>

int
DbEnv::mutex_unlock(db_mutex_t mutex);
```

The `DbEnv::mutex_unlock()` method unlocks the mutex locked by DbEnv::mutex_lock().

The `DbEnv::mutex_unlock()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

mutex

The **mutex** parameter is a mutex previously locked by DbEnv::mutex_lock().

Errors

The `DbEnv::mutex_unlock()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv

See Also

Mutex Methods

Chapter 11. Replication Methods

This chapter describes the APIs available to build Berkeley DB replicated applications. There are two different ways to build replication into a Berkeley DB application, and the APIs for both are described in this chapter.

For an overview of the two different ways to build a replicated application, see the *Berkeley DB Getting Started with Replicated Applications* guide.

The first, and simplest, way to build a replication Berkeley DB application is via the *Replication Manager*. If the Replication Manager does not meet your application's architectural requirements, they you can write your own replication implementation using the "Base APIs".

Note that the Replication Manager is written using the Base APIs.

Note, also, that applications which make use of the Replication Manager use many of the Base APIs as the situation warrants. That said, a few Base API methods cannot be used by applications that are making use of the Replication Manager. Where this is the case, this is noted in the following method descriptions.

Replication and Related Methods

Replication Manager Methods	Description
DbEnv::repmgr_add_remote_site()	Specify the Replication Manager's remote sites
DbEnv::repmgr_set_ack_policy(), DbEnv::repmgr_get_ack_policy()	Specify the Replication Manager's client acknowledgement policy
DbEnv::repmgr_set_local_site()	Specify the Replication Manager's local site
DbEnv::repmgr_site_list()	List the sites and their status
DbEnv::repmgr_start()	Start the Replication Manager
DbEnv::repmgr_stat()	Replication Manager statistics
DbEnv::repmgr_stat_print()	Print Replication Manager statistics
Base API Methods	
DbEnv::rep_elect()	Hold a replication election
DbEnv::rep_process_message()	Process a replication message
DbEnv::rep_set_transport()	Configure replication transport callback
DbEnv::rep_start()	Start replication
Additional Replication Methods	
DbEnv::rep_stat()	Replication statistics
DbEnv::rep_stat_print()	Print replication statistics
DbEnv::rep_sync()	Replication synchronization
Replication Configuration	
DbEnv::rep_set_clockskew(), DbEnv::rep_get_clockskew()	Configure master lease clock adjustment
DbEnv::rep_set_config(), DbEnv::rep_get_config()	Configure the replication subsystem
DbEnv::rep_set_limit(), DbEnv::rep_get_limit()	Limit data sent in response to a single message
DbEnv::rep_set_nsites(), DbEnv::rep_get_nsites()	Configure replication group site count
DbEnv::rep_set_priority(), DbEnv::rep_get_priority()	Configure replication site priority
DbEnv::rep_set_request(), DbEnv::rep_get_request()	Configure replication client retransmission requests
DbEnv::rep_set_timeout(), DbEnv::rep_get_timeout()	Configure replication timeouts

DbEnv::rep_elect()

```
#include <db_cxx.h>

int
DbEnv::rep_elect(u_int32_t nsites, u_int32_t nvotes, u_int32_t flags);
```

The `DbEnv::rep_elect()` method holds an election for the master of a replication group.

The `DbEnv::rep_elect()` method is not called by most replication applications. It should only be called by Base API applications implementing their own network transport layer, explicitly holding replication group elections and handling replication messages outside of the Replication Manager framework.

If the election is successful, Berkeley DB will notify the application of the results of the election by means of either the DB_EVENT_REP_ELECTED or DB_EVENT_REP_NEWMASTER events (see DbEnv::set_event_notify() method for more information). The application is responsible for adjusting its relationship to the other database environments in the replication group, including directing all database updates to the newly selected master, in accordance with the results of the election.

The thread of control that calls the `DbEnv::rep_elect()` method must not be the thread of control that processes incoming messages; processing the incoming messages is necessary to successfully complete an election.

Before calling this method, the enclosing database environment must already have been opened by calling the DbEnv::open() method and must already have been configured to send replication messages by calling the DbEnv::rep_set_transport() method.

How Elections are Held

Elections are done in two parts: first, replication sites collect information from the other replication sites they know about, and second, replication sites cast their votes for a new master. The second phase is triggered by one of two things: either the replication site gets election information from **nsites** sites, or the election timeout expires. Once the second phase is triggered, the replication site will cast a vote for the new master of its choice if, and only if, the site has election information from at least **nvotes** sites. If a site receives **nvotes** votes for it to become the new master, then it will become the new master.

We recommend **nvotes** be set to at least:

```
        (sites participating in the election / 2) + 1
```

to ensure there are never more than two masters active at the same time even in the case of a network partition. When a network partitions, the side of the partition with more than half the environments will elect a new master and continue, while the environments communicating with fewer than half of the environments will fail to find a new master, as no site can get **nvotes** votes.

We recommend **nsites** be set to:

```
        number of sites in the replication group - 1
```

when choosing a new master after a current master fails. This allows the group to reach a consensus without having to wait for the timeout to expire.

When choosing a master from among a group of client sites all restarting at the same time, it makes more sense to set **nsites** to the total number of sites in the group, since there is no known missing site. Furthermore, in order to ensure the best choice from among sites that may take longer to boot than the local site, setting **nvotes** also to this same total number of sites will guarantee that every site in the group is considered. Alternatively, using the special timeout for full elections allows full participation on restart but allows election of a master if one site does not reboot and rejoin the group in a reasonable amount of time. (See the Elections section in the *Berkeley DB Programmer's Reference Guide* for more information.)

Setting **nsites** to lower values can increase the speed of an election, but can also result in election failure, and is usually not recommended.

Parameters

flags

The **flags** parameter is currently unused, and must be set to 0.

nsites

The **nsites** parameter specifies the number of replication sites expected to participate in the election. Once the current site has election information from that many sites, it will short-circuit the election and immediately cast its vote for a new master. The **nsites** parameter must be no less than **nvotes**, or 0 if the election should use the value previously set using the DbEnv::rep_set_nsites() method. If an application is using master leases, then the value **must** be 0 and the value from DbEnv::rep_set_nsites() method must be used.

nvotes

The **nvotes** parameter specifies the minimum number of replication sites from which the current site must have election information, before the current site will cast a vote for a new master. The **nvotes** parameter must be no greater than **nsites**, or 0 if the election should use the value ((**nsites** / 2) + 1) as the **nvotes** argument.

Errors

The DbEnv::rep_elect() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DB_REP_UNAVAIL

The replication group was unable to elect a master, or was unable to complete the election in the election timeout period (see DbEnv::rep_set_timeout() method for more information).

EINVAL

If the database environment was not already configured to communicate with a replication group by a call to DbEnv::rep_set_transport(); if the database environment was not already opened; if this

method is called from a Replication Manager application; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Replication and Related Methods

DbEnv::rep_get_clockskew()

```
#include <db_cxx.h>

DbEnv::rep_get_clockskew(u_int32_t *fast_clockp, u_in32_t *slow_clockp);
```

The `DbEnv::rep_get_clockskew()` method returns the current clock skew ratio values, as set by the DbEnv::rep_set_clockskew() method.

The `DbEnv::rep_get_clockskew()` method may be called at any time during the life of the application.

The `DbEnv::rep_get_clockskew()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

fast_clockp

The **fast_clockp** parameter references memory into which the value for the fastest clock in the group of sites is copied.

slow_clockp

The **slow_clockp** parameter references memory into which the value for the slowest clock in the group of sites is copied.

Class

DbEnv

See Also

Replication and Related Methods, DbEnv::rep_set_clockskew()

DbEnv::rep_get_config()

```
#include <db_cxx.h>

int
DbEnv::rep_get_config(u_int32_t which, int *onoffp);
```

The `DbEnv::rep_get_config()` method returns whether the specified **which** parameter is currently set or not. See the DbEnv::rep_set_config() method for the configuration flags that can be set for replication.

The `DbEnv::rep_get_config()` method may be called at any time during the life of the application.

The `DbEnv::rep_get_config()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

onoffp

The **onoffp** parameter references memory into which the configuration of the specified **which** parameter is copied.

If the returned **onoff** value is zero, the parameter is off; otherwise it is on.

which

The **which** parameter is the configuration flag which is being checked. See the DbEnv::rep_set_config() method for a list of configuration flags that you can provide to this parameter.

Class

DbEnv

See Also

Replication and Related Methods, DbEnv::rep_set_config()

DbEnv::rep_get_limit()

```
#include <db_cxx.h>

int
DbEnv::rep_get_limit(u_int32_t *gbytesp, u_int32_t *bytesp);
```

The `DbEnv::rep_get_limit()` method returns the byte-count limit on the amount of data that will be transmitted from a site in response to a single message processed by the DbEnv::rep_process_message() method. This value is configurable using the DbEnv::rep_set_limit() method.

The `DbEnv::rep_get_limit()` method may be called at any time during the life of the application.

The `DbEnv::rep_get_limit()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

bytesp

The **bytesp** parameter references memory into which the bytes component of the current transmission limit is copied.

gbytesp

The **gbytesp** parameter references memory into which the gigabytes component of the current transmission limit is copied.

Class

DbEnv

See Also

Replication and Related Methods, DbEnv::rep_set_limit()

DbEnv::rep_get_nsites()

```
#include <db_cxx.h>

int
DbEnv::rep_get_nsites(u_int32_t *nsitesp);
```

The `DbEnv::rep_get_nsites()` method returns the total number of sites in the replication group. This value is configurable using the DbEnv::rep_set_nsites() method.

The `DbEnv::rep_get_nsites()` method may be called at any time during the life of the application.

The `DbEnv::rep_get_nsites()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

nsitesp

The `DbEnv::rep_get_nsites()` method returns the total number of sites in the replication group in nsitesp.

Class

DbEnv

See Also

Replication and Related Methods, DbEnv::rep_set_nsites()

DbEnv::rep_get_priority()

```
#include <db_cxx.h>

int
DbEnv::rep_get_priority(u_int32_t *priorityp);
```

The `DbEnv::rep_get_priority()` method returns the database environment priority as configured using the DbEnv::rep_set_priority() method.

The `DbEnv::rep_get_priority()` method may be called at any time during the life of the application.

The `DbEnv::rep_get_priority()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

priorityp

The `DbEnv::rep_get_priority()` method returns the database environment priority in **priorityp**.

Class

DbEnv

See Also

Replication and Related Methods, DbEnv::rep_set_priority()

DbEnv::rep_get_request()

```
#include <db_cxx.h>

int
DbEnv::rep_get_request(u_int32_t *minp, u_int32_t *maxp);
```

The `DbEnv::rep_get_request()` method returns the minimum and maximum number of microseconds a client waits before requesting retransmission. These values can be configured using the DbEnv::rep_set_request() method.

The `DbEnv::rep_get_request()` method may be called at any time during the life of the application.

The `DbEnv::rep_get_request()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

minp

The **minp** parameter references memory into which the minimum number of microseconds a client will wait before requesting retransmission is copied.

maxp

The **maxp** parameter references memory into which the maximum number of microseconds a client will wait before requesting retransmission is copied.

Class

DbEnv

See Also

Replication and Related Methods, DbEnv::rep_set_request()

DbEnv::rep_get_timeout()

```
#include <db_cxx.h>

int
DbEnv::rep_get_timeout(int which, u_int32_t *timeoutp);
```

The `DbEnv::rep_get_timeout()` method returns the timeout value for the specified **which** parameter. Timeout values can be managed using the DbEnv::rep_set_timeout() method.

The `DbEnv::rep_get_timeout()` method may be called at any time during the life of the application.

The `DbEnv::rep_get_timeout()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

timeoutp

The **timeoutp** parameter references memory into which the timeout value of the specified **which** parameter is copied.

The returned timeout value is in microseconds.

which

The **which** parameter is the timeout for which the value is being returned. See the DbEnv::rep_set_timeout() method for a list of timeouts that you can provide to this parameter.

Errors

The `DbEnv::rep_get_timeout()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv

See Also

Replication and Related Methods, DbEnv::rep_set_timeout()

DbEnv::rep_process_message()

```
#include <db_cxx.h>

int
DbEnv::rep_process_message(Dbt *control, Dbt *rec, int envid, DbLsn *ret_lsnp)
```

The `DbEnv::rep_process_message()` method processes an incoming replication message sent by a member of the replication group to the local database environment.

The `DbEnv::rep_process_message()` method is not called by most replication applications. It should only be called by Base API applications implementing their own network transport layer, explicitly holding replication group elections and handling replication messages outside of the Replication Manager framework.

For implementation reasons, all incoming replication messages must be processed using the same DbEnv handle. It is not required that a single thread of control process all messages, only that all threads of control processing messages use the same handle.

Before calling this method, the enclosing database environment must already have been opened by calling the DbEnv::open() method and must already have been configured to send replication messages by calling the DbEnv::rep_set_transport() method.

The `DbEnv::rep_process_message()` method has additional return values:

- **DB_REP_DUPMASTER**

    The `DbEnv::rep_process_message()` method will return `DB_REP_DUPMASTER` if the replication group has more than one master. The application should reconfigure itself as a client by calling the DbEnv::rep_start() method, and then call for an election by calling DbEnv::rep_elect().

- **DB_REP_HOLDELECTION**

    The `DbEnv::rep_process_message()` method will return `DB_REP_HOLDELECTION` if an election is needed. The application should call for an election by calling DbEnv::rep_elect().

- **DB_REP_IGNORE**

    The `DbEnv::rep_process_message()` method will return `DB_REP_IGNORE` if this message cannot be processed. This is an indication that this message is irrelevant to the current replication state (for example, an old message from a previous generation arrives and is processed late).

- **DB_REP_ISPERM**

    The `DbEnv::rep_process_message()` method will return `DB_REP_ISPERM` if processing this message results in the processing of records that are permanent. The maximum LSN of the permanent records stored is returned.

- **DB_REP_JOIN_FAILURE**

The `DbEnv::rep_process_message()` method will return `DB_REP_JOIN_FAILURE` if a new master has been chosen but the client is unable to synchronize with the new master (possibly because the client has been configured with the DB_REP_CONF_NOAUTOINIT flag to turn off automatic internal initialization).

- **DB_REP_NEWSITE**

    The `DbEnv::rep_process_message()` method will return `DB_REP_NEWSITE` if the system received contact information from a new environment. The **rec** parameter contains the opaque data specified in the **cdata** parameter to the DbEnv::rep_start(). The application should take whatever action is needed to establish a communication channel with this new environment.

- **DB_REP_NOTPERM**

    The `DbEnv::rep_process_message()` method will return `DB_REP_NOTPERM` if a message carrying a DB_REP_PERMANENT flag was processed successfully, but was not written to disk. The LSN of this record is returned. The application should take whatever action is deemed necessary to retain its recoverability characteristics.

Unless otherwise specified, the `DbEnv::rep_process_message()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

control

The **control** parameter should reference a copy of the **control** parameter specified by Berkeley DB on the sending environment. See the DbEnv::rep_set_transport() method for more information.

envid

The **envid** parameter should contain the local identifier that corresponds to the environment that sent the message to be processed (see Replication environment IDs for more information).

rec

The **rec** parameter should reference a copy of the **rec** parameter specified by Berkeley DB on the sending environment. See the DbEnv::rep_set_transport() method for more information.

ret_lsnp

If `DbEnv::rep_process_message()` method returns `DB_REP_NOTPERM` then the **ret_lsnp** parameter will contain the log sequence number of this permanent log message that could not be written to disk. If `DbEnv::rep_process_message()` method returns `DB_REP_ISPERM` then the **ret_lsnp** parameter will contain largest log sequence number of the permanent records that are now written to disk as a result of processing this message. In all other cases the value of **ret_lsnp** is undefined.

Errors

The DbEnv::rep_process_message() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the database environment was not already configured to communicate with a replication group by a call to DbEnv::rep_set_transport(); if the database environment was not already opened; if this method is called from a Replication Manager application; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Replication and Related Methods

DbEnv::rep_set_clockskew()

```
#include <db_cxx.h>

int
DbEnv::rep_set_clockskew(u_int32_t fast_clock, u_int32_t slow_clock);
```

The `DbEnv::rep_set_clockskew()` method sets the clock skew ratio among replication group members based on the fastest and slowest measurements among the group for use with master leases. Calling this method is optional; the default values for clock skew assume no skew. The user must also configure leases via the DbEnv::rep_set_config() method. Additionally, the user must also set the master lease timeout via the DbEnv::rep_set_timeout() method and the number of sites in the replication group via the DbEnv::rep_set_nsites() method. These methods may be called in any order. For a description of the clock skew values, see Clock skew in the *Berkeley DB Programmer's Reference Guide*. For a description of master leases, see Master leases in the *Berkeley DB Programmer's Reference Guide*.

These arguments can be used to express either raw measurements of a clock timing experiment or a percentage across machines. For example, if a group of sites has a 2% variance, then **fast_clock** should be set to 102, and **slow_clock** should be set to 100. Or, for a 0.03% difference, you can use 10003 and 10000 respectively.

The database environment's replication subsystem may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "rep_set_clockskew", one or more whitespace characters, and the clockskew specified in two parts: the fast_clock and the slow_clock. For example, "rep_set_clockskew 102 100". Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The `DbEnv::rep_set_clockskew()` method configures a database environment, not only operations performed using the specified DbEnv handle.

The `DbEnv::rep_set_clockskew()` method may not be called after the DbEnv::repmgr_start() or DbEnv::rep_start() methods are called.

The `DbEnv::rep_set_clockskew()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

fast_clock

The value, relative to the **slow_clock**, of the fastest clock in the group of sites.

slow_clock

The value of the slowest clock in the group of sites.

Errors

The `DbEnv::rep_set_clockskew()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after replication is started with a call to the DbEnv::repmgr_start() or the DbEnv::rep_start() method; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Replication and Related Methods

DbEnv::rep_set_config()

```
#include <db_cxx.h>

int
DbEnv::rep_set_config(u_int32_t which, int onoff);
```

The `DbEnv::rep_set_config()` method configures the Berkeley DB replication subsystem.

The database environment's replication subsystem may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "rep_set_config", one or more whitespace characters, and the method **which** parameter as a string; for example, "rep_set_config DB_REP_CONF_NOWAIT". Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The `DbEnv::rep_set_config()` method configures a database environment, not only operations performed using the specified DbEnv handle.

The `DbEnv::rep_set_config()` method may not be called to set in-memory replication after the environment is opened using the DbEnv::open() method. This method may also not be called to set master leases after the DbEnv::rep_start() or DbEnv::repmgr_start() methods are called. For all other **which** parameters, this method may be called at any time during the life of the application.

The `DbEnv::rep_set_config()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

onoff

If the **onoff** parameter is zero, the configuration flag is turned off. Otherwise, it is turned on. All configuration flags are turned off by default.

which

The **which** parameter must be set to one of the following values:

- `DB_REP_CONF_BULK`

  The replication master sends groups of records to the clients in a single network transfer.

- `DB_REP_CONF_DELAYCLIENT`

  The client should delay synchronizing to a newly declared master. Clients configured in this way will remain unsynchronized until the application calls the DbEnv::rep_sync() method.

- `DB_REP_CONF_INMEM`

  Store internal replication information in memory only.

By default, replication creates files in the environment home directory to preserve some internal information. If this configuration flag is turned on, replication only stores this internal information in-memory and cannot keep persistent state across a site crash or reboot. This results in the following limitations:

- A master site should not reappoint itself master immediately after crashing or rebooting because the application would incur a slightly higher risk of client crashes. The former master site should rejoin the replication group as a client. The application should either hold an election or appoint a different site to be the next master.

- An application has a slightly higher risk that elections will fail or be unable to complete. Calling additional elections should eventually yield a winner.

- An application has a slight risk that the wrong site may win an election, resulting in the loss of some data. This is consistent with the general loss of data durability when running in-memory.

This configuration flag can only be turned on before the environment is opened with the DbEnv::open() method. Its value cannot be changed while the environment is open.

- DB_REP_CONF_LEASE

    Master leases will be used for this site.

    Configuring this option may result in DB_REP_LEASE_EXPIRED error returns from the Db::get() and Dbc::get() methods when attempting to read entries from a database after the site's master lease has expired.

    This configuration flag may not be set after the DbEnv::repmgr_start() method or the DbEnv::rep_start() method is called.

- DB_REP_CONF_NOAUTOINIT

    The replication master will not automatically re-initialize outdated clients.

- DB_REP_CONF_NOWAIT

    Berkeley DB method calls that would normally block while clients are in recovery will return errors immediately.

- DB_REPMGR_CONF_2SITE_STRICT

    Replication Manager observes the strict "majority" rule in managing elections, even in a group with only 2 sites. This means the client in a 2-site group will be unable to take over as master if the original master fails or becomes disconnected. (See the Elections section in the *Berkeley DB Programmer's Reference Guide* for more information.) Both sites in the replication group should have the same value for this configuration flag.

Errors

The DbEnv::rep_set_config() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If setting in-memory replication after the database environment is already opened; if setting master leases after replication is started; if setting the 2-site strict majority rule for a Base API application; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Replication and Related Methods

DbEnv::rep_set_limit()

```
#include <db_cxx.h>

int
DbEnv::rep_set_limit(u_int32_t gbytes, u_int32_t bytes);
```

The `DbEnv::rep_set_limit()` method sets record transmission throttling. This is a byte-count limit on the amount of data that will be transmitted from a site in response to a single message processed by the `DbEnv::rep_process_message()` method. The limit is not a hard limit, and the record that exceeds the limit is the last record to be sent.

Record transmission throttling is turned on by default with a limit of 10MB.

If the values passed to the `DbEnv::rep_set_limit()` method are both zero, then the transmission limit is turned off.

The database environment's replication subsystem may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "rep_set_limit", one or more whitespace characters, and the limit specified in two parts: the gigabytes and the bytes values. For example, "rep_set_limit 0 1048576" sets a 1 megabyte limit.. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The `DbEnv::rep_set_limit()` method configures a database environment, not only operations performed using the specified DbEnv handle.

The `DbEnv::rep_set_limit()` method may be called at any time during the life of the application.

The `DbEnv::rep_set_limit()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

bytes

The **bytes** parameter specifies the number of bytes which, when added to the **gbytes** parameter, specifies the maximum number of bytes that will be sent in a single call to the DbEnv::rep_process_message() method.

gbytes

The **gbytes** parameter specifies the number of gigabytes which, when added to the **bytes** parameter, specifies the maximum number of bytes that will be sent in a single call to the DbEnv::rep_process_message() method.

Class

DbEnv

See Also

Replication and Related Methods

DbEnv::rep_set_nsites()

```
#include <db_cxx.h>

int
DbEnv::rep_set_nsites(u_int32_t nsites);
```

The DbEnv::rep_set_nsites() method specifies the total number of sites in a replication group.

The DbEnv::rep_set_nsites() method is typically called by Replication Manager applications. (However, see also the DbEnv::rep_elect() method **nsites** parameter.)

The database environment's replication subsystem may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "rep_set_nsites", one or more whitespace characters, and the number of sites specified. For example, "rep_set_nsites 5" sets the number of sites to 5. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The DbEnv::rep_set_nsites() method configures a database environment, not only operations performed using the specified DbEnv handle.

If master leases are in use, the DbEnv::rep_set_nsites() method may not be called after the DbEnv::rep_start() or DbEnv::repmgr_start() methods are called. If master leases are not in use, the DbEnv::rep_set_nsites() method may be called at any time during the life of the application.

The DbEnv::rep_set_nsites() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

nsites

An integer specifying the total number of sites in the replication group.

Errors

The DbEnv::rep_set_nsites() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If master leases are in use and replication has already been started; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Replication and Related Methods

DbEnv::rep_set_priority()

```
#include <db_cxx.h>

int
DbEnv::rep_set_priority(u_int32_t priority);
```

The `DbEnv::rep_set_priority()` method specifies the database environment's priority in replication group elections. A special value of 0 indicates that this environment cannot be a replication group master.

The database environment's replication subsystem may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "rep_set_priority", one or more whitespace characters, and the priority of this site. For example, "rep_set_priority 1" sets the priority of this site to 1. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

Note that if the application never explicitly sets a priority, then a default value of 100 is used.

The `DbEnv::rep_set_priority()` method configures a database environment, not only operations performed using the specified DbEnv handle.

The `DbEnv::rep_set_priority()` method may be called at any time during the life of the application.

The `DbEnv::rep_set_priority()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

priority

The priority of this database environment in the replication group. The priority must be a non-zero integer, or 0 if this environment cannot be a replication group master. (See Replication environment priorities for more information).

Class

DbEnv

See Also

Replication and Related Methods

DbEnv::rep_set_request()

```
#include <db_cxx.h>

int
DbEnv::rep_set_request(u_int32_t min, u_int32_t max);
```

The `DbEnv::rep_set_request()` method sets a threshold for the minimum and maximum time that a client waits before requesting retransmission of a missing message. Specifically, if the client detects a gap in the sequence of incoming log records or database pages, Berkeley DB will wait for at least **min** microseconds before requesting retransmission of the missing record. Berkeley DB will double that amount before requesting the same missing record again, and so on, up to a maximum threshold of **max** microseconds.

These values are thresholds only. Since Berkeley DB has no thread available in the library as a timer, the threshold is only checked when a thread enters the Berkeley DB library to process an incoming replication message. Any amount of time may have passed since the last message arrived and Berkeley DB only checks whether the amount of time since a request was made is beyond the threshold value or not.

By default the minimum is 40000 and the maximum is 1280000 (1.28 seconds). These defaults are fairly arbitrary and the application likely needs to adjust these. The values should be based on expected load and performance characteristics of the master and client host platforms and transport infrastructure as well as round-trip message time.

The database environment's replication subsystem may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "rep_set_request", one or more whitespace characters, and the request times specified in two parts: the min and the max. For example, "rep_set_request 40000 1280000". Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The `DbEnv::rep_set_request()` method configures a database environment, not only operations performed using the specified DbEnv handle.

The `DbEnv::rep_set_request()` method may be called at any time during the life of the application.

The `DbEnv::rep_set_request()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

max

The maximum number of microseconds a client waits before requesting retransmission.

min

The minimum number of microseconds a client waits before requesting retransmission.

Errors

The `DbEnv::rep_set_request()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv

See Also

Replication and Related Methods

DbEnv::rep_set_timeout()

```
#include <db_cxx.h>

int
DbEnv::rep_set_timeout(int which, u_int32_t timeout);
```

The `DbEnv::rep_set_timeout()` method specifies a variety of replication timeout values.

The database environment's replication subsystem may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "rep_set_timeout", one or more whitespace characters, and the **which** parameter specified as a string and the timeout specified as two parts. For example, "rep_set_timeout DB_REP_CONNECTION_RETRY 15000000" specifies the connection retry timeout for 15 seconds. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The `DbEnv::rep_set_timeout()` method configures a database environment, not only operations performed using the specified DbEnv handle.

The `DbEnv::rep_set_timeout()` method may not be called to set the master lease timeout after the DbEnv::repmgr_start() method or the DbEnv::rep_start() method is called. For all other timeouts, the `DbEnv::rep_set_timeout()` method may be called at any time during the life of the application.

The `DbEnv::rep_set_timeout()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

timeout

The **timeout** parameter is the timeout value. It must be specified as an unsigned 32-bit number of microseconds, limiting the maximum timeout to roughly 71 minutes.

which

The **which** parameter must be set to one of the following values:

- DB_REP_ACK_TIMEOUT

  Configure the amount of time the Replication Manager's transport function waits to collect enough acknowledgments from replication group clients, before giving up and returning a failure indication. The default wait time is 1 second.

- DB_REP_CHECKPOINT_DELAY

  Configure the amount of time a master site will delay between completing a checkpoint and writing a checkpoint record into the log. This delay allows clients to complete their own checkpoints before the master requires completion of them. The default is 30 seconds. If all databases in the environment, and the environment's transaction log, are configured to reside in memory (never preserved to disk), then, although checkpoints are still necessary, the delay is not useful and should be set to 0.

- DB_REP_CONNECTION_RETRY

  Configure the amount of time the Replication Manager will wait before trying to re-establish a connection to another site after a communication failure. The default wait time is 30 seconds.

- DB_REP_ELECTION_TIMEOUT

  The timeout period for an election. The default timeout is 2 seconds.

- DB_REP_ELECTION_RETRY

  Configure the amount of time the Replication Manager will wait before retrying a failed election. The default wait time is 10 seconds.

- DB_REP_FULL_ELECTION_TIMEOUT

  An optional configuration timeout period to wait for full election participation the first time the replication group finds a master. By default this option is turned off and normal election timeouts are used. (See the Elections section in the *Berkeley DB Programmer's Reference Guide* for more information.)

- DB_REP_HEARTBEAT_MONITOR

  The amount of time the Replication Manager, running at a client site, waits for some message activity on the connection from the master (heartbeats or other messages) before concluding that the connection has been lost. This timeout should be of longer duration than the DB_REP_HEARTBEAT_SEND timeout to ensure that heartbeats are not missed. When 0 (the default), no monitoring is performed.

- DB_REP_HEARTBEAT_SEND

  The frequency at which the Replication Manager, running at a master site, broadcasts a heartbeat message in an otherwise idle system. When 0 (the default), no heartbeat messages will be sent.

- DB_REP_LEASE_TIMEOUT

  Configure the amount of time a client grants its master lease to a master. When using master leases all sites in a replication group must use the same lease timeout value. There is no default value. If leases are desired, this method must be called prior to calling DbEnv::rep_start() method. See also DbEnv::rep_set_clockskew() method, DbEnv::rep_set_config() method or Master leases.

Errors

The DbEnv::rep_set_timeout() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If setting the lease timeout and replication has already been started; if setting a Replication Manager timeout for a Base API application; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Replication and Related Methods

DbEnv::rep_set_transport()

```
#include <db_cxx.h>

int
DbEnv::rep_set_transport(int envid,
    int (*send)(DB_ENV *dbenv,
    const Dbt *control, const Dbt *rec, const DbLsn *lsnp,
    int envid, u_int32_t flags));
```

The `DbEnv::rep_set_transport()` method initializes the communication infrastructure for a database environment participating in a replicated application.

The `DbEnv::rep_set_transport()` method is not called by most replication applications. It should only be called by Base API applications implementing their own network transport layer, explicitly holding replication group elections and handling replication messages outside of the Replication Manager framework.

The `DbEnv::rep_set_transport()` method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The `DbEnv::rep_set_transport()` method may be called at any time during the life of the application.

The `DbEnv::rep_set_transport()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

☞ Berkeley DB is not re-entrant. The callback function for this method should not attempt to make library calls (for example, to release locks or close open handles). Re-entering Berkeley DB is not guaranteed to work correctly, and the results are undefined.

Parameters

envid

The **envid** parameter is the local environment's ID. It must be a non-negative integer and uniquely identify this Berkeley DB database environment (see Replication environment IDs for more information).

send

The **send** callback function is used to transmit data using the replication application's communication infrastructure. The parameters to **send** are as follows:

• dbenv

    The **dbenv** parameter is the enclosing database environment handle.

• control

    The **control** parameter is the first of the two data elements to be transmitted by the **send** function.

• rec

The **rec** parameter is the second of the two data elements to be transmitted by the **send** function.

- `lsnp`

  If the type of message to be sent has an LSN associated with it, then the **lsnp** parameter contains the LSN of the record being sent. This LSN can be used to determine that certain records have been processed successfully by clients.

- `envid`

  The **envid** parameter is a positive integer identifier that specifies the replication environment to which the message should be sent (see Replication environment IDs for more information).

  The special identifier `DB_EID_BROADCAST` indicates that a message should be broadcast to every environment in the replication group. The application may use a true broadcast protocol or may send the message in sequence to each machine with which it is in communication. In both cases, the sending site should not be asked to process the message.

  The special identifier DB_EID_INVALID indicates an invalid environment ID. This may be used to initialize values that are subsequently checked for validity.

- `flags`

  The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

  - `DB_REP_ANYWHERE`

    The message is a client request that can be satisfied by another client as well as by the master.

  - `DB_REP_NOBUFFER`

    The record being sent should be transmitted immediately and not buffered or delayed.

  - `DB_REP_PERMANENT`

    The record being sent is critical for maintaining database integrity (for example, the message includes a transaction commit). The application should take appropriate action to enforce the reliability guarantees it has chosen, such as waiting for acknowledgement from one or more clients.

  - `DB_REP_REREQUEST`

    The message is a client request that has already been made and to which no response was received.

It may sometimes be useful to pass application-specific data to the send function; see Environment FAQ for a discussion on how to do this.

The **send** function must return 0 on success and non-zero on failure. If the send function fails, the message being sent is necessary to maintain database integrity, and the local log is not configured for synchronous flushing, the local log will be flushed; otherwise, any error from the **send** function will be ignored.

Errors

The `DbEnv::rep_set_transport()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

The method is called from a Replication Manager application; or an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Replication and Related Methods

DbEnv::rep_start()

```
#include <db_cxx.h>

int
DbEnv::rep_start(Dbt *cdata, u_int32_t flags);
```

The `DbEnv::rep_start()` method configures the database environment as a client or master in a group of replicated database environments.

The `DbEnv::rep_start()` method is not called by most replication applications. It should only be called by Base API applications implementing their own network transport layer, explicitly holding replication group elections and handling replication messages outside of the Replication Manager framework.

Replication master environments are the only database environments where replicated databases may be modified. Replication client environments are read-only as long as they are clients. Replication client environments may be upgraded to be replication master environments in the case that the current master fails or there is no master present. If master leases are in use, this method cannot be used to appoint a master, and should only be used to configure a database environment as a master as the result of an election.

The enclosing database environment must already have been opened by calling the DbEnv::open() method and must already have been configured to send replication messages by calling the DbEnv::rep_set_transport() method.

The `DbEnv::rep_start()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

cdata

The **cdata** parameter is an opaque data item that is sent over the communication infrastructure when the client comes online (see Connecting to a new site for more information). If no such information is useful, **cdata** should be NULL.

flags

The **flags** parameter must be set to one of the following values:

- `DB_REP_CLIENT`

  Configure the environment as a replication client.

- `DB_REP_MASTER`

  Configure the environment as a replication master.

Errors

The `DbEnv::rep_start()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DB_REP_UNAVAIL

DB_REP_UNAVAIL

If the flags parameter was passed as DB_REP_MASTER but the database environment cannot currently become the replication master because it is temporarily initializing and is incomplete.

If the database environment was not already configured to communicate with a replication group by a call to DbEnv::rep_set_transport(); the database environment was not already opened; this method is called from a Replication Manager application; outstanding master leases are granted; this method is used to appoint a new master when master leases are in use; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Replication and Related Methods

DbEnv::rep_stat()

```
#include <db_cxx.h>

int
DbEnv::rep_stat(DB_REP_STAT **statp, u_int32_t flags);
```

The `DbEnv::rep_stat()` method returns the replication subsystem statistics.

The `DbEnv::rep_stat()` method creates a statistical structure of type `DB_REP_STAT` and copies a pointer to it into a user-specified memory location.

Statistical structures are stored in allocated memory. If application-specific allocation routines have been declared (see DbEnv::set_alloc() for more information), they are used to allocate the memory; otherwise, the standard C library **malloc**(3) is used. The caller is responsible for deallocating the memory. To deallocate the memory, free the memory reference; references inside the returned memory need not be individually freed.

The following `DB_REP_STAT` fields will be filled in:

- **uintmax_t st_bulk_fills;**

  The number of times the bulk buffer filled up, forcing the buffer content to be sent.

- **uintmax_t st_bulk_overflows;**

  The number of times a record was bigger than the entire bulk buffer, and therefore had to be sent as a singleton.

- **uintmax_t st_bulk_records;**

  The number of records added to a bulk buffer.

- **uintmax_t st_bulk_transfers;**

  The number of bulk buffers transferred (via a call to the application's **send** function).

- **uintmax_t st_client_rerequests;**

  The number of times this client site received a "re-request" message, indicating that a request it previously sent to another client could not be serviced by that client. (Compare to **st_client_svc_miss**.)

- **uintmax_t st_client_svc_miss;**

  The number of "request" type messages received by this client that could not be processed, forcing the originating requester to try sending the request to the master (or another client).

- **uintmax_t st_client_svc_req;**

The number of "request" type messages received by this client. ("Request" messages are usually sent from a client to the master, but a message marked with the DB_REP_ANYWHERE flag in the invocation of the application's **send** function may be sent to another client instead.)

- **u_int32_t st_dupmasters;**

  The number of duplicate master conditions originally detected at this site.

- **u_int32_t st_egen;**

  The election generation number for the current or next election.

- **int st_election_cur_winner;**

  The environment ID of the winner of the current or last election.

- **u_int32_t st_election_gen;**

  The master generation number of the winner of the current or last election.

- **DB_LSN st_election_lsn;**

  The maximum LSN of the winner of the current or last election.

- **u_int32_t st_election_nsites;**

  The number of sites responding to this site during the current election.

- **u_int32_t st_election_nvotes;**

  The number of votes required in the current or last election.

- **u_int32_t st_election_priority;**

  The priority of the winner of the current or last election.

- **u_int32_t st_election_sec;**

  The number of seconds the last election took (the total election time is **st_election_sec** plus **st_election_usec**).

- **int st_election_status;**

  The current election phase (0 if no election is in progress).

- **u_int32_t st_election_tiebreaker;**

  The tiebreaker value of the winner of the current or last election.

- **u_int32_t st_election_usec;**

  The number of microseconds the last election took (the total election time is **st_election_sec** plus **st_election_usec**).

- **u_int32_t st_election_votes;**

  The number of votes received during the current election.

- **uintmax_t st_elections;**

  The number of elections held.

- **uintmax_t st_elections_won;**

  The number of elections won.

- **int st_env_id;**

  The current environment ID.

- **u_int32_t st_env_priority;**

  The current environment priority.

- **u_int32_t st_gen;**

  The current master generation number.

- **uintmax_t st_log_duplicated;**

  The number of duplicate log records received.

- **uintmax_t st_log_queued;**

  The number of log records currently queued.

- **uintmax_t st_log_queued_max;**

  The maximum number of log records ever queued at once.

- **uintmax_t st_log_queued_total;**

  The total number of log records queued.

- **uintmax_t st_log_records;**

  The number of log records received and appended to the log.

- **uintmax_t st_log_requested;**

  The number of times log records were missed and requested.

- **int st_master;**

  The current master environment ID.

- **uintmax_t st_master_changes;**

The number of times the master has changed.

- **u_int32_t st_max_lease_sec;**

  The number of seconds of the longest lease (the total lease time is **st_max_lease_sec** plus **st_max_lease_usec**).

- **u_int32_t st_max_lease_usec;**

  The number of microseconds of the longest lease (the total lease time is **st_max_lease_sec** plus **st_max_lease_usec**).

- **DB_LSN st_max_perm_lsn;**

  The LSN of the maximum permanent log record, or 0 if there are no permanent log records.

- **uintmax_t st_msgs_badgen;**

  The number of messages received with a bad generation number.

- **uintmax_t st_msgs_processed;**

  The number of messages received and processed.

- **uintmax_t st_msgs_recover;**

  The number of messages ignored due to pending recovery.

- **uintmax_t st_msgs_send_failures;**

  The number of failed message sends.

- **uintmax_t st_msgs_sent;**

  The number of messages sent.

- **uintmax_t st_newsites;**

  The number of new site messages received.

- **DB_LSN st_next_lsn;**

  In replication environments configured as masters, the next LSN to be used. In replication environments configured as clients, the next LSN expected.

- **u_int32_t st_next_pg;**

  The next page number we expect to receive.

- **u_int32_t st_nsites;**

  The number of sites used in the last election.

- **uintmax_t st_nthrottles;**

  Transmission limited. This indicates the number of times that data transmission was stopped to limit the amount of data sent in response to a single call to DbEnv::rep_process_message().

- **uintmax_t st_outdated;**

  The number of outdated conditions detected.

- **uintmax_t st_pg_duplicated;**

  The number of duplicate pages received.

- **uintmax_t st_pg_records;**

  The number of pages received and stored.

- **uintmax_t st_pg_requested;**

  The number of pages missed and requested from the master.

- **uintmax_t st_startsync_delayed;**

  The number of times the client had to delay the start of a cache flush operation (initiated by the master for an impending checkpoint) because it was missing some previous log record(s).

- **u_int32_t st_startup_complete;**

  The client site has completed its startup procedures and is now handling live records from the master.

- **u_int32_t st_status;**

  The current replication mode. Set to DB_REP_MASTER if the environment is a replication master, DB_REP_CLIENT if the environment is a replication client, or 0 if replication is not configured.

- **uintmax_t st_txns_applied;**

  The number of transactions applied.

- **DB_LSN st_waiting_lsn;**

  The LSN of the first log record we have after missing log records being waited for, or 0 if no log records are currently missing.

- **u_int32_t st_waiting_pg;**

  The page number of the first page we have after missing pages being waited for, or 0 if no pages are currently missing.

The DbEnv::rep_stat() method may not be called before the DbEnv::open() method is called.

The `DbEnv::rep_stat()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or the following value:

- `DB_STAT_CLEAR`

  Reset statistics after returning their values.

statp

The **statp** parameter references memory into which a pointer to the allocated statistics structure is copied.

Errors

The `DbEnv::rep_stat()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the database environment was not already opened; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Replication and Related Methods

DbEnv::rep_stat_print()

```
#include <db_cxx.h>

int
DbEnv::rep_stat_print(u_int32_t flags);
```

The `DbEnv::rep_stat_print()` method displays the replication subsystem statistical information, as described for the `DbEnv::rep_stat()` method. The information is printed to a specified output channel (see the DbEnv::set_msgfile() method for more information), or passed to an application callback function (see the DbEnv::set_msgcall() method for more information).

The `DbEnv::rep_stat_print()` method may not be called before the DbEnv::open() method is called.

The `DbEnv::rep_stat_print()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

- `DB_STAT_ALL`

  Display all available information.

- `DB_STAT_CLEAR`

  Reset statistics after displaying their values.

Errors

The `DbEnv::rep_stat_print()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called before DbEnv::open() was called; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Replication and Related Methods

DbEnv::rep_sync()

```
#include <db_cxx.h>

int
DbEnv::rep_sync(u_int32_t flags);
```

The `DbEnv::rep_sync()` method forces master synchronization to begin for this client. This method is the other half of setting the DB_REP_CONF_DELAYCLIENT flag via the DbEnv::rep_set_config() method.

If an application has configured delayed master synchronization, the application must synchronize explicitly (otherwise the client will remain out-of-date and will ignore all database changes forwarded from the replication group master). The `DbEnv::rep_sync()` method may be called any time after the client application learns that the new master has been established (by receiving a DB_EVENT_REP_NEWMASTER event notification).

Before calling this method, the enclosing database environment must already have been opened by calling the DbEnv::open() method and must already have been configured to send replication messages by calling the DbEnv::rep_set_transport() method.

The `DbEnv::rep_sync()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter is currently unused, and must be set to 0.

Errors

The `DbEnv::rep_sync()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the database environment was not already configured to communicate with a replication group by a call to DbEnv::rep_set_transport(); the database environment was not already opened; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Replication and Related Methods

DbEnv::repmgr_add_remote_site()

```
#include <db_cxx.h>

int
DbEnv::repmgr_add_remote_site(const char *host, u_int port, int *eidp, u_int32_t flags);
```

The `DbEnv::repmgr_add_remote_site()` method adds a new replication site to the Replication Manager's list of known sites. It is not necessary for all sites in a replication group to know about all other sites in the group.

The `DbEnv::repmgr_add_remote_site()` method may be called at any time during the life of the application.

The `DbEnv::repmgr_add_remote_site()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

eidp

If **eidp** is non-NULL, and the database environment has already been opened by calling DbEnv::open(), then the environment ID assigned to the remote site is returned in the memory location referenced by **eidp**.

flags

The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

• `DB_REPMGR_PEER`

Specifying the `DB_REPMGR_PEER` flag configures client-to-client synchronization with the specified remote site.

Currently, the Replication Manager framework only supports a single client peer, and the last specified peer is used.

host

The site's host identification string, generally a TCP/IP host name.

port

The port number on which the remote site is receiving.

Errors

The `DbEnv::repmgr_add_remote_site()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If this method is called from a base replication API application; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Replication and Related Methods

DbEnv::repmgr_get_ack_policy()

```
#include <db_cxx.h>

int
DbEnv::repmgr_get_ack_policy(int *ack_policyp);
```

The `DbEnv::repmgr_get_ack_policy()` method returns the Replication Manager's client acknowledgment policy. This is configured using the DbEnv::repmgr_set_ack_policy() method.

The `DbEnv::repmgr_get_ack_policy()` method may be called at any time during the life of the application.

The `DbEnv::repmgr_get_ack_policy()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

ack_policyp

The **ack_policyp** parameter references memory into which the Replication Manager's client acknowledgement policy is copied.

Class

DbEnv

See Also

Replication and Related Methods, DbEnv::repmgr_set_ack_policy()

DbEnv::repmgr_set_ack_policy()

```
#include <db_cxx.h>

int
DbEnv::repmgr_set_ack_policy(int ack_policy);
```

The `DbEnv::repmgr_set_ack_policy()` method specifies how master and client sites will handle acknowledgment of replication messages which are necessary for "permanent" records. The current implementation requires all sites in a replication group configure the same acknowledgement policy.

The database environment's replication subsystem may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "repmgr_set_ack_policy", one or more whitespace characters, and the **ack_policy** parameter specified as a string. For example, "repmgr_set_ack_policy DB_REPMGR_ACKS_ALL". Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

Waiting for client acknowledgements is always limited by the DB_REP_ACK_TIMEOUT specified by the DbEnv::rep_set_timeout() method.

The `DbEnv::repmgr_set_ack_policy()` method configures a database environment, not only operations performed using the specified DbEnv handle.

The `DbEnv::repmgr_set_ack_policy()` method may be called at any time during the life of the application.

The `DbEnv::repmgr_set_ack_policy()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

ack_policy

Some acknowledgement policies use the concept of an electable peer, which is a client capable of being subsequently elected master of the replication group. The **ack_policy** parameter must be set to one of the following values:

- **DB_REPMGR_ACKS_ALL**

  The master should wait until all replication clients have acknowledged each permanent replication message.

- **DB_REPMGR_ACKS_ALL_PEERS**

  The master should wait until all electable peers have acknowledged each permanent replication message.

- **DB_REPMGR_ACKS_NONE**

  The master should not wait for any client replication message acknowledgments.

- **DB_REPMGR_ACKS_ONE**

  The master should wait until at least one client site has acknowledged each permanent replication message.

- **DB_REPMGR_ACKS_ONE_PEER**

  The master should wait until at least one electable peer has acknowledged each permanent replication message.

- **DB_REPMGR_ACKS_QUORUM**

  The master should wait until it has received acknowledgements from the minimum number of electable peers sufficient to ensure that the effect of the permanent record remains durable if an election is held. This is the default acknowledgement policy.

Errors

The `DbEnv::repmgr_set_ack_policy()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If this method is called from a base replication API application; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Replication and Related Methods

DbEnv::repmgr_set_local_site()

```
#include <db_cxx.h>

int
repmgr_set_local_site(const char * host, u_int16_t port, u_int32_t flags);
```

The `DbEnv::repmgr_set_local_site()` method specifies the host identification string and port number for the local system.

The `DbEnv::repmgr_set_local_site()` method may not be called after the DbEnv::repmgr_start() method is called.

The `DbEnv::repmgr_set_local_site()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter is currently unused, and must be set to 0.

host

The site's host identification string, generally a TCP/IP host name.

port

The port number on which the local site is listening.

Errors

The `DbEnv::repmgr_set_local_site()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If this method is called from a Base API application; if the method was called after replication is started with a call to the DbEnv::repmgr_start() method; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Replication and Related Methods

DbEnv::repmgr_site_list()

```
#include <db_cxx.h>

int
DbEnv::repmgr_site_list(u_int *countp, DB_REPMGR_SITE **listp);
```

The `DbEnv::repmgr_site_list()` method returns the status of the sites currently known by the Replication Manager.

The `DbEnv::repmgr_site_list()` method creates a statistical structure of type `DB_REPMGR_SITE` and copies a pointer to it into a user-specified memory location.

Statistical structures are stored in allocated memory. If application-specific allocation routines have been declared (see DbEnv::set_alloc() for more information), they are used to allocate the memory; otherwise, the standard C library **malloc**(3) is used. The caller is responsible for deallocating the memory. To deallocate the memory, free the memory reference; references inside the returned memory need not be individually freed.

The following `DB_REPMGR_SITE` fields will be filled in:

- **int eid;**

  Environment ID assigned by the Replication Manager. This is the same value that is passed to the application's event notification function for the DB_EVENT_REP_NEWMASTER event.

- **char host[];**

  Null-terminated host name.

- **u_int port;**

  TCP/IP port number.

- **u_int32_t status;**

  Zero (if unknown), or one of the following constants: `DB_REPMGR_CONNECTED`, `DB_REPMGR_DISCONNECTED`.

The `DbEnv::repmgr_site_list()` method may be called at any time during the life of the application.

The `DbEnv::repmgr_site_list()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

countp

A count of the returned structures will be stored into the memory referenced by **countp**.

listp

A reference to an array of structures will be stored into the memory referenced by **listp**.

Class

DbEnv

See Also

Replication and Related Methods

DbEnv::repmgr_start()

```
#include <db_cxx.h>

int
DbEnv::repmgr_start(int nthreads, u_int32_t flags);
```

The `DbEnv::repmgr_start()` method starts the Replication Manager.

There are two ways to build Berkeley DB replication applications: the most common approach is to use the Berkeley DB library Replication Manager, where the Berkeley DB library manages the replication group, including network transport, all replication message processing and acknowledgment, and group elections. Applications using the Replication Manager generally make the following calls:

1. Call DbEnv::repmgr_set_local_site() to configure the local site in the replication group.

2. Call DbEnv::repmgr_add_remote_site() to configure the remote site(s) in the replication group.

3. Call DbEnv::repmgr_set_ack_policy() to configure the message acknowledgment policy which best supports the replication group's transactional needs.

4. Call DbEnv::rep_set_priority() to configure the local site's election priority.

5. Call `DbEnv::repmgr_start()` to start the replication application.

For more information on building Replication Manager applications, please see the *Replication Getting Started Guide* included in the Berkeley DB documentation.

Applications with special needs (for example, applications using network protocols not supported by the Berkeley DB Replication Manager), must perform additional configuration and call other Berkeley DB replication Base API methods. For more information on building Base API applications, please see the Base API Methods section in the *Berkeley DB Programmer's Reference Guide*.

Starting the Replication Manager consists of opening the TCP/IP listening socket to accept incoming connections, and starting all necessary background threads. When multiple processes share a database environment, only one process can open the listening socket; the `DbEnv::repmgr_start()` method automatically opens the socket in the first process to call it, and skips this step in the later calls from other processes.

The `DbEnv::repmgr_start()` method may not be called before the DbEnv::open() method is called to open the local environment and the DbEnv::repmgr_set_local_site() method is called to configure the local site. In addition, the local environment must be opened with the DB_THREAD flag set.

The `DbEnv::repmgr_start()` method will return `DB_REP_IGNORE` as an informational, non-error return code, if another process has previously become the TCP/IP listener (though the current call has nevertheless successfully started Replication Manager's background threads). Unless otherwise specified, the `DbEnv::repmgr_start()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to one of the following values:

- DB_REP_MASTER

  Start as a master site, and do not call for an election. Note there must never be more than a single master in any replication group, and only one site at a time should ever be started with the DB_REP_MASTER flag specified.

- DB_REP_CLIENT

  Start as a client site, and do not call for an election.

- DB_REP_ELECTION

  Start as a client, and call for an election if no master is found.

nthreads

Specify the number of threads of control created and dedicated to processing replication messages. In addition to these message processing threads, the Replication Manager creates and manages a few of its own threads of control.

Errors

The `DbEnv::repmgr_start()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the database environment was not already opened or was opened without the `DB_THREAD` flag set; a local site has not already been configured, this method is called from a Base API application; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Replication and Related Methods

DbEnv::repmgr_stat()

```
#include <db_cxx.h>

int
DbEnv::repmgr_stat(DB_REPMGR_STAT **statp, u_int32_t flags);
```

The `DbEnv::repmgr_stat()` method returns the Replication Manager statistics.

The `DbEnv::repmgr_stat()` method creates a statistical structure of type `DB_REPMGR_STAT` and copies a pointer to it into a user-specified memory location.

Statistical structures are stored in allocated memory. If application-specific allocation routines have been declared (see DbEnv::set_alloc() for more information), they are used to allocate the memory; otherwise, the standard C library **malloc**(3) is used. The caller is responsible for deallocating the memory. To deallocate the memory, free the memory reference; references inside the returned memory need not be individually freed.

The following `DB_REPMGR_STAT` fields will be filled in:

- **uintmax_t st_connect_fail;**

    The number of times an attempt to open a new TCP/IP connection failed.

- **uintmax_t st_connection_drop;**

    The number of times an existing TCP/IP connection failed.

- **uintmax_t st_msgs_dropped;**

    The number of outgoing messages that were completely dropped, because the outgoing message queue was full. (Berkeley DB replication is tolerant of dropped messages, and will automatically request retransmission of any missing messages as needed.)

- **uintmax_t st_msgs_queued;**

    The number of outgoing messages which could not be transmitted immediately, due to a full network buffer, and had to be queued for later delivery.

- **uintmax_t st_perm_failed;**

    The number of times a message critical for maintaining database integrity (for example, a transaction commit), originating at this site, did not receive sufficient acknowledgement from clients, according to the configured acknowledgement policy and acknowledgement timeout.

The `DbEnv::repmgr_stat()` method may not be called before the DbEnv::open() method is called.

The `DbEnv::repmgr_stat()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or the following value:

• `DB_STAT_CLEAR`

  Reset statistics after returning their values.

statp

The **statp** parameter references memory into which a pointer to the allocated statistics structure is copied.

Errors

The `DbEnv::repmgr_stat()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called before DbEnv::open() was called; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Replication and Related Methods

DbEnv::repmgr_stat_print()

```
#include <db_cxx.h>

int
DbEnv::repmgr_stat_print(u_int32_t flags);
```

The `DbEnv::repmgr_stat_print()` method displays the Replication Manager statistical information, as described for the `DbEnv::repmgr_stat()` method. The information is printed to a specified output channel (see the DbEnv::set_msgfile() method for more information), or passed to an application callback function (see the DbEnv::set_msgcall() method for more information).

The `DbEnv::repmgr_stat_print()` method may not be called before the DbEnv::open() method is called.

The `DbEnv::repmgr_stat_print()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

- `DB_STAT_ALL`

  Display all available information.

- `DB_STAT_CLEAR`

  Reset statistics after displaying their values.

Errors

The `DbEnv::repmgr_stat_print()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called before DbEnv::open() was called; or if an invalid flag value or parameter was specified.

Class

DbEnv

See Also

Replication and Related Methods

Chapter 12. The DbSequence Handle

Sequences provide an arbitrary number of persistent objects that return an increasing or decreasing sequence of integers. Opening a sequence handle associates it with a record in a database. The handle can maintain a cache of values from the database so that a database update is not needed as the application allocates a value.

A sequence is stored as a record pair in a database. The database may be of any type, but must not have been configured to support duplicate data items. The sequence is referenced by the key used when the sequence is created, therefore the key must be compatible with the underlying access method. If the database stores fixed-length records, the record size must be at least 64 bytes long.

You create a sequence using the DbSequence method.

For more information on sequences, see the *Berkeley DB Programmer's Reference Guide* guide.

Sequences and Related Methods

Sequences and Related Methods	Description
DbSequence	Create a sequence handle
DbSequence::close()	Close a sequence
DbSequence::get()	Get the next sequence element(s)
DbSequence::get_dbp()	Return a handle for the underlying sequence database
DbSequence::get_key()	Return the key for a sequence
DbSequence::initial_value()	Set the initial value of a sequence
DbSequence::open()	Open a sequence
DbSequence::remove()	Remove a sequence
DbSequence::stat()	Return sequence statistics
DbSequence::stat_print()	Print sequence statistics
Sequences Configuration	
DbSequence::set_cachesize(), DbSequence::get_cachesize()	Set/get the cache size of a sequence
DbSequence::set_flags(), DbSequence::get_flags()	Set/get the flags for a sequence
DbSequence::set_range(), DbSequence::get_range()	Set/get the range for a sequence

DbSequence

```
#include <db_cxx.h>

class DbSequence {
public:
        DbSequence(Db *db, u_int32_t flags);
        ~DbSequence();

        DB_SEQUENCE *DbSequence::get_DB();
        const DB *DbSequence::get_const_DB() const;
        static DbSequence *DbSequence::get_DbSequence(DB *db);
        static const DbSequence *DbSequence::get_const_DbSequence(const DB *db);

        ...
};
```

The DbSequence handle is the handle used to manipulate a sequence object. A sequence object is stored in a record in a database.

DbSequence handles are free-threaded if the DB_THREAD flag is specified to the DbSequence::open() method when the sequence is opened. Once the DbSequence::close() or DbSequence::remove() methods are called, the handle can not be accessed again, regardless of the method's return.

Each handle opened on a sequence may maintain a separate cache of values which are returned to the application using the DbSequence::get() method either singly or in groups depending on its **delta** parameter.

Calling the DbSequence::close() or DbSequence::remove() methods discards this handle.

Parameters

db

The **db** parameter is an open database handle which holds the persistent data for the sequence. The database may be of any type, but must not have been configured to support duplicate data items.

flags

The **flags** parameter is currently unused, and must be set to 0.

Errors

The db_sequence_create method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbSequence

See Also

Sequences and Related Methods

DbSequence::close()

```
#include <db_cxx.h>

int
DbSequence::close(u_int32_t flags);
```

The `DbSequence::close()` method closes the sequence handle. Any unused cached values are lost.

The DbSequence handle may not be accessed again after `DbSequence::close()` is called, regardless of its return.

The `DbSequence::close()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter is currently unused, and must be set to 0.

Errors

The `DbSequence::close()` method method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbSequence

See Also

Sequences and Related Methods

DbSequence::get()

```
#include <db_cxx.h>

int
DbSequence::get(DbTxn *txnid, int32_t delta, db_seq_t *retp, u_int32_t flags);
```

The `DbSequence::get()` method returns the next available element in the sequence and changes the sequence value by **delta**. The value of **delta** must be greater than zero. If there are enough cached values in the sequence handle then they will be returned. Otherwise the next value will be fetched from the database and incremented (decremented) by enough to cover the **delta** and the next batch of cached values.

For maximum concurrency a non-zero cache size should be specified prior to opening the sequence handle and DB_TXN_NOSYNC should be specified for each `DbSequence::get()` method call.

By default, sequence ranges do not wrap; to cause the sequence to wrap around the beginning or end of its range, specify the DB_SEQ_WRAP flag to the DbSequence::set_flags() method.

The `DbSequence::get()` method will return EINVAL if the record in the database is not a valid sequence record, or the sequence has reached the beginning or end of its range and is not configured to wrap.

Parameters

delta

Specifies the amount to increment or decrement the sequence.

flags

The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

- DB_TXN_NOSYNC

  If the operation is implicitly transaction protected (the **txnid** argument is NULL but the operation occurs to a transactional database), do not synchronously flush the log when the transaction commits.

retp

retp points to the memory to hold the return value from the sequence.

txnid

If the operation is part of an application-specified transaction, the **txnid** parameter is a transaction handle returned from DbEnv::txn_begin(); if the operation is part of a Berkeley DB Concurrent Data Store group, the **txnid** parameter is a handle returned from DbEnv::cdsgroup_begin(); otherwise NULL. If no transaction handle is specified, but the operation occurs in a transactional database, the operation will be implicitly transaction protected. No **txnid** handle may be specified if the sequence handle was opened with a non-zero cache size.

If the underlying database handle was opened in a transaction, calling `DbSequence::get()` may result in changes to the sequence object; these changes will be automatically committed in a transaction internal to the Berkeley DB library. If the thread of control calling `DbSequence::get()` has an active transaction, which holds locks on the same database as the one in which the sequence object is stored, it is possible for a thread of control calling `DbSequence::get()` to self-deadlock because the active transaction's locks conflict with the internal transaction's locks. For this reason, it is often preferable for sequence objects to be stored in their own database.

Class

DbSequence

See Also

Sequences and Related Methods

DbSequence::get_cachesize()

```
#include <db_cxx.h>

int DbSequence::get_cachesize(u_int32_t *sizep);
```

The `DbSequence::get_cachesize()` method returns the current cache size.

The `DbSequence::get_cachesize()` method may be called at any time during the life of the application.

The `DbSequence::get_cachesize()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

sizep

The `DbSequence::get_cachesize()` method returns the current cache size in **sizep**.

Class

DbSequence

See Also

Sequences and Related Methods

DbSequence::get_dbp()

```
#include <db_cxx.h>

int
DbSequence::get_dbp(Db **dbp);
```

The `DbSequence::get_dbp()` method returns the database handle used by the sequence.

The `DbSequence::get_dbp()` method may be called at any time during the life of the application.

The `DbSequence::get_dbp()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

dbp

The **dbp** parameter references memory into which a pointer to the database handle is copied.

Class

DbSequence

See Also

Sequences and Related Methods

DbSequence::get_flags()

```
#include <db_cxx.h>

int DbSequence::get_flags(u_int32_t *flagsp);
```

The `DbSequence::get_flags()` method returns the current flags.

The `DbSequence::get_flags()` method may be called at any time during the life of the application.

The `DbSequence::get_flags()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flagsp

The `DbSequence::get_flags()` method returns the current flags in **flagsp**.

Class

DbSequence

See Also

Sequences and Related Methods

DbSequence::get_key()

```
#include <db_cxx.h>

int
DbSequence::get_key(Dbt *key);
```

The `DbSequence::get_key()` method returns the key for the sequence.

The `DbSequence::get_key()` method may be called at any time during the life of the application.

The `DbSequence::get_key()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

key

The **key** parameter references memory into which a pointer to the key data is copied.

Class

DbSequence

See Also

Sequences and Related Methods

DbSequence::get_range()

```
#include <db_cxx.h>

int DbSequence::get_range(u_int32_t, db_seq_t *minp, db_seq_t *maxp);
```

The `DbSequence::get_range()` method returns the range of values in the sequence.

The `DbSequence::get_range()` method may be called at any time during the life of the application.

The `DbSequence::get_range()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

minp

The `DbSequence::get_range()` method returns the minimum value in **minp**.

maxp

The `DbSequence::get_range()` method returns the maximum value in **maxp**.

Class

DbSequence

See Also

Sequences and Related Methods

DbSequence::initial_value()

```
#include <db_cxx.h>

int
DbSequence::initial_value(db_seq_t value);
```

Set the initial value for a sequence. This call is only effective when the sequence is being created.

The DbSequence::initial_value() method may not be called after the DbSequence::open() method is called.

The DbSequence::initial_value() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

value

The initial value to set.

Errors

The DbSequence::initial_value() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbSequence

See Also

Sequences and Related Methods

DbSequence::open()

```
#include <db_cxx.h>

int
DbSequence::open(DbTxn *txnid, Dbt *key, u_int32_t flags);
```

The `DbSequence::open()` method opens the sequence represented by the **key**. The key must be compatible with the underlying database specified in the corresponding call to DbSequence.

The `DbSequence::open()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

key

The **key** specifies which record in the database stores the persistent sequence data.

flags

The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

- `DB_CREATE`

  Create the sequence. If the sequence does not already exist and the `DB_CREATE` flag is not specified, the `DbSequence::open()` method will fail.

- `DB_EXCL`

  Return an error if the sequence already exists. This flag is only meaningful when specified with the `DB_CREATE` flag.

- `DB_THREAD`

  Cause the DbSequence handle returned by `DbSequence::open()` to be *free-threaded*; that is, usable by multiple threads within a single address space. Note that if multiple threads create multiple sequences using the same database handle that handle must have been opened specifying this flag.

txnid

If the operation is part of an application-specified transaction, the **txnid** parameter is a transaction handle returned from DbEnv::txn_begin(); if the operation is part of a Berkeley DB Concurrent Data Store group, the **txnid** parameter is a handle returned from DbEnv::cdsgroup_begin(); otherwise NULL. If no transaction handle is specified, but the operation occurs in a transactional database, the operation will be implicitly transaction protected. Transactionally protected operations on a DbSequence handle require the DbSequence handle itself be transactionally protected during its open if the open creates the sequence.

Class

DbSequence

See Also

Sequences and Related Methods

DbSequence::remove()

```
#include <db_cxx.h>

int
DbSequence::remove(u_int32_t flags);
```

The `DbSequence::remove()` method removes the sequence from the database. This method should not be called if there are other open handles on this sequence.

The [DbSequence](#) handle may not be accessed again after `DbSequence::remove()` is called, regardless of its return.

The `DbSequence::remove()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

• `DB_TXN_NOSYNC`

 If the operation is implicitly transaction protected (the **txnid** argument is NULL but the operation occurs to a transactional database), do not synchronously flush the log when the transaction commits.

txnid

If the operation is part of an application-specified transaction, the **txnid** parameter is a transaction handle returned from [DbEnv::txn_begin()](#); if the operation is part of a Berkeley DB Concurrent Data Store group, the **txnid** parameter is a handle returned from [DbEnv::cdsgroup_begin()](#); otherwise NULL. If no transaction handle is specified, but the operation occurs in a transactional database, the operation will be implicitly transaction protected.

Errors

The `DbSequence::remove()` method may fail and throw a [DbException](#) exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

[DbSequence](#)

See Also

Sequences and Related Methods

DbSequence::set_cachesize()

```
#include <db_cxx.h>

int
DbSequence::set_cachesize(int32_t size);
```

Configure the number of elements cached by a sequence handle.

The `DbSequence::set_cachesize()` method may not be called after the DbSequence::open() method is called.

The `DbSequence::set_cachesize()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

size

The number of elements in the cache.

Errors

The `DbSequence::set_cachesize()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbSequence

See Also

Sequences and Related Methods

DbSequence::set_flags()

```
#include <db_cxx.h>

int
DbSequence::set_flags(u_int32_t flags);
```

Configure a sequence. The flags are only effective when creating a sequence. Calling `DbSequence::set_flags()` is additive; there is no way to clear flags.

The `DbSequence::set_flags()` method may not be called after the DbSequence::open() method is called.

The `DbSequence::set_flags()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

- `DB_SEQ_DEC`

  Specify that the sequence should be decremented.

- `DB_SEQ_INC`

  Specify that the sequence should be incremented. This is the default.

- `DB_SEQ_WRAP`

  Specify that the sequence should wrap around when it is incremented (decremented) past the specified maximum (minimum) value.

Errors

The `DbSequence::set_flags()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbSequence

See Also

Sequences and Related Methods

DbSequence::set_range()

```
#include <db_cxx.h>

int
DbSequence::set_range(db_seq_t min, db_seq_t max);
```

Configure a sequence range. This call is only effective when the sequence is being created. The range is limited to a signed 64 bit integer.

The `DbSequence::set_range()` method may not be called after the DbSequence::open() method is called.

The `DbSequence::set_range()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

min

Specifies the minimum value for the sequence.

max

Specifies the maximum value for the sequence.

Errors

The `DbSequence::set_range()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbSequence

See Also

Sequences and Related Methods

DbSequence::stat()

```
#include <db_cxx.h>

int
DbSequence::stat(DB_SEQUENCE_STAT **spp, u_int32_t flags);
```

The `DbSequence::stat()` method creates a statistical structure and copies a pointer to it into user-specified memory locations. Specifically, if **spp** is non-NULL, a pointer to the statistics for the database are copied into the memory location to which it refers.

Statistical structures are stored in allocated memory. If application-specific allocation routines have been declared (see DbEnv::set_alloc() for more information), they are used to allocate the memory; otherwise, the standard C library **malloc**(3) is used. The caller is responsible for deallocating the memory. To deallocate the memory, free the memory reference; references inside the returned memory need not be individually freed.

In the presence of multiple threads or processes accessing an active sequence, the information returned by `DbSequence::stat()` may be out-of-date.

The `DbSequence::stat()` method cannot be transaction-protected. For this reason, it should be called in a thread of control that has no open cursors or active transactions.

The `DbSequence::stat()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

The statistics are stored in a structure of type DB_SEQUENCE_STAT. The following fields will be filled in:

- **uintmax_t st_wait;**

  The number of times a thread of control was forced to wait on the handle mutex.

- **uintmax_t st_nowait;**

  The number of times that a thread of control was able to obtain handle mutex without waiting.

- **db_seq_t st_current;**

  The current value of the sequence in the database.

- **db_seq_t st_value;**

  The current cached value of the sequence.

- **db_seq_t st_last_value;**

  The last cached value of the sequence.

- **db_seq_t st_min;**

  The minimum permitted value of the sequence.

- **db_seq_t st_max;**

  The maximum permitted value of the sequence.

- **int32_t st_cache_size;**

  The number of values that will be cached in this handle.

- **u_int32_t st_flags;**

  The flags value for the sequence.

Parameters

flags

The **flags** parameter must be set by bitwise inclusively **OR**'ing together one or more of the following values:

- DB_STAT_CLEAR

  Reset statistics after printing their values.

Class

DbSequence

See Also

Sequences and Related Methods

DbSequence::stat_print()

```
#include <db_cxx>

int
DbSequence::stat_print(u_int32_t flags);
```

The `DbSequence::stat_print()` method prints diagnostic information to the output channel described by the DbEnv::set_msgfile() method.

The `DbSequence::stat_print()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set by bitwise inclusively **OR**'ing together one or more of the following values:

- `DB_STAT_CLEAR`

  Reset statistics after printing their values.

Class

DbSequence

See Also

Sequences and Related Methods

Chapter 13. The DbTxn Handle

```
#include <db_cxx.h>

class DbTxn {
public:
    DB_TXN *DbTxn::get_DB_TXN();
    const DB_TXN *DbTxn::get_const_DB_TXN() const;
    static DbTxn *DbTxn::get_DbTxn(DB_TXN *txn);
    static const DbTxn *DbTxn::get_const_DbTxn(const DB_TXN *txn);
    ...

};
```

The `DbTxn` object is the handle for a transaction. Methods of the `DbTxn` handle are used to configure, abort and commit the transaction. `DbTxn` handles are provided to Db methods in order to transactionally protect those database operations.

`DbTxn` handles are not free-threaded; transactions handles may be used by multiple threads, but only serially, that is, the application must serialize access to the `DbTxn` handle. Once the DbTxn::abort() or DbTxn::commit() methods are called, the handle may not be accessed again, regardless of the method's return. In addition, parent transactions may not issue any Berkeley DB operations while they have active child transactions (child transactions that have not yet been committed or aborted) except for DbEnv::txn_begin(), DbTxn::abort() and DbTxn::commit().

Each `DbTxn` object has an associated `DB_TXN` struct, which is used by the underlying implementation of Berkeley DB and its C++ language API. The `DbTxn::get_DB_TXN()` method returns a pointer to this struct. Given a const `DbTxn` object, `txnMget_const_DB_TXN()` returns a const pointer to the same struct.

Given a `DB_TXN` struct, the `DbTxn::get_DbTxn()` method returns the corresponding `DbTxn` object, if there is one. If the `DB_TXN` object was not associated with a `DbTxn` (that is, it was not returned from a call to `DbTxn::get_DB_TXN()`), then the result of `DbTxn::get_DbTxn` is undefined. Given a const `DB_TXN` struct, `DbTxn::get_const_DbTxn()` returns the associated const `DbTxn` object, if there is one.

These methods may be useful for Berkeley DB applications including both C and C++ language software. It should not be necessary to use these calls in a purely C++ application.

Transaction Subsystem and Related Methods

Transaction Subsystem and Related Methods	Description
DbEnv::txn_checkpoint()	Checkpoint the transaction subsystem
DbEnv::txn_recover()	Distributed transaction recovery
DbTxn::set_timeout()	Set transaction timeout
DbEnv::txn_stat()	Return transaction subsystem statistics
DbEnv::txn_stat_print()	Print transaction subsystem statistics
Transaction Subsystem Configuration	
DbEnv::set_timeout(), DbEnv::get_timeout()	Set/get lock and transaction timeout
Db::get_transactional()	Does the Db have transaction support
DbEnv::cdsgroup_begin()	Get a locker ID in Berkeley DB Concurrent Data Store
DbEnv::set_tx_max(), DbEnv::get_tx_max()	Set/get maximum number of transactions
DbEnv::set_tx_timestamp(), DbEnv::get_tx_timestamp()	Set/get recovery timestamp
Transaction Operations	
DbTxn::abort()	Abort a transaction
DbEnv::txn_begin()	Begin a transaction
DbTxn::commit()	Commit a transaction
DbTxn::discard()	Discard a prepared but not resolved transaction handle
DbTxn::id()	Return a transaction's ID
DbTxn::prepare()	Prepare a transaction for commit
DbTxn::set_name(), DbTxn::get_name()	Associate a string with a transaction

Db::get_transactional()

```
#include <db_cxx.h>

int
Db::get_transactional()
```

The `Db::get_transactional()` method returns non-zero if the Db handle has been opened in a transactional mode.

The `Db::get_transactional()` method may be called at any time during the life of the application.

The `Db::get_transactional()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Class

Db

See Also

Transaction Subsystem and Related Methods

DbEnv::cdsgroup_begin()

```
#include <db_cxx.h>

int
DbEnv::cdsgroup_begin(DbTxn **tid);
```

The `DbEnv::cdsgroup_begin()` method allocates a locker ID in an environment configured for Berkeley DB Concurrent Data Store applications. It copies a pointer to a DbTxn that uniquely identifies the locker ID into the memory to which **tid** refers. Calling the DbTxn::commit() method will discard the allocated locker ID.

See Berkeley DB Concurrent Data Store applications for more information about when this is required.

The `DbEnv::cdsgroup_begin()` method may be called at any time during the life of the application.

The `DbEnv::cdsgroup_begin()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Errors

The `DbEnv::cdsgroup_begin()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

ENOMEM

The maximum number of lockers has been reached.

Class

DbEnv, DbTxn

See Also

Transaction Subsystem and Related Methods

DbEnv::get_tx_max()

```
#include <db_cxx.h>

int
DbEnv::get_tx_max(u_int32_t *tx_maxp);
```

The `DbEnv::get_tx_max()` method returns the maximum number of active transactions currently configured for the environment. You can manage this value using the DbEnv::set_tx_max() method.

The `DbEnv::get_tx_max()` method may be called at any time during the life of the application.

The `DbEnv::get_tx_max()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

tx_maxp

The `DbEnv::get_tx_max()` method returns the number of active transactions in **tx_maxp**.

Class

DbEnv, DbTxn

See Also

Transaction Subsystem and Related Methods, DbEnv::set_tx_max()

DbEnv::get_tx_timestamp()

```
#include <db_cxx.h>

int
DbEnv::get_tx_timestamp(time_t *timestampp);
```

The `DbEnv::get_tx_timestamp()` method returns the recovery timestamp. This value can be modified using the DbEnv::set_tx_timestamp() method.

The `DbEnv::get_tx_timestamp()` method may be called at any time during the life of the application.

The `DbEnv::get_tx_timestamp()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

timestampp

The `DbEnv::get_tx_timestamp()` method returns the recovery timestamp in **timestampp**.

Class

DbEnv, DbTxn

See Also

Transaction Subsystem and Related Methods, DbEnv::set_tx_timestamp()

DbEnv::set_tx_max()

```
#include <db_cxx.h>

int
DbEnv::set_tx_max(u_int32_t max);
```

Configure the Berkeley DB database environment to support at least **max** active transactions. This value bounds the size of the memory allocated for transactions. Child transactions are counted as active until they either commit or abort.

Transactions that update multiversion databases are not freed until the last page version that the transaction created is flushed from cache. This means that applications using multi-version concurrency control may need a transaction for each page in cache, in the extreme case.

When all of the memory available in the database environment for transactions is in use, calls to DbEnv::txn_begin() will fail (until some active transactions complete). If DbEnv::set_tx_max() is never called, the database environment is configured to support at least 100 active transactions.

The database environment's number of active transactions may also be configured using the environment's DB_CONFIG file. The syntax of the entry in that file is a single line with the string "set_tx_max", one or more whitespace characters, and the number of transactions. Because the DB_CONFIG file is read when the database environment is opened, it will silently overrule configuration done before that time.

The DbEnv::set_tx_max() method configures a database environment, not only operations performed using the specified DbEnv handle.

The DbEnv::set_tx_max() method may not be called after the DbEnv::open() method is called. If the database environment already exists when DbEnv::open() is called, the information specified to DbEnv::set_tx_max() will be ignored.

The DbEnv::set_tx_max() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

max

The **max** parameter configures the minimum number of simultaneously active transactions supported by Berkeley DB database environment.

Errors

The DbEnv::set_tx_max() method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the method was called after DbEnv::open() was called; or if an invalid flag value or parameter was specified.

Class

DbEnv, DbTxn

See Also

Transaction Subsystem and Related Methods

DbEnv::set_tx_timestamp()

```
#include <db_cxx.h>

int
DbEnv::set_tx_timestamp(time_t *timestamp);
```

Recover to the time specified by **timestamp** rather than to the most current possible date.

Once a database environment has been upgraded to a new version of Berkeley DB involving a log format change (see Upgrading Berkeley DB installations), it is no longer possible to recover to a specific time before that upgrade.

The `DbEnv::set_tx_timestamp()` method configures operations performed using the specified DbEnv handle, not all operations performed on the underlying database environment.

The `DbEnv::set_tx_timestamp()` method may not be called after the DbEnv::open() method is called.

The `DbEnv::set_tx_timestamp()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

timestamp

The **timestamp** parameter references the memory location where the recovery timestamp is located.

The **timestamp** parameter should be the number of seconds since 0 hours, 0 minutes, 0 seconds, January 1, 1970, Coordinated Universal Time; that is, the Epoch.

Errors

The `DbEnv::set_tx_timestamp()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If it is not possible to recover to the specified time using the log files currently present in the environment; or if an invalid flag value or parameter was specified.

Class

DbEnv, DbTxn

See Also

Transaction Subsystem and Related Methods

DbTxn::abort()

```
#include <db_cxx.h>

int
DbTxn::abort();
```

The `DbTxn::abort()` method causes an abnormal termination of the transaction. The log is played backward, and any necessary undo operations are done through the **tx_recover** function specified to DbEnv::set_app_dispatch(). Before `DbTxn::abort()` returns, any locks held by the transaction will have been released.

In the case of nested transactions, aborting a parent transaction causes all children (unresolved or not) of the parent transaction to be aborted.

All cursors opened within the transaction must be closed before the transaction is aborted.

After `DbTxn::abort()` has been called, regardless of its return, the DbTxn handle may not be accessed again.

The `DbTxn::abort()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Class

DbEnv, DbTxn

See Also

Transaction Subsystem and Related Methods

DbEnv::txn_begin()

```
#include <db_cxx.h>

int
DbEnv::txn_begin(DbTxn *parent, DbTxn **tid, u_int32_t flags);
```

The `DbEnv::txn_begin()` method creates a new transaction in the environment and copies a pointer to a DbTxn that uniquely identifies it into the memory to which **tid** refers. Calling the DbTxn::abort(), DbTxn::commit() or DbTxn::discard() methods will discard the returned handle.

☞ Transactions may only span threads if they do so serially; that is, each transaction must be active in only a single thread of control at a time. This restriction holds for parents of nested transactions as well; no two children may be concurrently active in more than one thread of control at any one time.

☞ Cursors may not span transactions; that is, each cursor must be opened and closed within a single transaction.

☞ A parent transaction may not issue any Berkeley DB operations — except for `DbEnv::txn_begin()`, DbTxn::abort() and DbTxn::commit() — while it has active child transactions (child transactions that have not yet been committed or aborted).

The `DbEnv::txn_begin()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

- `DB_READ_COMMITTED`

  This transaction will have degree 2 isolation. This provides for cursor stability but not repeatable reads. Data items which have been previously read by this transaction may be deleted or modified by other transactions before this transaction completes.

- `DB_READ_UNCOMMITTED`

  This transaction will have degree 1 isolation. Read operations performed by the transaction may read modified but not yet committed data. Silently ignored if the `DB_READ_UNCOMMITTED` flag was not specified when the underlying database was opened.

- `DB_TXN_NOSYNC`

  Do not synchronously flush the log when this transaction commits or prepares. This means the transaction will exhibit the ACI (atomicity, consistency, and isolation) properties, but not D (durability); that is, database integrity will be maintained but it is possible that this transaction may be undone during recovery.

This behavior may be set for a Berkeley DB environment using the DbEnv::set_flags() method. Any value specified to this method overrides that setting.

- DB_TXN_NOWAIT

  If a lock is unavailable for any Berkeley DB operation performed in the context of this transaction, cause the operation to return DB_LOCK_DEADLOCK (or DB_LOCK_NOTGRANTED if the database environment has been configured using the DB_TIME_NOTGRANTED flag).

  This behavior may be set for a Berkeley DB environment using the DbEnv::set_flags() method. Any value specified to this method overrides that setting.

- DB_TXN_SNAPSHOT

  This transaction will execute with snapshot isolation. For databases with the DB_MULTIVERSION flag set, data values will be read as they are when the transaction begins, without taking read locks. Silently ignored for operations on databases with DB_MULTIVERSION not set on the underlying database (read locks are acquired).

  The error DB_LOCK_DEADLOCK will be returned from update operations if a snapshot transaction attempts to update data which was modified after the snapshot transaction read it.

- DB_TXN_SYNC

  Synchronously flush the log when this transaction commits or prepares. This means the transaction will exhibit all of the ACID (atomicity, consistency, isolation, and durability) properties.

  This behavior is the default for Berkeley DB environments unless the DB_TXN_NOSYNC flag was specified to the DbEnv::set_flags() method. Any value specified to this method overrides that setting.

- DB_TXN_WAIT

  If a lock is unavailable for any Berkeley DB operation performed in the context of this transaction, wait for the lock.

  This behavior is the default for Berkeley DB environments unless the DB_TXN_NOWAIT flag was specified to the DbEnv::set_flags() method. Any value specified to this method overrides that setting.

- DB_TXN_WRITE_NOSYNC

  Write, but do not synchronously flush, the log when this transaction commits. This means the transaction will exhibit the ACI (atomicity, consistency, and isolation) properties, but not D (durability); that is, database integrity will be maintained, but if the system fails, it is possible some number of the most recently committed transactions may be undone during recovery. The number of transactions at risk is governed by how often the system flushes dirty buffers to disk and how often the log is flushed or checkpointed.

  This behavior may be set for a Berkeley DB environment using the DbEnv::set_flags() method. Any value specified to this method overrides that setting.

parent

If the **parent** parameter is non-NULL, the new transaction will be a nested transaction, with the transaction indicated by **parent** as its parent. Transactions may be nested to any level. In the presence of distributed transactions and two-phase commit, only the parental transaction, that is a transaction without a **parent** specified, should be passed as an parameter to DbTxn::prepare().

Errors

The `DbEnv::txn_begin()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

DbMemoryException or ENOMEM

The maximum number of concurrent transactions has been reached.

DbMemoryException is thrown if your Berkeley DB API is configured to throw exceptions. Otherwise, `ENOMEM` is returned.

Class

DbEnv, DbTxn

See Also

Transaction Subsystem and Related Methods

DbEnv::txn_checkpoint()

```
#include <db_cxx.h>

int
DbEnv::txn_checkpoint(u_int32_t kbyte, u_int32_t min, u_int32_t flags) const;
```

If there has been any logging activity in the database environment since the last checkpoint, the
DbEnv::txn_checkpoint() method flushes the underlying memory pool, writes a checkpoint record to
the log, and then flushes the log.

The DbEnv::txn_checkpoint() method returns a non-zero error value on failure and 0 on success.

The DbEnv::txn_checkpoint() method is the underlying method used by the db_checkpoint utility. See
the db_checkpoint utility source code for an example of using DbEnv::txn_checkpoint() in a IEEE/ANSI
Std 1003.1 (POSIX) environment.

Parameters

flags

The **flags** parameter must be set to 0 or the following value:

- DB_FORCE

  Force a checkpoint record, even if there has been no activity since the last checkpoint.

kbyte

If the **kbyte** parameter is non-zero, a checkpoint will be done if more than **kbyte** kilobytes of log data
have been written since the last checkpoint.

min

If the **min** parameter is non-zero, a checkpoint will be done if more than **min** minutes have passed
since the last checkpoint.

Errors

The DbEnv::txn_checkpoint() method may fail and throw a DbException exception, encapsulating one
of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv, DbTxn

See Also

Transaction Subsystem and Related Methods

DbTxn::commit()

```
#include <db_cxx.h>

int
DbTxn::commit(u_int32_t flags);
```

The `DbTxn::commit()` method ends the transaction.

In the case of nested transactions, if the transaction is a parent transaction, committing the parent transaction causes all unresolved children of the parent to be committed. In the case of nested transactions, if the transaction is a child transaction, its locks are not released, but are acquired by its parent. Although the commit of the child transaction will succeed, the actual resolution of the child transaction is postponed until the parent transaction is committed or aborted; that is, if its parent transaction commits, it will be committed; and if its parent transaction aborts, it will be aborted.

All cursors opened within the transaction must be closed before the transaction is committed.

After `DbTxn::commit()` has been called, regardless of its return, the DbTxn handle may not be accessed again. If `DbTxn::commit()` encounters an error, the transaction and all child transactions of the transaction are aborted.

The `DbTxn::commit()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or one of the following values:

• `DB_TXN_NOSYNC`

  Do not synchronously flush the log. This means the transaction will exhibit the ACI (atomicity, consistency, and isolation) properties, but not D (durability); that is, database integrity will be maintained, but it is possible that this transaction may be undone during recovery.

  This behavior may be set for a Berkeley DB environment using the DbEnv::set_flags() method or for a single transaction using the DbEnv::txn_begin() method. Any value specified to this method overrides both of those settings.

• `DB_TXN_SYNC`

  Synchronously flush the log. This means the transaction will exhibit all of the ACID (atomicity, consistency, isolation, and durability) properties.

  This behavior is the default for Berkeley DB environments unless the DB_TXN_NOSYNC flag was specified to the DbEnv::set_flags() method. This behavior may also be set for a single transaction using the DbEnv::txn_begin() method. Any value specified to this method overrides both of those settings.

Class

DbEnv, DbTxn

See Also

Transaction Subsystem and Related Methods

DbTxn::discard()

```
#include <db_cxx.h>

int
DbTxn::discard(u_int32_t flags);
```

The `DbTxn::discard()` method frees up all the per-process resources associated with the specified `DbTxn` handle, neither committing nor aborting the transaction. This call may be used only after calls to `DbEnv::txn_recover()` when there are multiple global transaction managers recovering transactions in a single Berkeley DB environment. Any transactions returned by `DbEnv::txn_recover()` that are not handled by the current global transaction manager should be discarded using `DbTxn::discard()`.

The `DbTxn::discard()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

After `DbTxn::discard()` has been called, regardless of its return, the `DbTxn` handle may not be accessed again.

Parameters

flags

The **flags** parameter is currently unused, and must be set to 0.

Errors

The `DbTxn::discard()` method may fail and throw a `DbException` exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

If the transaction handle does not refer to a transaction that was recovered into a prepared but not yet completed state; or if an invalid flag value or parameter was specified.

Class

DbEnv, DbTxn

See Also

Transaction Subsystem and Related Methods

DbTxn::get_name()

```
#include <db_cxx.h>

int
DbTxn::get_name(const char **namep);
```

The `DbTxn::get_name()` method returns the string associated with the transaction.

The `DbTxn::get_name()` method may be called at any time during the life of the application.

The `DbTxn::get_name()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

namep

The `DbTxn::get_name()` method returns a reference to the string associated with the transaction in namep.

Class

DbEnv, DbTxn

See Also

Transaction Subsystem and Related Methods

DbTxn::id()

```
#include <db_cxx.h>

u_int32_t
DbTxn::id();
```

The `DbTxn::id()` method returns the unique transaction id associated with the specified transaction. Locking calls made on behalf of this transaction should use the value returned from `DbTxn::id()` as the locker parameter to the DbEnv::lock_get() or DbEnv::lock_vec() calls.

Class

DbEnv, DbTxn

See Also

Transaction Subsystem and Related Methods

DbTxn::prepare()

```
#include <db_cxx.h>

int
DbTxn::prepare(u_int8_t gid[DB_XIDDATASIZE]);
```

The `DbTxn::prepare()` method initiates the beginning of a two-phase commit.

In a distributed transaction environment, Berkeley DB can be used as a local transaction manager. In this case, the distributed transaction manager must send *prepare* messages to each local manager. The local manager must then issue a `DbTxn::prepare()` and await its successful return before responding to the distributed transaction manager. Only after the distributed transaction manager receives successful responses from all of its *prepare* messages should it issue any *commit* messages.

In the case of nested transactions, preparing the parent causes all unresolved children of the parent transaction to be committed. Child transactions should never be explicitly prepared. Their fate will be resolved along with their parent's during global recovery.

The `DbTxn::prepare()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

gid

The **gid** parameter specifies the global transaction ID by which this transaction will be known. This global transaction ID will be returned in calls to DbEnv::txn_recover() telling the application which global transactions must be resolved.

Class

DbEnv, DbTxn

See Also

Transaction Subsystem and Related Methods

DbEnv::txn_recover()

```
#include <db_cxx.h>

int
DbEnv::txn_recover(DB_PREPLIST preplist[],
    long count, long *retp, u_int32_t flags);
```

Database environment recovery restores transactions that were prepared, but not yet resolved at the time of the system shut down or crash, to their state prior to the shut down or crash, including any locks previously held. The DbEnv::txn_recover() method returns a list of those prepared transactions.

The DbEnv::txn_recover() method should only be called after the environment has been recovered.

Multiple threads of control may call DbEnv::txn_recover(), but only one thread of control may resolve each returned transaction, that is, only one thread of control may call DbTxn::commit() or DbTxn::abort() on each returned transaction. Callers of DbEnv::txn_recover() must call DbTxn::discard() to discard each transaction they do not resolve.

On return from DbEnv::txn_recover(), the **preplist** parameter will be filled in with a list of transactions that must be resolved by the application (committed, aborted or discarded). The **preplist** parameter is a structure of type DB_PREPLIST; the following DB_PREPLIST fields will be filled in:

- **DB_TXN * txn;**

    The transaction handle for the transaction.

- **u_int8_t gid[DB_XIDDATASIZE];**

    The global transaction ID for the transaction. The global transaction ID is the one specified when the transaction was prepared. The application is responsible for ensuring uniqueness among global transaction IDs.

The DbEnv::txn_recover() method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

count

The **count** parameter specifies the number of available entries in the passed-in **preplist** array. The **retp** parameter returns the number of entries DbEnv::txn_recover() has filled in, in the array.

flags

The **flags** parameter must be set to one of the following values:

- DB_FIRST

    Begin returning a list of prepared, but not yet resolved transactions. Specifying this flag begins a new pass over all prepared, but not yet completed transactions, regardless of whether they have

already been returned in previous calls to `DbEnv::txn_recover.()` Calls to `DbEnv::txn_recover()` from different threads of control should not be intermixed in the same environment.

- `DB_NEXT`

Continue returning a list of prepared, but not yet resolved transactions, starting where the last call to `DbEnv::txn_recover()` left off.

preplist

The **preplist** parameter references memory into which the list of transactions to be resolved by the application is copied.

Class

DbEnv, DbTxn

See Also

Transaction Subsystem and Related Methods

DbTxn::set_name()

```
#include <db_cxx.h>

int
DbTxn::set_name(const char *name);
```

The `DbTxn::set_name()` method associates the specified string with the transaction. The string is returned by DbEnv::txn_stat() and displayed by DbEnv::txn_stat_print().

If the database environment has been configured for logging and the Berkeley DB library was configured with --enable-diagnostic, a debugging log record is written including the transaction ID and the name.

The `DbTxn::set_name()` method may be called at any time during the life of the application.

The `DbTxn::set_name()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

name

The **name** parameter is the string to associate with the transaction.

Class

DbEnv, DbTxn

See Also

Transaction Subsystem and Related Methods

DbTxn::set_timeout()

```
#include <db_cxx.h>

u_int32_t
DbTxn::set_timeout(db_timeout_t timeout, u_int32_t flags);
```

The `DbTxn::set_timeout()` method sets timeout values for locks or transactions for the specified transaction.

Timeouts are checked whenever a thread of control blocks on a lock or when deadlock detection is performed. In the case of `DB_SET_LOCK_TIMEOUT`, the timeout is for any single lock request. In the case of `DB_SET_TXN_TIMEOUT`, the timeout is for the life of the transaction. As timeouts are only checked when the lock request first blocks or when deadlock detection is performed, the accuracy of the timeout depends on how often deadlock detection is performed.

Timeout values may be specified for the database environment as a whole. See DbEnv::set_timeout() and for more information.

The `DbTxn::set_timeout()` method configures operations performed on the underlying transaction, not only operations performed using the specified DbTxn handle.

The `DbTxn::set_timeout()` method may be called at any time during the life of the application.

The `DbTxn::set_timeout()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to one of the following values:

- `DB_SET_LOCK_TIMEOUT`

  Set the timeout value for locks in this transaction.

- `DB_SET_TXN_TIMEOUT`

  Set the timeout value for this transaction.

timeout

The **timeout** parameter is specified as an unsigned 32-bit number of microseconds, limiting the maximum timeout to roughly 71 minutes. A value of 0 disables timeouts for the transaction.

Errors

The `DbTxn::set_timeout()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv, DbTxn

See Also

Transaction Subsystem and Related Methods

DbEnv::txn_stat()

```
#include <db_cxx.h>

int
DbEnv::txn_stat(DB_TXN_STAT **statp, u_int32_t flags);
```

The `DbEnv::txn_stat()` method returns the transaction subsystem statistics.

The `DbEnv::txn_stat()` method creates a statistical structure of type `DB_TXN_STAT` and copies a pointer to it into a user-specified memory location.

Statistical structures are stored in allocated memory. If application-specific allocation routines have been declared (see DbEnv::set_alloc() for more information), they are used to allocate the memory; otherwise, the standard C library **malloc**(3) is used. The caller is responsible for deallocating the memory. To deallocate the memory, free the memory reference; references inside the returned memory need not be individually freed.

The following DB_TXN_STAT fields will be filled in:

- **DB_LSN st_last_ckp;**

  The LSN of the last checkpoint.

- **time_t st_time_ckp;**

  The time the last completed checkpoint finished (as the number of seconds since the Epoch, returned by the IEEE/ANSI Std 1003.1 (POSIX) **time** function).

- **u_int32_t st_last_txnid;**

  The last transaction ID allocated.

- **u_int32_t st_maxtxns;**

  The maximum number of active transactions configured.

- **u_int32_t st_nactive;**

  The number of transactions that are currently active.

- **u_int32_t st_nsnapshot;**

  The number of transactions on the snapshot list. These are transactions which modified a database opened with DB_MULTIVERSION, and which have committed or aborted, but the copies of pages they created are still in the cache.

- **u_int32_t st_maxnactive;**

  The maximum number of active transactions at any one time.

- **u_int32_t st_maxnsnapshot;**

The maximum number of transactions on the snapshot list at any one time.

- **uintmax_t st_nbegins;**

  The number of transactions that have begun.

- **uintmax_t st_naborts;**

  The number of transactions that have aborted.

- **uintmax_t st_ncommits;**

  The number of transactions that have committed.

- **u_int32_t st_nrestores;**

  The number of transactions that have been restored.

- **roff_t st_regsize;**

  The size of the transaction region, in bytes.

- **uintmax_t st_region_wait;**

  The number of times that a thread of control was forced to wait before obtaining the transaction region mutex.

- **uintmax_t st_region_nowait;**

  The number of times that a thread of control was able to obtain the transaction region mutex without waiting.

- **DB_TXN_ACTIVE *st_txnarray;**

  A pointer to an array of **st_nactive** DB_TXN_ACTIVE structures, describing the currently active transactions. The following fields of the DB_TXN_ACTIVE structure will be filled in:

  - **u_int32_t txnid;**

    The transaction ID of the transaction.

  - **u_int32_t parentid;**

    The transaction ID of the parent transaction (or 0, if no parent).

  - **pid_t pid;**

    The process ID of the originator of the transaction.

  - **db_threadid_t tid;**

    The thread of control ID of the originator of the transaction.

- **DB_LSN lsn;**

  The current log sequence number when the transaction was begun.

- **DB_LSN read_lsn;**

  The log sequence number of reads for snapshot transactions.

- **u_int32_t mvcc_ref;**

  The number of buffer copies created by this transaction that remain in cache.

- **u_int32_t status;**

  One of the following list of constants: TXN_ABORTED, TXN_COMMITTED, TXN_PREPARED, TXN_RUNNING.

- **u_int8_t gid[DB_GID_SIZE];**

  If the transaction was prepared using DbTxn::prepare(), then **gid** contains the transaction's Global ID. Otherwise, **gid's** contents are undefined.

- **char name[];**

  If a name was specified for the transaction, up to the first 50 bytes of that name, followed by a nul termination byte.

The `DbEnv::txn_stat()` method may not be called before the DbEnv::open() method is called.

The `DbEnv::txn_stat()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or the following value:

- `DB_STAT_CLEAR`

  Reset statistics after returning their values.

statp

The **statp** parameter references memory into which a pointer to the allocated statistics structure is copied.

Errors

The `DbEnv::txn_stat()` method may fail and throw a DbException exception, encapsulating one of the following non-zero errors, or return one of the following non-zero errors:

EINVAL

An invalid flag value or parameter was specified.

Class

DbEnv, DbTxn

See Also

Transaction Subsystem and Related Methods

DbEnv::txn_stat_print()

```
#include <db_cxx.h>

int
DbEnv::txn_stat_print(u_int32_t flags);
```

The `DbEnv::txn_stat_print()` method displays the transaction subsystem statistical information, as described for the `DbEnv::txn_stat()` method. The information is printed to a specified output channel (see the DbEnv::set_msgfile() method for more information), or passed to an application callback function (see the DbEnv::set_msgcall() method for more information).

The `DbEnv::txn_stat_print()` method may not be called before the DbEnv::open() method is called.

The `DbEnv::txn_stat_print()` method either returns a non-zero error value or throws an exception that encapsulates a non-zero error value on failure, and returns 0 on success.

Parameters

flags

The **flags** parameter must be set to 0 or by bitwise inclusively **OR**'ing together one or more of the following values:

- `DB_STAT_ALL`

  Display all available information.

- `DB_STAT_CLEAR`

  Reset statistics after displaying their values.

Class

DbEnv, DbTxn

See Also

Transaction Subsystem and Related Methods

Appendix A. Berkeley DB Command Line Utilities

The following describes the command line utilities that are available for Berkeley DB.

Utilities

Utility	Description
db_archive	Archival utility
db_checkpoint	Transaction checkpoint utility
db_deadlock	Deadlock detection utility
db_dump	Database dump utility
db_hotbackup	Hot backup utility
db_load	Database load utility
db_printlog	Transaction log display utility
db_recover	Recovery utility
db_sql	SQL schema to Berkeley DB code in C
db_stat	Statistics utility
db_upgrade	Database upgrade utility
db_verify	Verification utility

db_archive

```
db_archive [-adlsVv] [-h home] [-P password]
```

The **db_archive** utility writes the pathnames of log files that are no longer in use (for example, no longer involved in active transactions), to the standard output, one pathname per line. These log files should be written to backup media to provide for recovery in the case of catastrophic failure (which also requires a snapshot of the database files), but they may then be deleted from the system to reclaim disk space.

The options are as follows:

- **-a**

  Write all pathnames as absolute pathnames, instead of relative to the database home directory.

- **-d**

  Remove log files that are no longer needed; no filenames are written. This automatic log file removal is likely to make catastrophic recovery impossible.

- **-h**

  Specify a home directory for the database environment; by default, the current working directory is used.

- **-l**

  Write out the pathnames of all the database log files, whether or not they are involved in active transactions.

- **-P**

  Specify an environment password. Although Berkeley DB utilities overwrite password strings as soon as possible, be aware there may be a window of vulnerability on systems where unprivileged users can see command-line arguments or where utilities are not able to overwrite the memory containing the command-line arguments.

- **-s**

  Write the pathnames of all the database files that need to be archived in order to recover the database from catastrophic failure. If any of the database files have not been accessed during the lifetime of the current log files, **db_archive** will not include them in this output.

  It is possible that some of the files to which the log refers have since been deleted from the system. In this case, **db_archive** will ignore them. When db_recover is run, any files to which the log refers that are not present during recovery are assumed to have been deleted and will not be recovered.

- **-V**

  Write the library version number to the standard output, and exit.

- **-v**

    Run in verbose mode.

Log cursor handles (returned by the DbEnv::log_cursor() method) may have open file descriptors for log files in the database environment. Also, the Berkeley DB interfaces to the database environment logging subsystem (for example, DbEnv::log_put() and DbTxn::abort() may allocate log cursors and have open file descriptors for log files as well. On operating systems where filesystem related system calls (for example, rename and unlink on Windows/NT) can fail if a process has an open file descriptor for the affected file, attempting to move or remove the log files listed by **db_archive** may fail. All Berkeley DB internal use of log cursors operates on active log files only and furthermore, is short-lived in nature. So, an application seeing such a failure should be restructured to close any open log cursors it may have, and otherwise to retry the operation until it succeeds. (Although the latter is not likely to be necessary; it is hard to imagine a reason to move or rename a log file in which transactions are being logged or aborted.)

The **db_archive** utility uses a Berkeley DB environment (as described for the **-h** option, the environment variable **DB_HOME**, or because the utility was run in a directory containing a Berkeley DB environment). In order to avoid environment corruption when using a Berkeley DB environment, **db_archive** should always be given the chance to detach from the environment and exit gracefully. To cause **db_archive** to release all environment resources and exit cleanly, send it an interrupt signal (SIGINT).

The **db_archive** utility exits 0 on success, and >0 if an error occurs.

Environment Variables

DB_HOME

If the **-h** option is not specified and the environment variable DB_HOME is set, it is used as the path of the database home, as described in the DbEnv::open() method.

db_checkpoint

```
db_checkpoint [-1Vv] [-h home]
    [-k kbytes] [-L file] [-P password] [-p min]
```

The **db_checkpoint** utility is a daemon process that monitors the database log, and periodically calls DbEnv::txn_checkpoint() to checkpoint it.

The options are as follows:

- **-1**

  Force a single checkpoint of the log (regardless of whether or not there has been activity since the last checkpoint), and then exit.

  When the **-1** flag is specified, the **db_checkpoint** utility will checkpoint the log even if unable to find an existing database environment. This functionality is useful when upgrading database environments from one version of Berkeley DB to another.

- **-h**

  Specify a home directory for the database environment; by default, the current working directory is used.

- **-k**

  Checkpoint the database at least as often as every **kbytes** of log file are written.

- **-L**

  Log the execution of the **db_checkpoint** utility to the specified file in the following format, where ### is the process ID, and the date is the time the utility was started.

  ```
  db_checkpoint: ### Wed Jun 15 01:23:45 EDT 1995
  ```

  This file will be removed if the **db_checkpoint** utility exits gracefully.

- **-P**

  Specify an environment password. Although Berkeley DB utilities overwrite password strings as soon as possible, be aware there may be a window of vulnerability on systems where unprivileged users can see command-line arguments or where utilities are not able to overwrite the memory containing the command-line arguments.

- **-p**

  Checkpoint the database at least every **min** minutes if there has been any activity since the last checkpoint.

- **-V**

  Write the library version number to the standard output, and exit.

- **-v**

  Write the time of each checkpoint attempt to the standard output.

At least one of the **-1**, **-k**, and **-p** options must be specified.

The **db_checkpoint** utility uses a Berkeley DB environment (as described for the **-h** option, the environment variable **DB_HOME**, or because the utility was run in a directory containing a Berkeley DB environment). In order to avoid environment corruption when using a Berkeley DB environment, **db_checkpoint** should always be given the chance to detach from the environment and exit gracefully. To cause **db_checkpoint** to release all environment resources and exit cleanly, send it an interrupt signal (SIGINT).

The **db_checkpoint** utility does not attempt to create the Berkeley DB shared memory regions if they do not already exist. The application that creates the region should be started first, and once the region is created, the **db_checkpoint** utility should be started.

The **db_checkpoint** utility exits 0 on success, and >0 if an error occurs.

Environment Variables

DB_HOME

If the **-h** option is not specified and the environment variable DB_HOME is set, it is used as the path of the database home, as described in the DbEnv::open() method.

db_deadlock

```
db_deadlock [-Vv]
    [-a e | m | n | o | W | w | y] [-h home] [-L file] [-t sec.usec]
```

The **db_deadlock** utility traverses the database environment lock region, and aborts a lock request each time it detects a deadlock or a lock request that has timed out. By default, in the case of a deadlock, a random lock request is chosen to be aborted.

This utility should be run as a background daemon, or the underlying Berkeley DB deadlock detection interfaces should be called in some other way, whenever there are multiple threads or processes accessing a database and at least one of them is modifying it.

The options are as follows:

- **-a**

  When a deadlock is detected, abort the locker:

  - **m**

    with the most locks

  - **n**

    with the fewest locks

  - **o**

    with the oldest locks

  - **W**

    with the most write locks

  - **w**

    with the fewest write locks

  - **y**

    with the youngest locks

  - **e**

    When lock or transaction timeouts have been specified, abort any lock request that has timed out. Note that this option does not perform the entire deadlock detection algorithm, but instead only checks for timeouts.

- **-h**

Specify a home directory for the database environment; by default, the current working directory is used.

- **-L**

  Log the execution of the **db_deadlock** utility to the specified file in the following format, where ### is the process ID, and the date is the time the utility was started.

  ```
          db_deadlock: ### Wed Jun 15 01:23:45 EDT 1995
  ```

  This file will be removed if the **db_deadlock** utility exits gracefully.

- **-t**

  Check the database environment every **sec** seconds plus **usec** microseconds to see if a process has been forced to wait for a lock; if one has, review the database environment lock structures.

- **-V**

  Write the library version number to the standard output, and exit.

- **-v**

  Run in verbose mode, generating messages each time the detector runs.

If the **-t** option is not specified, **db_deadlock** will run once and exit.

The **db_deadlock** utility uses a Berkeley DB environment (as described for the **-h** option, the environment variable **DB_HOME**, or because the utility was run in a directory containing a Berkeley DB environment). In order to avoid environment corruption when using a Berkeley DB environment, **db_deadlock** should always be given the chance to detach from the environment and exit gracefully. To cause **db_deadlock** to release all environment resources and exit cleanly, send it an interrupt signal (SIGINT).

The **db_deadlock** utility does not attempt to create the Berkeley DB shared memory regions if they do not already exist. The application which creates the region should be started first, and then, once the region is created, the **db_deadlock** utility should be started.

The **db_deadlock** utility exits 0 on success, and >0 if an error occurs.

Environment Variables

DB_HOME

If the **-h** option is not specified and the environment variable DB_HOME is set, it is used as the path of the database home, as described in the DbEnv::open() method.

db_dump

```
db_dump [-klNpRrV] [-d ahr]
 [-f output] [-h home] [-P password] [-s database] file

db_dump [-kNpV] [-d ahr] [-f output] [-h home] -m database

db_dump185 [-p] [-f output] file
```

The **db_dump** utility reads the database file **file** and writes it to the standard output using a portable flat-text format understood by the db_load utility. The **file** argument must be a file produced using the Berkeley DB library functions.

The **db_dump185** utility is similar to the **db_dump** utility, except that it reads databases in the format used by Berkeley DB versions 1.85 and 1.86.

The options are as follows:

- **-d**

  Dump the specified database in a format helpful for debugging the Berkeley DB library routines.

  - **a**

    Display all information.

  - **h**

    Display only page headers.

  - **r**

    Do not display the free-list or pages on the free list. This mode is used by the recovery tests.

  The output format of the -d option is not standard and may change, without notice, between releases of the Berkeley DB library.

- **-f**

  Write to the specified file instead of to the standard output.

- **-h**

  Specify a home directory for the database environment; by default, the current working directory is used.

- **-k**

  Dump record numbers from Queue and Recno databases as keys.

- **-l**

List the databases stored in the file.

- **-m**

  Specify a named in-memory database to dump. In this case the **file** argument must be omitted.

- **-N**

  Do not acquire shared region mutexes while running. Other problems, such as potentially fatal errors in Berkeley DB, will be ignored as well. This option is intended only for debugging errors, and should not be used under any other circumstances.

- **-P**

  Specify an environment password. Although Berkeley DB utilities overwrite password strings as soon as possible, be aware there may be a window of vulnerability on systems where unprivileged users can see command-line arguments or where utilities are not able to overwrite the memory containing the command-line arguments.

- **-p**

  If characters in either the key or data items are printing characters (as defined by **isprint**(3)), use printing characters in **file** to represent them. This option permits users to use standard text editors and tools to modify the contents of databases.

  Note: different systems may have different notions about what characters are considered *printing characters*, and databases dumped in this manner may be less portable to external systems.

- **-R**

  Aggressively salvage data from a possibly corrupt file. The **-R** flag differs from the **-r** option in that it will return all possible data from the file at the risk of also returning already deleted or otherwise nonsensical items. Data dumped in this fashion will almost certainly have to be edited by hand or other means before the data is ready for reload into another database

- **-r**

  Salvage data from a possibly corrupt file. When used on a uncorrupted database, this option should return equivalent data to a normal dump, but most likely in a different order.

- **-s**

  Specify a single database to dump. If no database is specified, all databases in the database file are dumped.

- **-V**

  Write the library version number to the standard output, and exit.

Dumping and reloading Hash databases that use user-defined hash functions will result in new databases that use the default hash function. Although using the default hash function may not be optimal for the new database, it will continue to work correctly.

Dumping and reloading Btree databases that use user-defined prefix or comparison functions will result in new databases that use the default prefix and comparison functions. **In this case, it is quite likely that the database will be damaged beyond repair permitting neither record storage or retrieval.**

The only available workaround for either case is to modify the sources for the db_load utility to load the database using the correct hash, prefix, and comparison functions.

The **db_dump185** utility may not be available on your system because it is not always built when the Berkeley DB libraries and utilities are installed. If you are unable to find it, see your system administrator for further information.

The **db_dump** and **db_dump185** utility output formats are documented in the Dump Output Formats section of the Berkeley DB Reference Guide.

The **db_dump** utility may be used with a Berkeley DB environment (as described for the **-h** option, the environment variable **DB_HOME**, or because the utility was run in a directory containing a Berkeley DB environment). In order to avoid environment corruption when using a Berkeley DB environment, **db_dump** should always be given the chance to detach from the environment and exit gracefully. To cause **db_dump** to release all environment resources and exit cleanly, send it an interrupt signal (SIGINT).

Even when using a Berkeley DB database environment, the db_dump utility does not use any kind of database locking if it is invoked with the **-d**, **-R**, or **-r** arguments. If used with one of these arguments, the **db_dump** utility may only be safely run on databases that are not being modified by any other process; otherwise, the output may be corrupt.

The **db_dump** utility exits 0 on success, and >0 if an error occurs.

The **db_dump185** utility exits 0 on success, and >0 if an error occurs.

Environment Variables

DB_HOME

If the **-h** option is not specified and the environment variable DB_HOME is set, it is used as the path of the database home, as described in the DbEnv::open() method.

G E B K R W E D F N X S

```
db_hotbackup [-cDuVv] [-d data_dir ...] [-h home]
    [-l log_dir] [-P password] -b backup_dir
```

7 K H G E B K R W E D F N X S X W L O L W B F U H D W H V D K R W E D F N X S R I B R X U O L Y H H Q Y L U R Q P H Q W

7 K H G E B K R W E D F N X S X W L O L W B S H U I R U P V W K H I R O O R Z L Q J V W H S V

, I W K H R S W L R Q L V Q R W V S H F L I L H G F K H F N S R L Q W B R X U K R P H G D W D E D V H H Q Y L U R Q P H Q W I R U X Q Q H F H V V D U B I L O H V

, I W K H D U G H O H W H F R S B W K H D F N X S R S H Q R Q O B L I L W V F U H D W L R Q W L P H L V R O G H U W K D Q D Q B Z U L W H R U H [H F X W W K R Z Q H U

, I W K H D U G H O H W H F R S B W K H D F N X S R S H Q R Q O B L I L W V F U H D W L R Q W L P H L V R O G H U W K D Q D Q B Z U L W H R U H [H F X W W K H D U G H O H W H F D W D O R J H U U H P R Y H G

, I W K H X R S W L R Q L V Q R W V S H F L I L H G D S S O L F D W L R Q O R J V K D W D E D V H H Q Y L U R Q P H Q W K R P H G L U H F W R U B D Q G G L U H F W R U B S H U L V Z K R G R S W L R Q W R W K H D U G H O H W H F D W D O R J H E D F N X S

& R S B D O O R I L I O H R X O G V W K H G L U H F W R U B S H U F L E L O L W B H G R S W L R Q W R R L Q W R W K H G D W D E D V H H Q Y L U R Q P H Q W K R P H G L U H F W R U B R S W L R Q V S H F L I L H G W R W K H D U G H O H W H F D W D O R J H E D F N X S

3 H U I R U P D W D V W U H H R S R K B L O H V W K H D U G H O H W H F D W D O R J H E D F N X S

5 H P R Y D O O X Q Q H F H V V D U B O R J I L O H V I U R P W K H D U G H O H W H F D W D O R J H E D F N X S

7 K H G E B K R W E D F N X S X W L O L W B S X W V R R U W H V R S O H Y H Q G V W D J Q V D V W F K D L U R U L H Q W V K H U H S D V V U H G D W H $ S S O L F D W L R Q W K D W F U H D W H E B 7 [Q B S U H S B P X W H S H F L B L B 5 (& 2 9 (5 B) Z K H R S H Q W L G H H Q Y L U R Q P H Q W D Q G X ' E (Q Y B W [Q B U B W R G H Y R D O O S H H Q G V W D J Q V D Z K H L G Q L Q O R L Y H U V R W K H E D F N X S

7 K H R S W L R Q D U H D V I R O O R Z V

‡ E

6 S H F L I B W K H D U G H O H W H F D W D E D F N X S

‡ F

% H I R U H S H U I R U P W K E D F N R K H F N S W R W K H V R X U G F D W D E O H Q Y H U R Q P L G G P D W R D Y H O R J I L O W K H D U G H O R Q O U H T X L U W K H G E O L Y L U R Q B W R R E N Q R P G N L Q J D V W U H R S K B L F U B L P S R V V O E O L K H U H P R Y P X O L H L Q W H J Z D W K R G L D O H F K L Y D O

‡ '

8 V H W K H G D W D D Q G O R G G L U H F W O R V U L W L D V % B & 2 1) F R Q I L J X U D W L O W K H R X G G L H F W R U B 7 K L R S W K L R O R X U I H F W V

- The specified data and log directories will be created relative to the target directory, with mode read-write-execute owner, if they do not already exist.

- In step #3 above, all files in any source data directories specified in the DB_CONFIG file will be copied to the target data directories.

- In step #4 above, log files will be copied from any log directory specified in the DB_CONFIG file, instead of from the default locations.

- The DB_CONFIG configuration file will be copied from the source directory to the target directory, and subsequently used for configuration if recovery is run in the target directory.

Care should be taken with the **-D** option where data and log directories are named relative to the source directory but are not subdirectories (that is, the name includes the element "..") Specifically, the constructed target directory names must be meaningful and distinct from the source directory names, otherwise running recovery in the target directory might corrupt the source data files.

It is an error to use absolute pathnames for data or log directories in this mode, as the DB_CONFIG configuration file copied into the target directory would then point at the source directories and running recovery would corrupt the source data files.

- **-d**

  Specify one or more directories that contain data files to be copied to the target directory.

  As all database files are copied into a single target directory, files named the same, stored in different source directories, would overwrite each other when copied to the target directory.

  Please note the database environment recovery log references database files as they are named by the application program. **If the application uses absolute or relative pathnames to name database files, (rather than filenames and the DbEnv::set_data_dir() method or the DB_CONFIG configuration file to specify filenames), running recovery in the target directory may not properly find the copies of the files or might even find the source files, potentially resulting in corruption.**

- **-h**

  Specify the source directory for the backup. That is, the database environment home directory.

- **-l**

  Specify a source directory that contains log files; if none is specified, the database environment home directory will be searched for log files.

- **-P**

  Specify an environment password. Although Berkeley DB utilities overwrite password strings as soon as possible, be aware there may be a window of vulnerability on systems where unprivileged users can see command-line arguments or where utilities are not able to overwrite the memory containing the command-line arguments.

- **-u**

Update a pre-existing hot backup snapshot by copying in new log files. If the **-u** option is specified, no databases will be copied into the target directory.

- **-V**

  Write the library version number to the standard output, and exit.

- **-v**

  Run in verbose mode, listing operations as they are done.

The **db_hotbackup** utility uses a Berkeley DB environment (as described for the **-h** option, the environment variable **DB_HOME**, or because the utility was run in a directory containing a Berkeley DB environment). In order to avoid environment corruption when using a Berkeley DB environment, **db_hotbackup** should always be given the chance to detach from the environment and exit gracefully. To cause **db_hotbackup** to release all environment resources and exit cleanly, send it an interrupt signal (SIGINT).

The **db_hotbackup** utility exits 0 on success, and >0 if an error occurs.

Environment Variables

DB_HOME

If the **-h** option is not specified and the environment variable DB_HOME is set, it is used as the path of the database home, as described in the DbEnv::open() method.

db_load

```
db_load [-nTV] [-c name=value] [-f file]
    [-h home] [-P password] [-t btree | hash | queue | recno] file

db_load [-r lsn | fileid] [-h home] [-P password] file
```

The **db_load** utility reads from the standard input and loads it into the database **file**. The database **file** is created if it does not already exist.

The input to **db_load** must be in the output format specified by the db_dump utility, utilities, or as specified for the **-T** below.

The options are as follows:

- **-c**

  Specify configuration options ignoring any value they may have based on the input. The command-line format is **name=value**. See the Supported Keywords section below for a list of keywords supported by the **-c** option.

- **-f**

  Read from the specified **input** file instead of from the standard input.

- **-h**

  Specify a home directory for the database environment.

  If a home directory is specified, the database environment is opened using the DB_INIT_LOCK, DB_INIT_LOG, DB_INIT_MPOOL, DB_INIT_TXN, and DB_USE_ENVIRON flags to DbEnv::open() (This means that **db_load** can be used to load data into databases while they are in use by other processes.) If the DbEnv::open() call fails, or if no home directory is specified, the database is still updated, but the environment is ignored; for example, no locking is done.

- **-n**

  Do not overwrite existing keys in the database when loading into an already existing database. If a key/data pair cannot be loaded into the database for this reason, a warning message is displayed on the standard error output, and the key/data pair are skipped.

- **-P**

  Specify an environment password. Although Berkeley DB utilities overwrite password strings as soon as possible, be aware there may be a window of vulnerability on systems where unprivileged users can see command-line arguments or where utilities are not able to overwrite the memory containing the command-line arguments.

- **-r**

  Reset the database's file ID or log sequence numbers (LSNs).

All database pages in transactional environments contain references to the environment's log records. In order to copy a database into a different database environment, database page references to the old environment's log records must be reset, otherwise data corruption can occur when the database is modified in the new environment. The **-r lsn** option resets a database's log sequence numbers.

All databases contain an ID string used to identify the database in the database environment cache. If a database is copied, and used in the same environment as another file with the same ID string, corruption can occur. The **-r fileid** option resets a database's file ID to a new value.

In both cases, the physical file specified by the file argument is modified in-place.

- **-T**

  The **-T** option allows non-Berkeley DB applications to easily load text files into databases.

  If the database to be created is of type Btree or Hash, or the keyword **keys** is specified as set, the input must be paired lines of text, where the first line of the pair is the key item, and the second line of the pair is its corresponding data item. If the database to be created is of type Queue or Recno and the keyword **keys** is not set, the input must be lines of text, where each line is a new data item for the database.

  A simple escape mechanism, where newline and backslash (\) characters are special, is applied to the text input. Newline characters are interpreted as record separators. Backslash characters in the text will be interpreted in one of two ways: If the backslash character precedes another backslash character, the pair will be interpreted as a literal backslash. If the backslash character precedes any other character, the two characters following the backslash will be interpreted as a hexadecimal specification of a single character; for example, \0a is a newline character in the ASCII character set.

  For this reason, any backslash or newline characters that naturally occur in the text input must be escaped to avoid misinterpretation by **db_load**.

  If the **-T** option is specified, the underlying access method type must be specified using the **-t** option.

- **-t**

  Specify the underlying access method. If no **-t** option is specified, the database will be loaded into a database of the same type as was dumped; for example, a Hash database will be created if a Hash database was dumped.

  Btree and Hash databases may be converted from one to the other. Queue and Recno databases may be converted from one to the other. If the **-k** option was specified on the call to db_dump then Queue and Recno databases may be converted to Btree or Hash, with the key being the integer record number.

- **-V**

  Write the library version number to the standard output, and exit.

The **db_load** utility may be used with a Berkeley DB environment (as described for the **-h** option, the environment variable **DB_HOME**, or because the utility was run in a directory containing a Berkeley

DB environment). In order to avoid environment corruption when using a Berkeley DB environment, **db_load** should always be given the chance to detach from the environment and exit gracefully. To cause **db_load** to release all environment resources and exit cleanly, send it an interrupt signal (SIGINT).

The **db_load** utility exits 0 on success, 1 if one or more key/data pairs were not loaded into the database because the key already existed, and >1 if an error occurs.

Examples

The **db_load** utility can be used to load text files into databases. For example, the following command loads the standard UNIX */etc/passwd* file into a database, with the login name as the key item and the entire password entry as the data item:

```
awk -F: '{print $1; print $0}' < /etc/passwd |
        sed 's/\\/\\\\/g' | db_load -T -t hash passwd.db
```

Note that backslash characters naturally occurring in the text are escaped to avoid interpretation as escape characters by **db_load**.

Environment Variables

DB_HOME

If the **-h** option is not specified and the environment variable DB_HOME is set, it is used as the path of the database home, as described in the DbEnv::open() method.

Supported Keywords

The following keywords are supported for the **-c** command-line option to the **db_load** utility. See the DbEnv::open() method for further discussion of these keywords and what values should be specified.

The parenthetical listing specifies how the value part of the **name=value** pair is interpreted. Items listed as (boolean) expect value to be **1** (set) or **0** (unset). Items listed as (number) convert value to a number. Items listed as (string) use the string value without modification.

- **bt_minkey (number)**

  The minimum number of keys per page.

- **chksum (boolean)**

  Enable page checksums.

- **database (string)**

  The database to load.

- **db_lorder (number)**

  The byte order for integers in the stored database metadata. For big endian systems, the order should be 4,321 while for little endian systems is should be 1,234.

- **db_pagesize (number)**

  The size of database pages, in bytes.

- **duplicates (boolean)**

  The value of the DB_DUP flag.

- **dupsort (boolean)**

  The value of the DB_DUPSORT flag.

- **extentsize (number)**

  The size of database extents, in pages, for Queue databases configured to use extents.

- **h_ffactor (number)**

  The density within the Hash database.

- **h_nelem (number)**

  The size of the Hash database.

- **keys (boolean)**

  Specify whether keys are present for Queue or Recno databases.

- **re_len (number)**

  Specify the length for fixed-length records. This number represents different things, depending on the access method the database is using. See the Db::set_re_len() method for details on what this number represents.

- **re_pad (string)**

  Specify the fixed-length record pad character.

- **recnum (boolean)**

  The value of the DB_RECNUM flag.

- **renumber (boolean)**

  The value of the DB_RENUMBER flag.

- **subdatabase (string)**

  The subdatabase to load.

db_printlog

```
db_printlog [-NrV] [-b start-LSN] [-e stop-LSN] [-h home] [-P password]
```

The **db_printlog** utility is a debugging utility that dumps Berkeley DB log files in a human-readable format.

The options are as follows:

- **-b**

  Display log records starting at log sequence number (LSN) **start-LSN**; **start-LSN** is specified as a file number, followed by a slash (/) character, followed by an offset number, with no intervening whitespace.

- **-e**

  Stop displaying log records at log sequence number (LSN) **stop-LSN**; **stop-LSN** is specified as a file number, followed by a slash (/) character, followed by an offset number, with no intervening whitespace.

- **-h**

  Specify a home directory for the database environment; by default, the current working directory is used.

- **-N**

  Do not acquire shared region mutexes while running. Other problems, such as potentially fatal errors in Berkeley DB, will be ignored as well. This option is intended only for debugging errors, and should not be used under any other circumstances.

- **-P**

  Specify an environment password. Although Berkeley DB utilities overwrite password strings as soon as possible, be aware there may be a window of vulnerability on systems where unprivileged users can see command-line arguments or where utilities are not able to overwrite the memory containing the command-line arguments.

- **-r**

  Read the log files in reverse order.

- **-V**

  Write the library version number to the standard output, and exit.

For more information on the **db_printlog** output and using it to debug applications, see Reviewing Berkeley DB log files.

The **db_printlog** utility uses a Berkeley DB environment (as described for the **-h** option, the environment variable **DB_HOME**, or because the utility was run in a directory containing a Berkeley DB environment).

In order to avoid environment corruption when using a Berkeley DB environment, **db_printlog** should always be given the chance to detach from the environment and exit gracefully. To cause **db_printlog** to release all environment resources and exit cleanly, send it an interrupt signal (SIGINT).

The **db_printlog** utility exits 0 on success, and >0 if an error occurs.

Environment Variables

DB_HOME

If the **-h** option is not specified and the environment variable DB_HOME is set, it is used as the path of the database home, as described in the DbEnv::open() method.

db_recover

```
db_recover [-cefVv] [-h home] [-P password] [-t [[CC]YY]MMDDhhmm[.SS]]]
```

The **db_recover** utility must be run after an unexpected application, Berkeley DB, or system failure to restore the database to a consistent state. All committed transactions are guaranteed to appear after **db_recover** has run, and all uncommitted transactions will be completely undone.

Note that this utility performs the same action as if the environment is opened with the DB_RECOVER flag. If DB_RECOVER is specified on environment open, then use of this utility is not necessary.

The options are as follows:

- **-c**

  Perform catastrophic recovery instead of normal recovery.

- **-e**

  Retain the environment after running recovery. This option will rarely be used unless a DB_CONFIG file is present in the home directory. If a DB_CONFIG file is not present, then the regions will be created with default parameter values.

- **-f**

  Display a message on the standard output showing the percent of recovery completed.

- **-h**

  Specify a home directory for the database environment; by default, the current working directory is used.

- **-P**

  Specify an environment password. Although Berkeley DB utilities overwrite password strings as soon as possible, be aware there may be a window of vulnerability on systems where unprivileged users can see command-line arguments or where utilities are not able to overwrite the memory containing the command-line arguments.

- **-t**

  Recover to the time specified rather than to the most current possible date. The timestamp argument should be in the form [[CC]YY]MMDDhhmm[.SS] where each pair of letters represents the following:

  - **CC**

    The first two digits of the year (the century).

  - **YY**

    The second two digits of the year. If "YY" is specified, but "CC" is not, a value for "YY" between 69 and 99 results in a "CC" value of 19. Otherwise, a "CC" value of 20 is used.

- **MM**

  The month of the year, from 1 to 12.

- **DD**

  The day of the month, from 1 to 31.

- **hh**

  The hour of the day, from 0 to 23.

- **mm**

  The minute of the hour, from 0 to 59.

- **SS**

  The second of the minute, from 0 to 61.

  If the "CC" and "YY" letter pairs are not specified, the values default to the current year. If the "SS" letter pair is not specified, the value defaults to 0.

- **-V**

  Write the library version number to the standard output, and exit.

- **-v**

  Run in verbose mode.

In the case of catastrophic recovery, an archival copy — or *snapshot* — of all database files must be restored along with all of the log files written since the database file snapshot was made. (If disk space is a problem, log files may be referenced by symbolic links). For further information on creating a database snapshot, see Archival Procedures. For further information on performing recovery, see Recovery Procedures.

If the failure was not catastrophic, the files present on the system at the time of failure are sufficient to perform recovery.

If log files are missing, **db_recover** will identify the missing log file(s) and fail, in which case the missing log files need to be restored and recovery performed again.

The **db_recover** utility uses a Berkeley DB environment (as described for the **-h** option, the environment variable **DB_HOME**, or because the utility was run in a directory containing a Berkeley DB environment). In order to avoid environment corruption when using a Berkeley DB environment, **db_recover** should always be given the chance to detach from the environment and exit gracefully. To cause **db_recover** to release all environment resources and exit cleanly, send it an interrupt signal (SIGINT).

The **db_recover** utility exits 0 on success, and >0 if an error occurs.

Environment Variables

DB_HOME

If the **-h** option is not specified and the environment variable DB_HOME is set, it is used as the path of the database home, as described in the DbEnv::open() method.

db_sql

```
db_sql [-i <ddl input file>] [-o <output C code file>]
    [-h <output header file>] [-t <test output file>]
```

Db_sql is a utility program that translates a schema description written in a SQL Data Definition Language dialect into C code that implements the schema using Berkeley DB. It is intended to provide a quick and easy means of getting started with Berkeley DB for users who are already conversant with SQL. It also introduces a convenient way to express a Berkeley DB schema in a format that is both external to the program that uses it and compatible with relational databases.

The **db_sql** command reads DDL from an input stream, and writes C code to an output stream. With no command line options, it will read from stdin and write to stdout. A more common usage mode would be to supply the DDL in a named input file (-i option). With only the -i option, **db_sql** will produce two files: a C-language source code (.c) file and a C-language header (.h) file, with names that are derived from the name of the input file. You can also control the names of these output files with the -o and -h options. Finally, the -t option will produce a simple application that invokes the generated function API. This is a C-language source file that includes a main function, and serves the dual purposes of providing a simple test for the generated C code, and of being an example of how to use the generated API.

The options are as follows:

- **-i**<ddl input file>

  Names the input file containing SQL DDL.

- **-o** <output C code file>

  Names the output C-language source code file.

- **-h** <output header file>

  Names the output C-language header file.

- **-t** <test output file>

  Names the output C-langage test file.

The **db_sql** utility exits 0 on success, and >0 if an error occurs.

Input Syntax

The input file can contain the following SQL DDL statements.

- **CREATE DATABASE**

  The DDL must contain a CREATE DATABASE statement. The syntax is simply

  ```
  CREATE DATABASE name;
  ```

. The name given here is used as the name of the Berkeley DB environment in which the Berkeley DB databases are created.

- **CREATE TABLE**

  Each CREATE TABLE statement produces functions to create and delete a primary Berkeley DB database. Also produced are functions to perform record insertion, retrieval and deletion on this database.

  CREATE TABLE establishes the field set of records that can be stored in the Berkeley DB database. Every CREATE TABLE statement must identify a primary key to be used as the lookup key in the Berkeley DB database.

  Here is an example to illustrate the syntax of CREATE TABLE that is accepted by **db_sql**:

```
CREATE TABLE person (person_id INTEGER PRIMARY KEY,
                     name VARCHAR(64),
                     age INTEGER);
```

  This results in the creation of functions to manage a database in which every record is an instance of the following C language data structure:

```
typedef struct _person_data {
    int person_id;
    char name[PERSON_DATA_NAME_LENGTH];
    int age;
} person_data;
```

- **CREATE INDEX** You can create secondary Berkeley DB databases to be used as indexes into a primary database. For example, to make an index on the "name" field of the "person" table mentioned above, the SQL DDL would be:

```
CREATE INDEX name_index ON person(name);
```

  This causes **db_sql** to emit functions to manage creation and deletion of a secondary database called "name_index," which is associated with the "person" database and is set up to perform lookups on the "name" field.

Hint Comments

The SQL DDL input may contain comments. Two types of comments are recognized. C-style comments begin with "/*" and end with "*/". These comments may extend over multiple lines.

Single line comments begin with "--" and run to the end of the line.

If the first character of a comment is "+" then the comment is interpreted as a "hint comment." Hint comments can be used to configure Berkeley DB features that cannot be represented in SQL DDL.

Hint comments are comma-separated lists of property assignments of the form "property=value." Hint comments apply to the SQL DDL statement that immediately precedes their appearance in the input. For example:

```
CREATE DATABASE peopledb; /*+ CACHESIZE = 16m */
```

This causes the generated environment creation function to set the cache size to sixteen megabytes.

In addition to the CACHESIZE example above, there is only one other hint comment that is currently recognized: After a CREATE TABLE statement, you may set the database type by assigning the DBTYPE property in a hint comment. Possible values for DBTYPE are BTREE and HASH.

Type Mapping

db_sql must map the schema expressed as SQL types into C language types. It implements the following mappings:

```
BIN         char[]
VARBIN      char[]
CHAR        char[]
VARCHAR     char[]
VARCHAR2    char[]
BIT         char
TINYINT     char
SMALLINT    short
INTEGER     int
INT         int
BIGINT      long
REAL        float
DOUBLE      double
FLOAT       double
DECIMAL     double
NUMERIC     double
NUMBER(p,s) int, long, float, or double
```

While BIN/VARBIN and CHAR/VARCHAR are both represented as char arrays, the latter are treated as null-terminated C strings, while the former are treated as binary data.

The Oracle type NUMBER is mapped to different C types, depending on its precision and scale values. If scale is 0, then it is mapped to an integer type (long if precision is greater than 9). Otherwise it is mapped to a floating point type (float if precision is less than 7, otherwise double).

Output

Depending on the options given on the command line, **db_sql** can produce three separate files: a .c file containing function definitions that implement the generated API; a .h file containing constants, data structures and prototypes of the generated functions; and a second .c file that contains a sample program that invokes the generated API. The latter program is usually referred to as a smoke test.

Given the following sample input in a file named "people.sql":

```
CREATE DATABASE peopledb;
CREATE TABLE person (person_id INTEGER PRIMARY KEY,
                name VARCHAR(64),
```

```
                       age INTEGER);
CREATE INDEX name_index ON person(name);
```

The command

```
db_sql -i people.sql -t test_people.c
```

Will produce files named people.h, people.c, and test_people.c.

The file people.h will contain the information needed to use the generated API. Among other things, an examination of the generated .h file will reveal:

```
#define PERSON_DATA_NAME_LENGTH 63
```

This is just a constant for the length of the string mapped from the VARCHAR field.

```
typedef struct _person_data {
    int    person_id;
    char   name[PERSON_DATA_NAME_LENGTH];
    int    age;
} person_data;
```

This is the data structure that represents the record type that is stored in the person database. There's that constant being used.

```
int create_peopledb_env(DB_ENV **envpp);
int create_person_database(DB_ENV *envp, DB **dbpp);
int create_name_index_secondary(DB_ENV *envp, DB *primary_dbp,
                          DB **secondary_dbpp);
```

These functions must be invoked to initialize the Berkeley DB environment. However, see the next bit:

```
extern DB_ENV * peopledb_envp;
extern DB *person_dbp;
extern DB *name_index_dbp;

int initialize_peopledb_environment();
```

For convenience, **db_sql** provides global variables for the environment and database, and a single initialization function that sets up the environment for you. You may choose to use the globals and the single initialization function, or you may declare your own DB_ENV and DB pointers, and invoke the individual create_* functions yourself.

The word "create" in these function names might be confusing. It means "create the environment/database if it doesn't already exist; otherwise open it."

All of the functions in the generated API return Berkeley DB error codes. If the return value is non-zero, there was an error of some kind, and an explanatory message should have been printed on stderr.

```
int person_insert_struct(DB *dbp, person_data *personp);
int person_insert_fields(DB * dbp,
```

```
        int person_id,
        char *name,
        int age);
```

These are the functions that you'd use to store a record in the database. The first form takes a pointer to the data structure that represents this record. The second form takes each field as a separate argument.

If two records with the same primary key value are stored, the first one is lost.

```
int get_person_data(DB *dbp, int person_key, person_data *data);
```

This function retrieves a record from the database. It seeks the record with the supplied key, and populates the supplied structure with the contents of the record. If no matching record is found, the function returns DB_NOTFOUND.

```
int delete_person_key(DB *dbp, int person_key);
```

This function removes the record matching the given key.

```
typedef void (*person_iteration_callback)(void *user_data,
                                          person_data *personp);

int person_full_iteration(DB *dbp,
    person_iteration_callback user_func,
    void *user_data);
```

This function performs a complete iteration over every record in the person table. The user must provide a callback function which is invoked once for every record found. The user's callback function must match the prototype provided in the typedef "person_iteration_callback." In the callback, the "user_data" argument is passed unchanged from the "user_data" argument given to person_full_iteration. This is provided so that the caller of person_full_iteration can communicate some context information to the callback function. The "personp" argument to the callback is a pointer to the record that was retrieved from the database. Personp points to data that is valid only for the duration of the callback invocation.

```
int name_index_query_iteration(DB *secondary_dbp,
    char *name_index_key,
    person_iteration_callback user_func,
    void *user_data);
```

This function performs lookups through the secondary index database. Because duplicate keys are allowed in secondary indexes, this query might return multiple instances. This function takes as an argument a pointer to a user-written callback function, which must match the function prototype typedef mentioned above (person_iteration_callback). The callback is invoked once for each record that matches the secondary key.

Test output

The test output file is useful as an example of how to invoke the generated API. It will contain calls to the functions mentioned above, to store a single record and retrieve it by primary key and through the secondary index.

To compile the test, you would issue a command such as

```
cc -I$BDB_INSTALL/include -L$BDB_INSTALL/lib -o test_people people.c \
    test_people.c -ldb-4.8
```

This will produce the executable file test_people, which can be run to exercise the generated API. The program generated from people.sql will create a database environment in a directory named "peopledb." This directory must be created before the program is run.

DB C++ API

db_stat

```
db_stat -d file [-fN] [-h home] [-P password] [-s database]

db_stat [-cEelmNrtVxZ] [-C Aclop] [-h home] [-L A] [-M Ah] [-R A]
        [-P password]
```

The **db_stat** utility displays statistics for Berkeley DB environments.

The options are as follows:

- **-C**

  Display detailed information about the locking subsystem.

  - **A**

    Display all information.

  - **c**

    Display lock conflict matrix.

  - **l**

    Display lockers within hash chains.

  - **o**

    Display lock objects within hash chains.

  - **p**

    Display locking subsystem parameters.

- **-c**

  Display locking subsystem statistics, as described in the DbEnv::lock_stat() method.

- **-d**

  Display database statistics for the specified file, as described in the Db::stat() method.

  If the database contains multiple databases and the **-s** flag is not specified, the statistics are for the internal database that describes the other databases the file contains, and not for the file as a whole.

- **-E**

  Display detailed information about the database environment.

- **-e**

Display information about the database environment, including all configured subsystems of the database environment.

- **-f**

  Display only those database statistics that can be acquired without traversing the database.

- **-h**

  Specify a home directory for the database environment; by default, the current working directory is used.

- **-l**

  Display logging subsystem statistics, as described in the DbEnv::log_stat() method.

- **-L**

  Display all logging subsystem statistics.

  - **A**

    Display all information.

- **-M**

  Display detailed information about the cache.

  - **A**

    Display all information.

  - **h**

    Display buffers within hash chains.

- **-m**

  Display cache statistics, as described in the DbEnv::memp_stat() method.

- **-N**

  Do not acquire shared region mutexes while running. Other problems, such as potentially fatal errors in Berkeley DB, will be ignored as well. This option is intended only for debugging errors, and should not be used under any other circumstances.

- **-P**

  Specify an environment password. Although Berkeley DB utilities overwrite password strings as soon as possible, be aware there may be a window of vulnerability on systems where unprivileged users can see command-line arguments or where utilities are not able to overwrite the memory containing the command-line arguments.

- **-R**

  Display detailed information about the replication subsystem.

  - **A**

    Display all information.

- **-r**

  Display replication statistics, as described in in the DbEnv::rep_stat() method.

- **-s**

  Display statistics for the specified database contained in the file specified with the **-d** flag.

- **-t**

  Display transaction subsystem statistics, as described in the DbEnv::txn_stat() method.

- **-V**

  Write the library version number to the standard output, and exit.

- **-x**

  Display mutex subsystem statistics, as described in the DbEnv::mutex_stat() method.

- **-Z**

  Reset the statistics after reporting them; valid only with the **-C**, **-c**, **-E**, **-e**, **-L**, **-l**, **-M**, **-m**, **-R**, **-r**, and **-t** options.

Values normally displayed in quantities of bytes are displayed as a combination of gigabytes (GB), megabytes (MB), kilobytes (KB), and bytes (B). Otherwise, values smaller than 10 million are displayed without any special notation, and values larger than 10 million are displayed as a number followed by "M".

The **db_stat** utility may be used with a Berkeley DB environment (as described for the **-h** option, the environment variable **DB_HOME**, or because the utility was run in a directory containing a Berkeley DB environment). In order to avoid environment corruption when using a Berkeley DB environment, **db_stat** should always be given the chance to detach from the environment and exit gracefully. To cause **db_stat** to release all environment resources and exit cleanly, send it an interrupt signal (SIGINT).

The **db_stat** utility exits 0 on success, and >0 if an error occurs.

Environment Variables

DB_HOME

If the **-h** option is not specified and the environment variable DB_HOME is set, it is used as the path of the database home, as described in the DbEnv::open() method.

db_upgrade

```
db_upgrade [-NsVv] [-h home] [-P password] file ...
```

The **db_upgrade** utility upgrades the Berkeley DB version of one or more files and the databases they contain to the current release version.

The options are as follows:

- **-h**

  Specify a home directory for the database environment; by default, the current working directory is used.

- **-N**

  Do not acquire shared region mutexes while running. Other problems, such as potentially fatal errors in Berkeley DB, will be ignored as well. This option is intended only for debugging errors, and should not be used under any other circumstances.

- **-P**

  Specify an environment password. Although Berkeley DB utilities overwrite password strings as soon as possible, be aware there may be a window of vulnerability on systems where unprivileged users can see command-line arguments or where utilities are not able to overwrite the memory containing the command-line arguments.

- **-s**

  This flag is only meaningful when upgrading databases from releases before the Berkeley DB 3.1 release.

  As part of the upgrade from the Berkeley DB 3.0 release to the 3.1 release, the on-disk format of duplicate data items changed. To correctly upgrade the format requires that applications specify whether duplicate data items in the database are sorted or not. Specifying the **-s** flag means that the duplicates are sorted; otherwise, they are assumed to be unsorted. Incorrectly specifying the value of this flag may lead to database corruption.

  Because the **db_upgrade** utility upgrades a physical file (including all the databases it contains), it is not possible to use **db_upgrade** to upgrade files where some of the databases it includes have sorted duplicate data items, and some of the databases it includes have unsorted duplicate data items. If the file does not have more than a single database, if the databases do not support duplicate data items, or if all the databases that support duplicate data items support the same style of duplicates (either sorted or unsorted), **db_upgrade** will work correctly as long as the **-s** flag is correctly specified. Otherwise, the file cannot be upgraded using **db_upgrade**, and must be upgraded manually using the db_dump and db_load utilities.

- **-V**

  Write the library version number to the standard output, and exit.

- **-v**

   Run in verbose mode, displaying a message for each successful upgrade.

It is important to realize that Berkeley DB database upgrades are done in place, and so are potentially destructive. This means that if the system crashes during the upgrade procedure, or if the upgrade procedure runs out of disk space, the databases may be left in an inconsistent and unrecoverable state. See Upgrading databases for more information.

The **db_upgrade** utility may be used with a Berkeley DB environment (as described for the **-h** option, the environment variable **DB_HOME**, or because the utility was run in a directory containing a Berkeley DB environment). In order to avoid environment corruption when using a Berkeley DB environment, **db_upgrade** should always be given the chance to detach from the environment and exit gracefully. To cause **db_upgrade** to release all environment resources and exit cleanly, send it an interrupt signal (SIGINT).

The **db_upgrade** utility exits 0 on success, and >0 if an error occurs.

Environment Variables

DB_HOME

If the **-h** option is not specified and the environment variable DB_HOME is set, it is used as the path of the database home, as described in the DbEnv::open() method.

db_verify

```
db_verify [-NoqV] [-h home] [-P password] file ...
```

The **db_verify** utility verifies the structure of one or more files and the databases they contain.

The options are as follows:

- **-h**

  Specify a home directory for the database environment; by default, the current working directory is used.

- **-o**

  Skip the database checks for btree and duplicate sort order and for hashing.

  If the file being verified contains databases with non-default comparison or hashing configurations, calling the **db_verify** utility without the **-o** flag will usually return failure. The **-o** flag causes **db_verify** to ignore database sort or hash ordering and allows **db_verify** to be used on these files. To fully verify these files, verify them explicitly using the Db::verify() method, after configuring the correct comparison or hashing functions.

- **-N**

  Do not acquire shared region mutexes while running. Other problems, such as potentially fatal errors in Berkeley DB, will be ignored as well. This option is intended only for debugging errors, and should not be used under any other circumstances.

- **-P**

  Specify an environment password. Although Berkeley DB utilities overwrite password strings as soon as possible, be aware there may be a window of vulnerability on systems where unprivileged users can see command-line arguments or where utilities are not able to overwrite the memory containing the command-line arguments.

- **-q**

  Suppress the printing of any error descriptions, simply exit success or failure.

- **-V**

  Write the library version number to the standard output, and exit.

The db_verify utility does not perform any locking, even in Berkeley DB environments that are configured with a locking subsystem. As such, it should only be used on files that are not being modified by another thread of control.

The **db_verify** utility may be used with a Berkeley DB environment (as described for the **-h** option, the environment variable **DB_HOME**, or because the utility was run in a directory containing a Berkeley DB environment). In order to avoid environment corruption when using a Berkeley DB environment, **db_verify** should always be given the chance to detach from the environment and exit gracefully. To

cause **db_verify** to release all environment resources and exit cleanly, send it an interrupt signal (SIGINT).

The **db_verify** utility exits 0 on success, and >0 if an error occurs.

Environment Variables

DB_HOME

If the **-h** option is not specified and the environment variable DB_HOME is set, it is used as the path of the database home, as described in the DbEnv::open() method.